[EVOP2016]
Introduction to **Visual Data Analysis**

Prof Jan Aerts
Visual Data Analysis lab, ESAT/STADIUS
Faculty of Engineering
KU Leuven

@jandot - jan.aerts@kuleuven.be - http://vda-lab.be

Overview

A. **Why visual analytics?**

B. Data visualization

- Data foundations
- Human perception foundations
- Visualization foundations and examples

C. Visualization evaluation

D. Tools of the trade

A. What’s the problem?

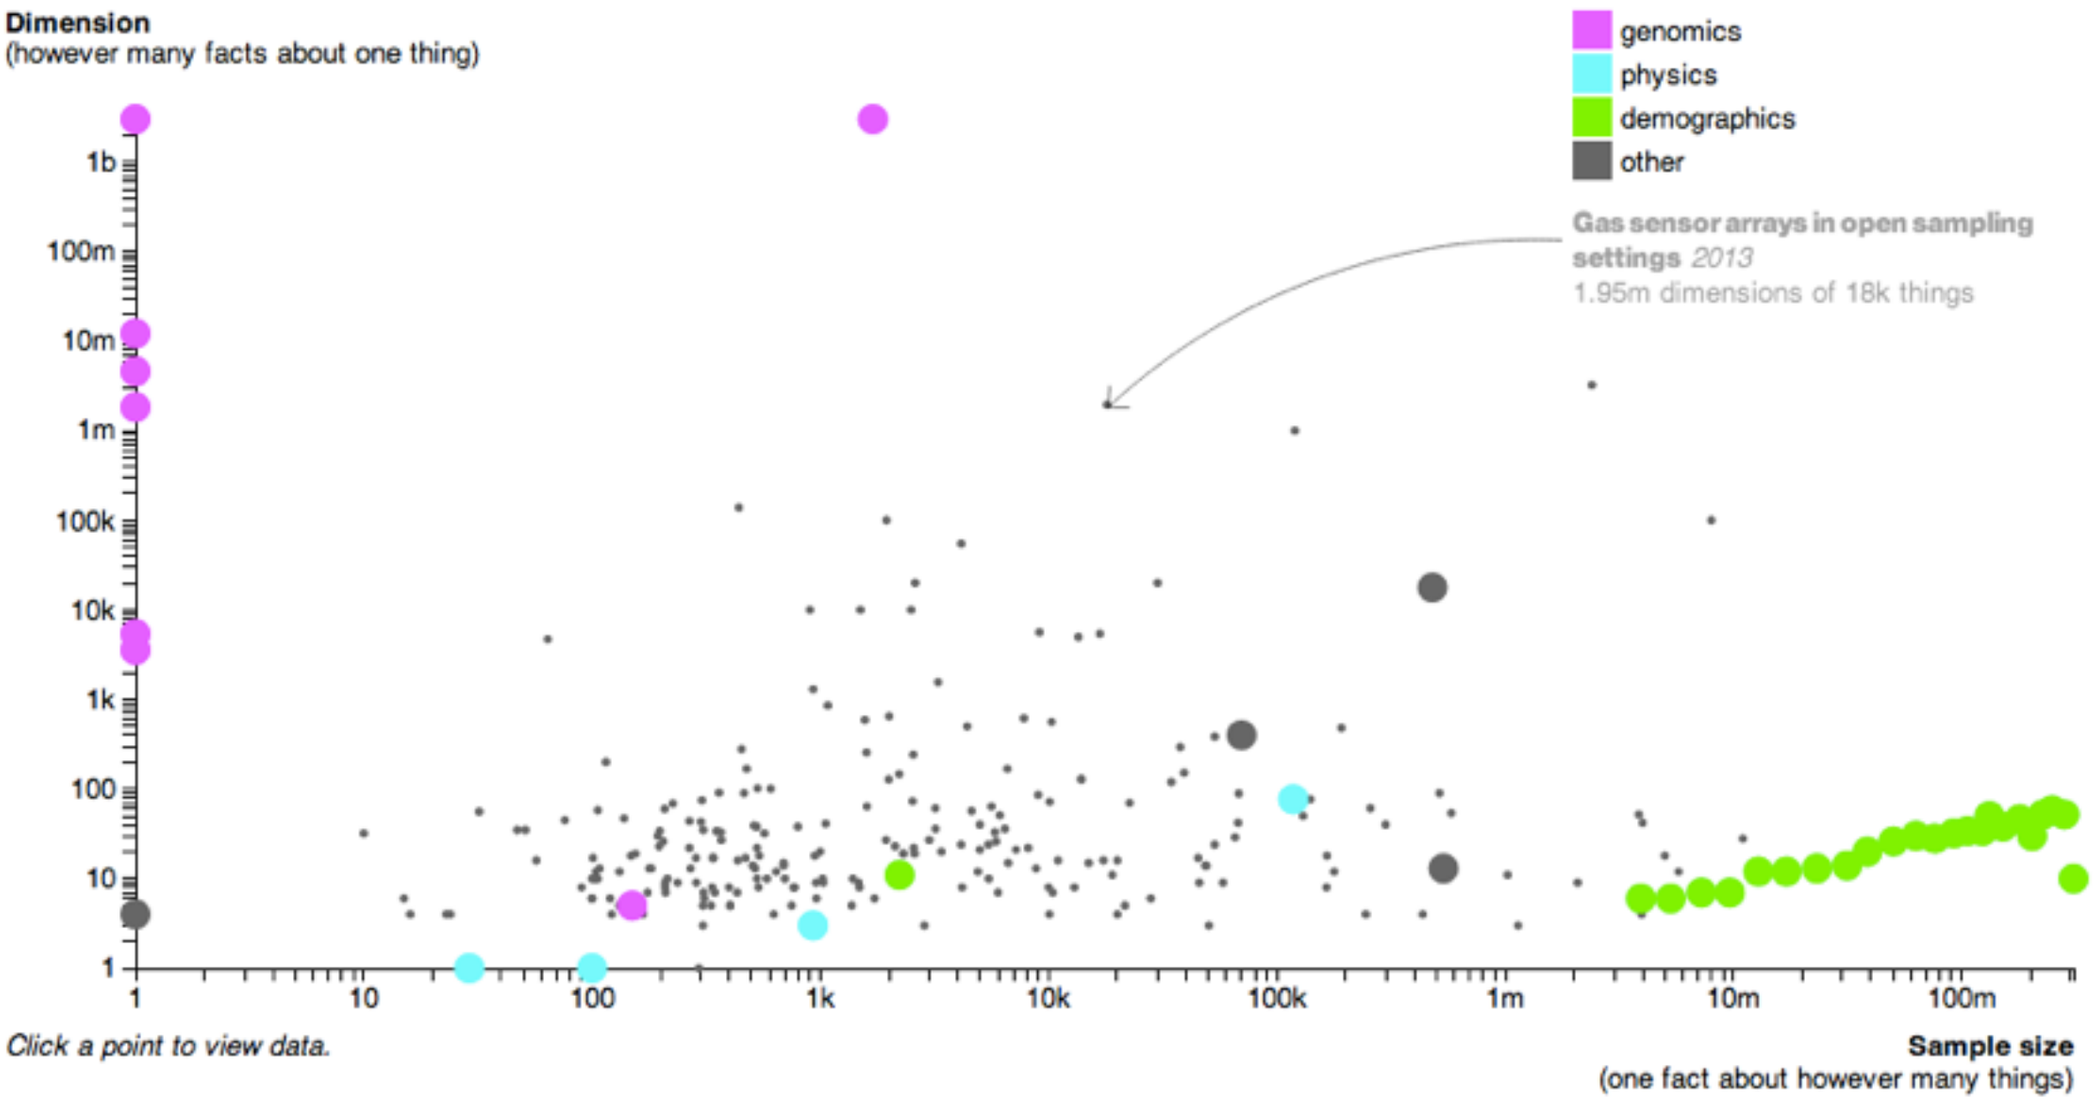
Dimension
(however many facts about one thing)
1b
100m
10m
1m
100k
10k
1k
100
10
1
1
10
100
1k
10k
100k
1m
10m
100m
genomics
physics
demographics
other
Gas sensor arrays in open sampling settings 2013
1.95m dimensions of 18k things
Click a point to view data.
Sample size
(one fact about however many things)

hypothesis-driven -> data-driven
hunting down unknown unknowns

Scientific Research Paradigms (Jim Gray, Microsoft)

1st	1,000s years ago	empirical
2nd	100s years ago	theoretical
3rd	last few decades	computational
4rd	today	data exploration

then:

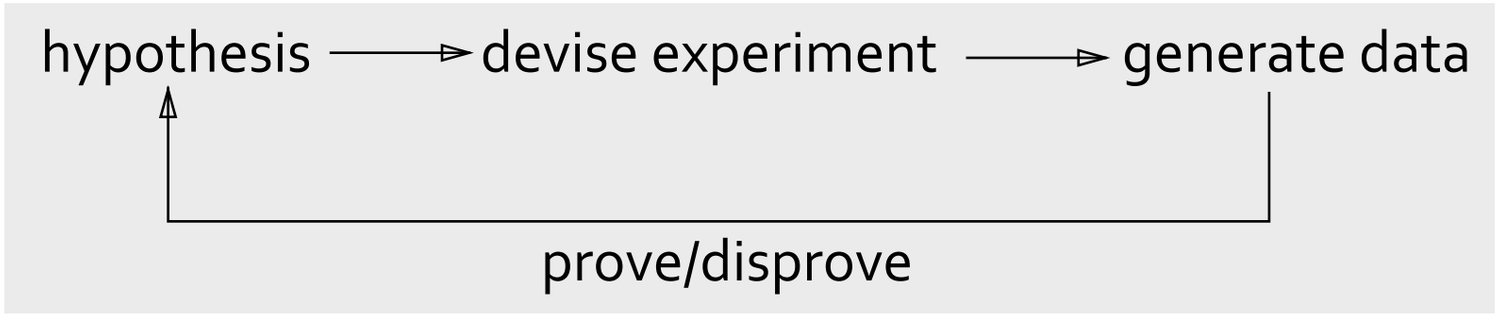

now:

have data → what's my hypothesis?

Data analysis
moving from
hypothesis-first
to
data-first

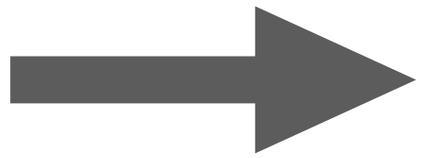

Challenge
moving from
finding the right answer to a question
to
finding the right question given the data

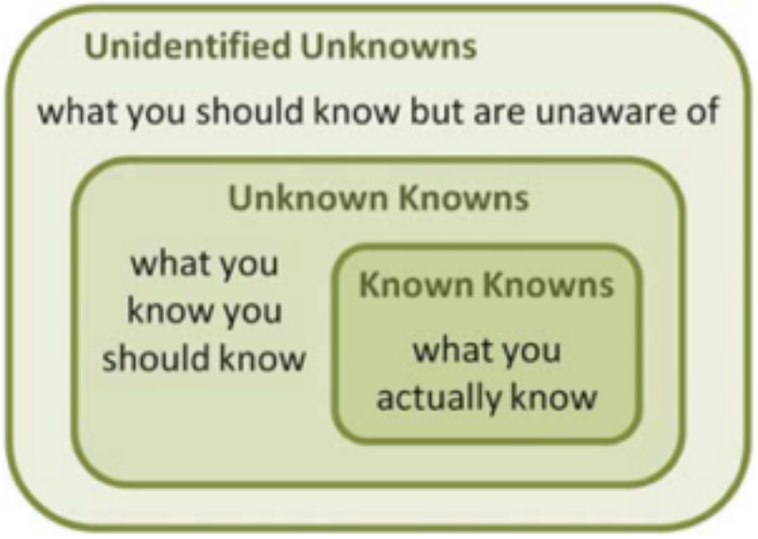

known unknowns	unknown unknowns
automated analysis (machine learning, statistics, …)	visual data analysis
replace the human	augment the human

Opening the black box

complex algorithms
in data analysis

obvious **link between input and output** difficult to see

user needs to **blindly trust** data analyst (and data analyst needs to trust their own skills)

result:
algorithm acts as a **black box**

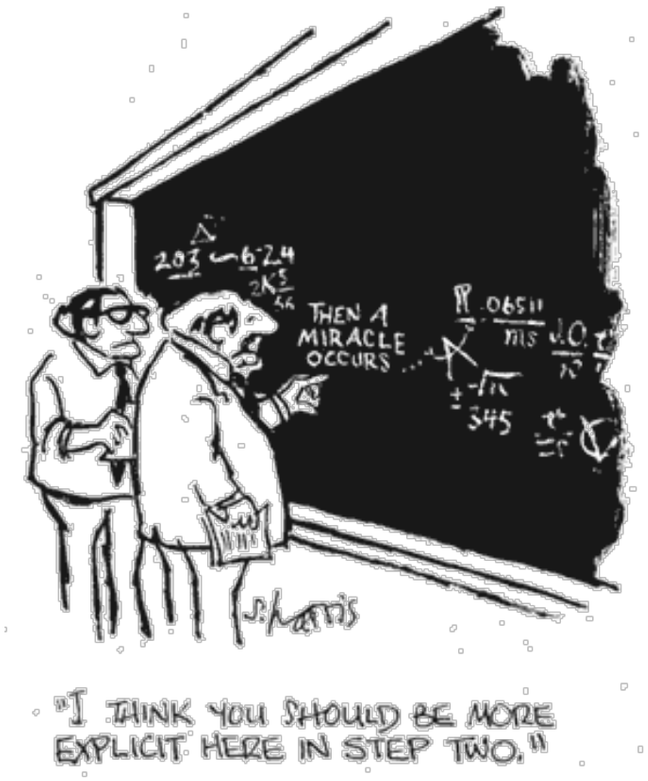

insight
in data
in algorithms

high-dimensional data
cell line + treatment + target + compound + …

intermediate feedback: e.g. intermediate result, goodness-of-fit for best subset so far

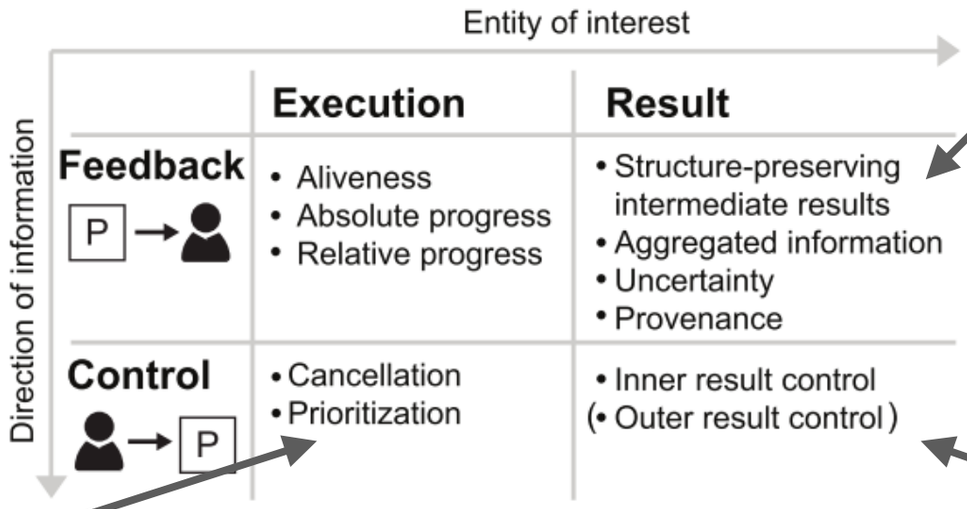

Muhlbacher et al, 2014

steer the final result: early validation of intermediate results, guided feature selection, weighting, avoiding being stuck in local minima, …

prioritization of remaining work: alter sequence of intermediate results to generate presumably more interesting ones earlier

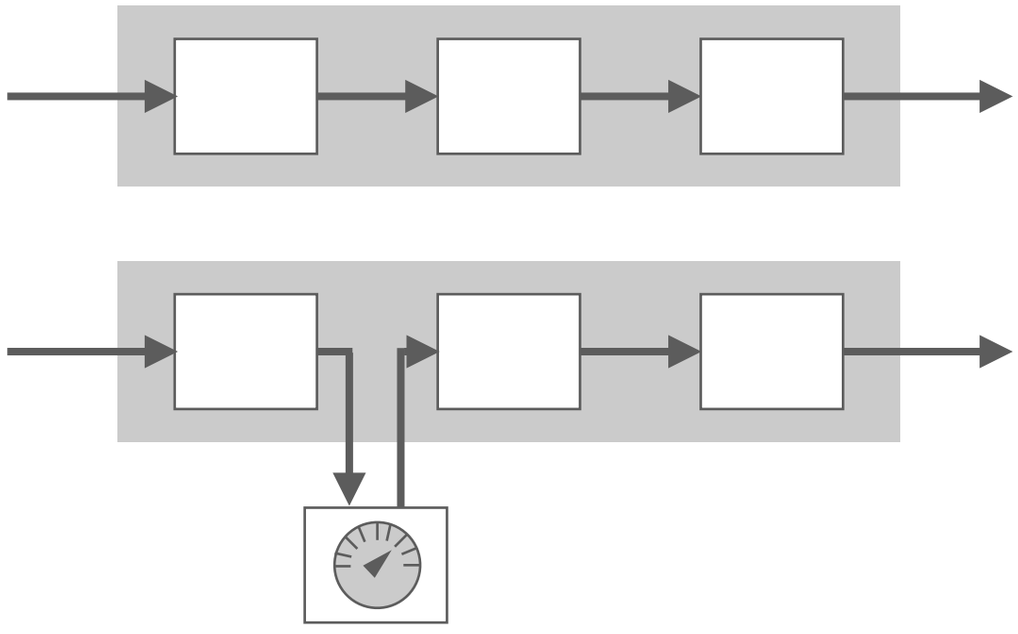

Example: mutation filtering

going from 5-6 million to 20,000

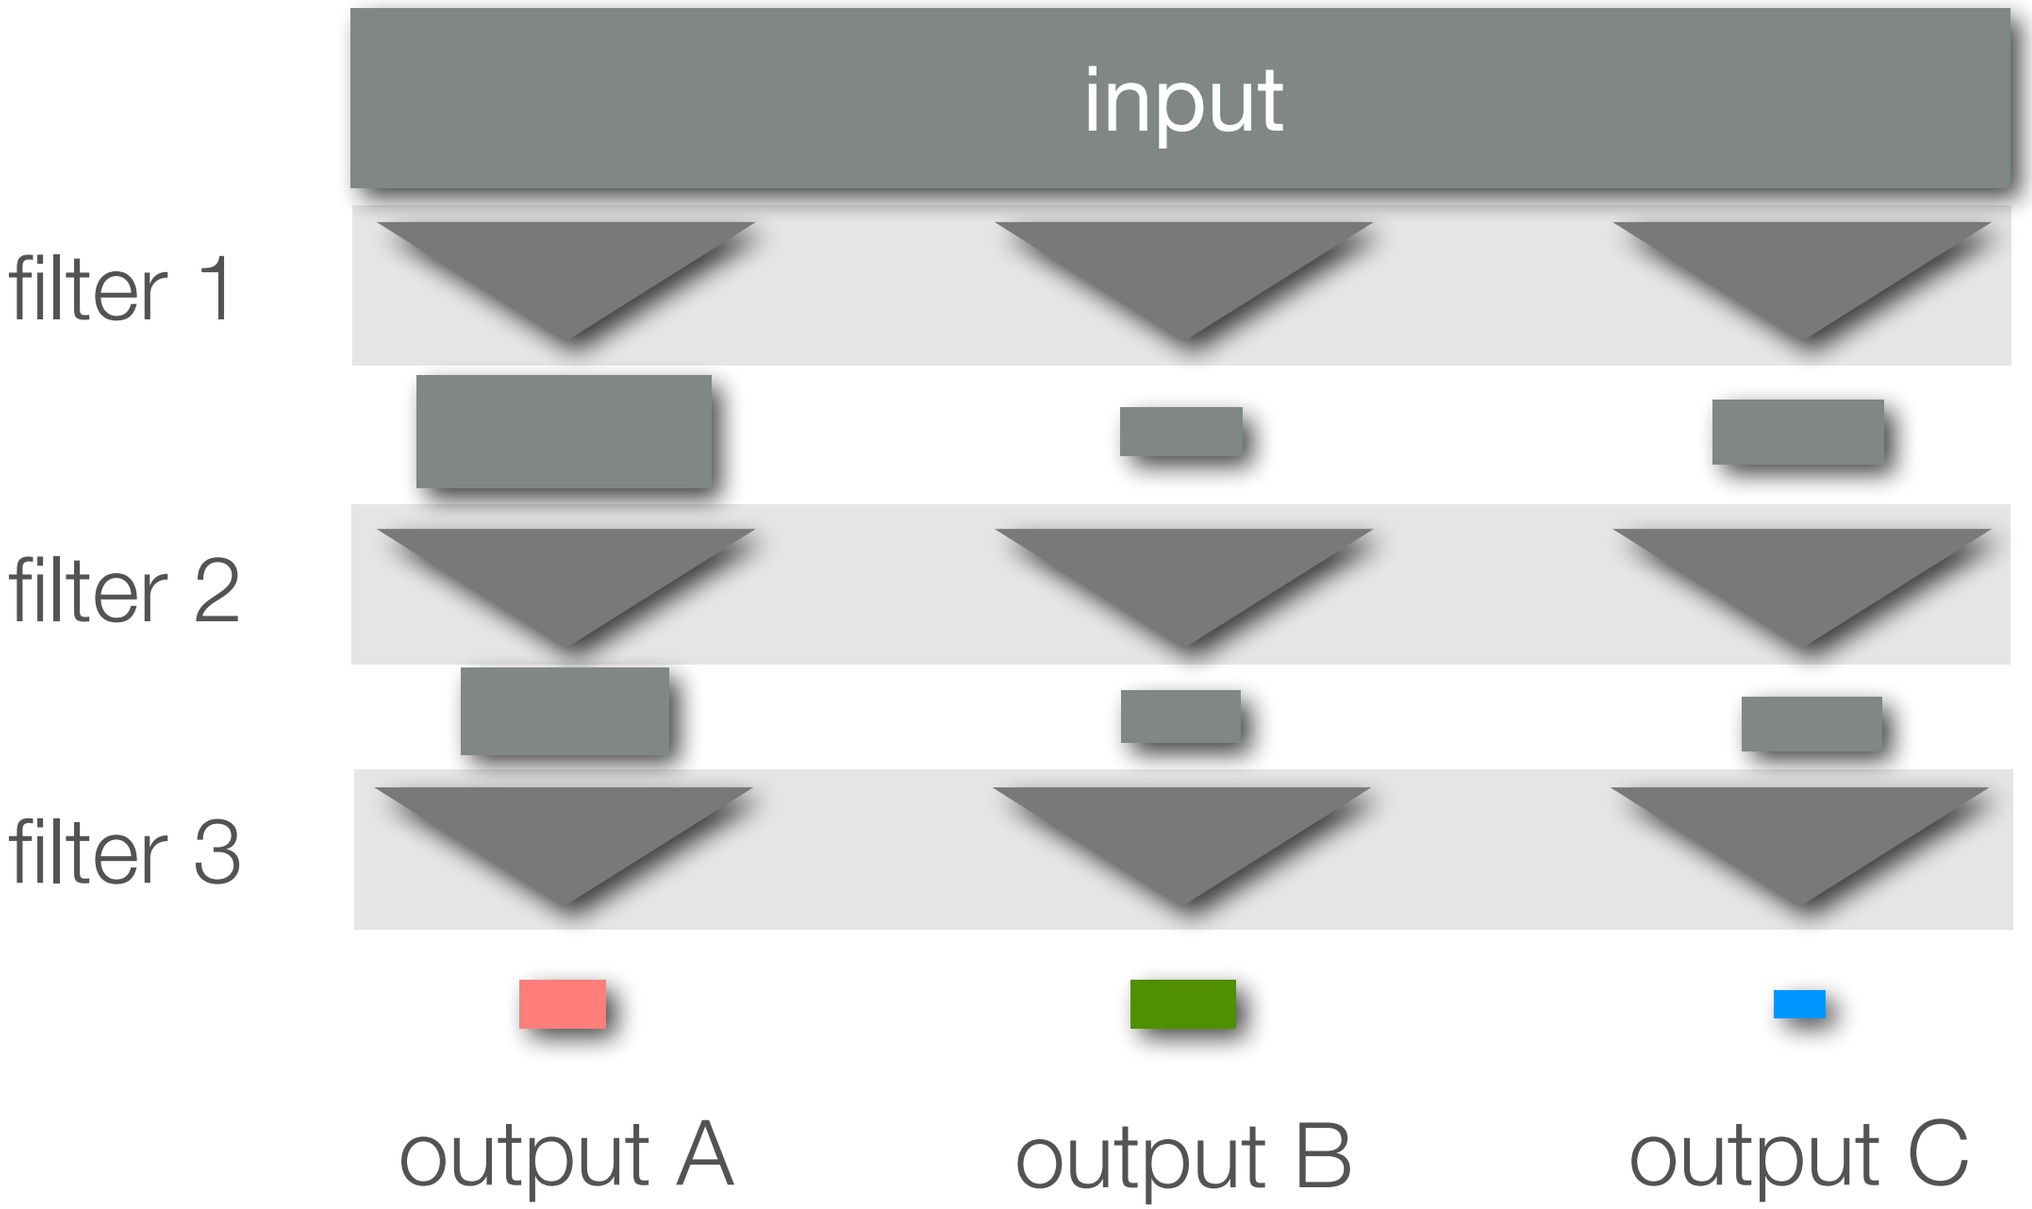

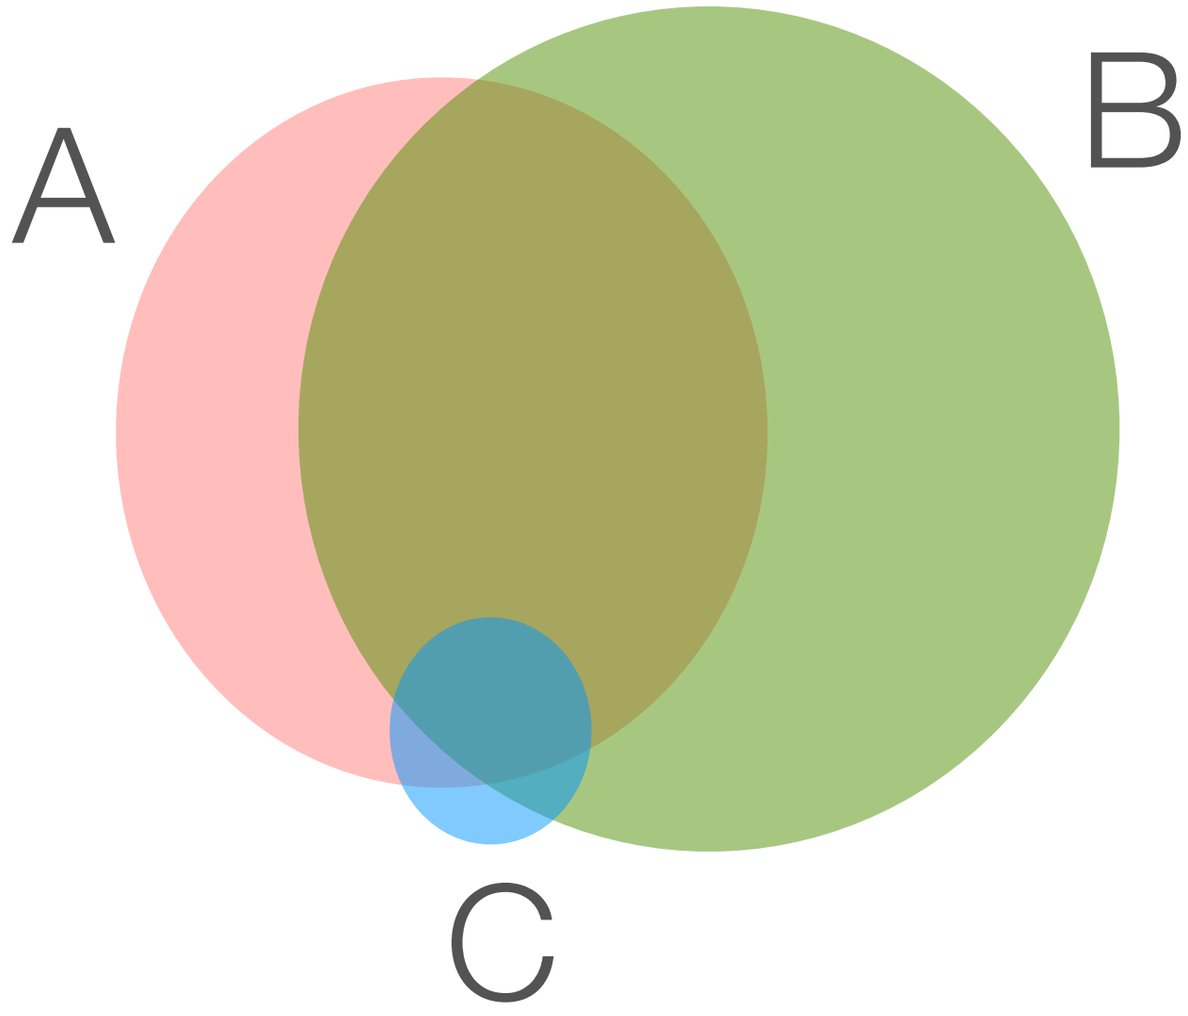
A
B
C

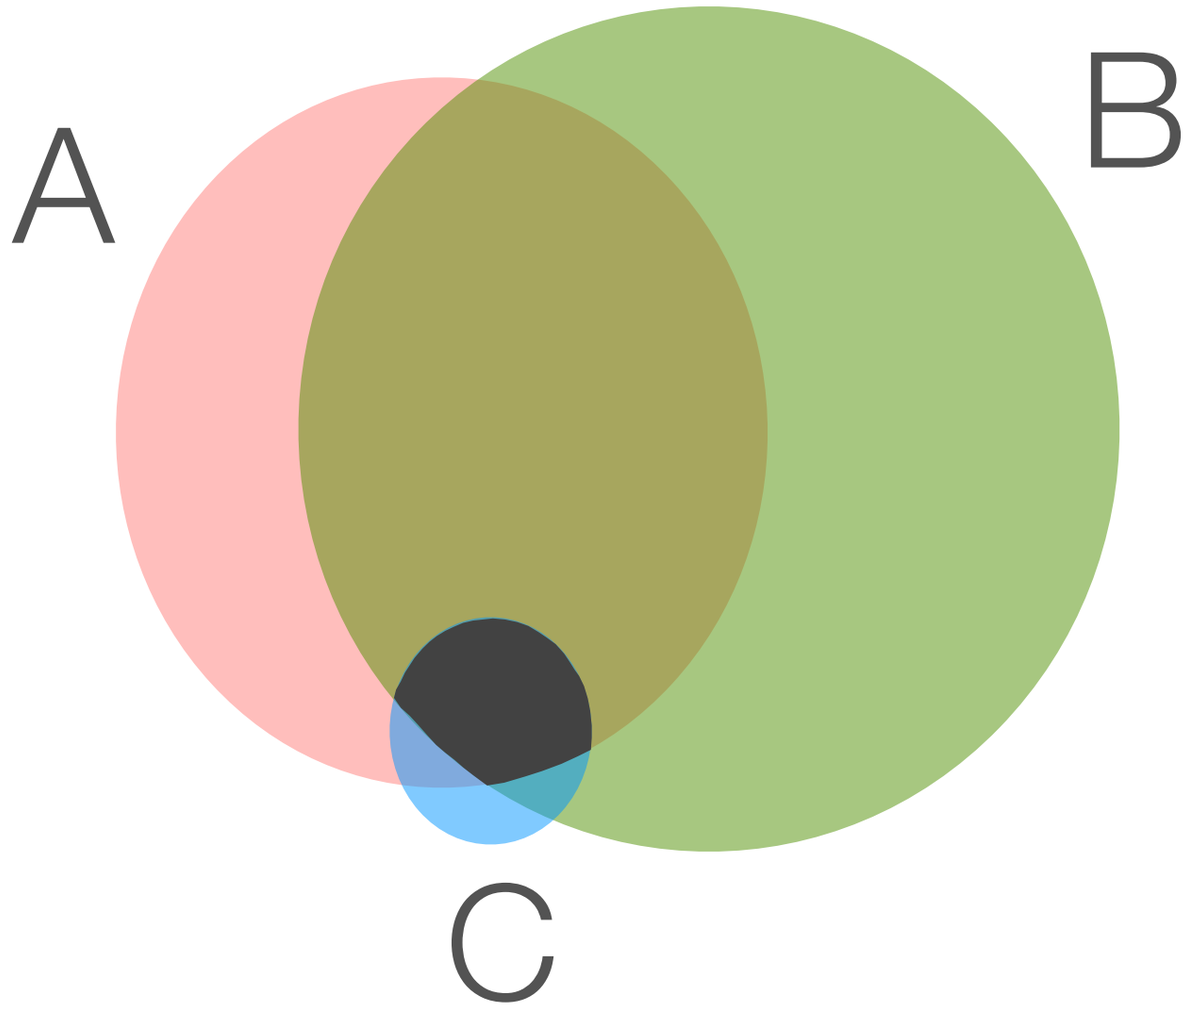
A
B
C

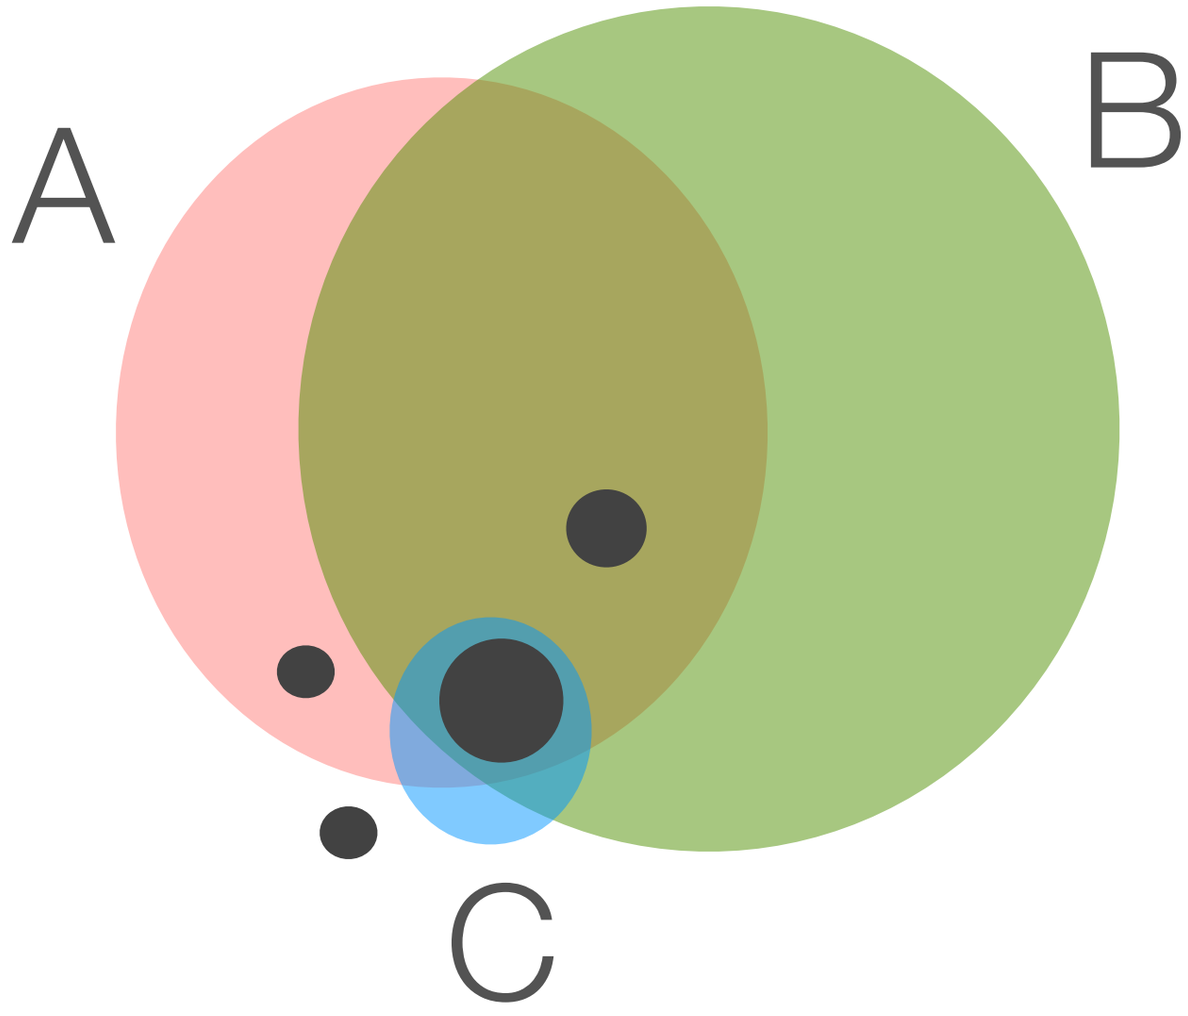
A
B
C

Visual Analytics to the rescue

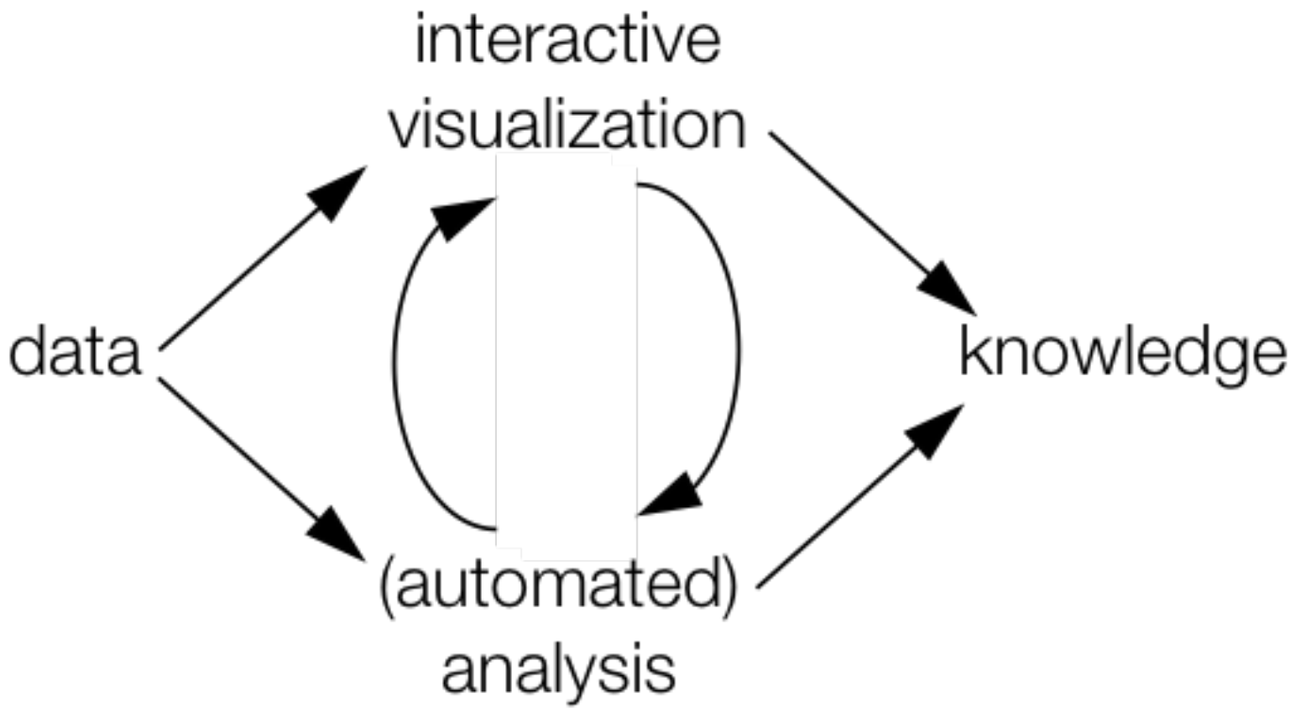

Avoid…

all taken from Genome Research

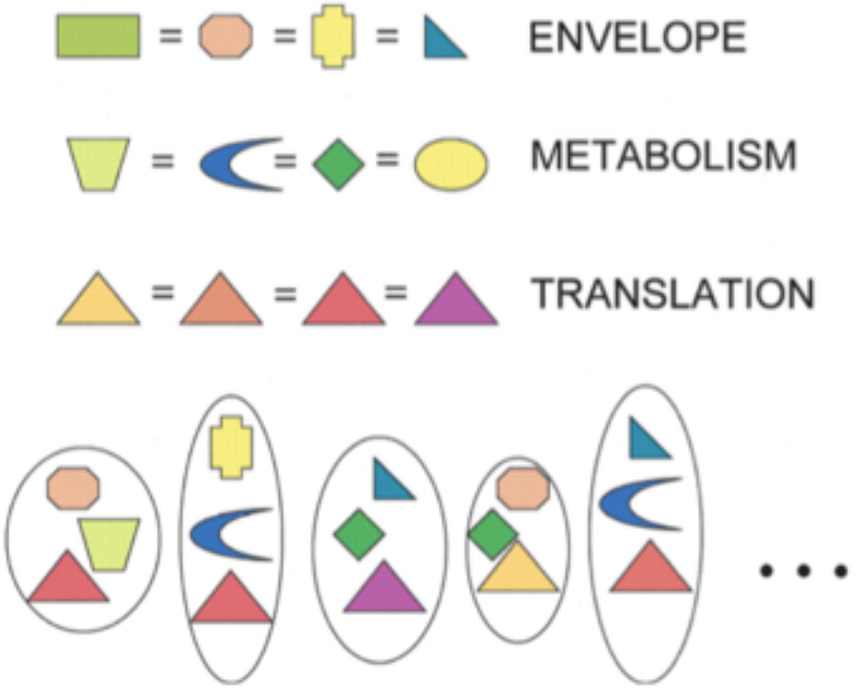

Charlebois & Doolittle, 2004

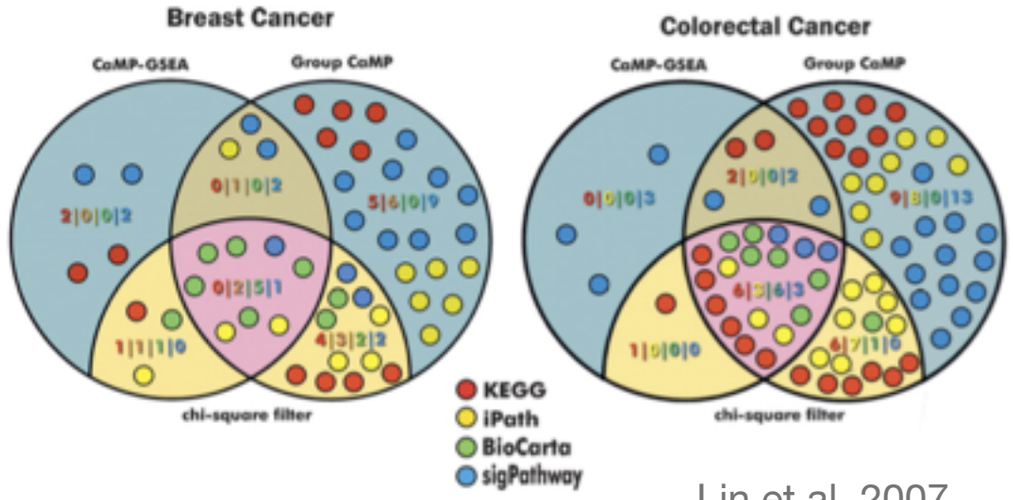

Lin et al, 2007

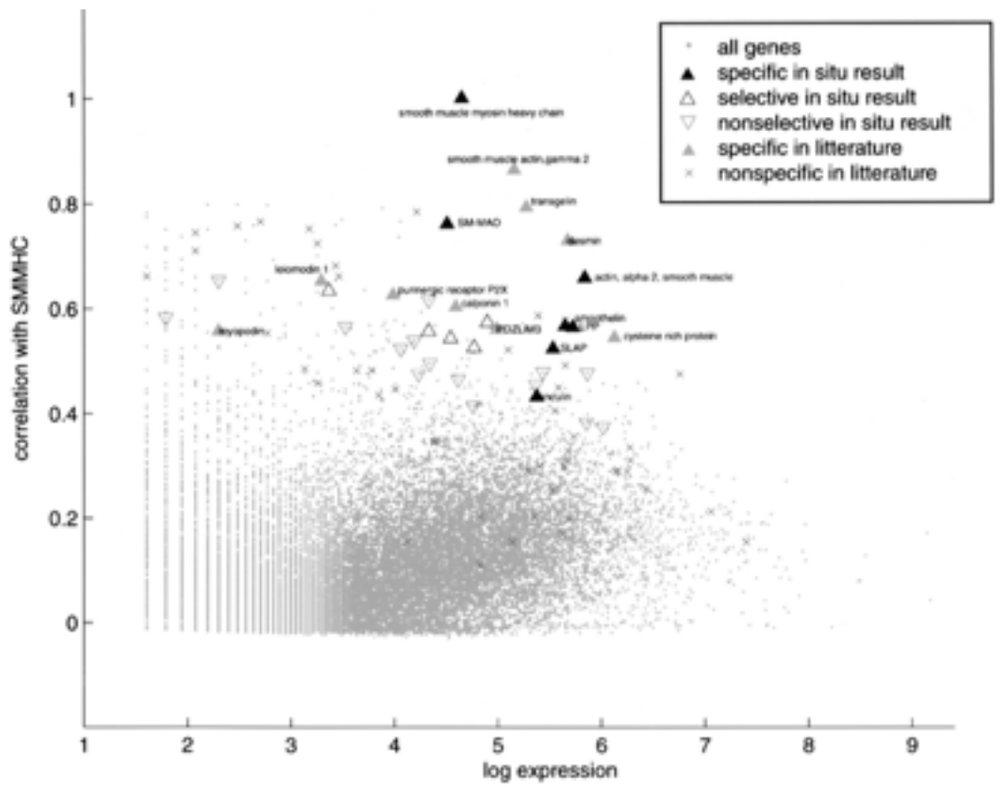

Nelander et al, 2003

Overview

A. Why visual analytics?

B. Data visualization

- Data foundations
- Human perception foundations
- Visualization foundations and examples

C. Visualization evaluation

D. Tools of the trade

B. Data Visualization

Historical perspective

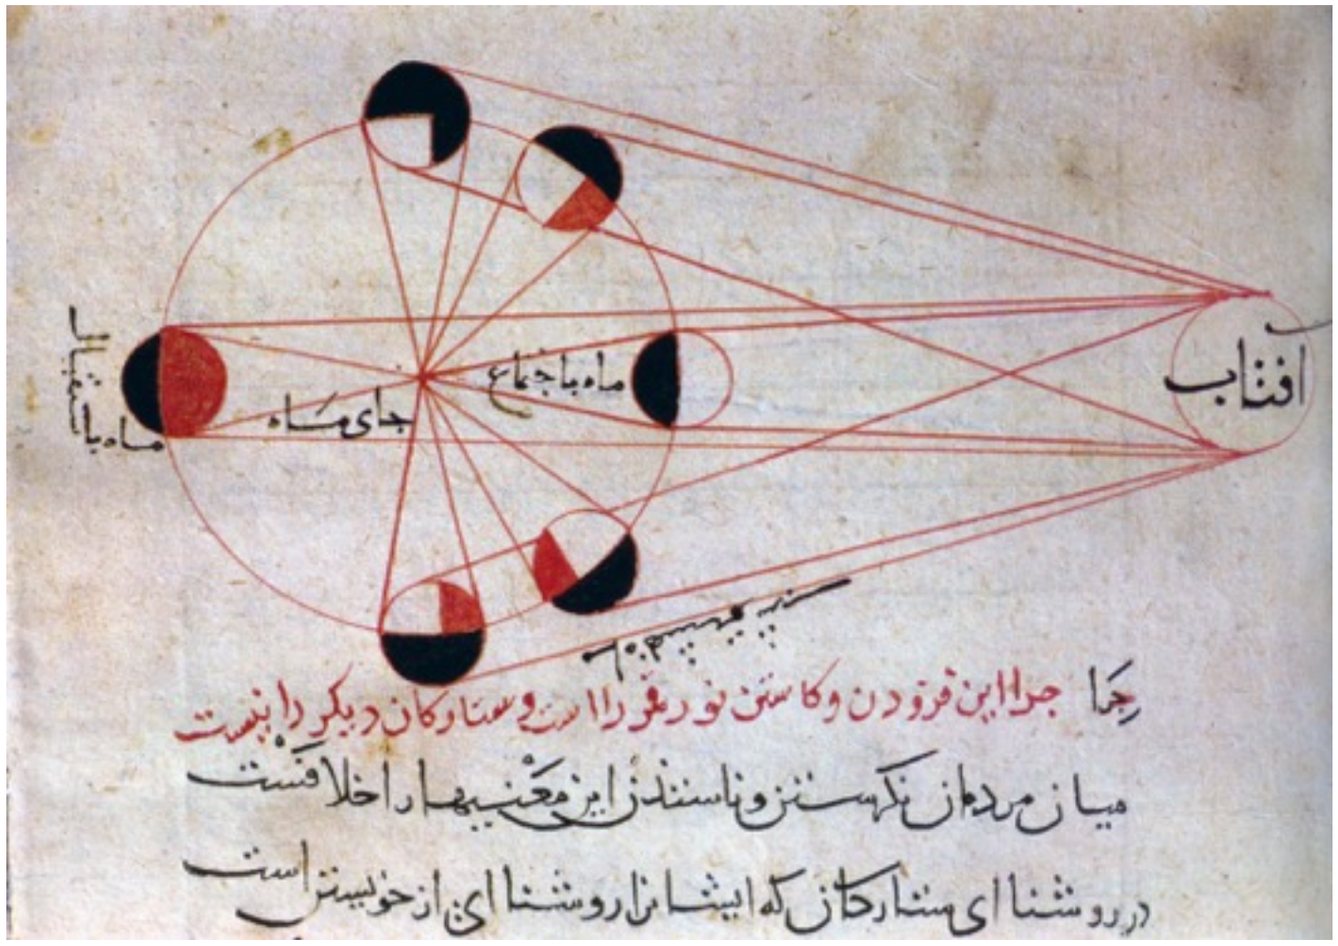

Al-Biruni - time series visualization: phases of the moon in orbit (circa 1030)

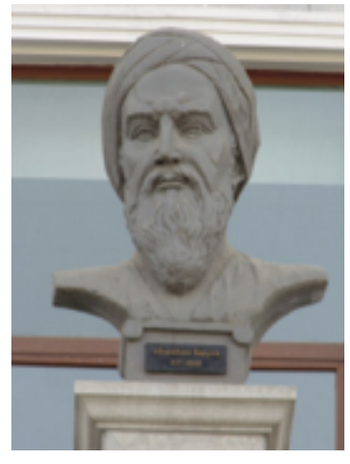

Hereford map - largest surviving map of the Middle Ages (1280s)

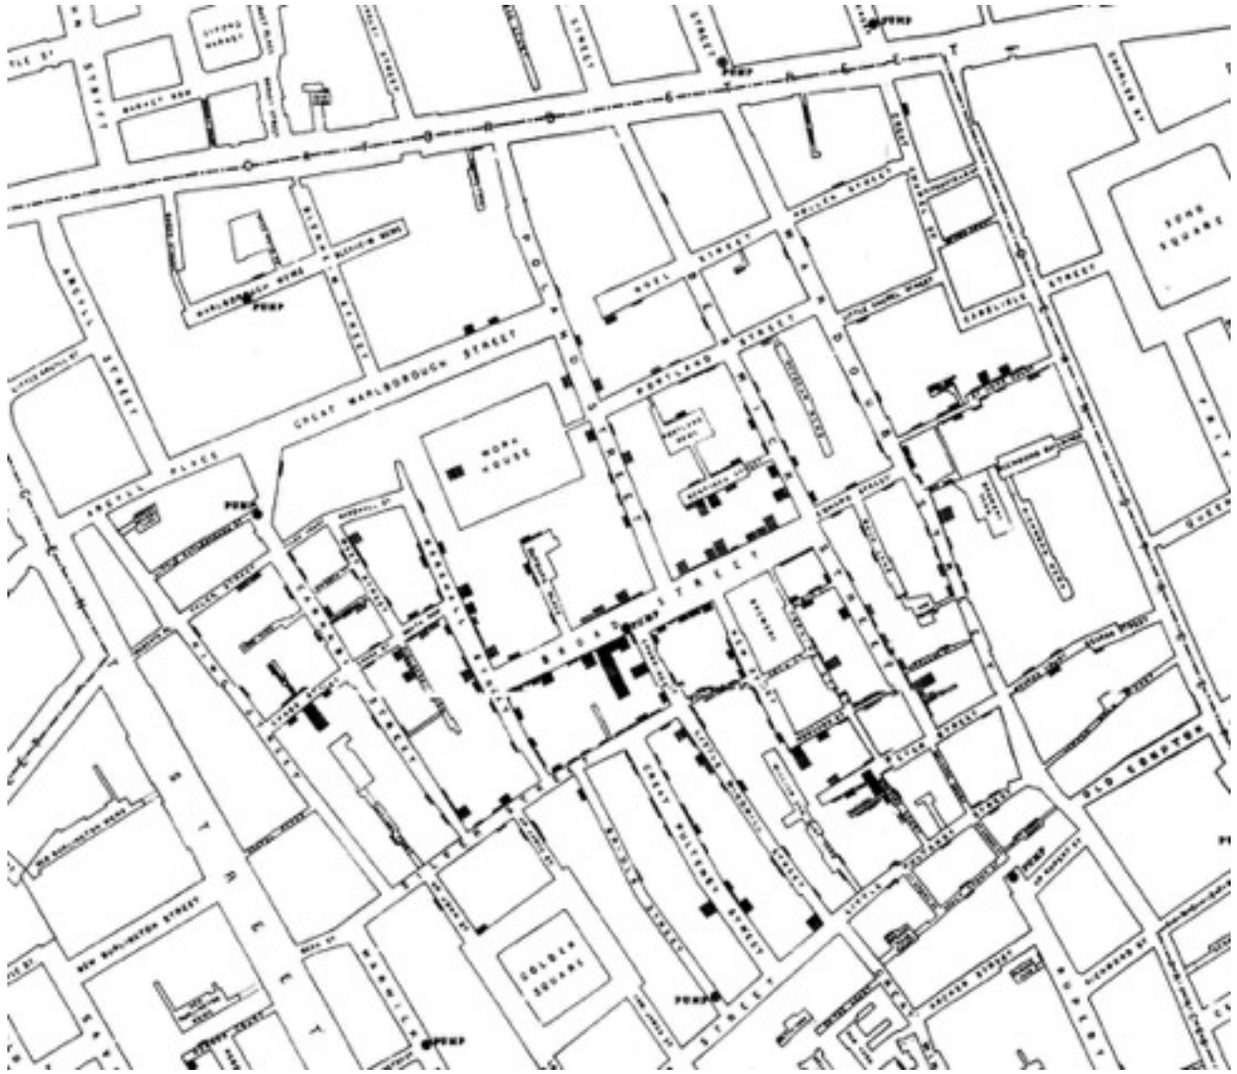

John Snow - cases of cholera in London (1663)

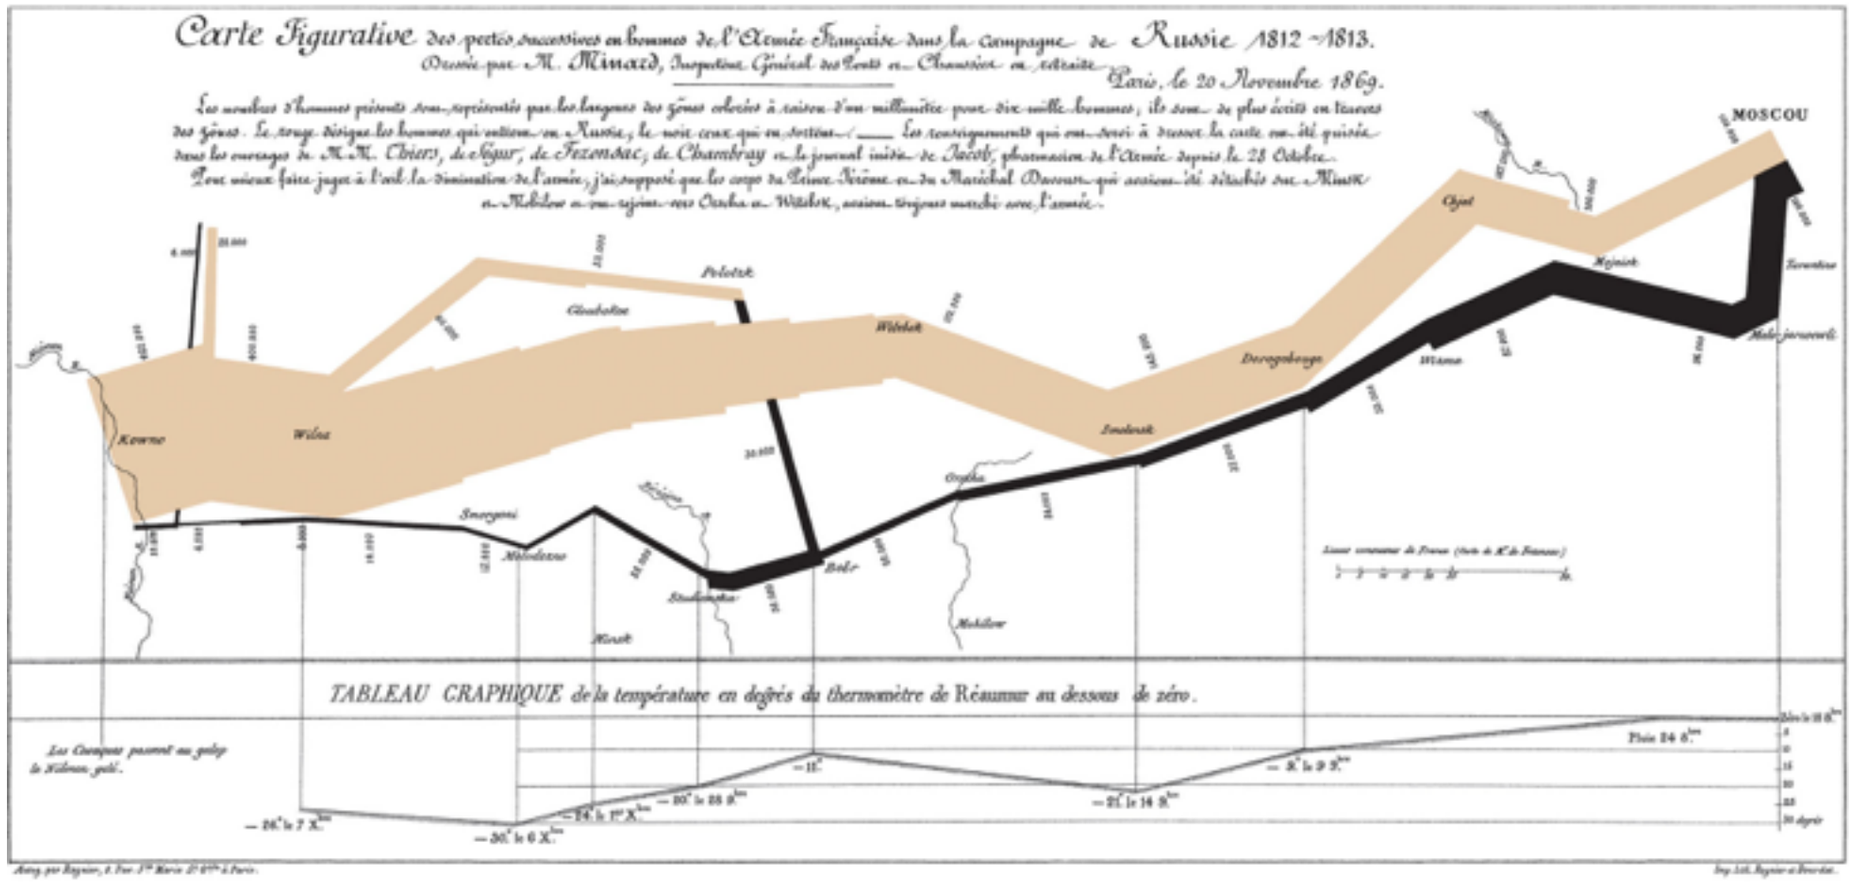

Charles Joseph Minard (1781-1870) - Napoleon's march on Moscow

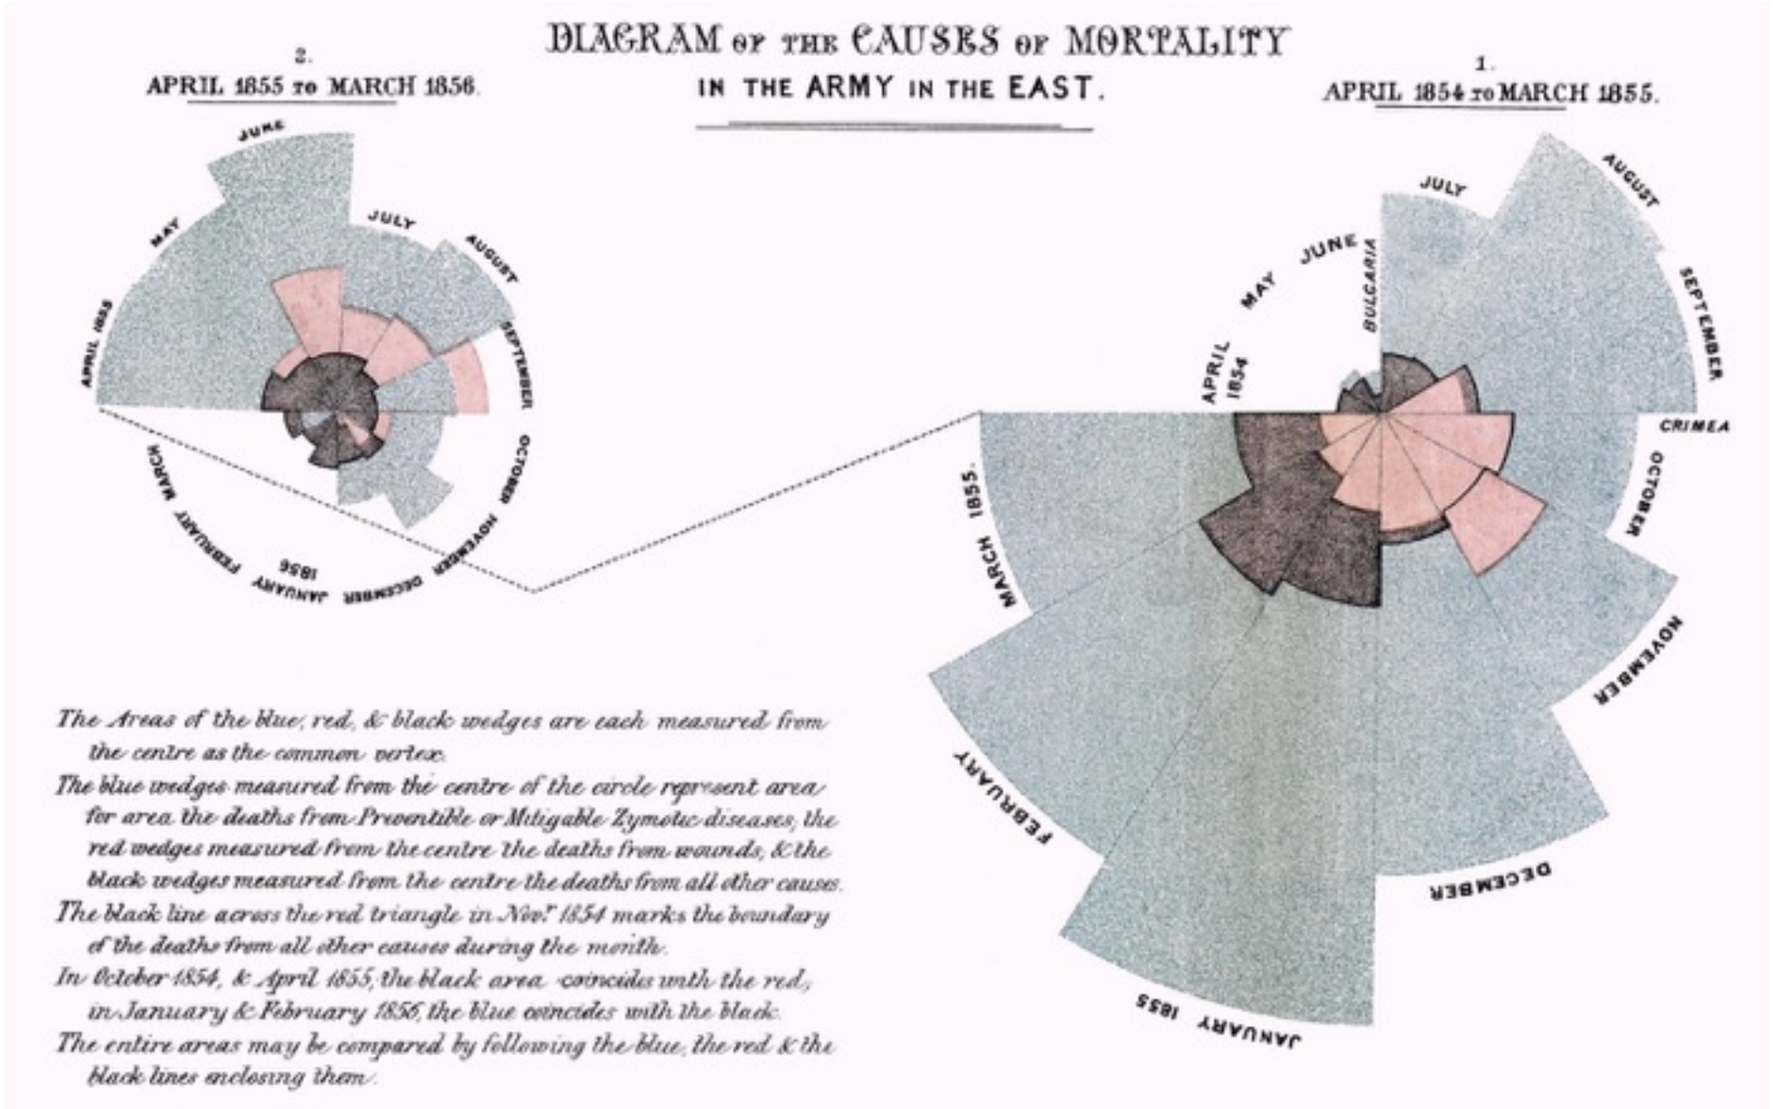

Florence Nightingale (1820-1910)
coxcomb chart monthly deaths from battle and other causes

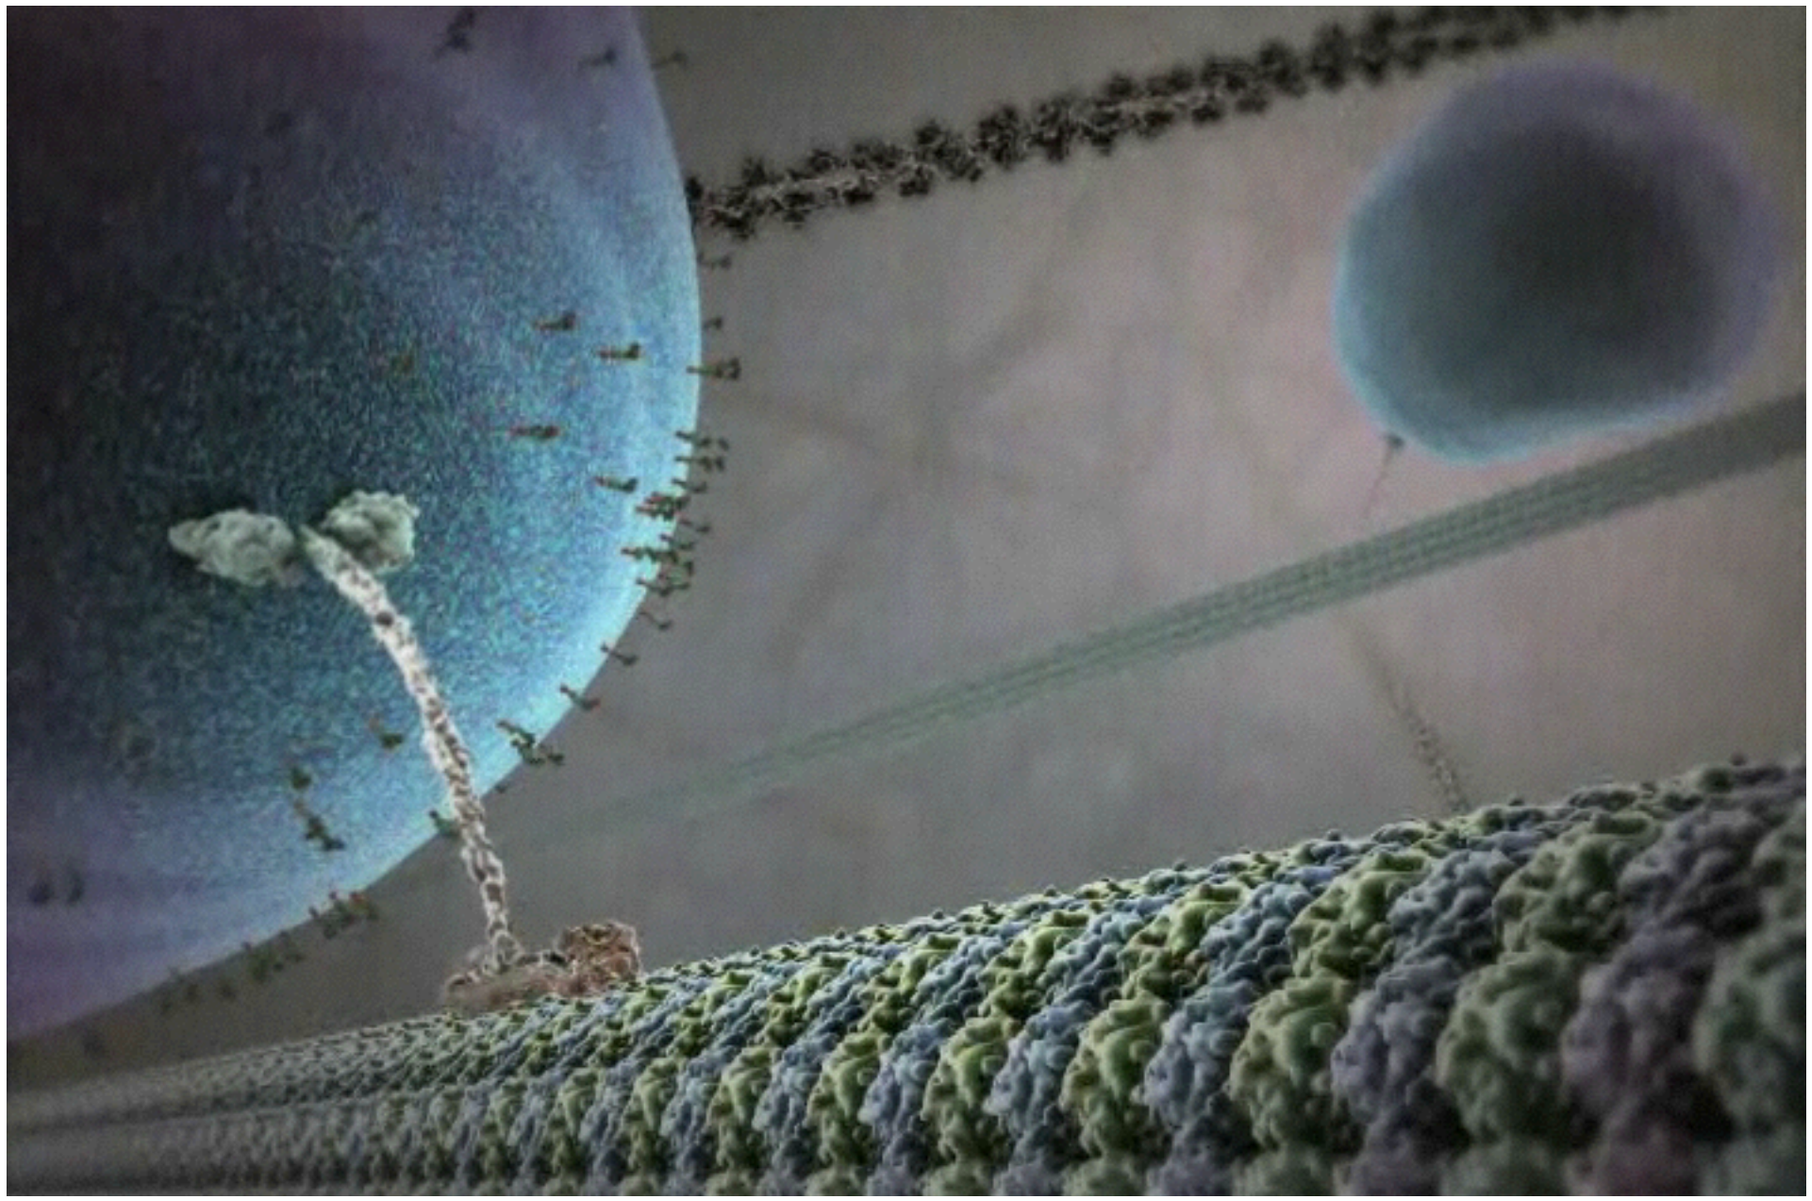

http://multimedia.mcb.harvard.edu/anim_innerlife.html

What is data visualization?

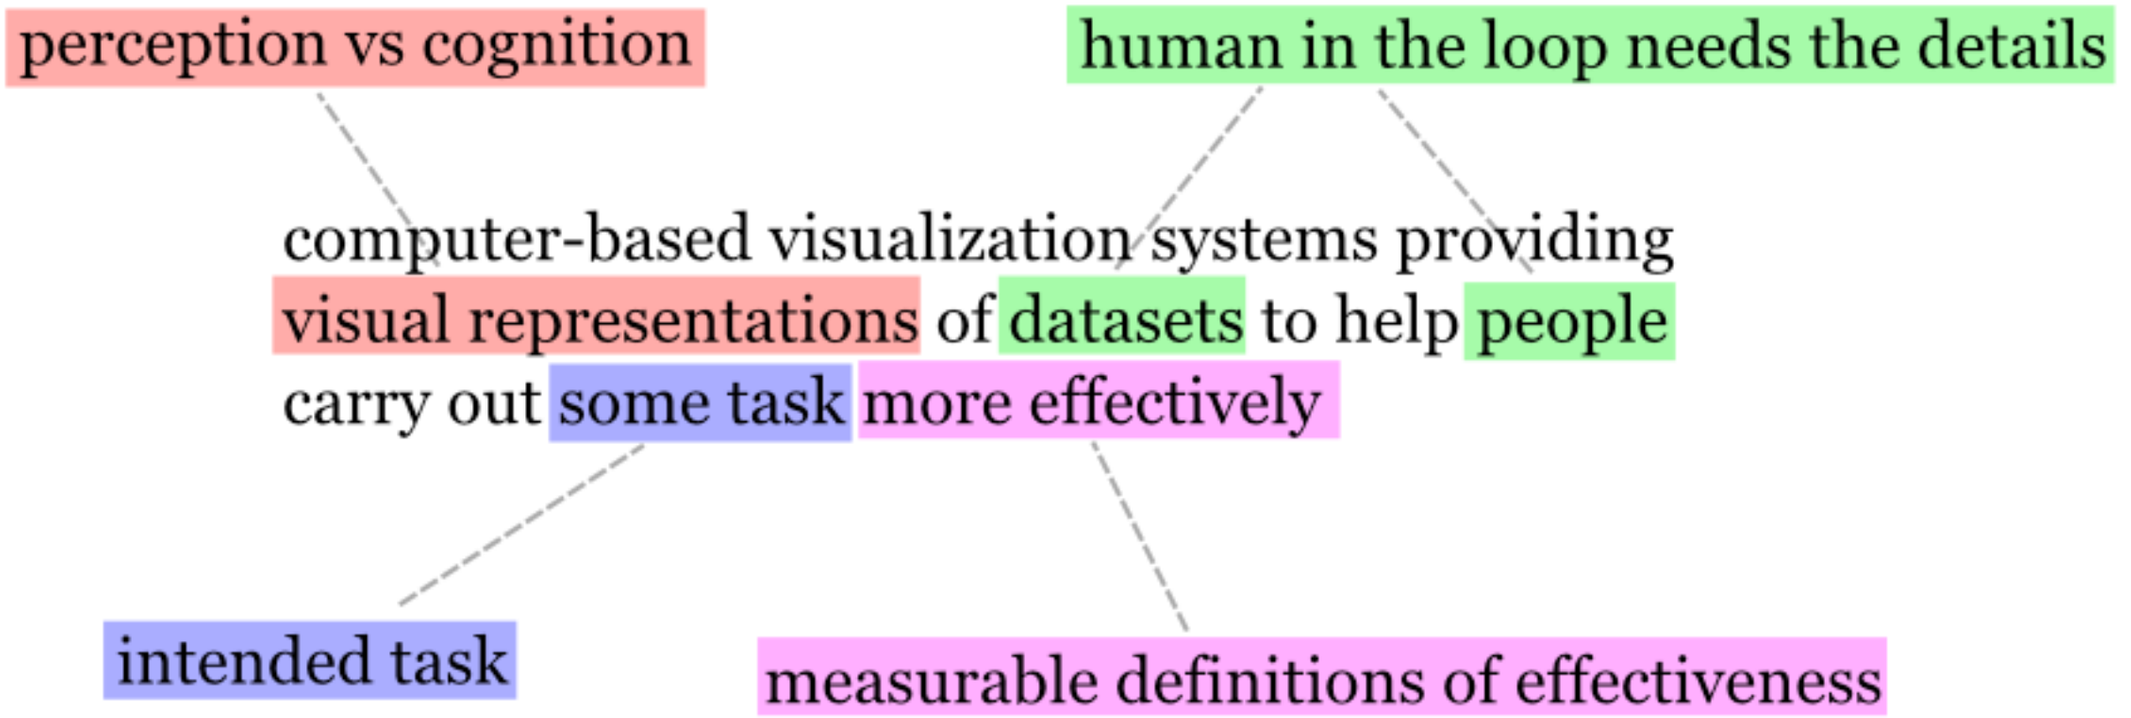
perception vs cognition
human in the loop needs the details
computer-based visualization systems providing visual representations of datasets to help people carry out some task more effectively
intended task
measurable definitions of effectiveness

T. Munzner

perception vs cognition

human in the loop needs the details

cognition <=> perception

cognitive task => perceptive task

identify anomalies, clusters, trends

T. Munzner

I		II		III	
x	y	x	y	x	y
10.0	8.04	10.0	9.14	10.0	7.46
8.0	6.95	8.0	8.14	8.0	6.77
13.0	7.58	13.0	8.74	13.0	12.74
9.0	8.81	9.0	8.77	9.0	7.11
11.0	8.33	11.0	9.26	11.0	7.81
14.0	9.96	14.0	8.10	14.0	8.84
6.0	7.24	6.0	6.13	6.0	6.08
4.0	4.26	4.0	3.10	4.0	5.39
12.0	10.84	12.0	9.13	12.0	8.15
7.0	4.82	7.0	7.26	7.0	6.42
5.0	5.68	5.0	4.74	5.0	5.73

n = 11

mean x = 9.0
mean y = 7.5

variance x = 11.0
variance y = 4.12

correlation x & y = 0.816
regression line: y = 3+0.5x

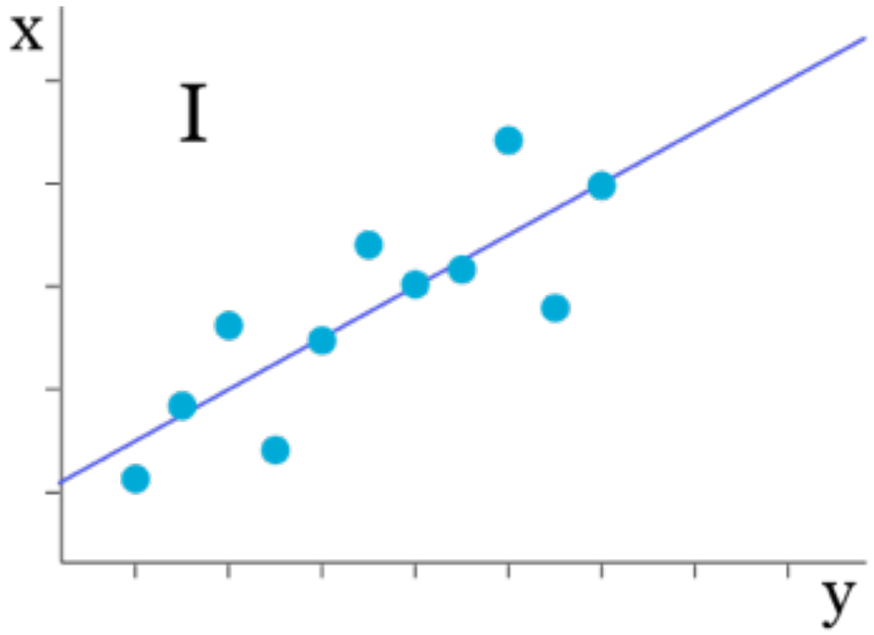
x
I
y

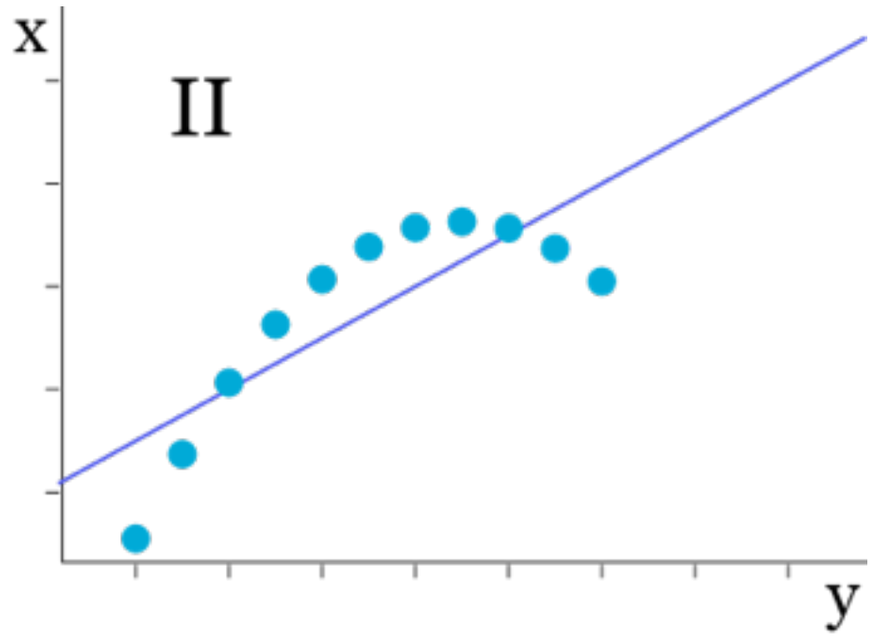
x
II
y

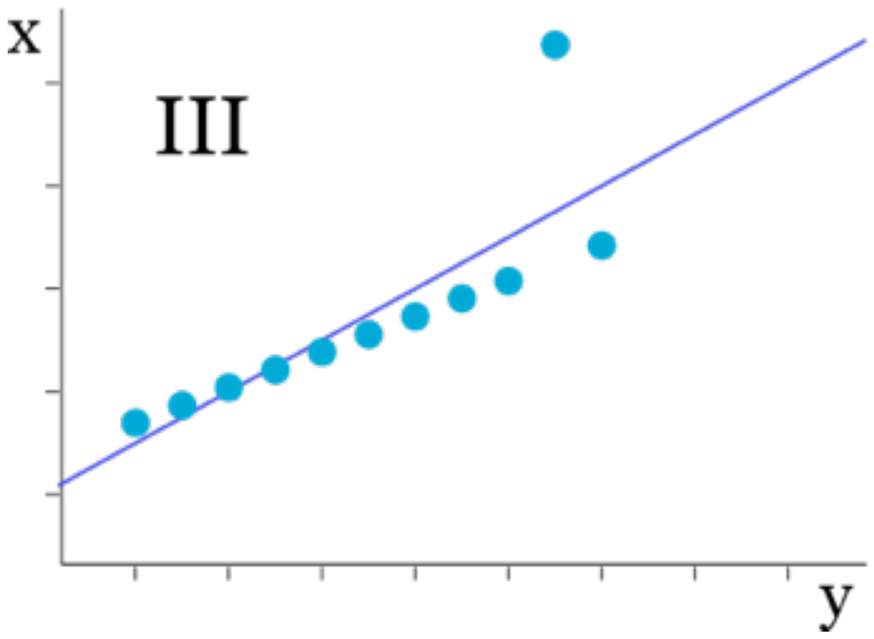
x
III
y

Components

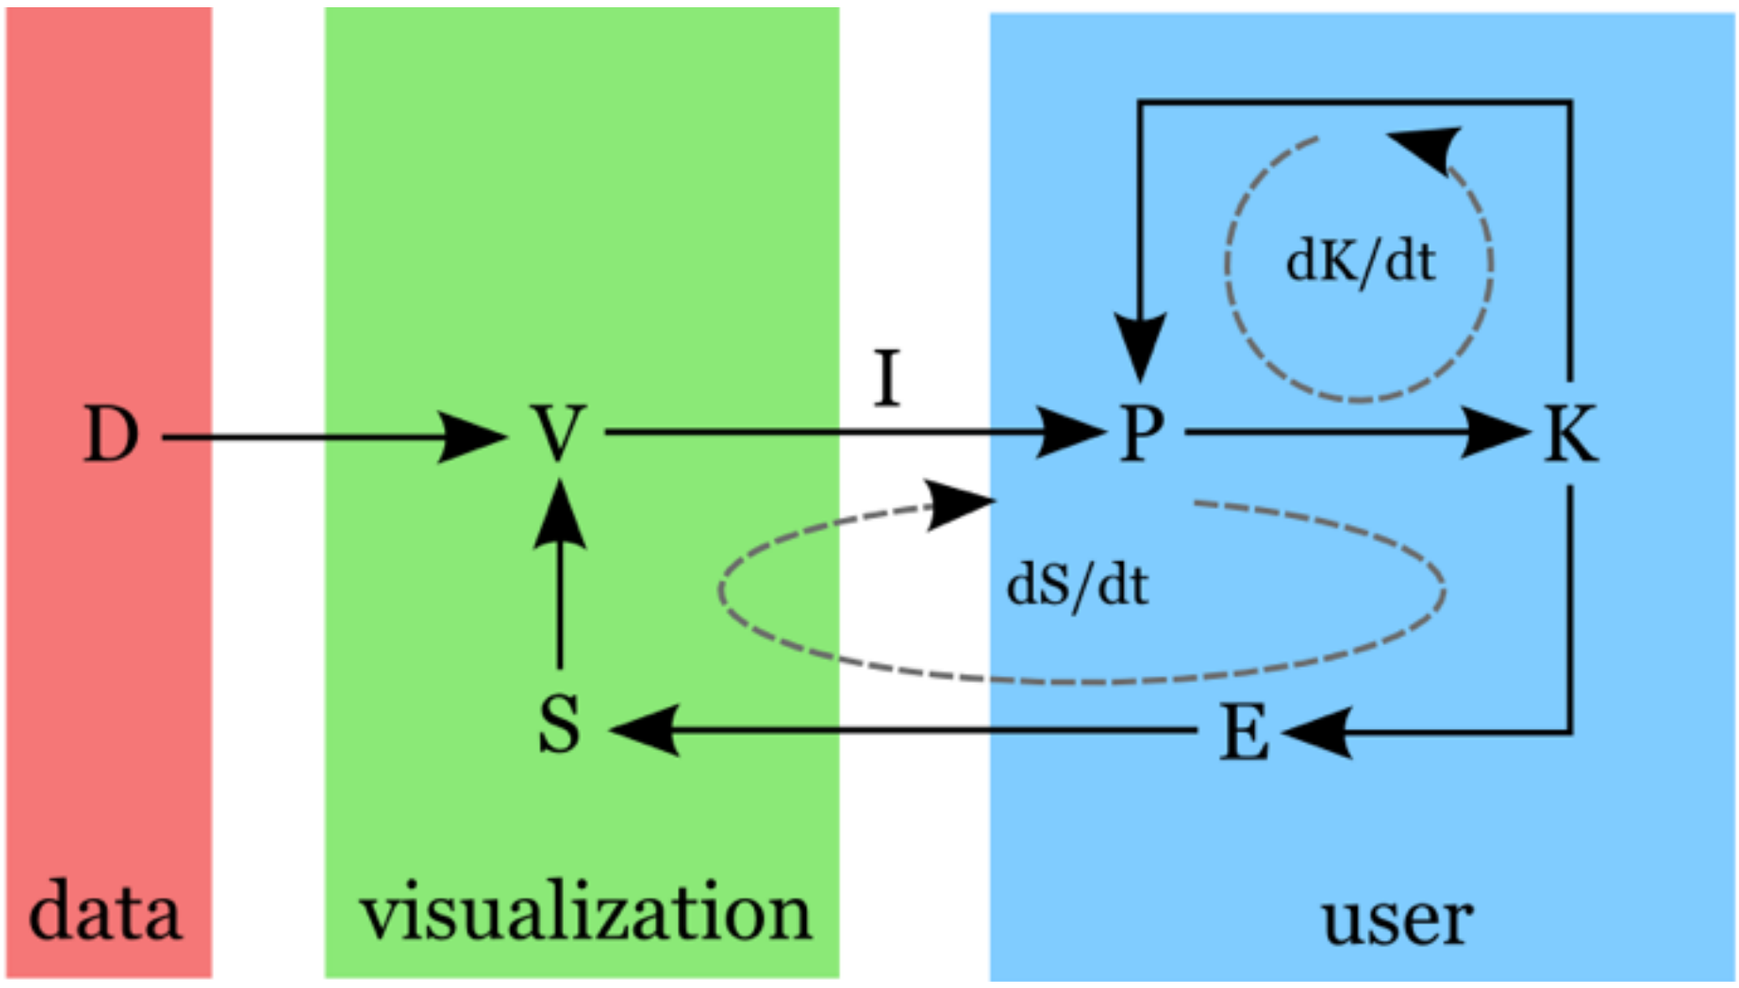

Data foundations

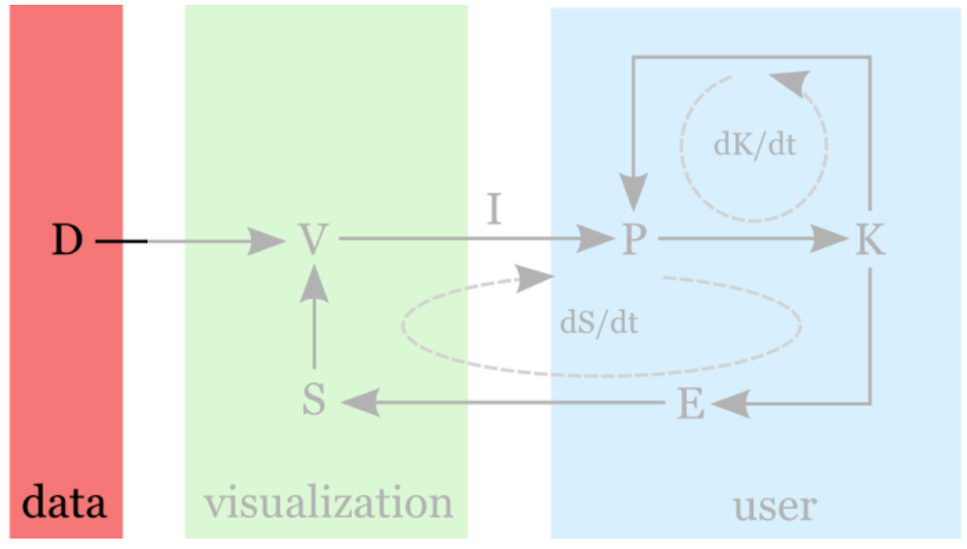

data attributes
dimensions
features
properties

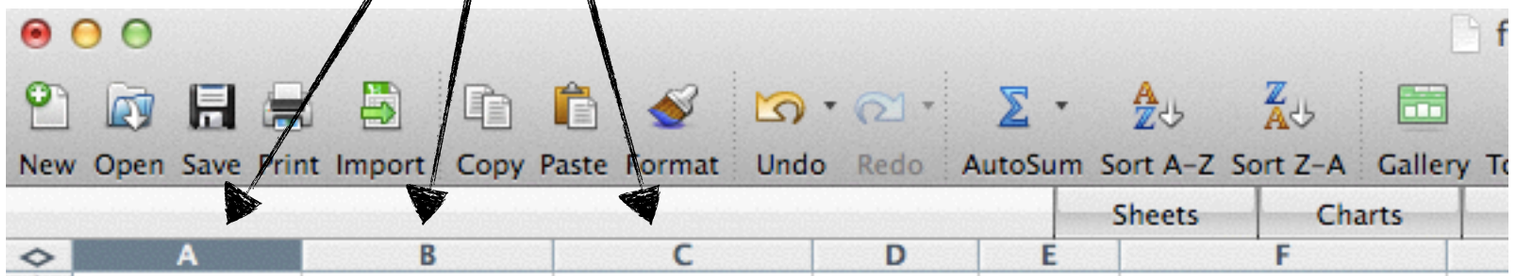

data items
datapoints

	from_airport	from_city	from_country	from_long	from_lat	to_airport	to_c
2	Balandino	Chelyabinsk	Russia	61.838	55.509	Domododevo	Mosc
3	Balandino	Chelyabinsk	Russia	61.838	55.509	Kazan	Kaza
4	Balandino	Chelyabinsk	Russia	61.838	55.509	Tolmachevo	Novc
5	Domododevo	Moscow	Russia	38.51	55.681	Balandino	Chel
6	Domododevo	Moscow	Russia	38.51	55.681	Khrabrovo	Kalir
7	Domododevo	Moscow	Russia	38.51	55.681	Kazan	Kaza
8	Domododevo	Moscow	Russia	38.51	55.681	Beaufort Mcas	Beau
9	Domododevo	Moscow	Russia	38.51	55.681	Penza Airport	Penz
10	Domododevo	Moscow	Russia	38.51	55.681	Bugulma Airport	Bugu
11	Heydar Aliyev	Baku	Azerbaijan	50.077	40.779	Beaufort Mcas	Beau
12	Khrabrovo	Kaliningrad	Russia	20.987	55.483	Domododevo	Mosc
13	Kazan	Kazan	Russia	49.464	56.01	Balandino	Chel
14	Kazan	Kazan	Russia	49.464	56.01	Domododevo	Mosc
15	Kazan	Kazan	Russia	49.464	56.01	Pulkovo	St. P
16	Kazan	Kazan	Russia	49.464	56.01	Franz Josef Strauss	Mun
17	Kazan	Kazan	Russia	49.464	56.01	Bugulma Airport	Bugu

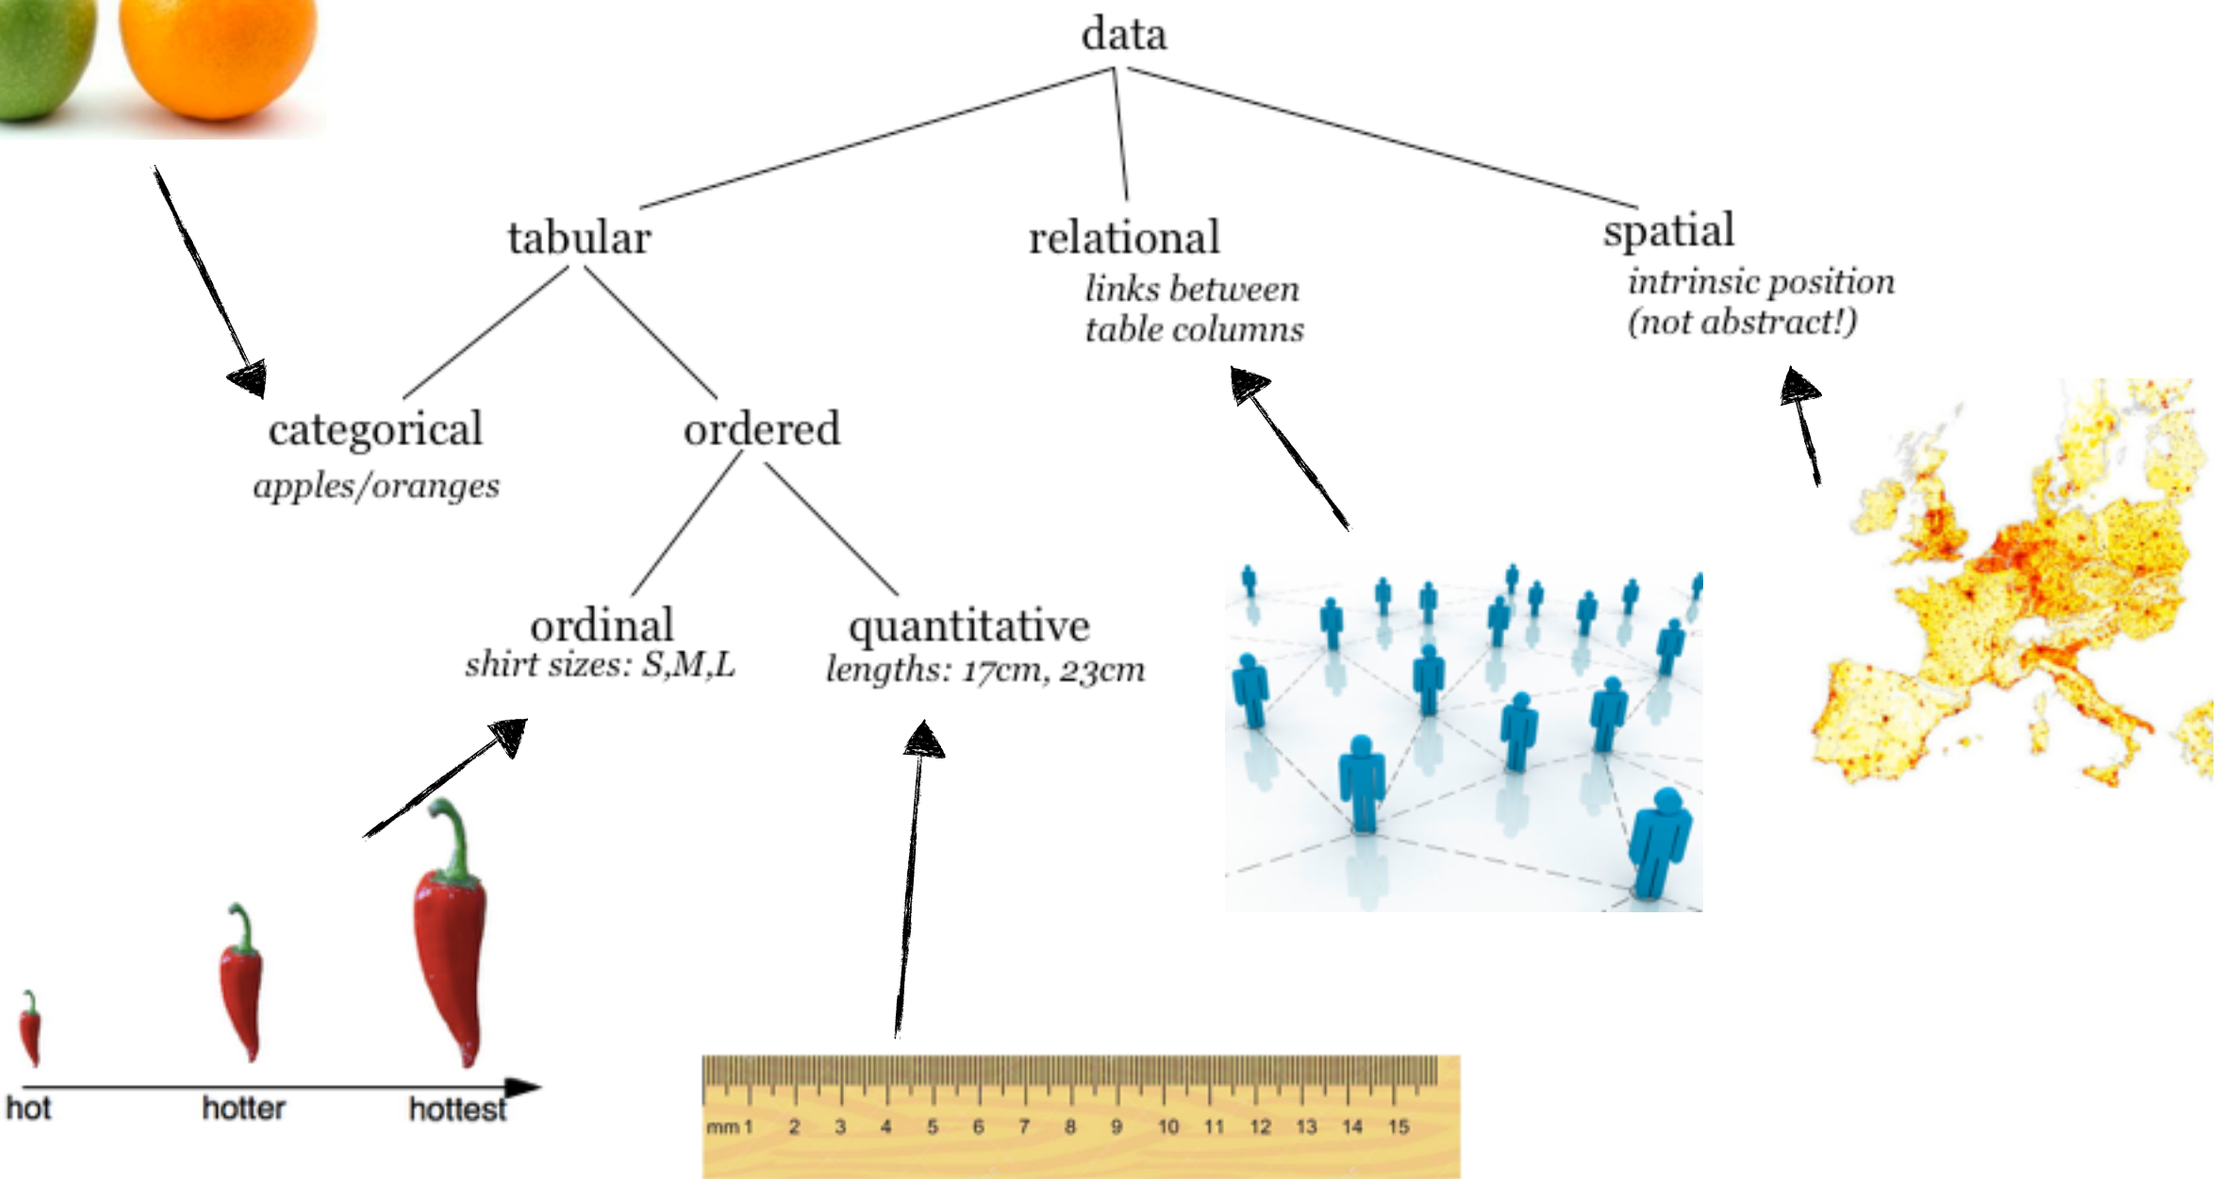

S Stevens “On the theory of scales and measurements” (1946)

Human perception foundations

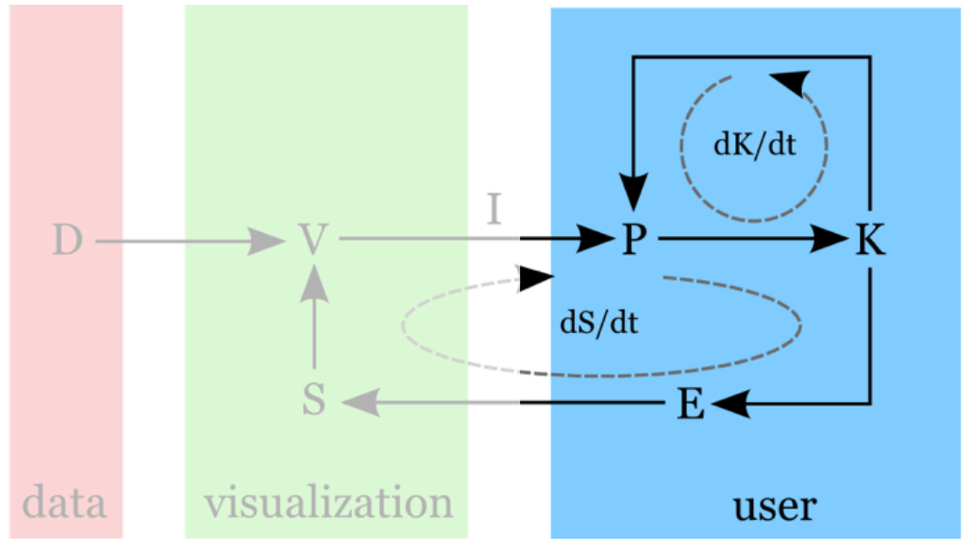

Human perception

- Gestalt laws
- pre-attentive vision
- selective attention
- colour
- What does this mean for visual design?

Gestalt laws - interplay between parts and the whole (Kurt Koffka)

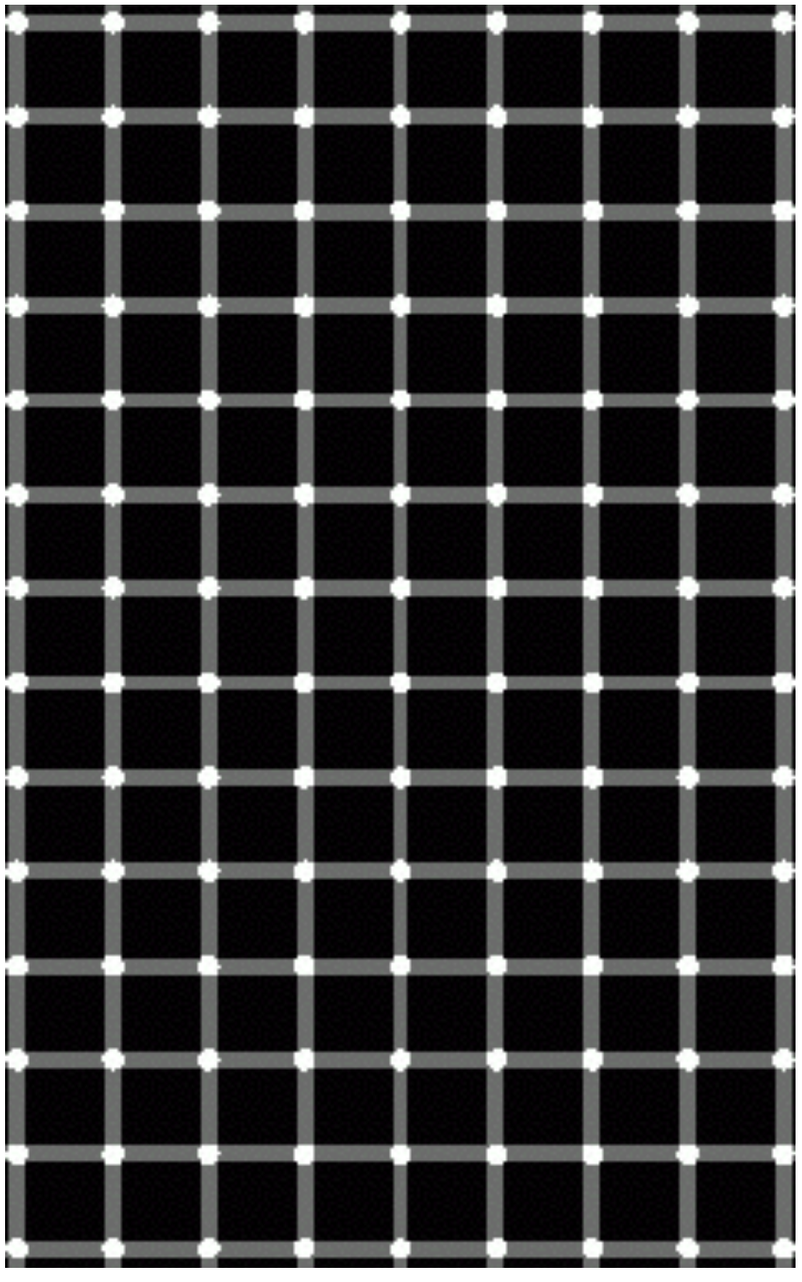

series of principles

Election results Florida:

- black = Bush
- white = Gore

Hermann grid illusion

Simplicity

Connectedness

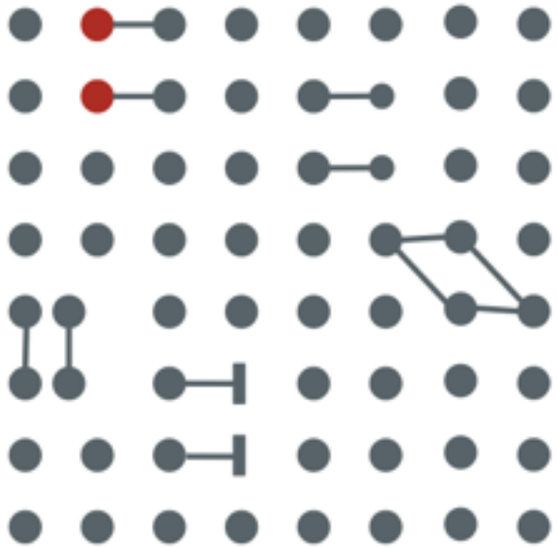

Proximity

Similarity

Closure

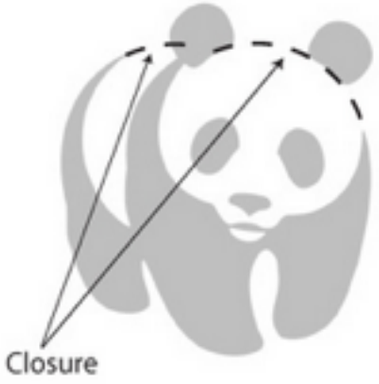

Kanizsa illusion

Symmetry, Familiarity, Continuity
Figure and Ground

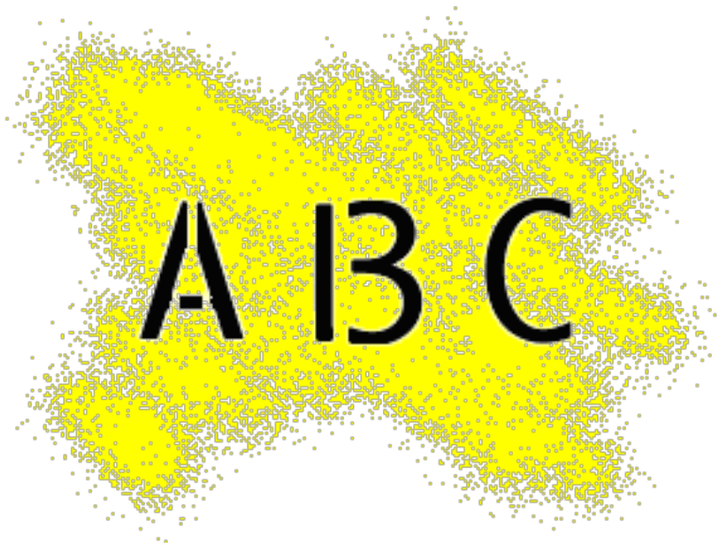
A B C


12 13 14

Figure and Ground

Pre-attentive vision

= ability of low-level human visual system to rapidly identify certain basic visual properties

- some features “pop out”
- used for:
  - target detection
  - boundary detection
  - counting/estimation
  - ...
- visual system takes over => all cognitive power available for interpreting the figure, rather than needing part of it for processing the figure

5

5

5

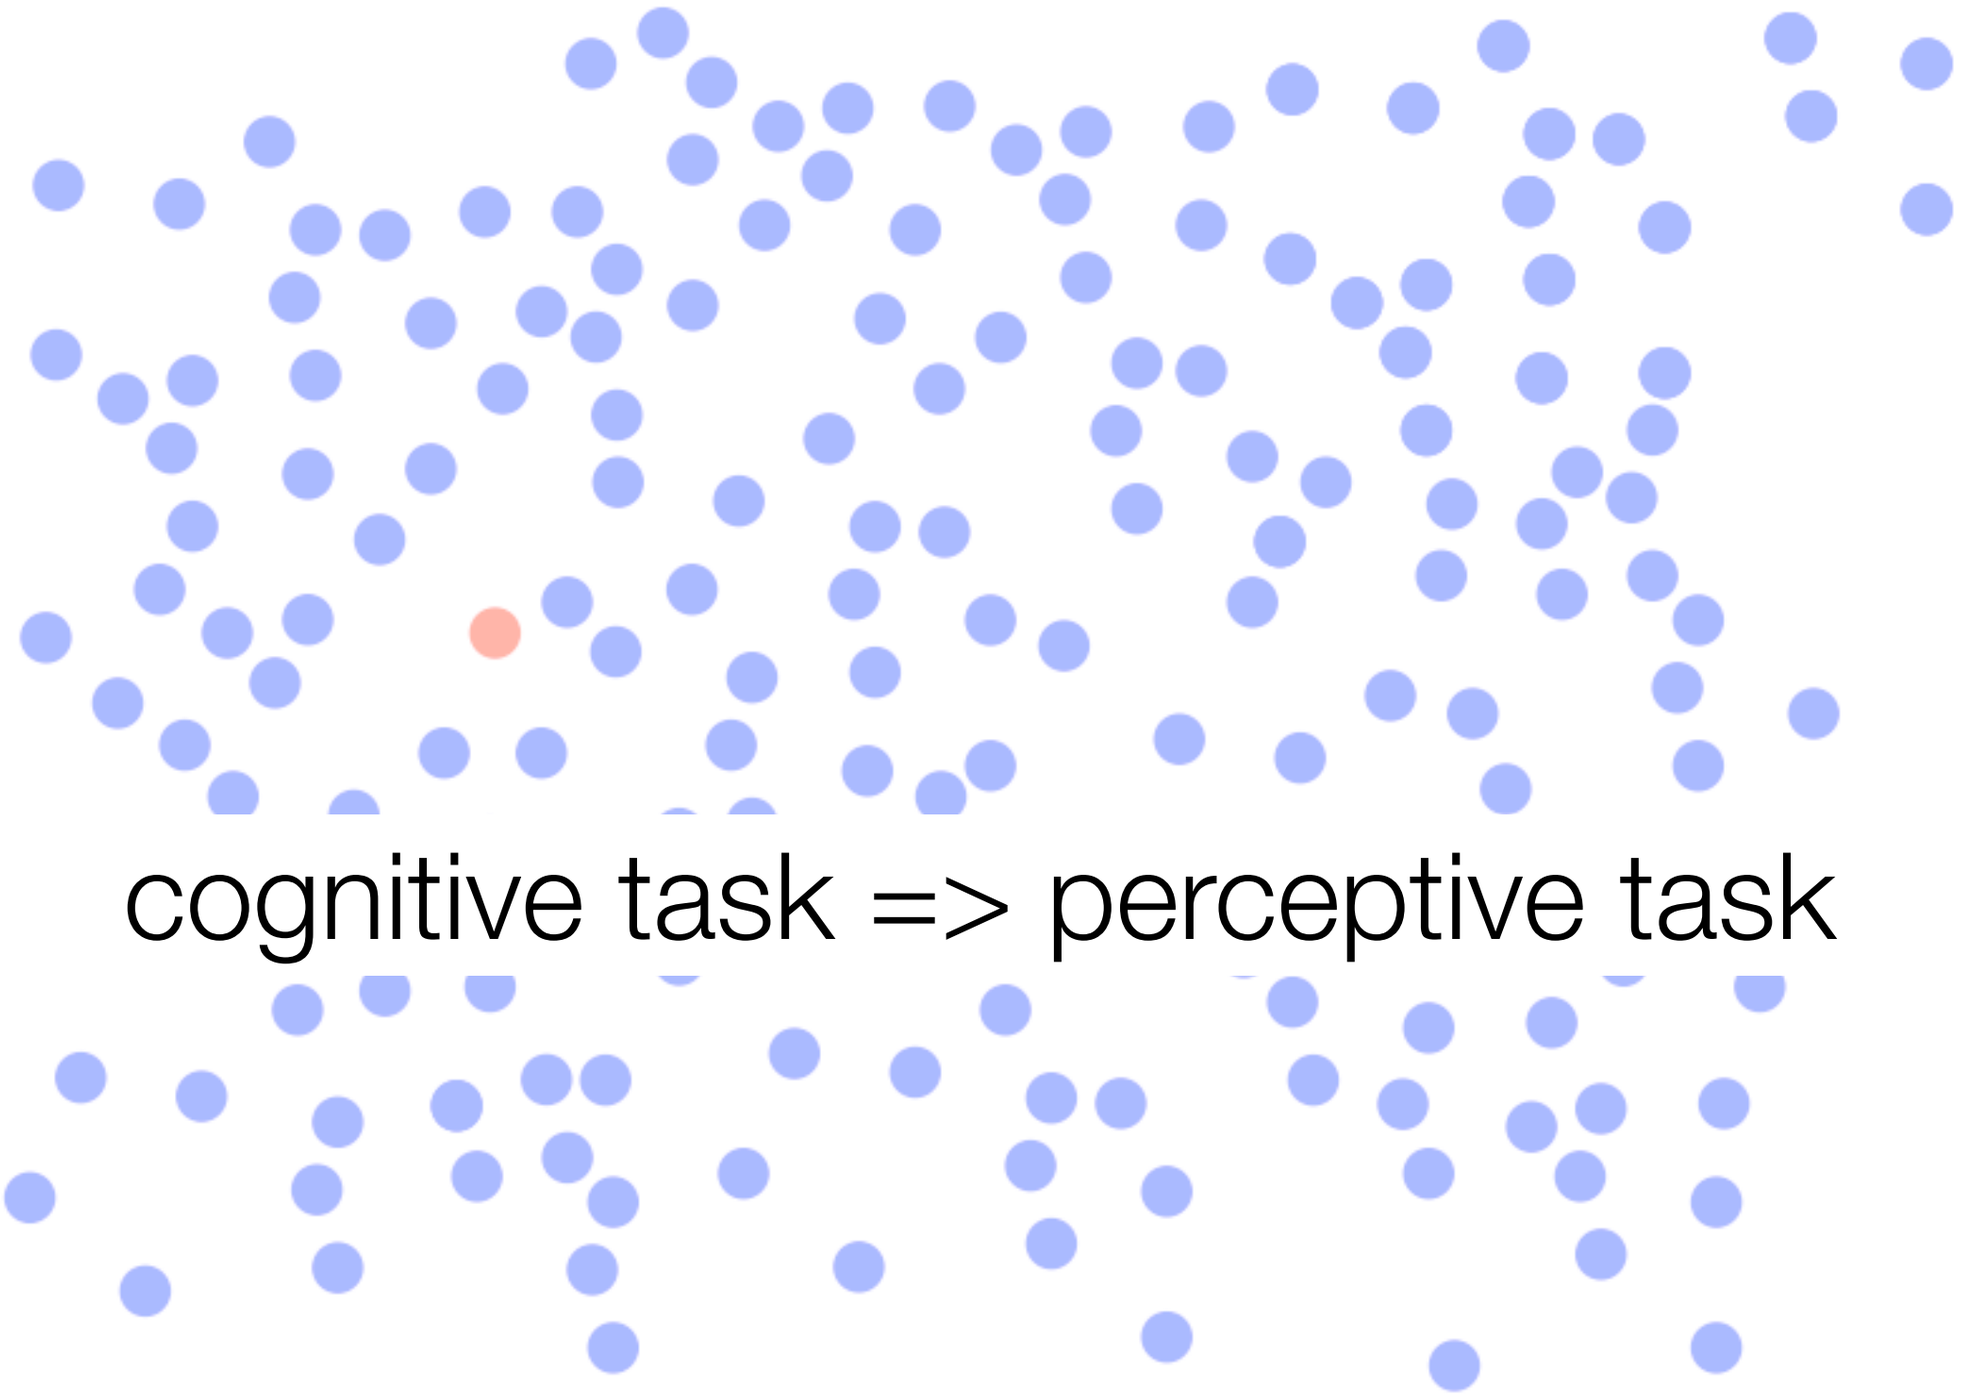

cognitive task => perceptive task

Limitations of preattentive vision

1. Speed depends on **which channel** (use one that is good for categorical)

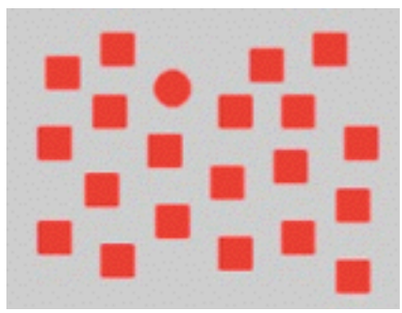

2. **Combining** pre-attentive features does *not* always work => would need to resort to "**serial search**" (most channel pairs; all channel triplets)

e.g. is there a red square in this picture

I		II		III		IV	
x	y	x	y	x	y	x	y
10.0	8.04	10.0	9.14	10.0	7.46	8.0	6.58
8.0	6.95	8.0	8.14	8.0	6.77	8.0	5.76
13.0	7.58	13.0	8.74	13.0	12.74	8.0	7.71
9.0	8.81	9.0	8.77	9.0	7.11	8.0	8.84
11.0	8.33	11.0	9.26	11.0	7.81	8.0	8.47
14.0	9.96	14.0	8.10	14.0	8.84	8.0	7.04
6.0	7.24	6.0	6.13	6.0	6.08	8.0	5.25
4.0	4.26	4.0	3.10	4.0	5.39	19.0	12.50
12.0	10.84	12.0	9.13	12.0	8.15	8.0	5.56
7.0	4.82	7.0	7.26	7.0	6.42	8.0	7.91
5.0	5.68	5.0	4.74	5.0	5.73	8.0	6.80

n = 11

mean x = 9.0 variance x = 11.0 correlation x & y = 0.816
mean y = 7.5 variance y = 4.12 regression line: y = 3+0.5x

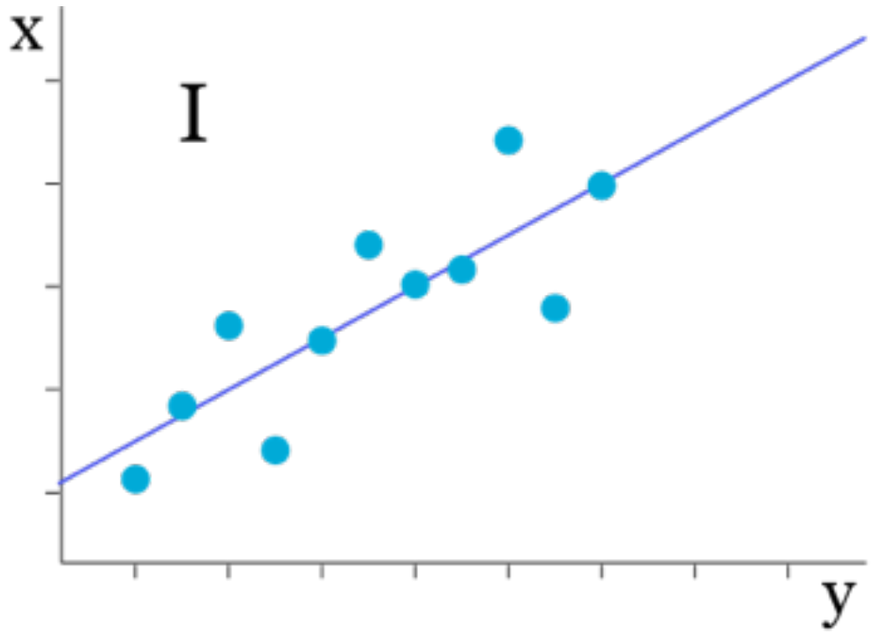
x
I
y

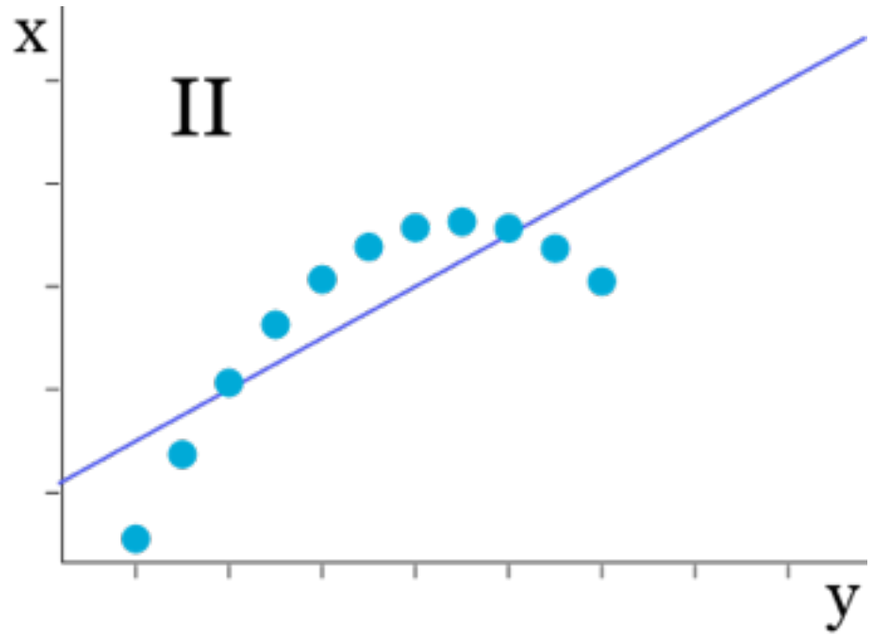
x
II
y

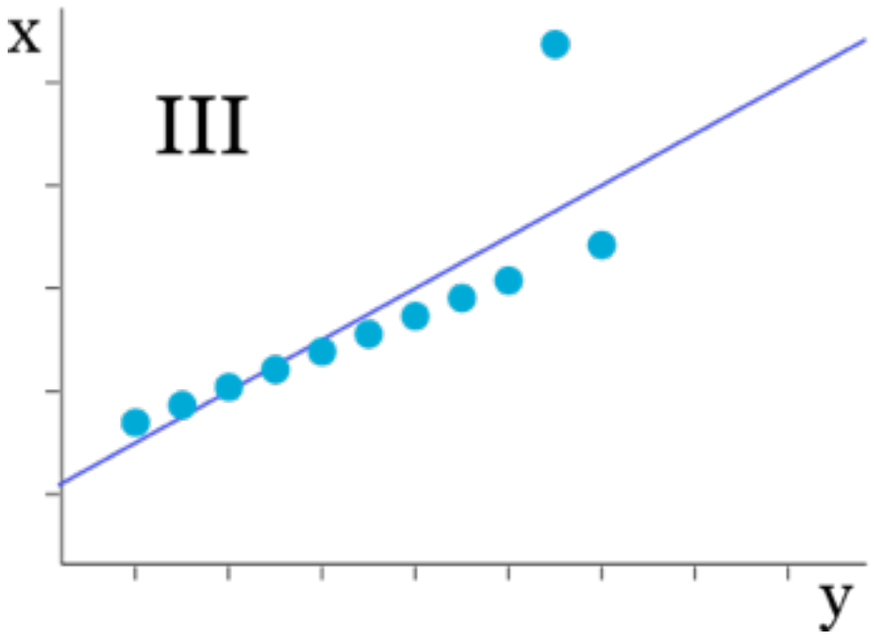
x
III
y

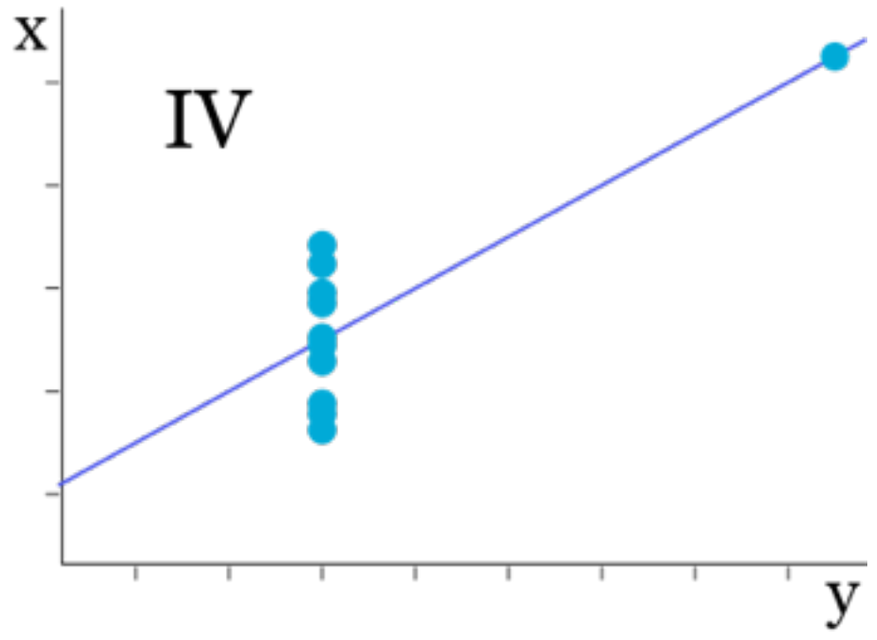
x
IV
y

Selective attention

https://www.youtube.com/watch?v=vJG698U2Mvo

About colour

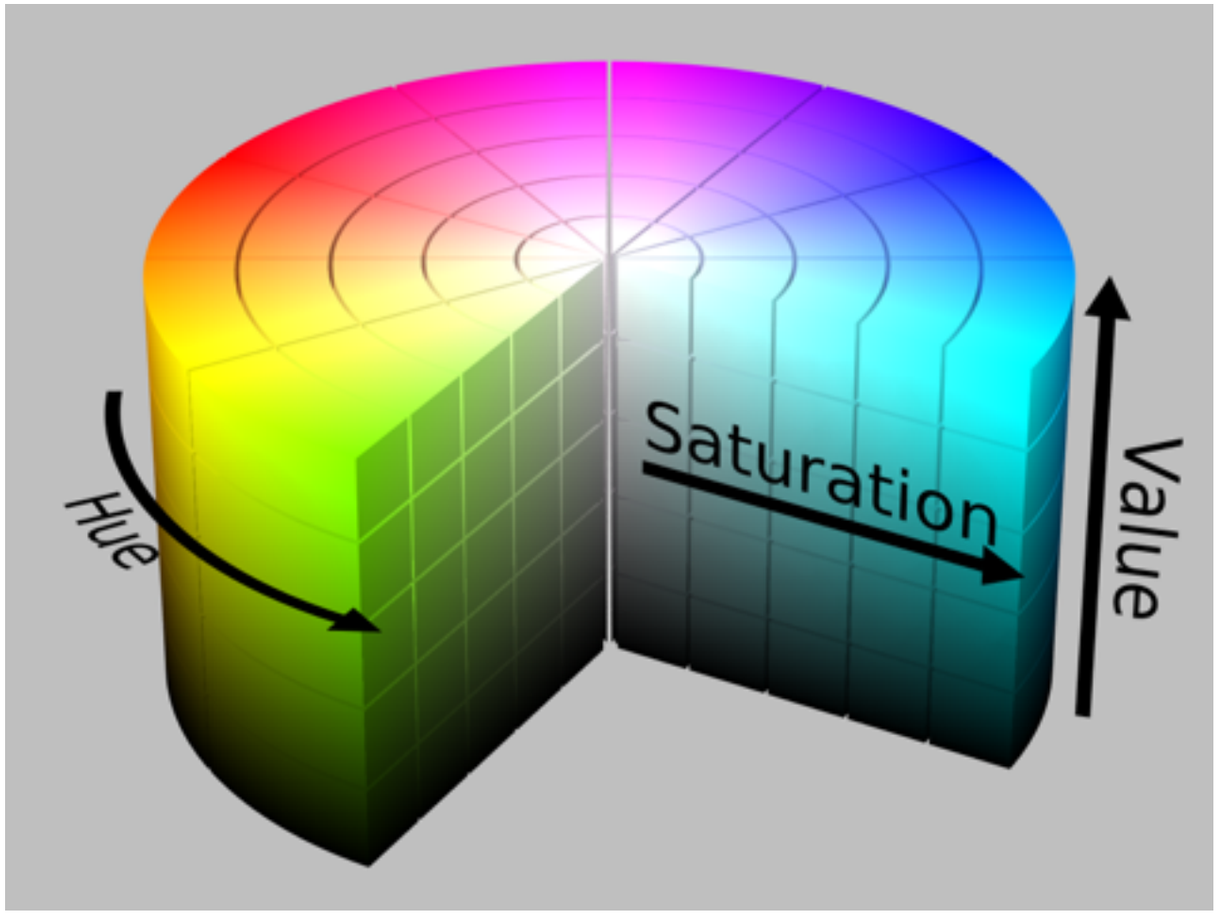

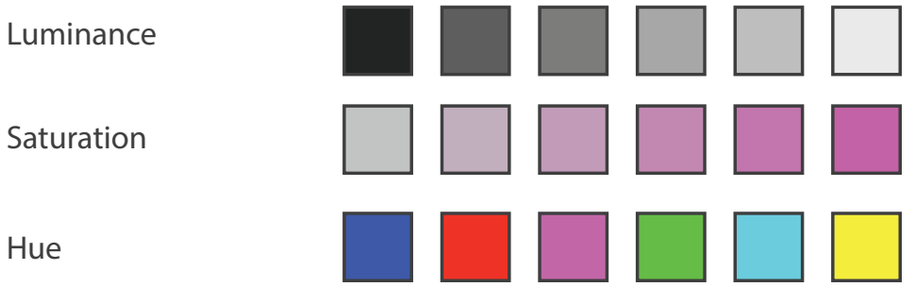

Use **HSV** (hue - saturation - value) instead of RGB

HSV: separates *luma* (intensity) from *chroma* (colour)

RGB has to do with *implementation*; HSV has to do with *perception*

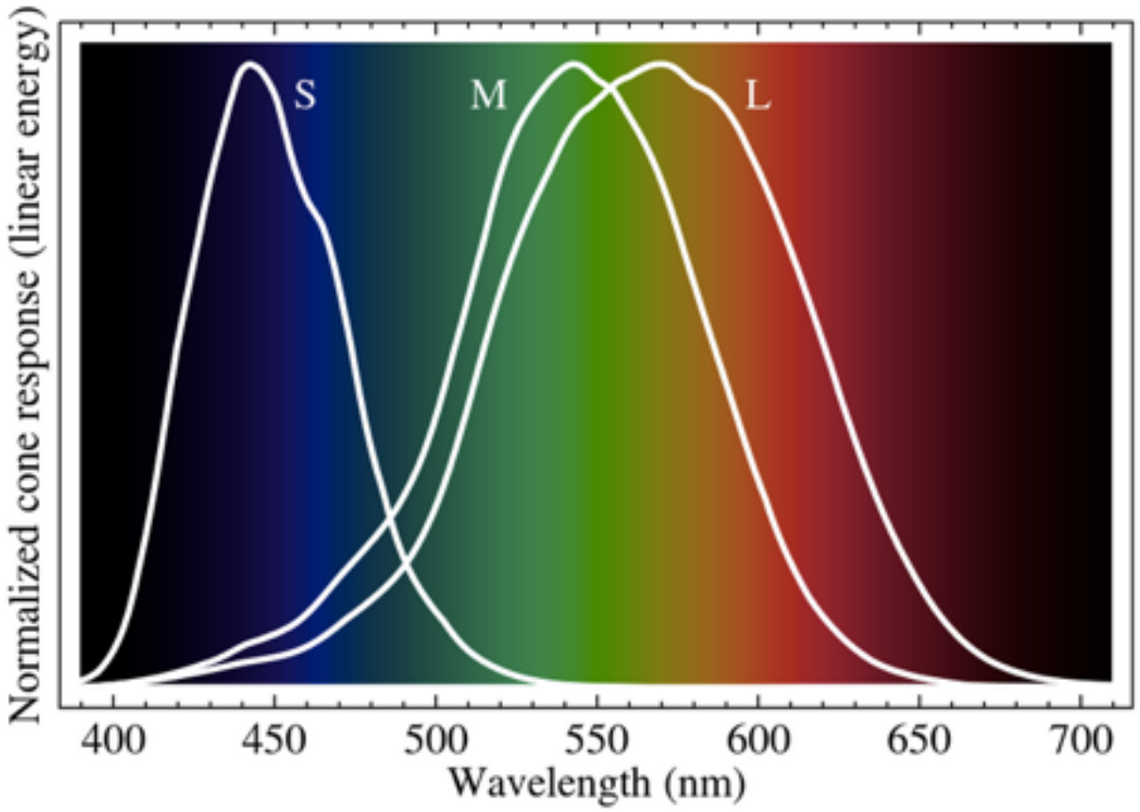

sensitivity of S/M/L cones to colour wavelengths

http://leaverou.github.io/whathecolor/

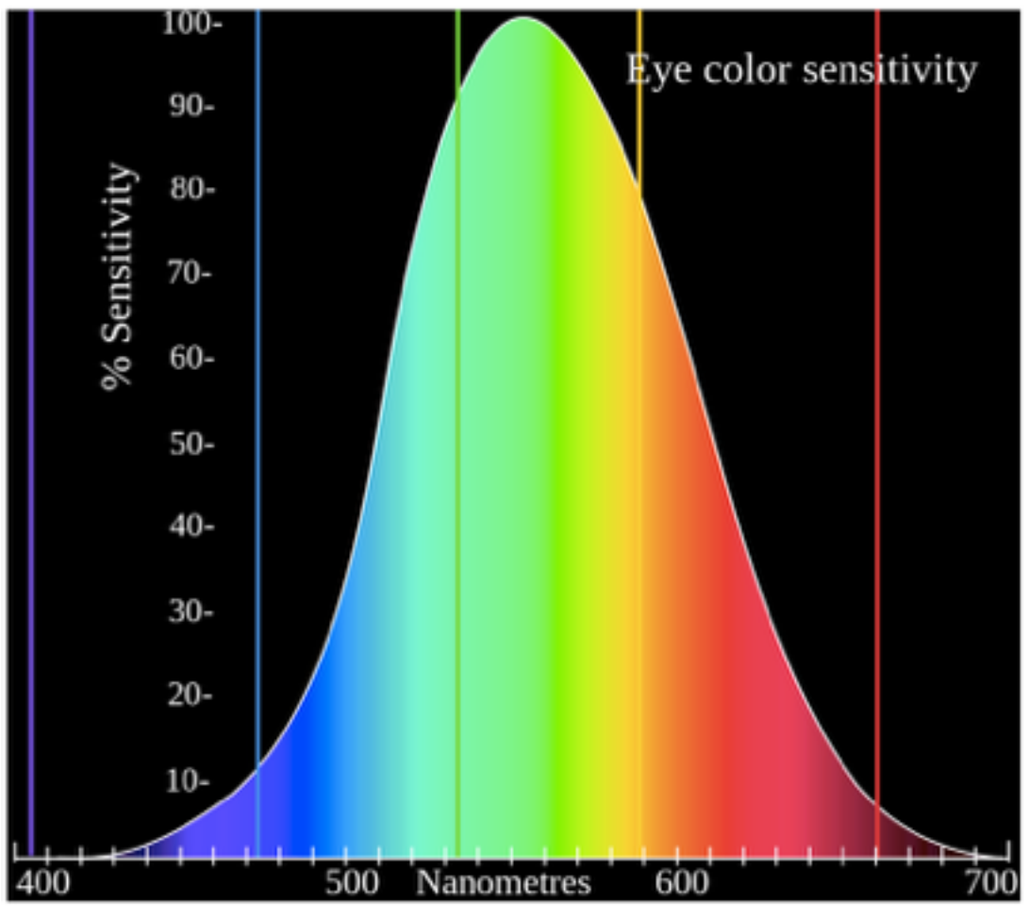

colorbrewer2.org - Never pick your own colour

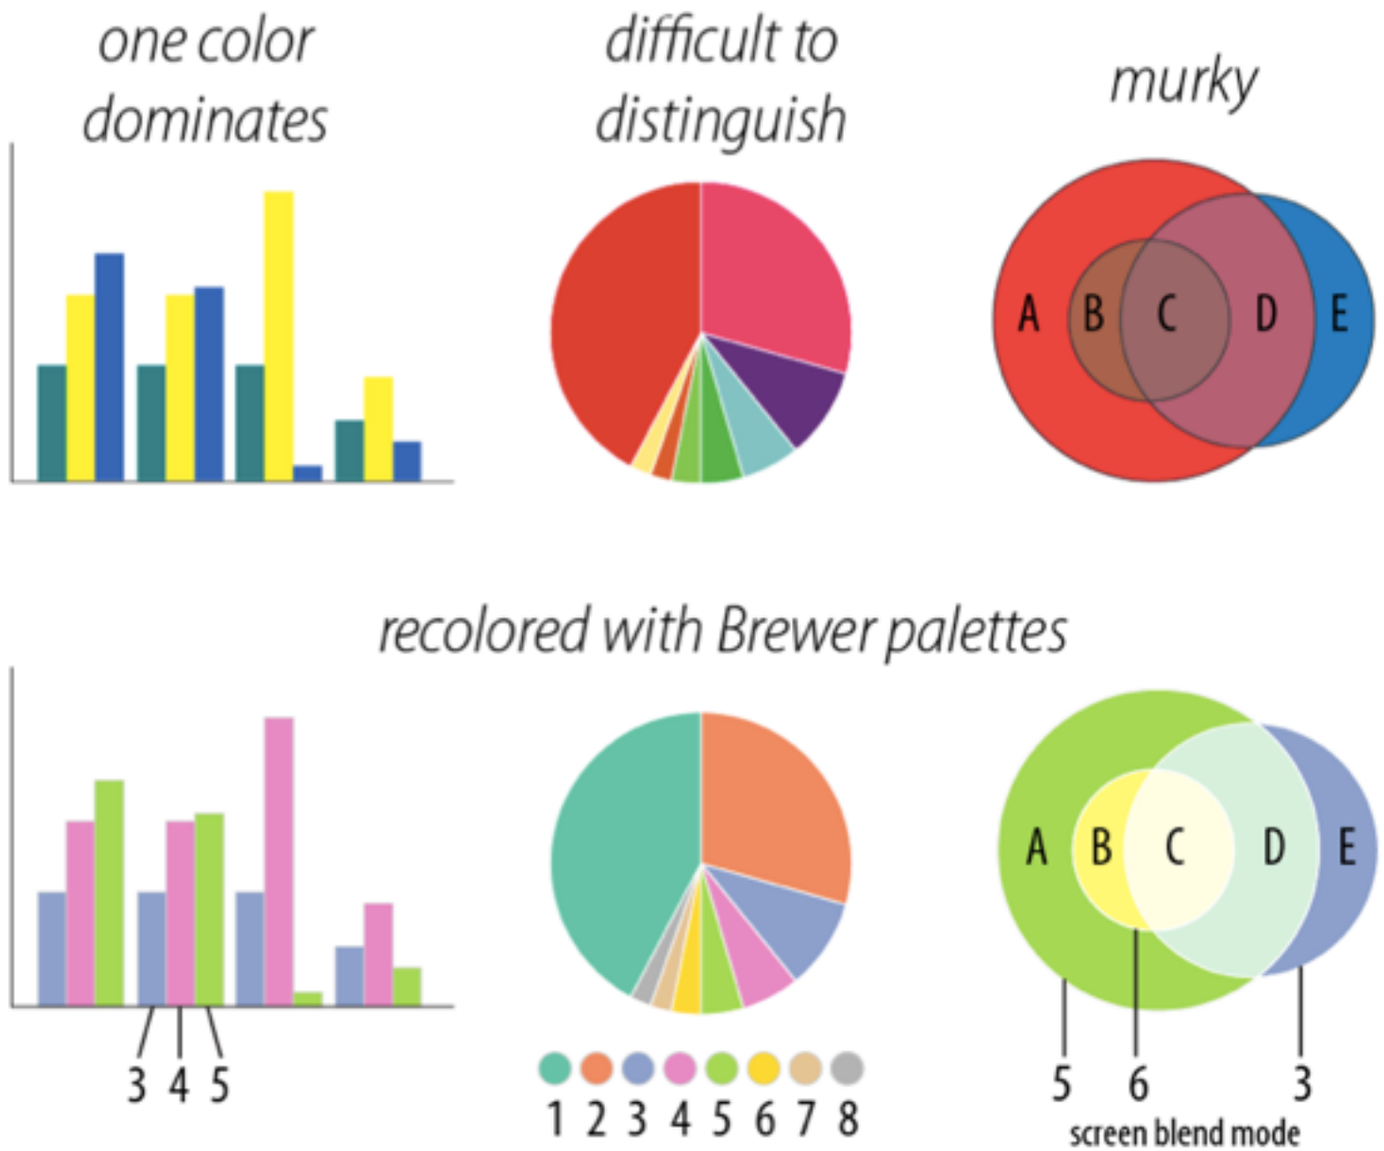

in R: please use RColorBrewer!

Steven's psychophysical law

= proposed relationship between the magnitude of a physical stimulus and its perceived intensity or strength

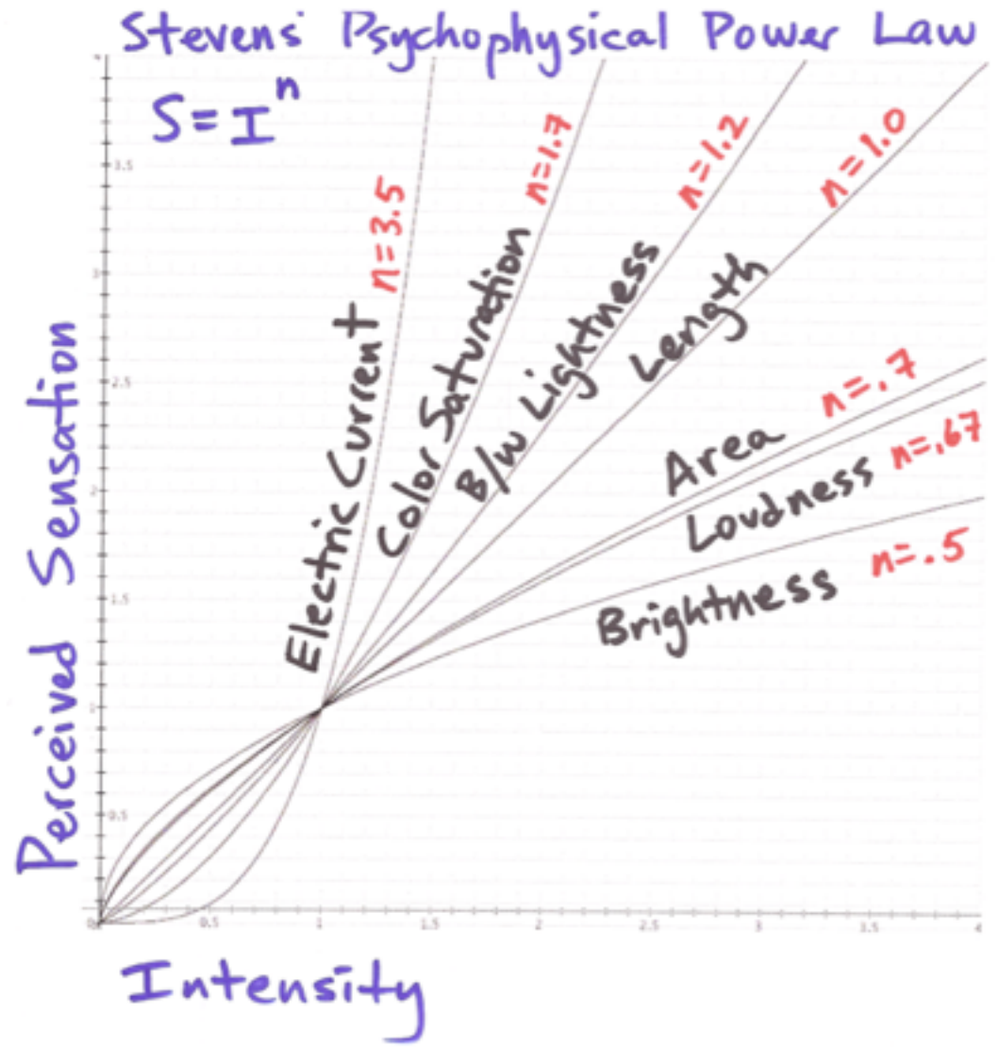

Semiology of graphics

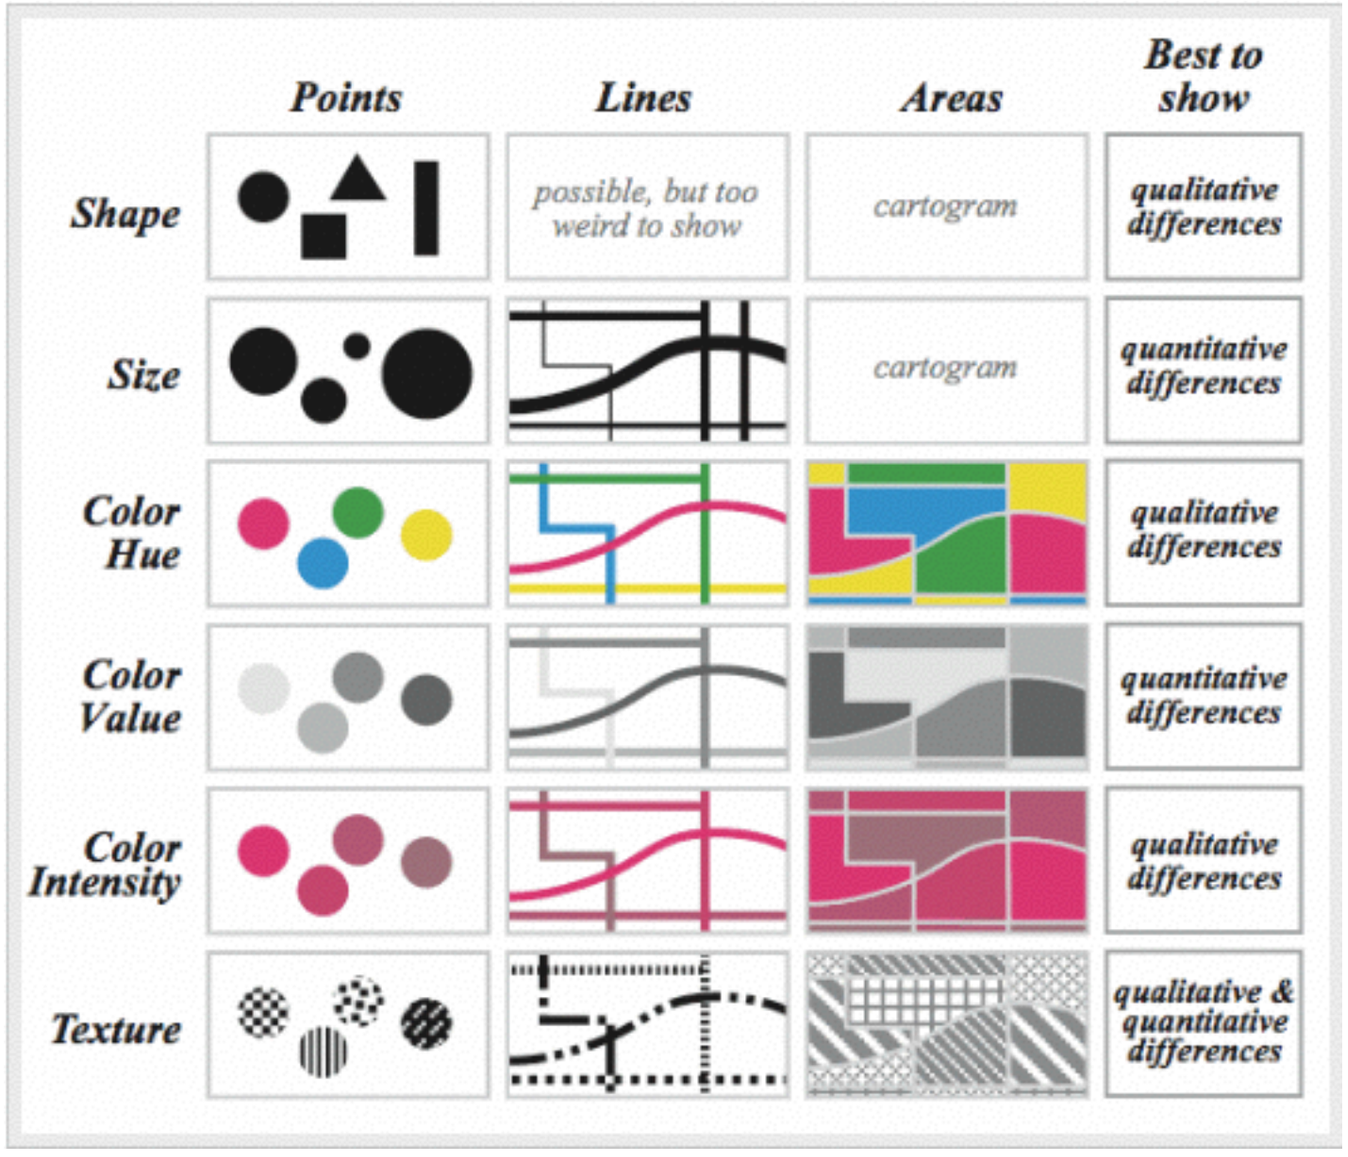

Accuracy of quantitative perceptual tasks

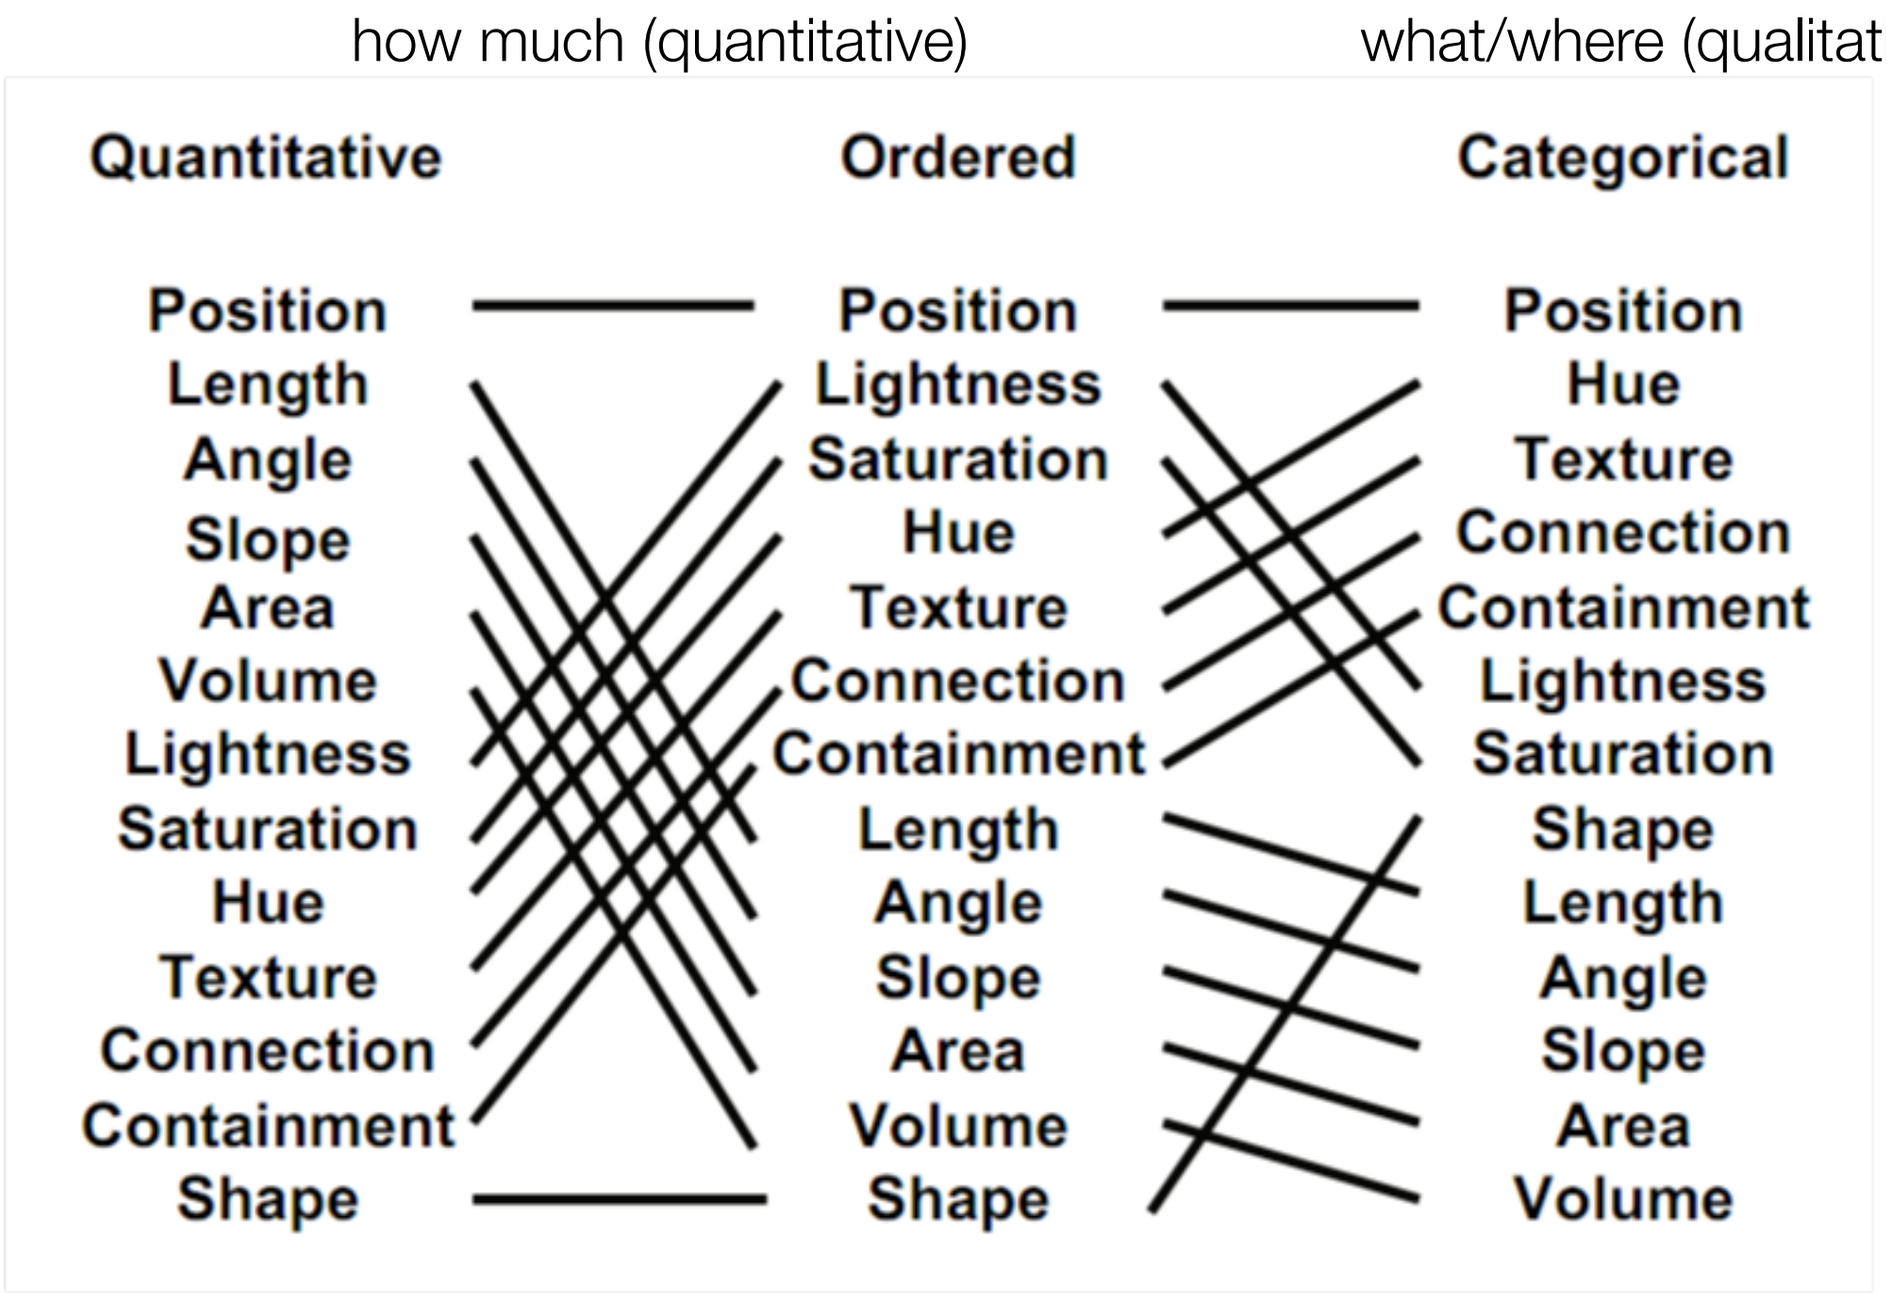

McKinlay

Accuracy of quantitative perceptual tasks

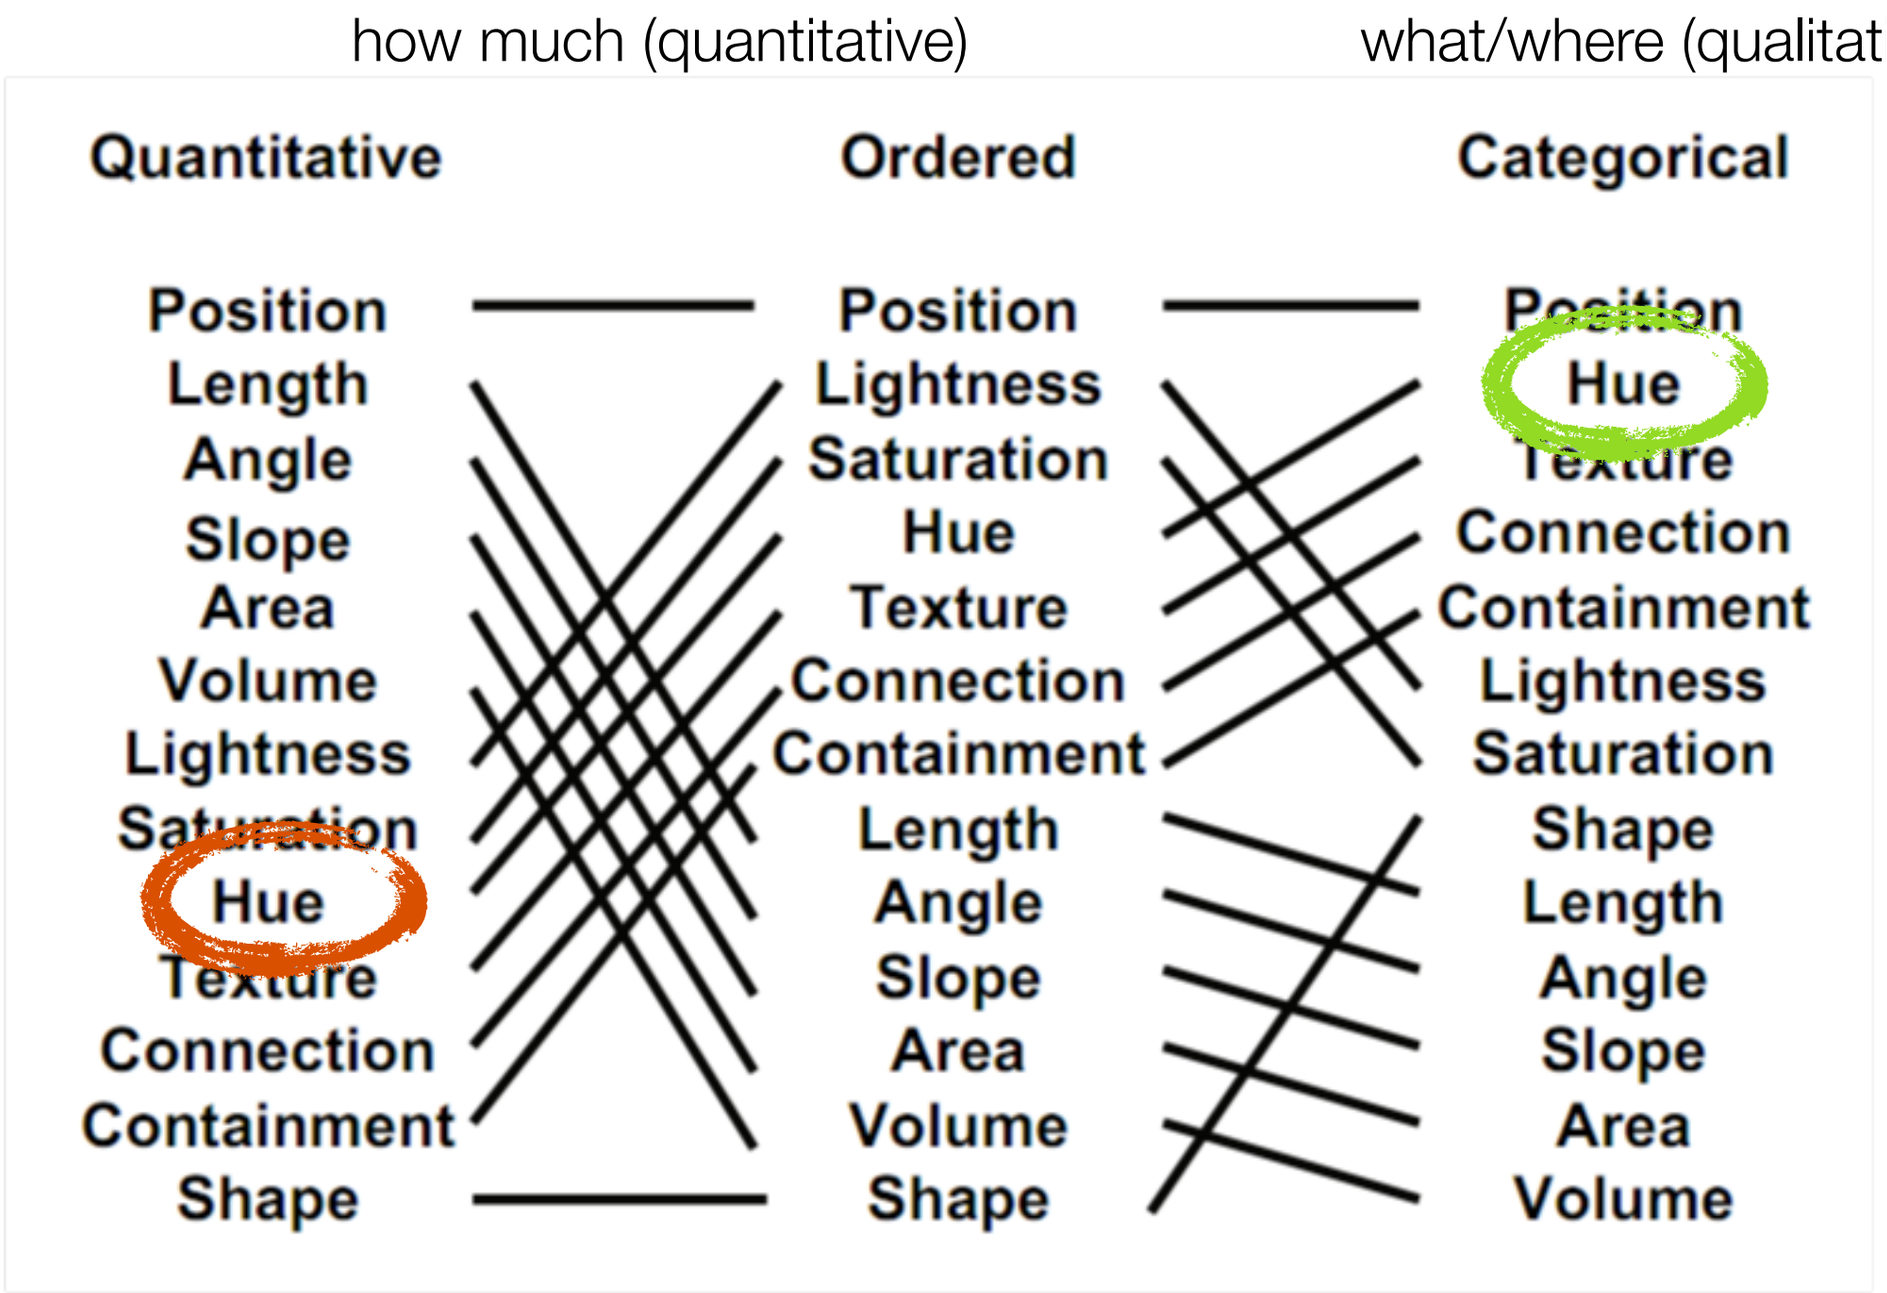

McKinlay

Accuracy of quantitative perceptual tasks

how much (quantitative) what/where (qualitative)

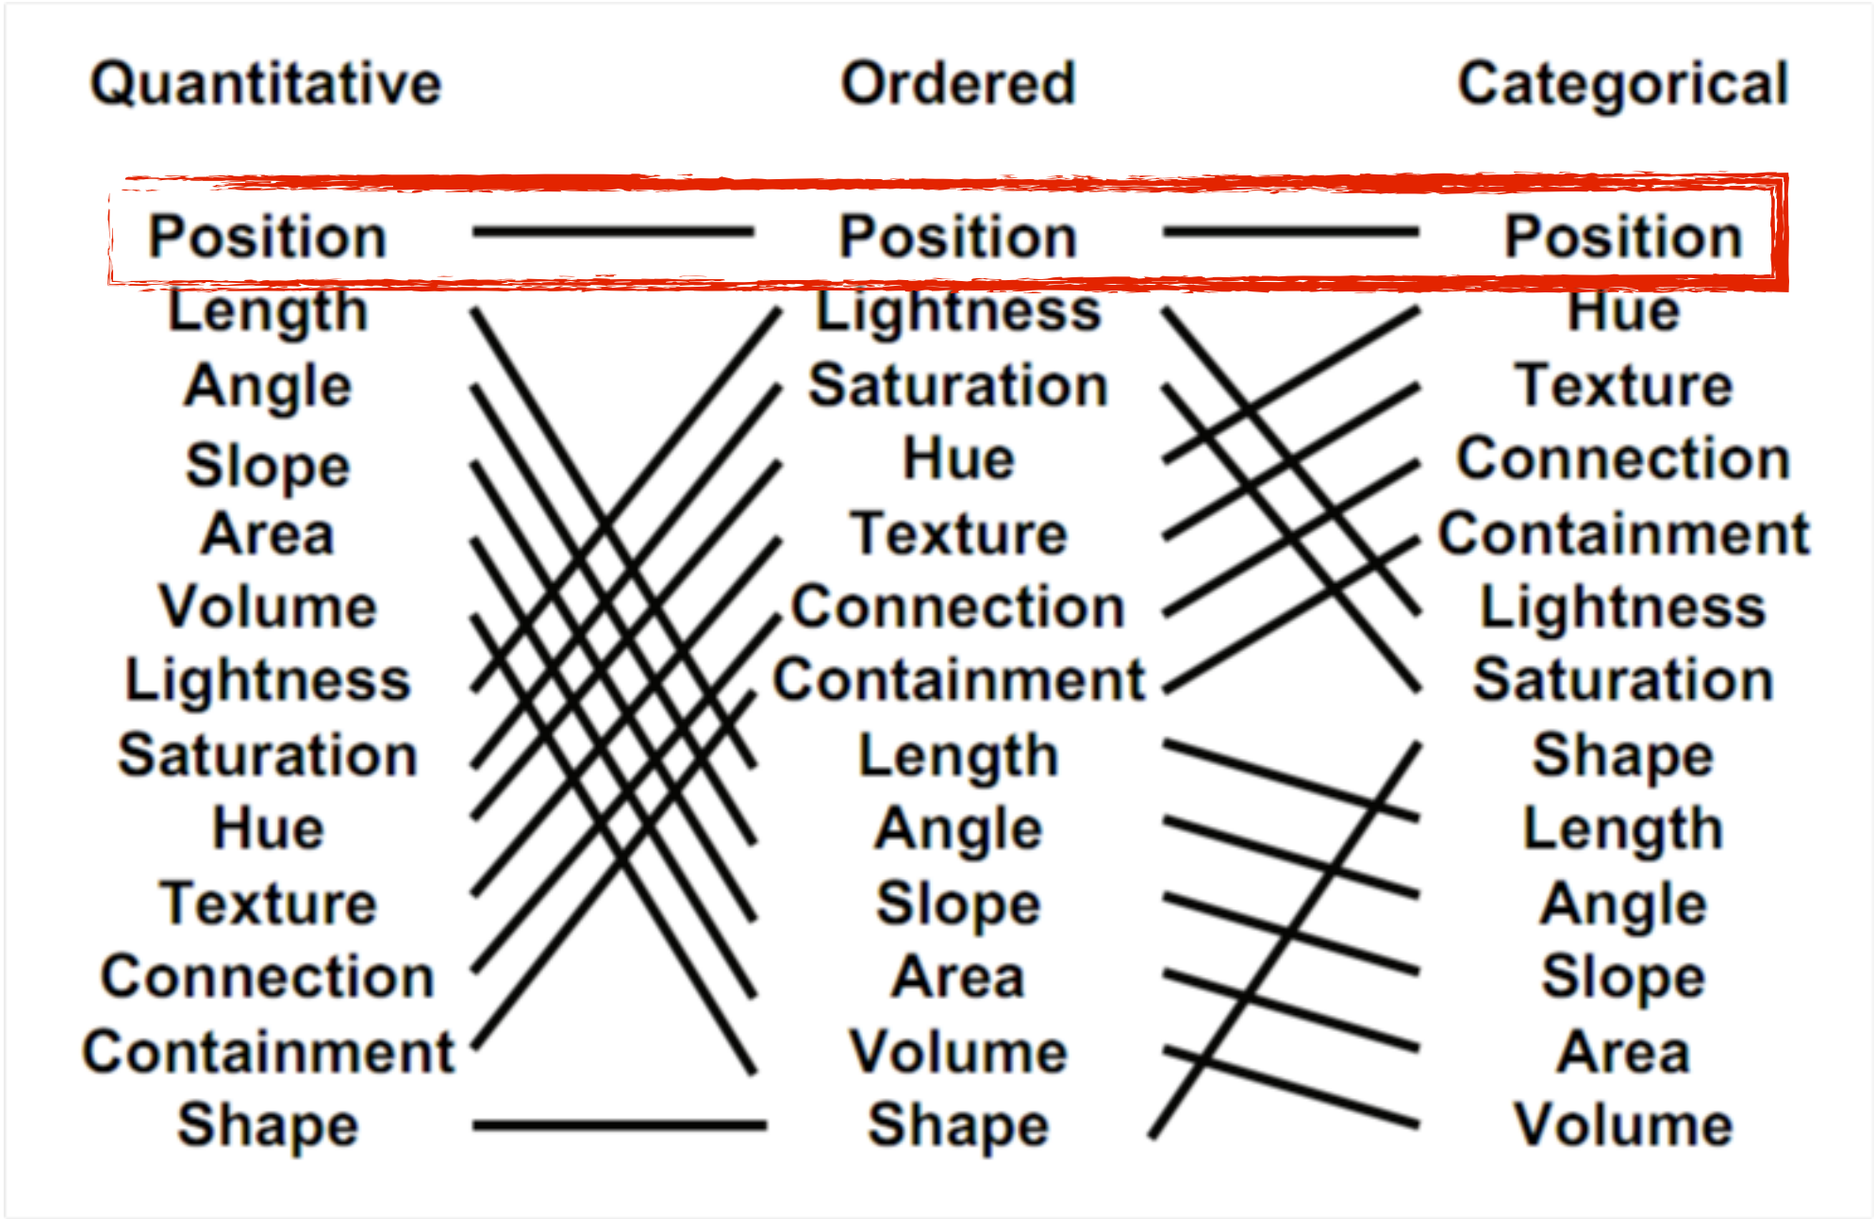

McKinlay

"power of the plane"

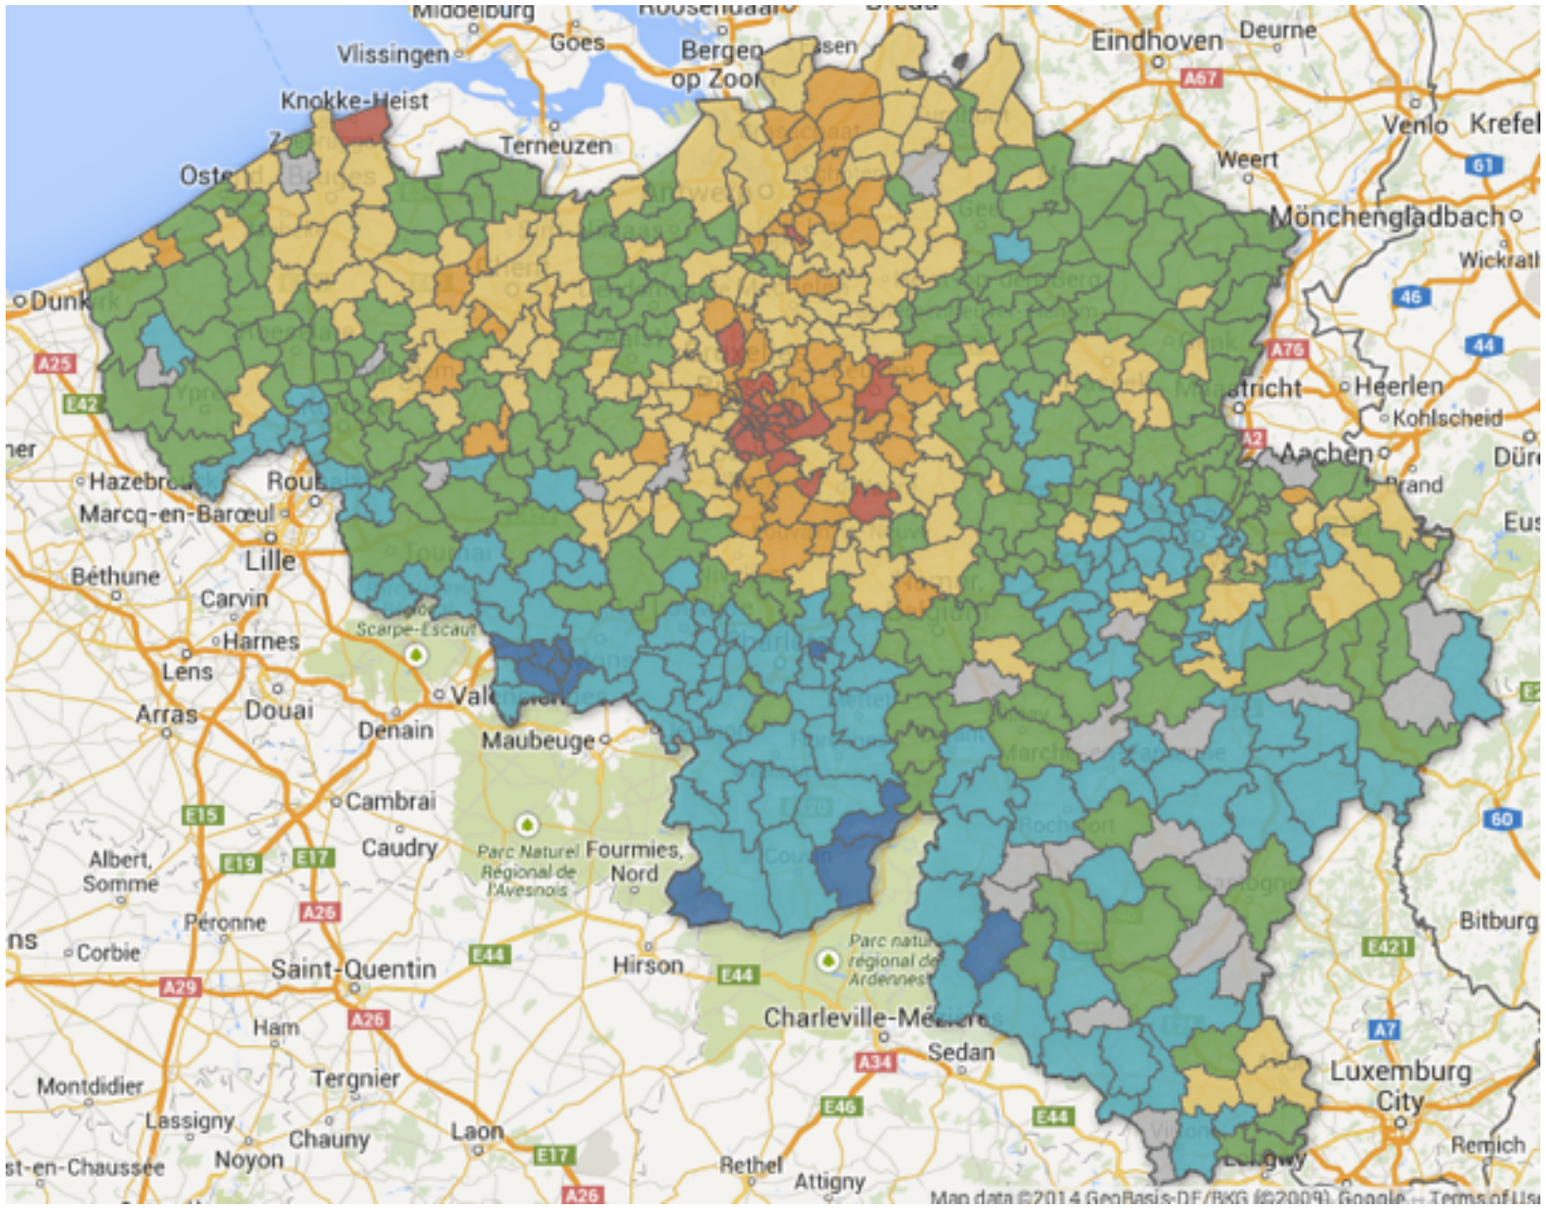

average prices per
municipality in Belgium
(De Standaard, 24/8/2014)

Rules of thumb

(partly based on book Munzner)

1. maximize data-to-ink ratio (Edward Tufte)

2. beware of the lie-factor

3. no unjustified 3D

4. eyes beat memory

5. overview first, zoom & filter, details on demand

6. don’t overengineer

1. maximize data-to-ink ratio

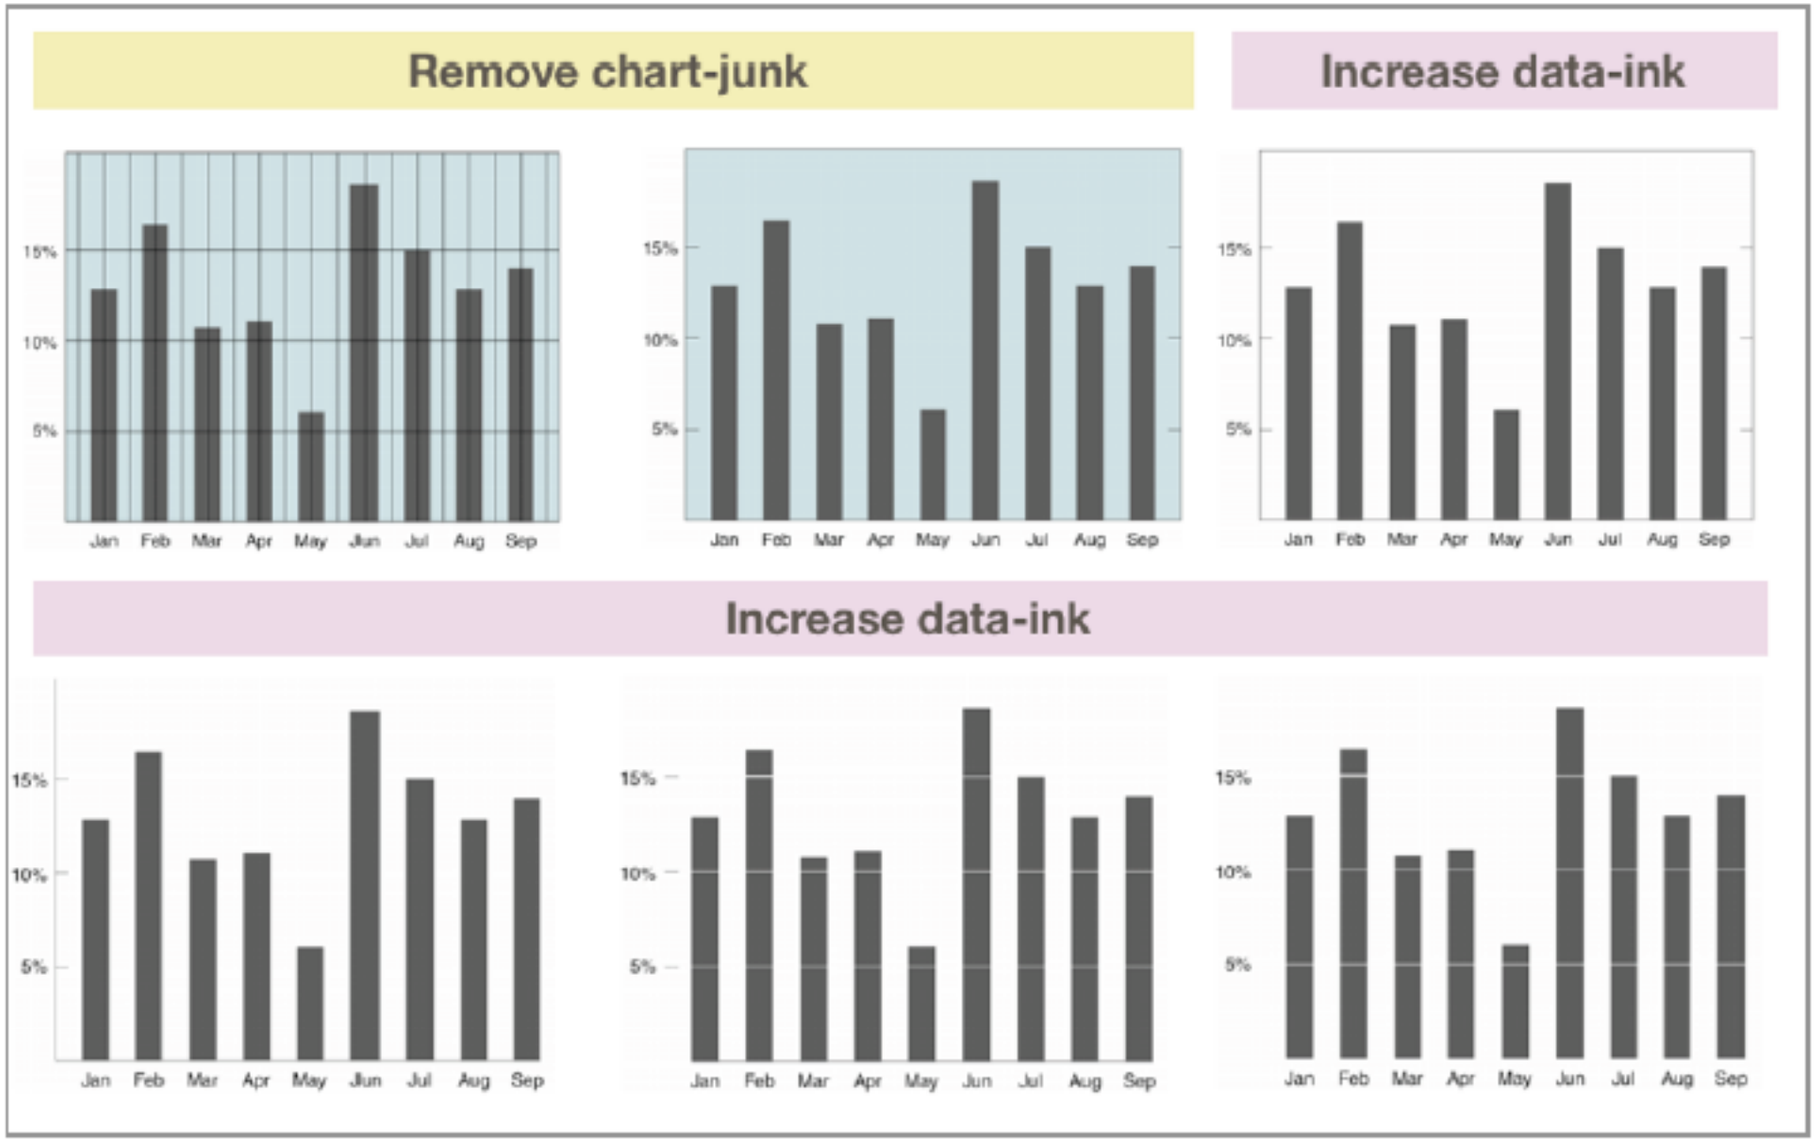

“data-to-ink ratio” = $\frac{\text{data-ink}}{\text{total ink}}$ = 1 - proportion of graphic that can be erased

2. beware of the lie-factor

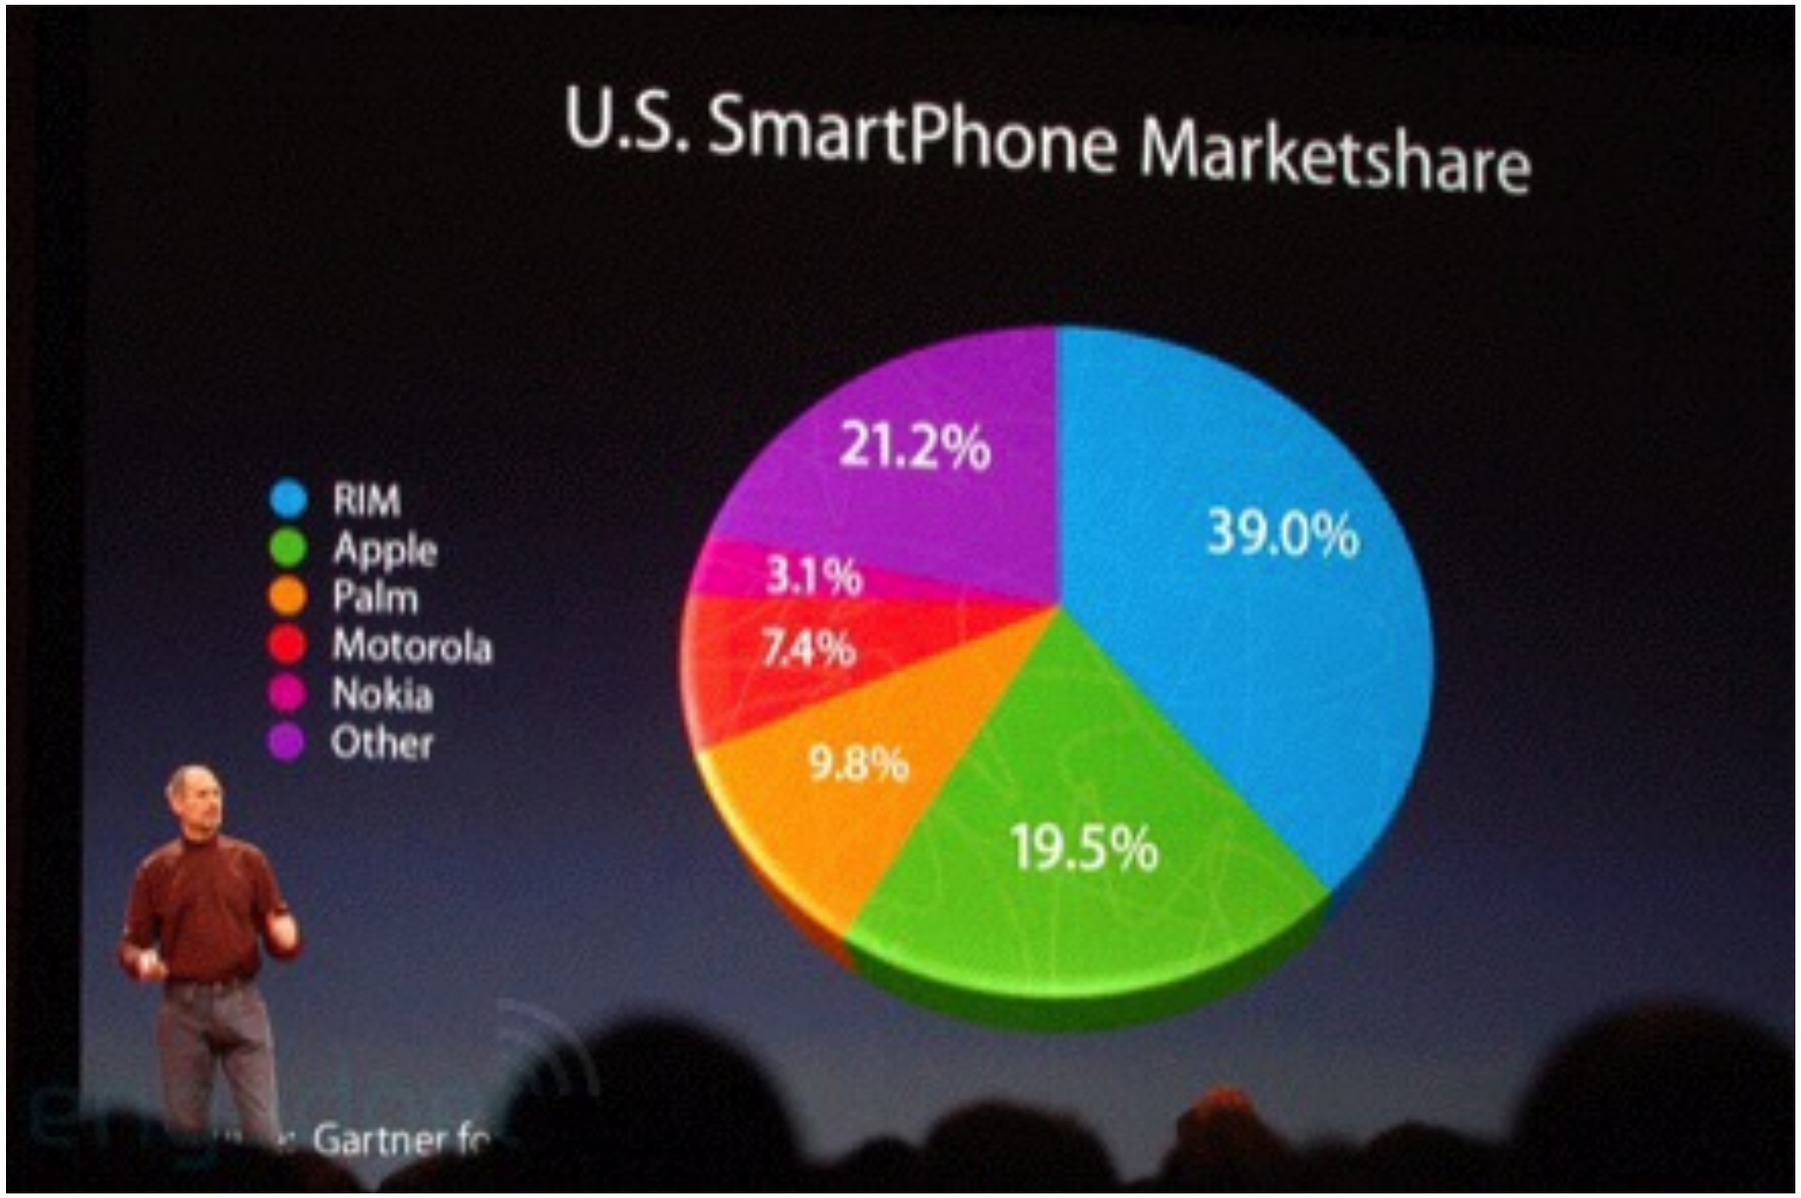

3D Charts!

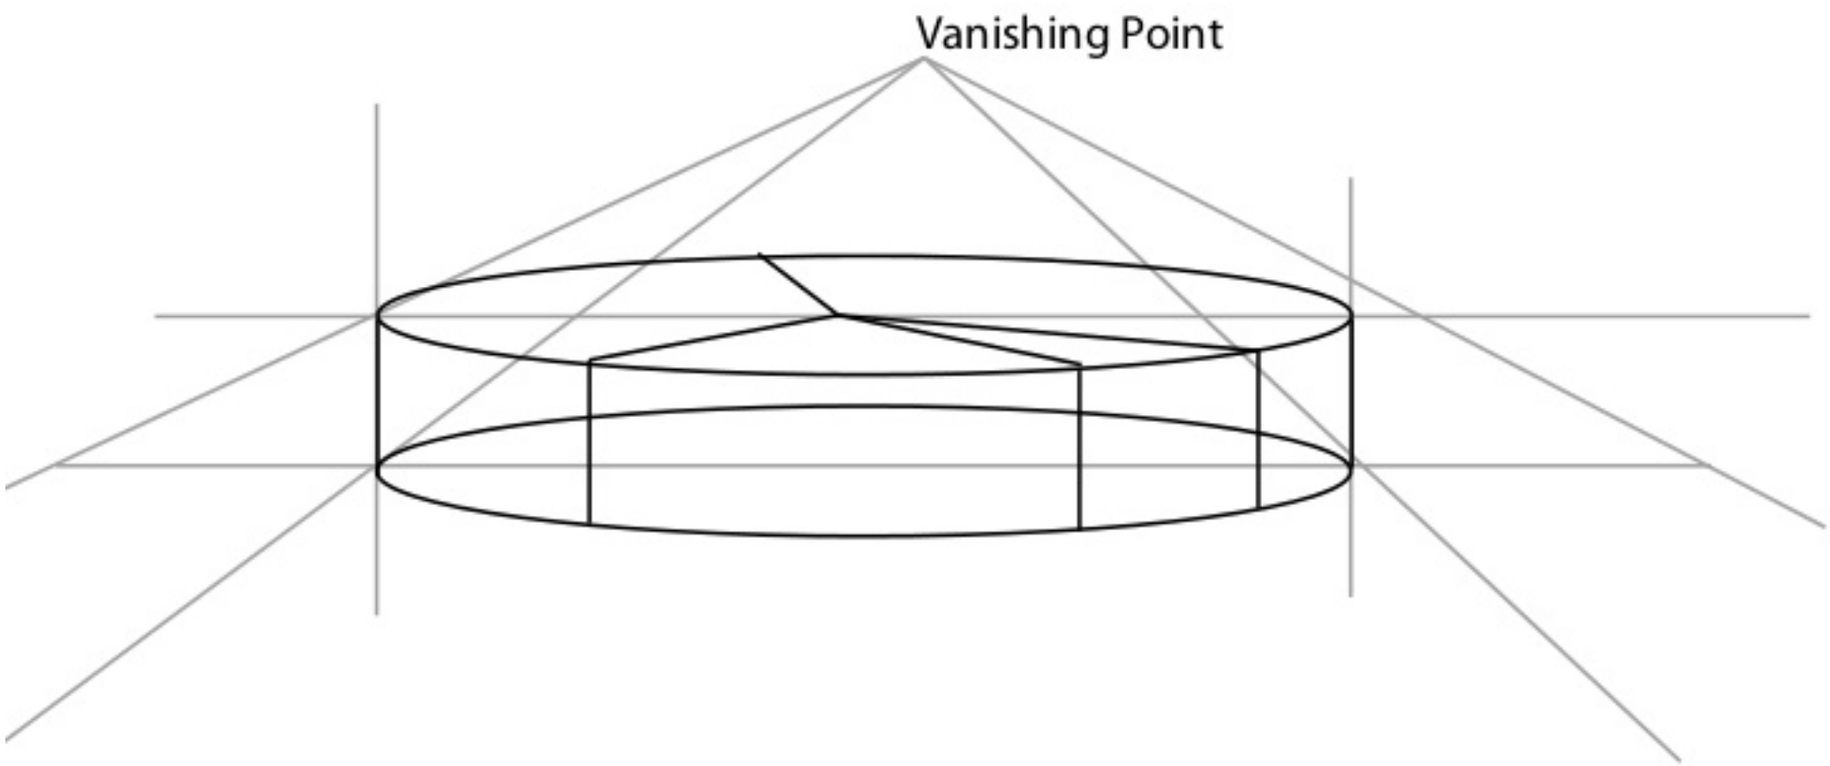

$$\textbf{“lie factor”} = \frac{\text{size of effect shown in graphic}}{\text{size of effect in data}}$$

3. no unjustified 3D

issues with occlusion, perspective distortion, text legibility, …

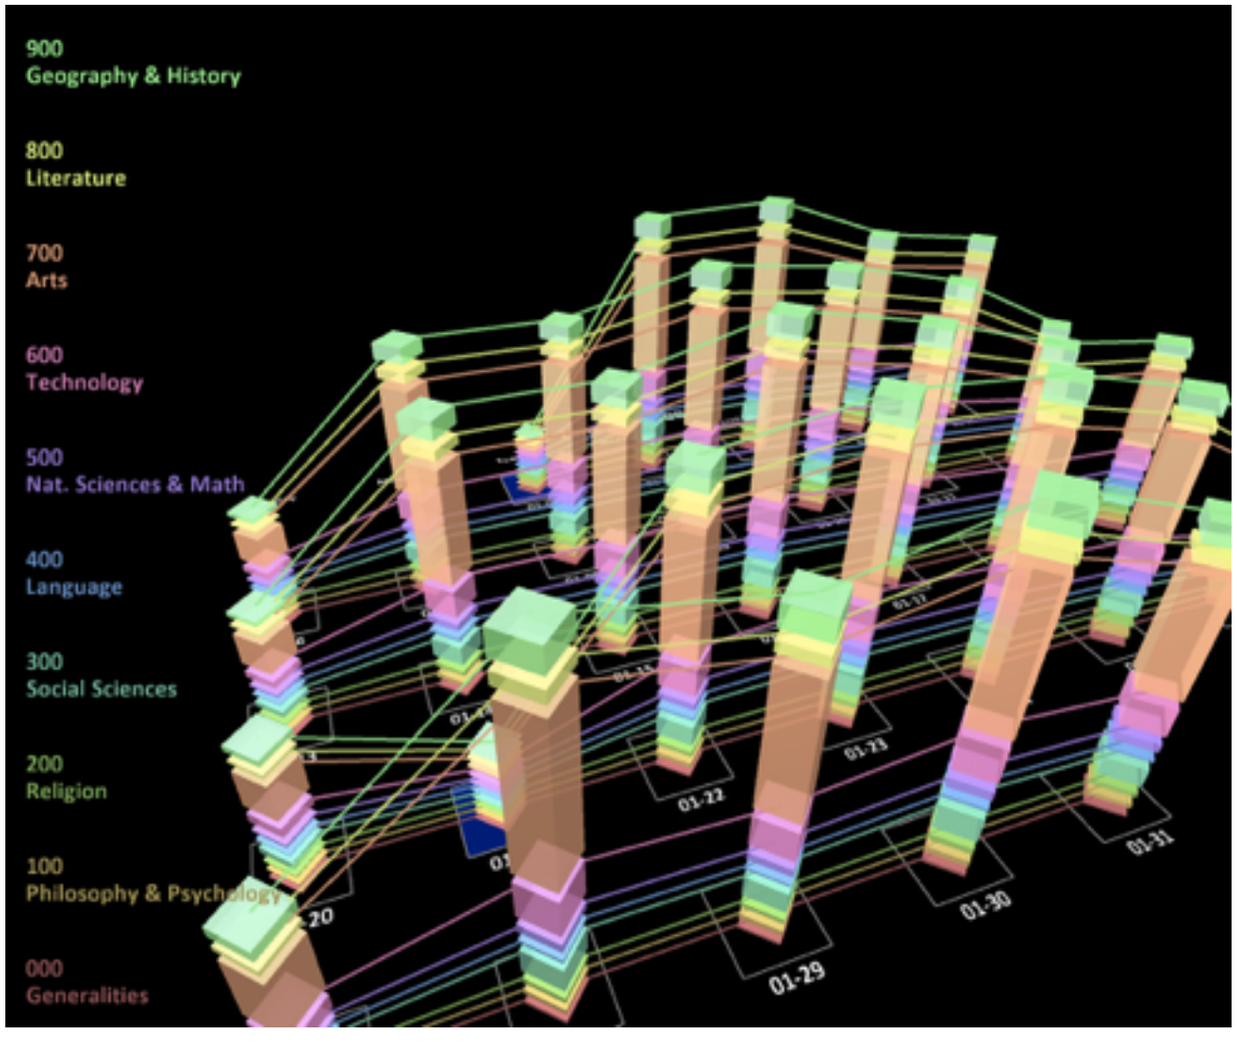

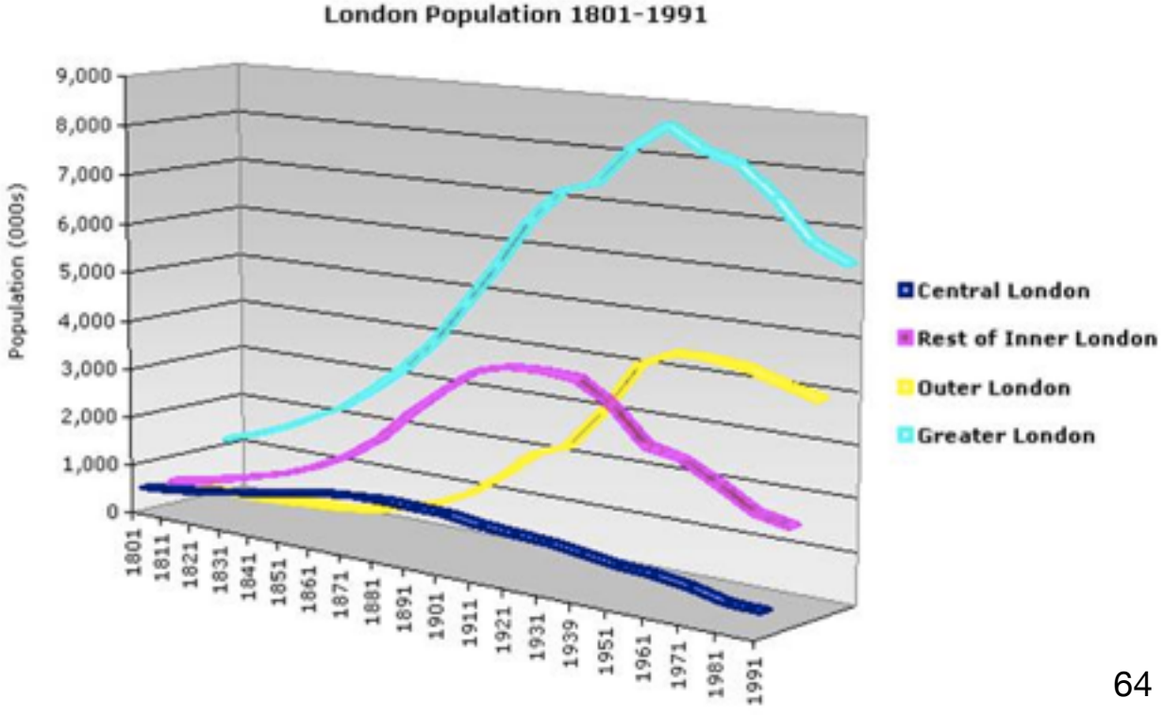

4. eyes beat memory

animation vs side-by-side views

- switch between different views that are visible at same time = lower cognitive load than consulting memory to compare current view with what was seen before

=> try to represent dynamic processes in a static way

5. overview first, zoom & filter, details on demand

task taxonomy Ben Schneiderman:

- *Overview*: see overall patterns in data
- *Zoom*: see a subset of data
- *Filter*: see a subset based on values
- *Detail on demand*: see values of items
- *Relate*: compare values
- *History*: keep track of actions
- *Extract*: mark and captur

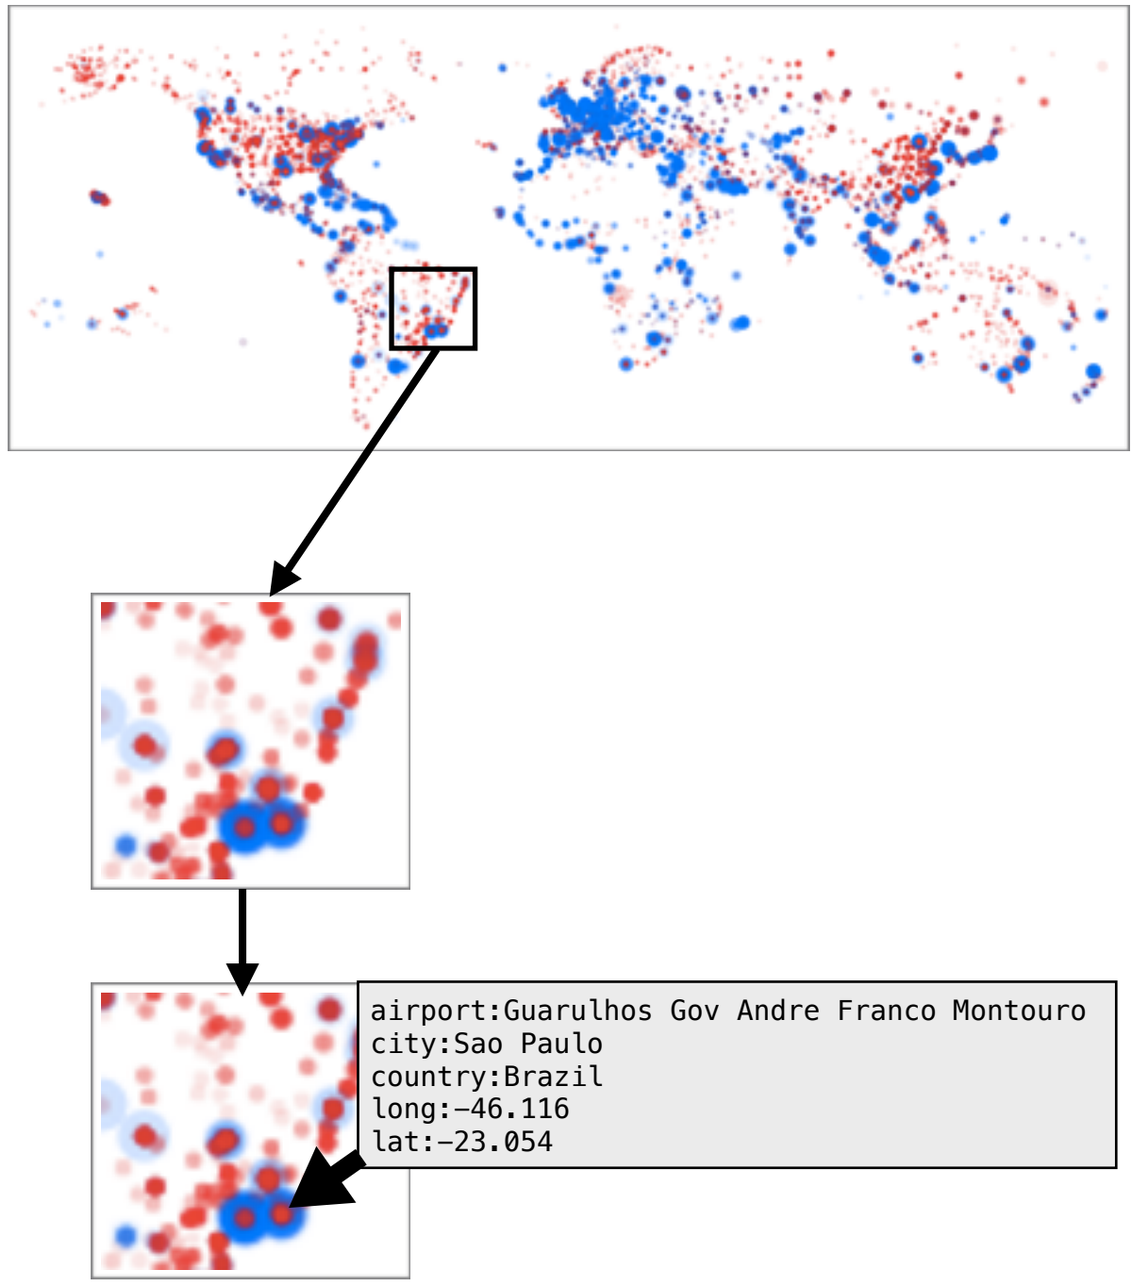

6. “underengineer”, if possible

important advantage of human in the loop vs algorithm: you can take shortcuts

- keep interaction simple (see mouse position & data filter flight patterns; vda-lab.be blog - hands-on visualisation using p5)
- might be OK to not handle edge cases in fast prototyping
- simple raw data visualization can have emergent properties

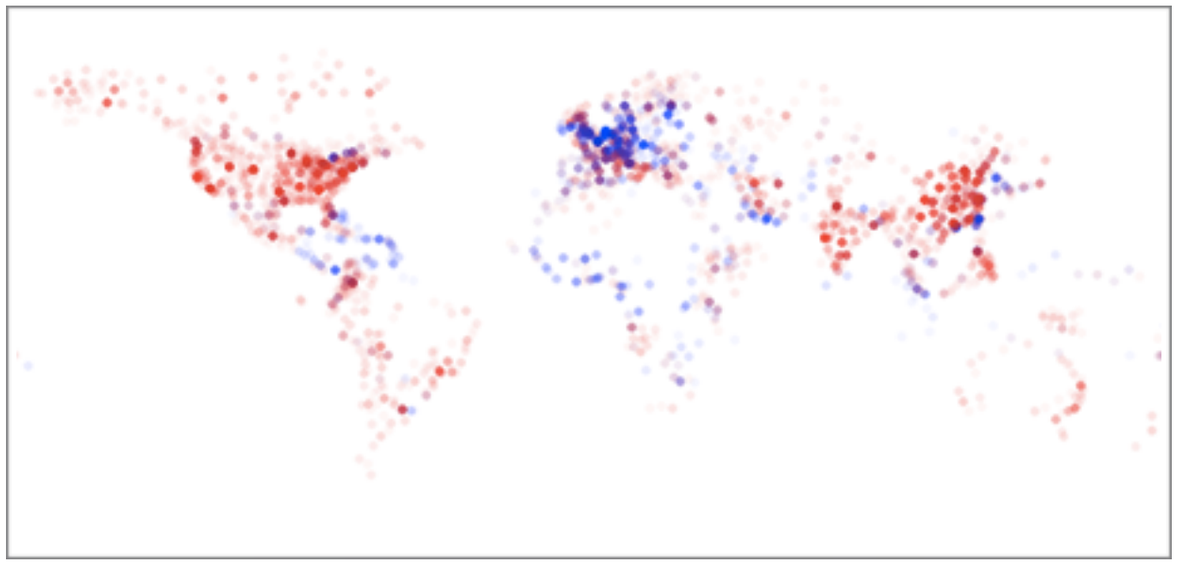

what is drawn: roads
what we see: cities, mountain ranges, population density, …
=> we observe higher-level patterns; not the raw data

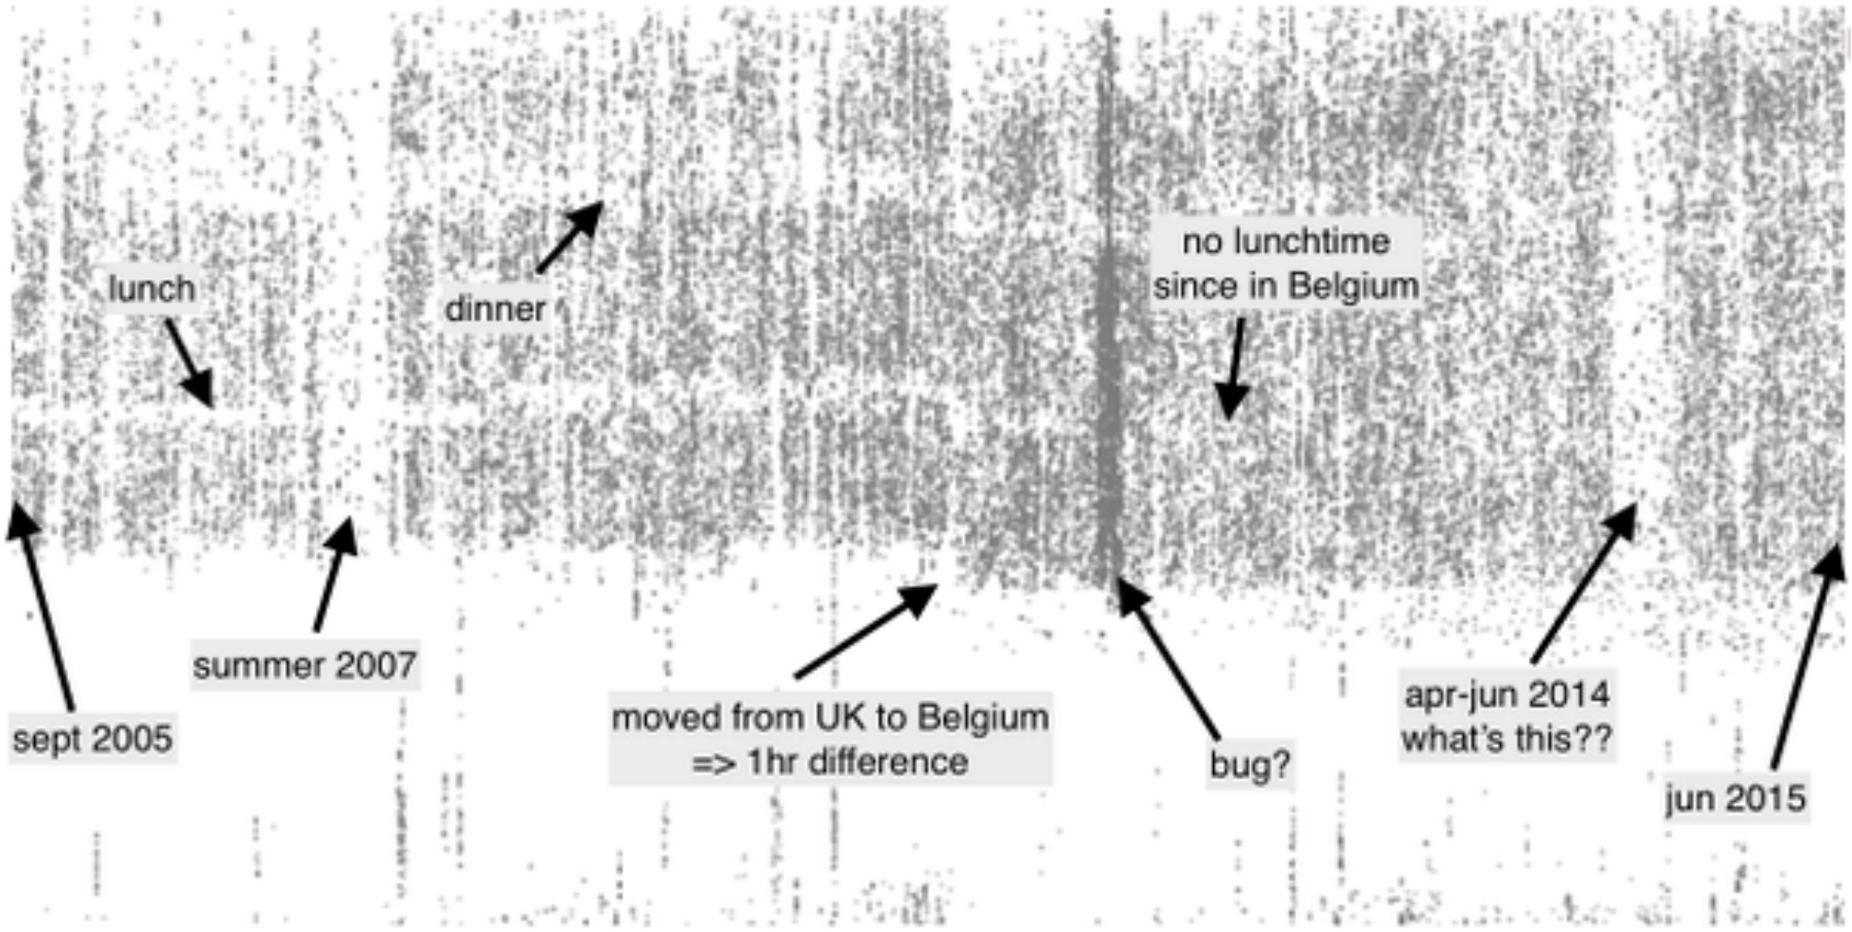

"send" date & time for all my outgoing email since 2005

Visualization foundations and techniques

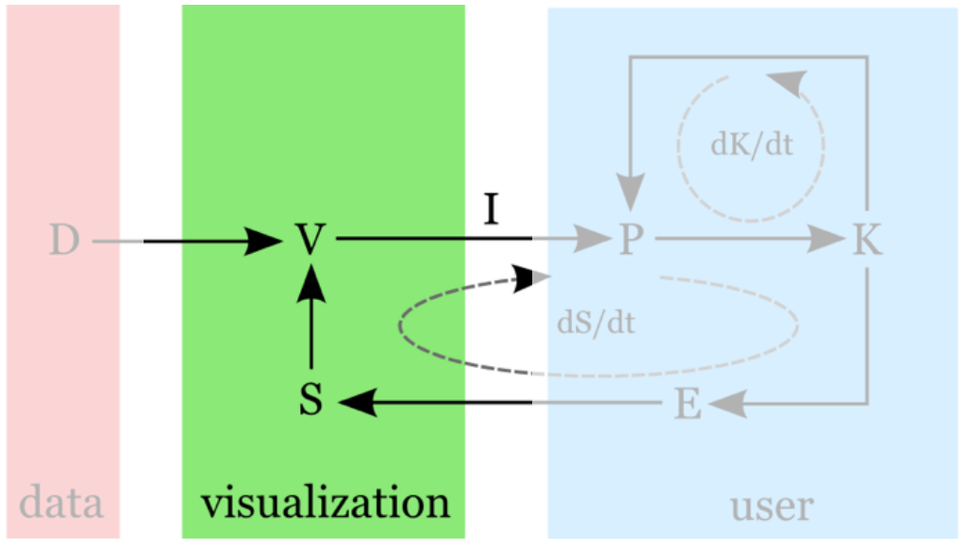

Mapping data to marks with properties

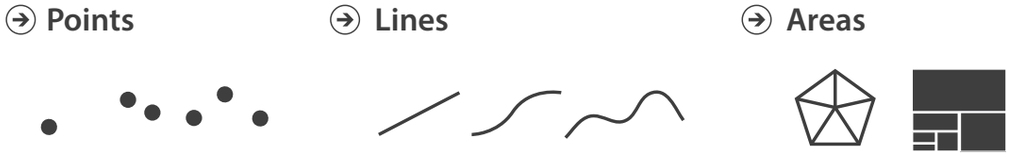

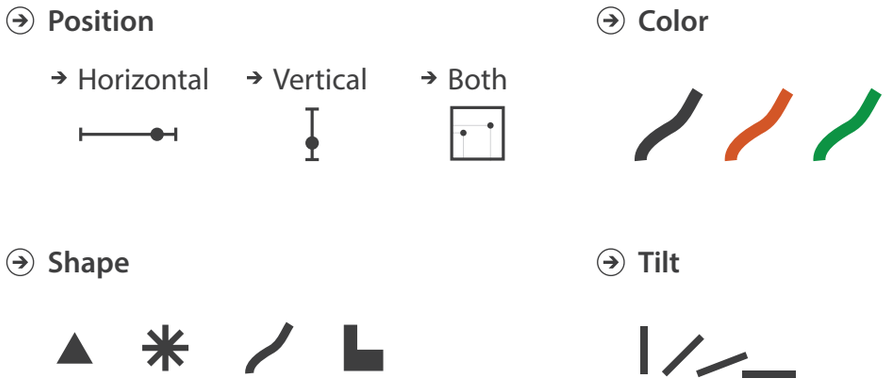

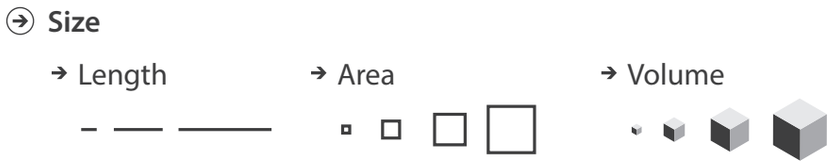

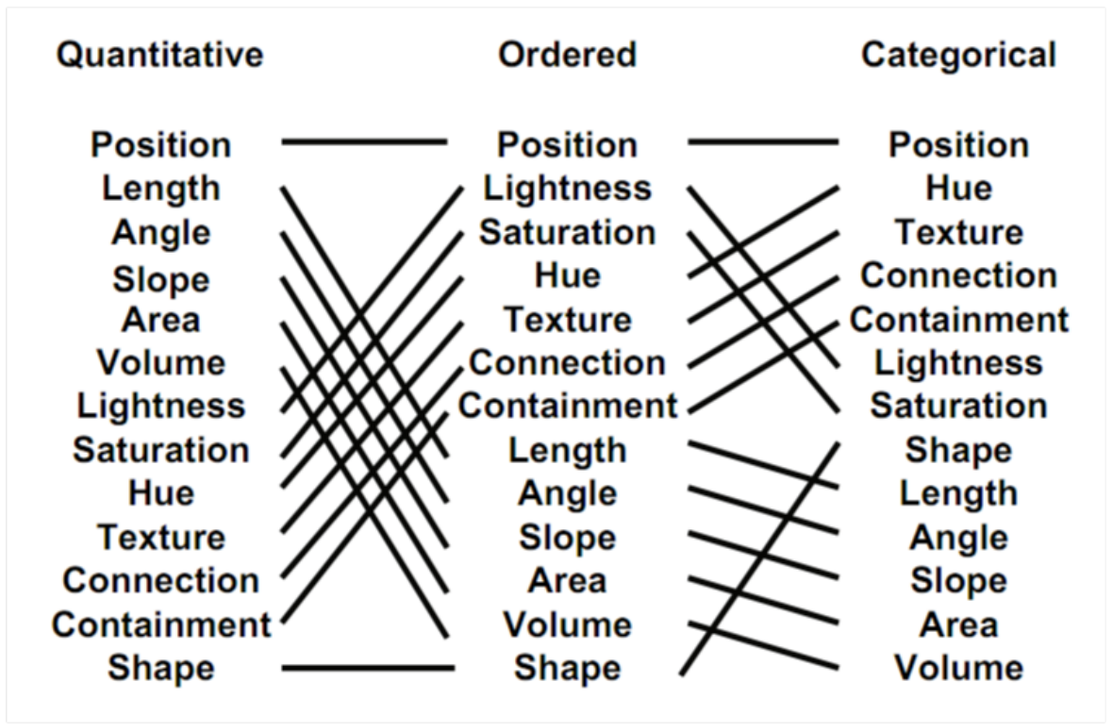

Combining visuals

- **juxtapose**/facet, e.g. *small multiples*
  - easier to see trends
  - pop-out effect
  - very powerful!

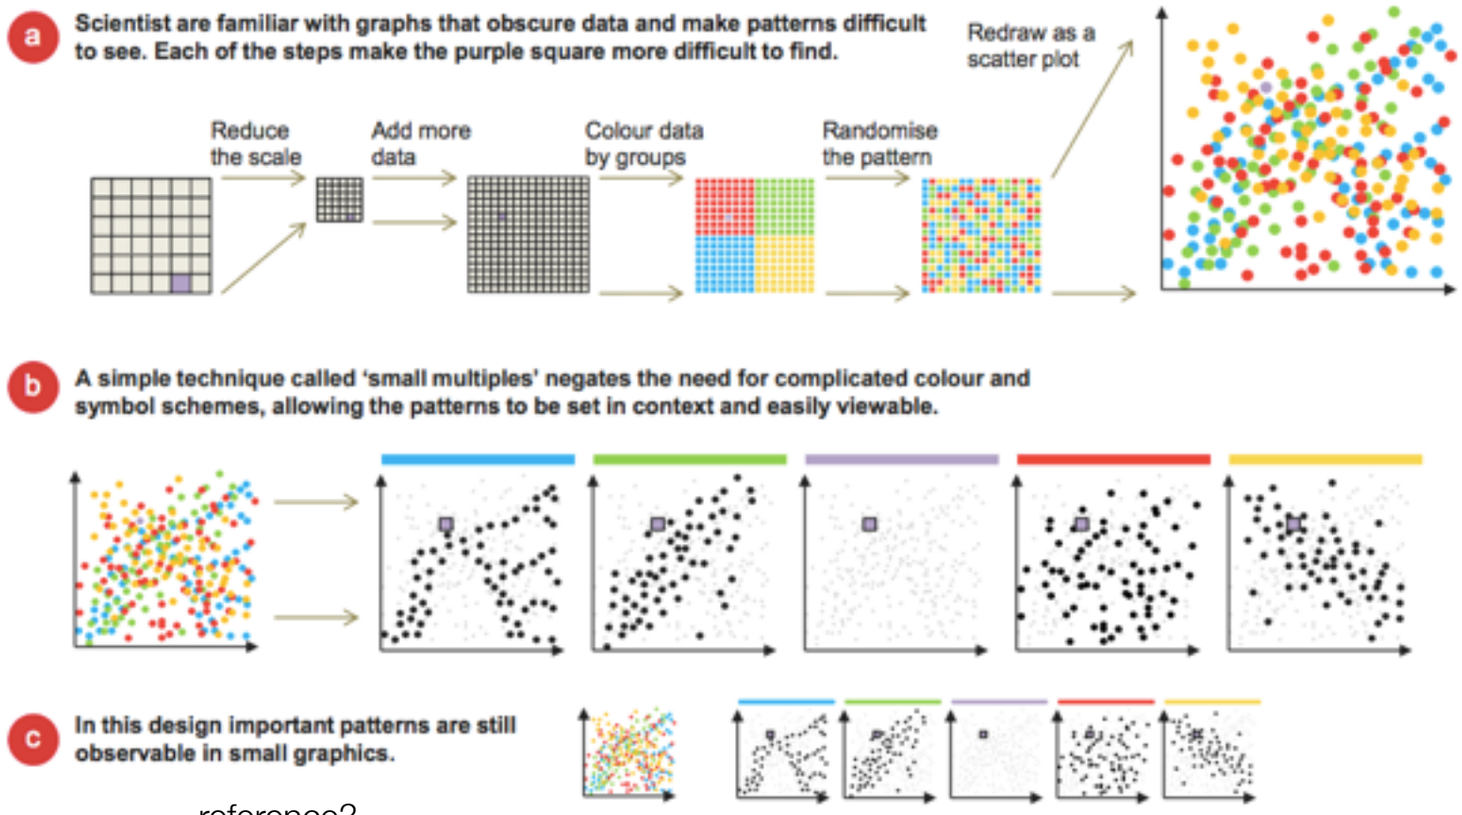

reference?

L1000 data

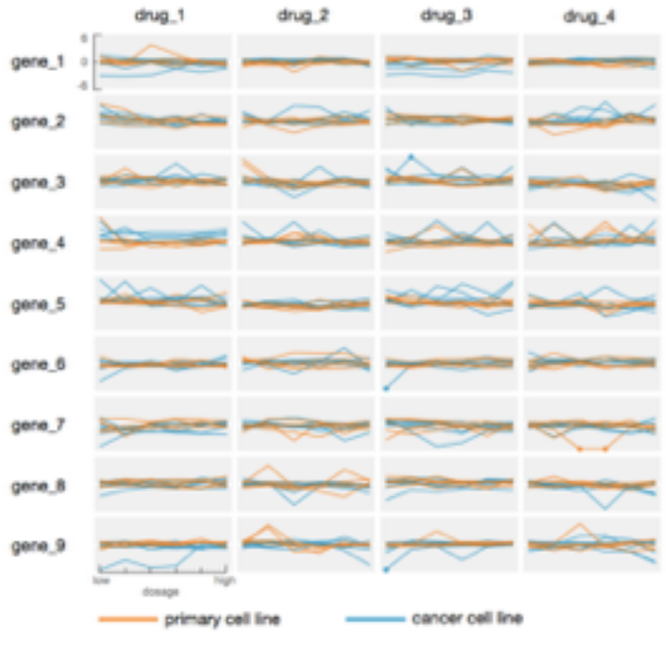

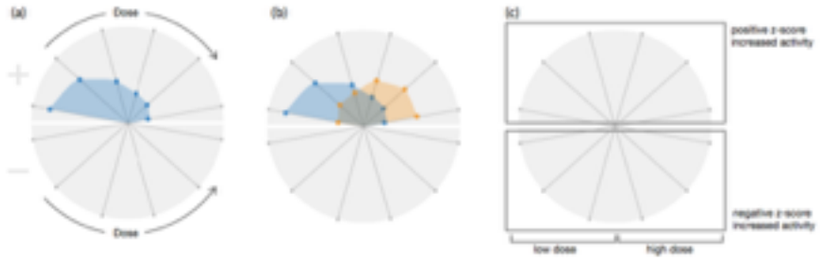

- **superimpose**

Hadlak et al, 2015

- **embed**/nest
- e.g. network + map: each can be “base”

Hadlak et al, 2015

Creating data visualisations - the process

1. Get to know the user
2. Get to know the data
3. Create visual designs
4. Iterate

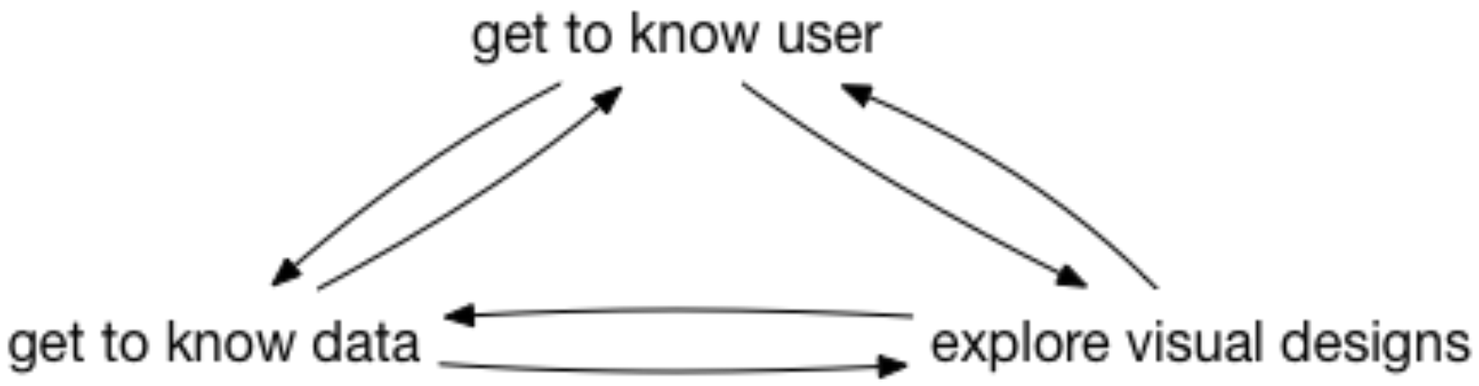

Design decision styles

1. **Unintended design** - Design decisions based on what's easiest to implement. Developer focuses on development and deployment without any consideration of what will happen when people use the tool.

2. **Self design** - Design decisions based on by developer's own use.

3. **Genius design** - Developer still does not look beyond own experience, but that experience is extensive.

4. **Task-focused design** - Developer investigates which actions the user wants/needs to perform.

5. **Goal-focused design** - Developer goes further than activities and investigates goals, needs and contexts of the user.

1. Get to know the user

1. Talk to the user
   - what they *want* != what the *need* => need to find **underlying goals**
   - e.g. let them **imagine** what they could do if some technologies were available that are (still) science-fiction (e.g. nanobots in blood; Gaviscon commercial http://m.youtube.com/watch?v=_skKmcLdyVQ)
   - additional methods, e.g. **card sorting**
   - if possible: tape the discussion (w/ agreement)

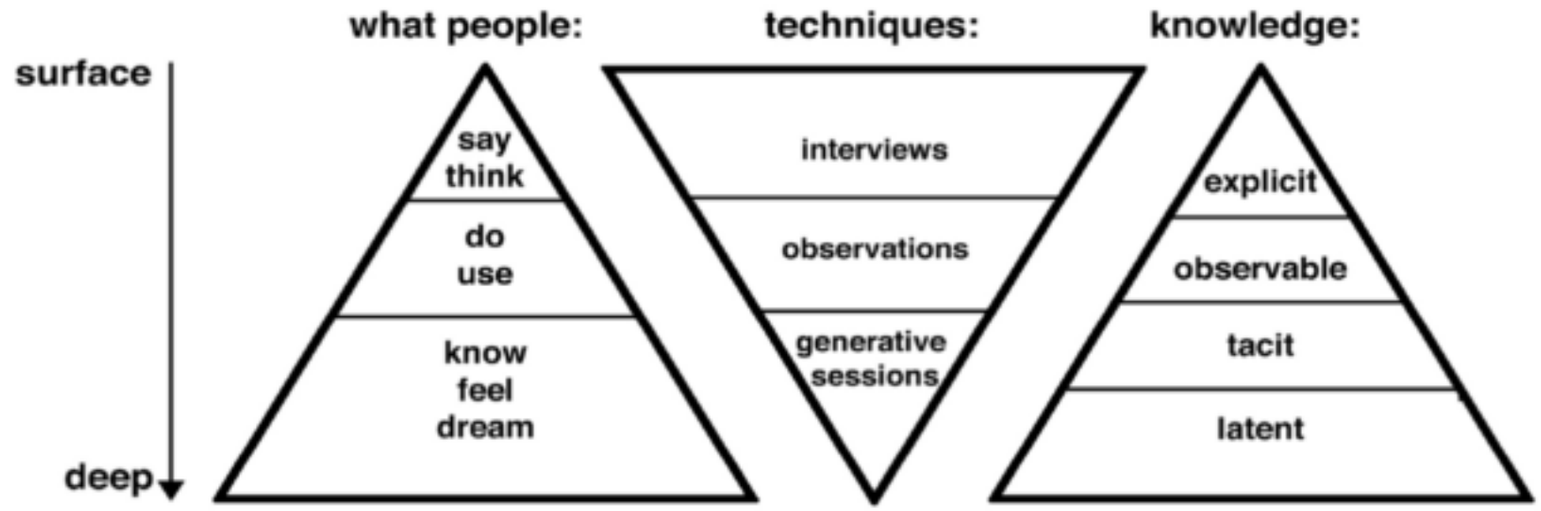

Different levels of knowledge about experience are accessed by different techniques

Visser et al, 2005

2. Post-hoc analysis of interview (e.g. using Munzner taxonomies)

- data abstraction
- task abstraction

Actions

Analyze

Consume

- *Discover*
- *Present*
- *Enjoy*

Produce

- *Annotate*
- *Record*
- *Derive*

Search

	Target known	Target unknown
Location known	*Lookup*	*Browse*
Location unknown	*Locate*	*Explore*

Query

- Identify
- Compare
- Summarise

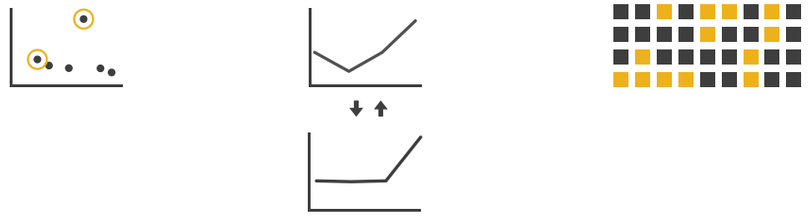

2. Get to know the data

Questions to answer:

- how many **dimensions**? What are the **types** of dimensions (categorical, numerical, geo-spatial, …)?
- For each dimension: what is **distribution** of datapoints?
- Are there any **correlations** between dimensions?
- What does **principal component analysis** or **singular value decomposition** reveal?
- What does **hierarchical clustering** show?
- Are there any **local clusters**? E.g. use **topological data analysis**

3. Create visual designs

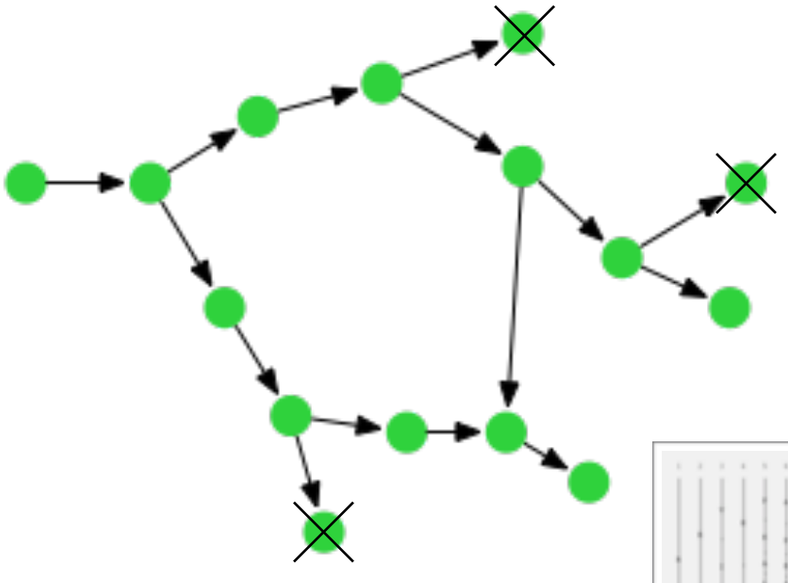

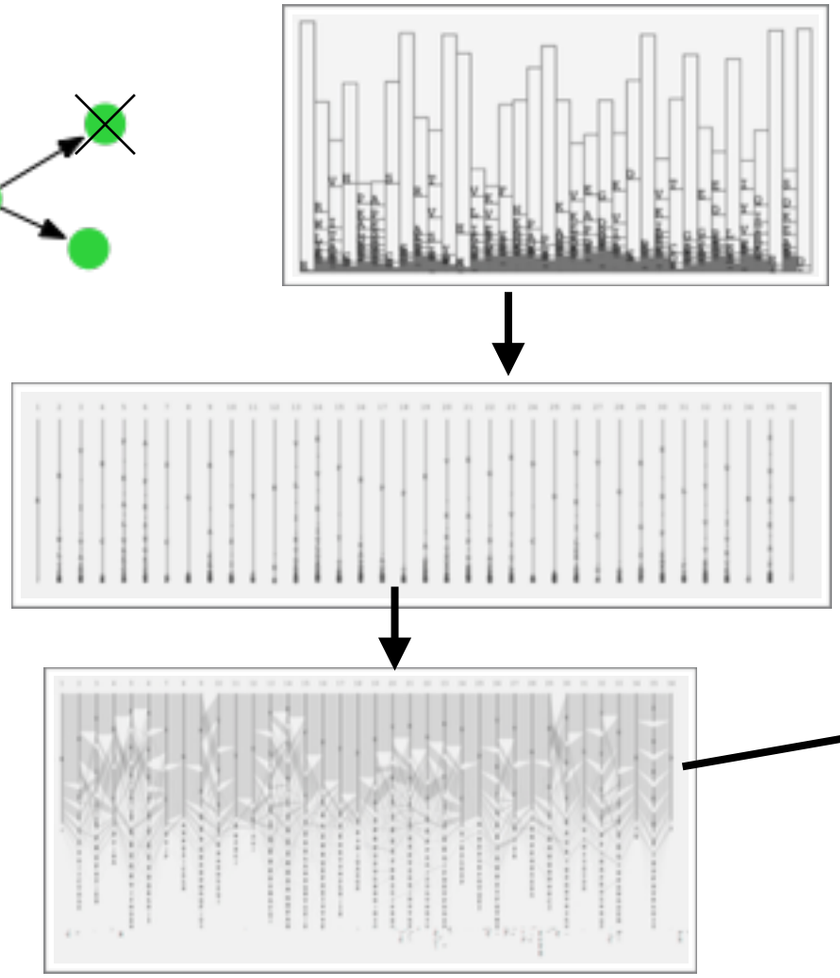

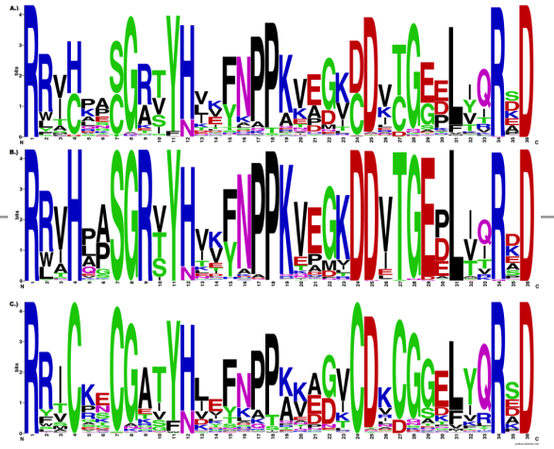

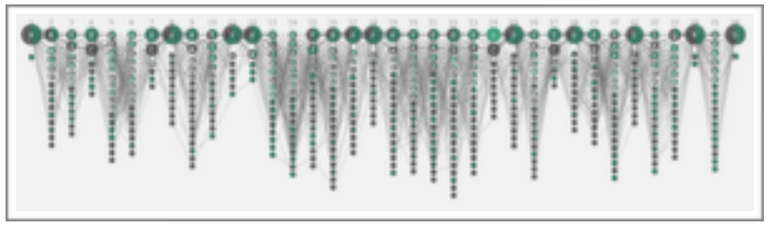

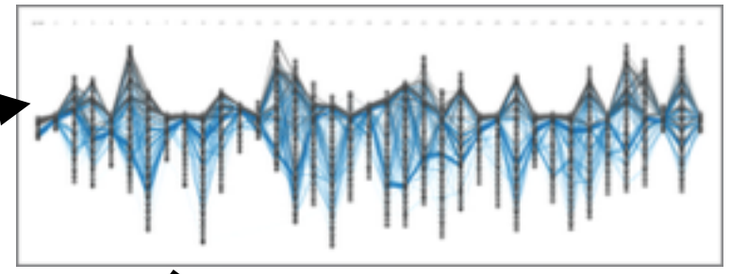

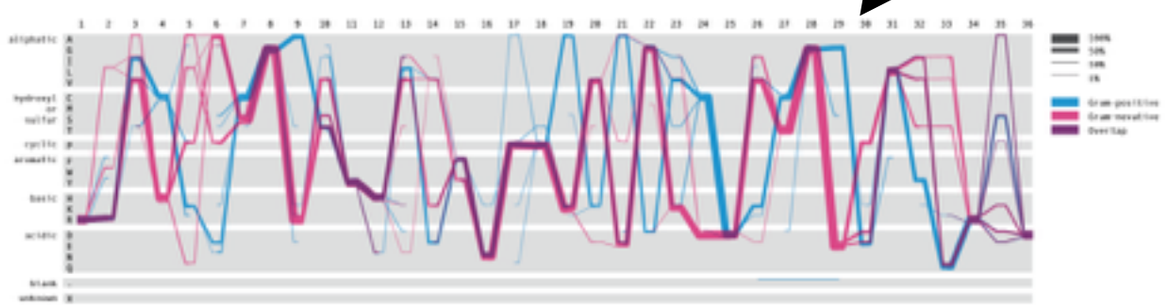

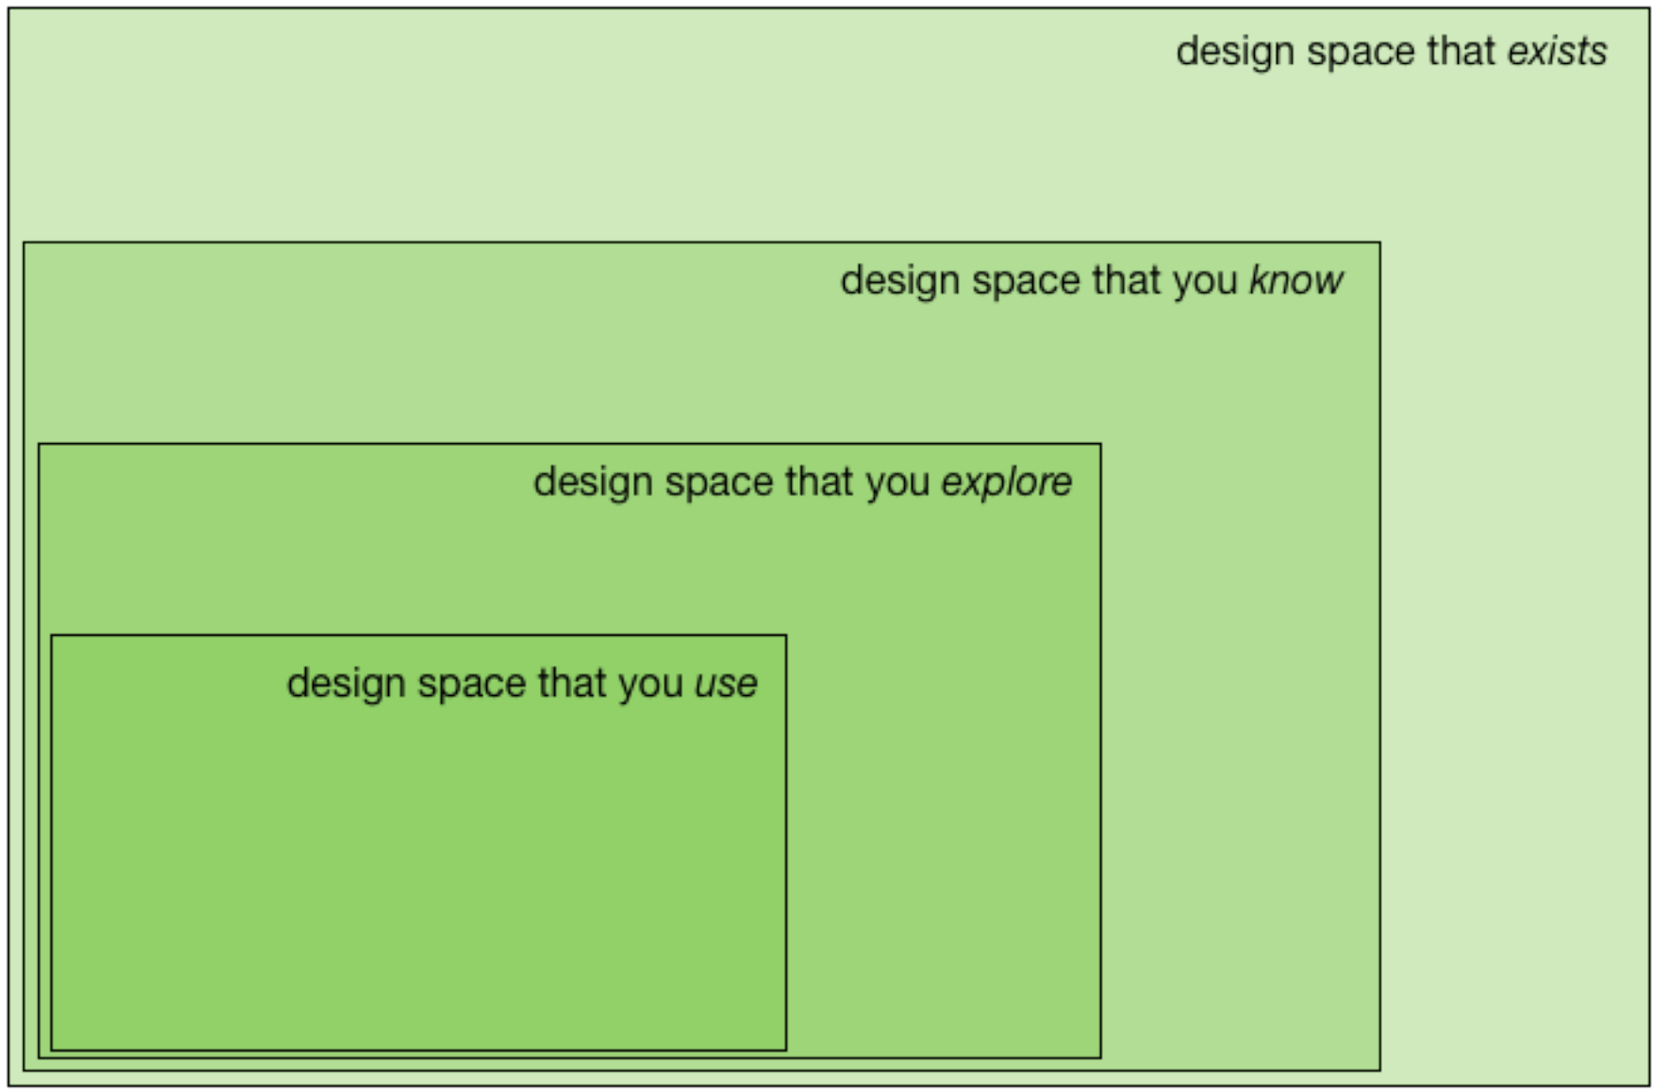
design space that *exists*
design space that you *know*
design space that you *explore*
design space that you *use*

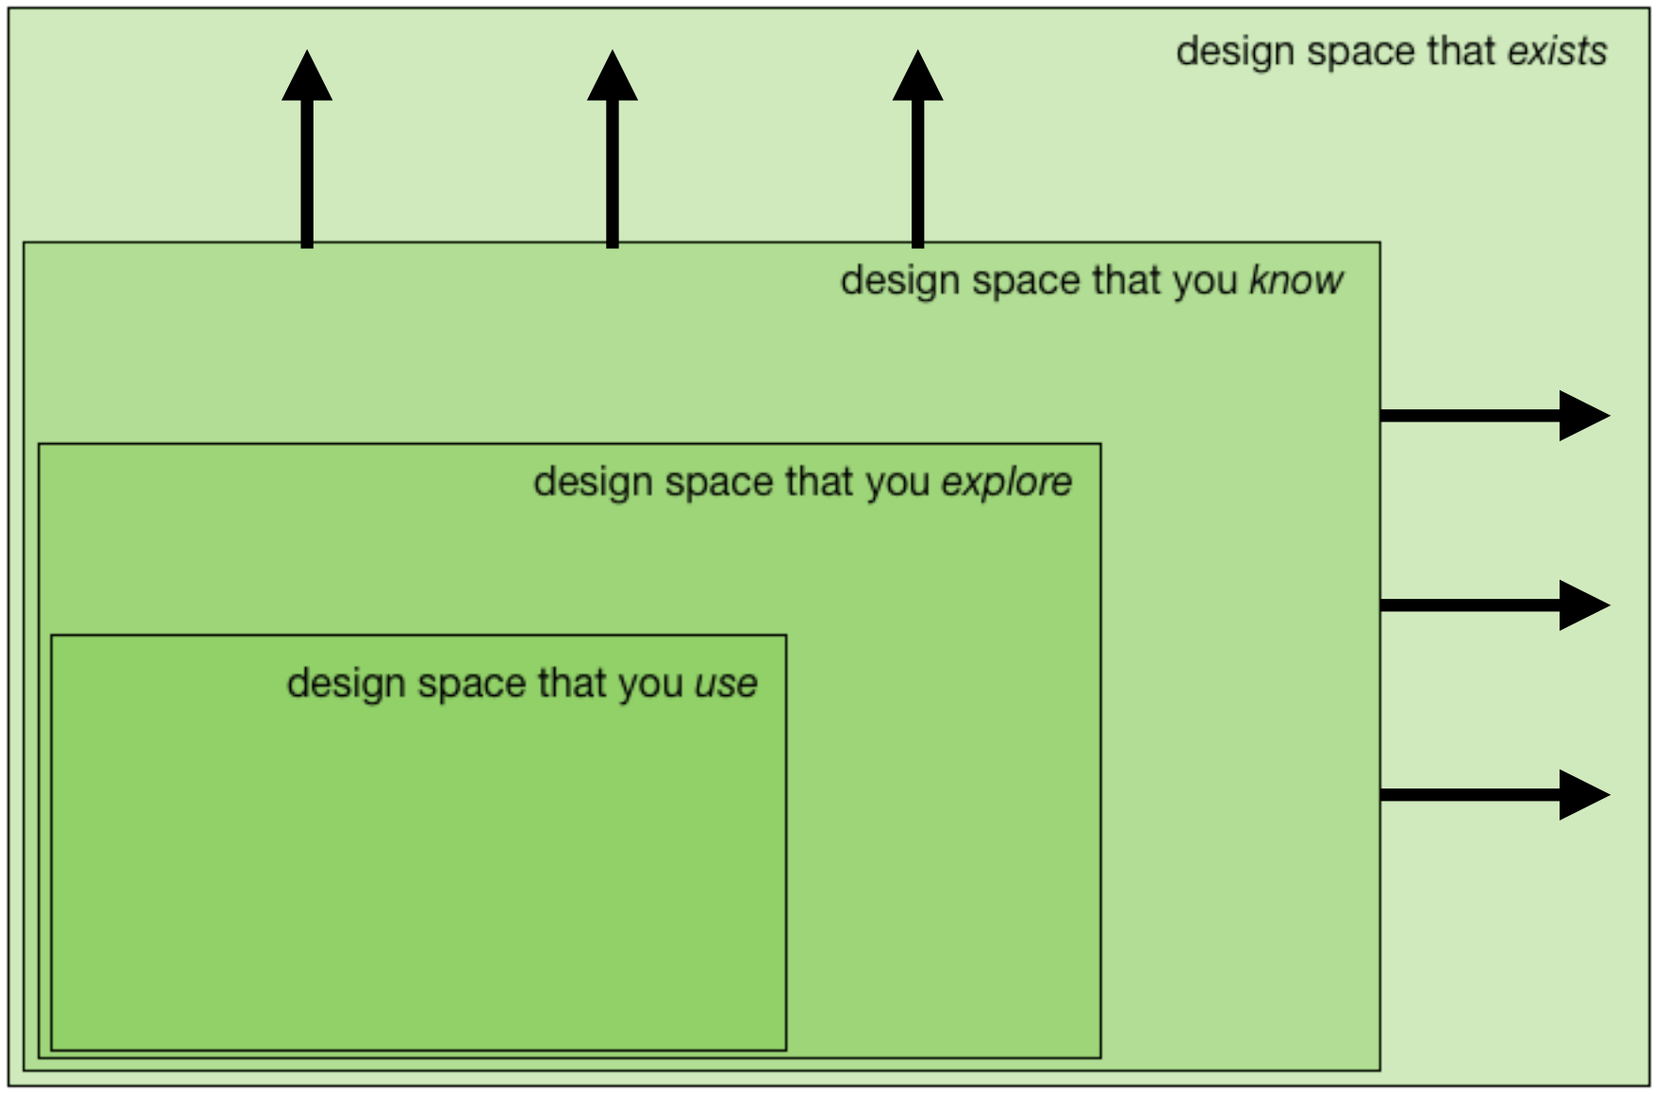

problem: initial design space that you know is small => how to expand?

= “design fixation”

Tips to explore design space:

- Use systematic approach such as 5-design-sheet (see paper)

- Make very “cheap” sketches: pen & paper => don’t get attached to a design because you spent a lot of time on it

- Extend your visual library: collect interesting visuals (useful *and* useless), e.g. https://www.pinterest.com/aertsjan/data-visualizationart/

- Can you go to a higher abstraction level?

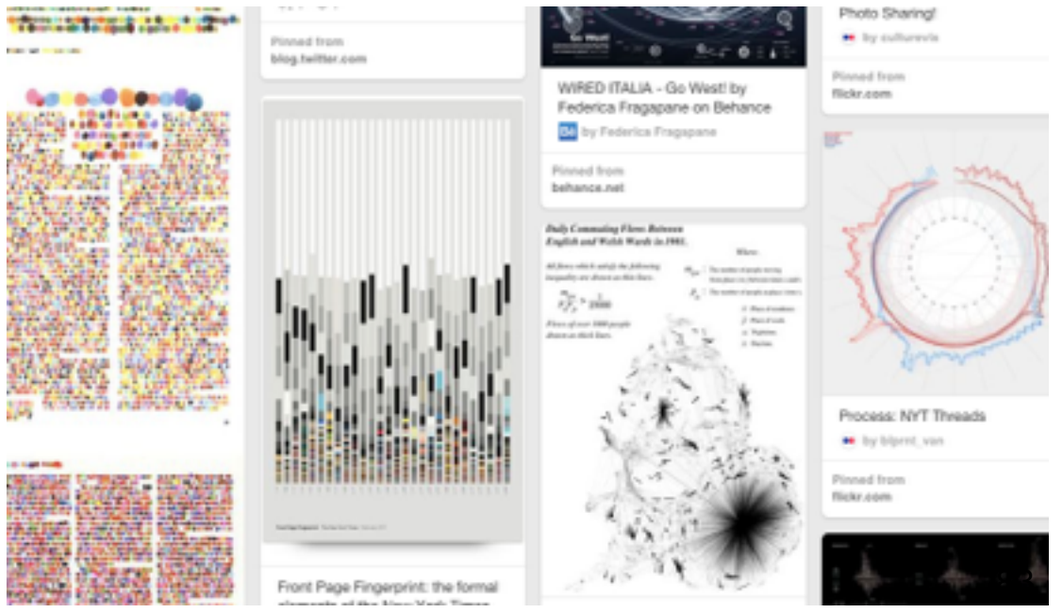

5-design sheet methodology

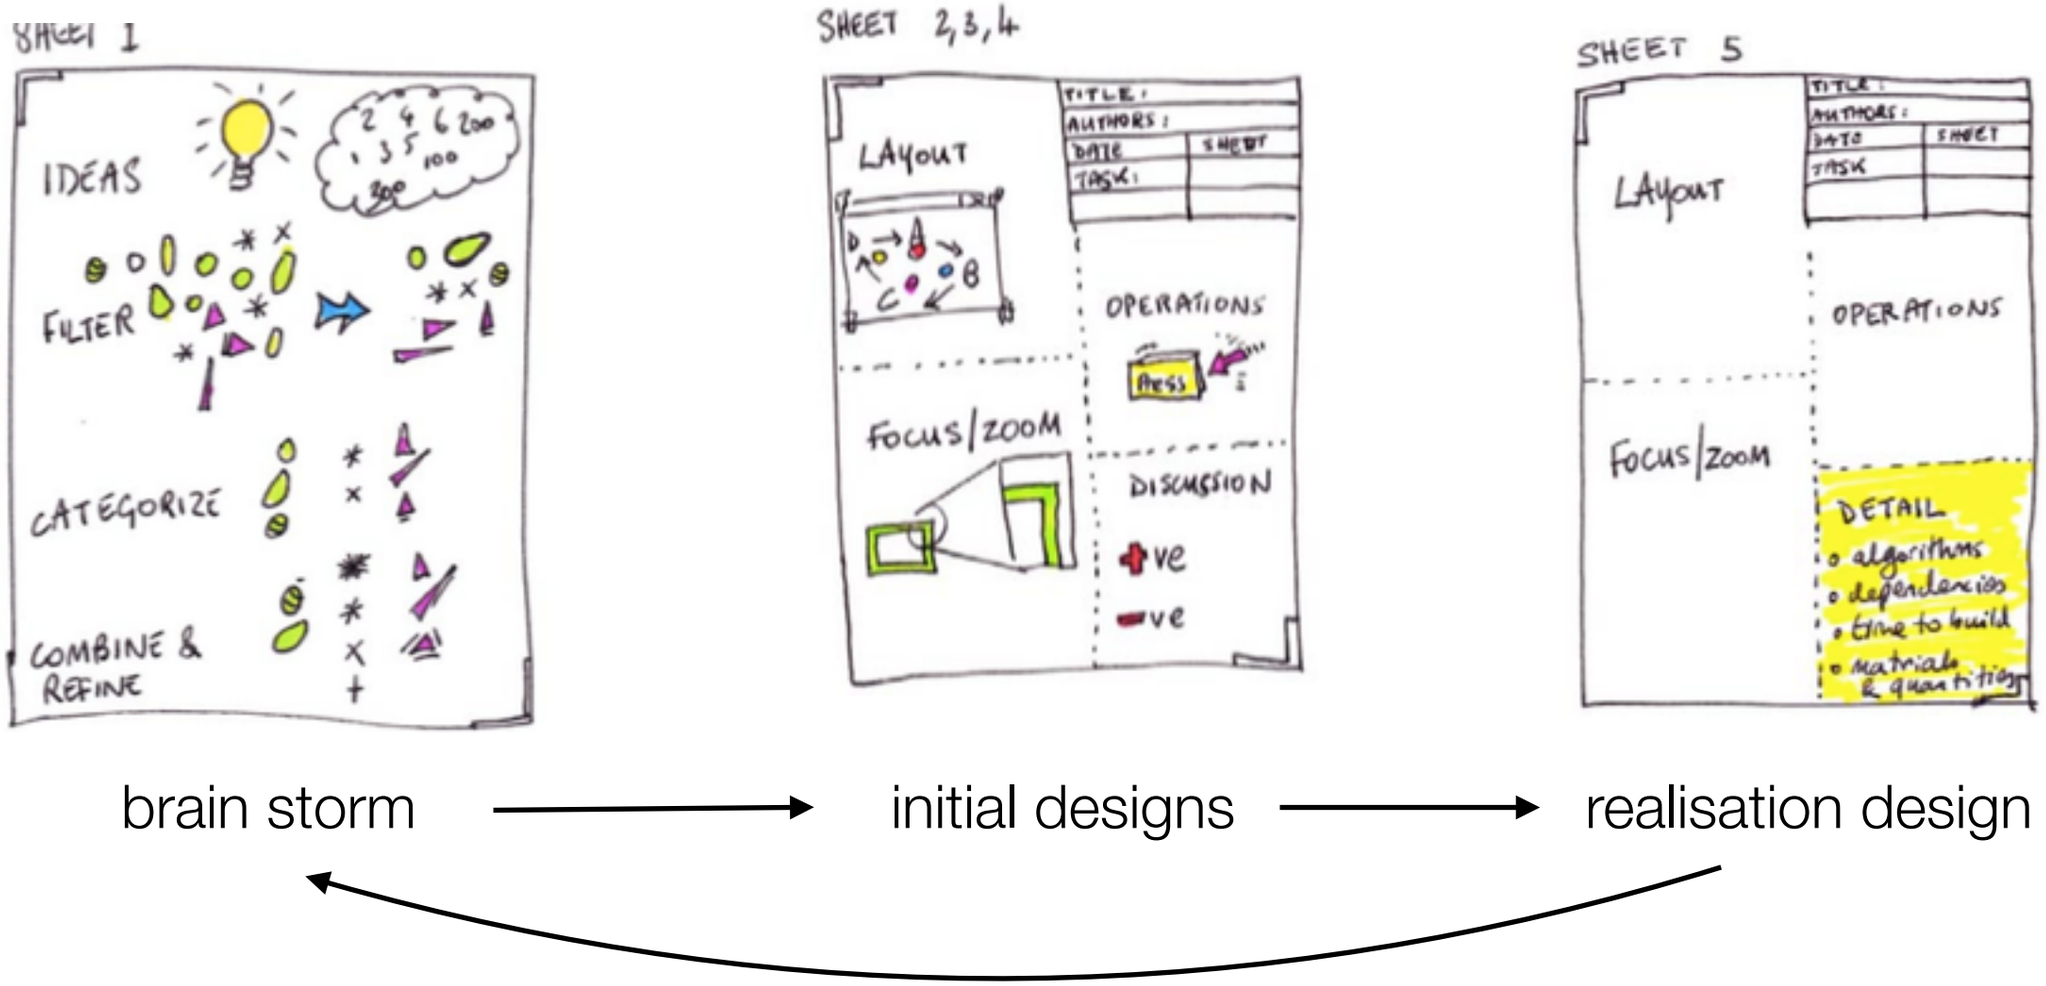

Data visualization framework

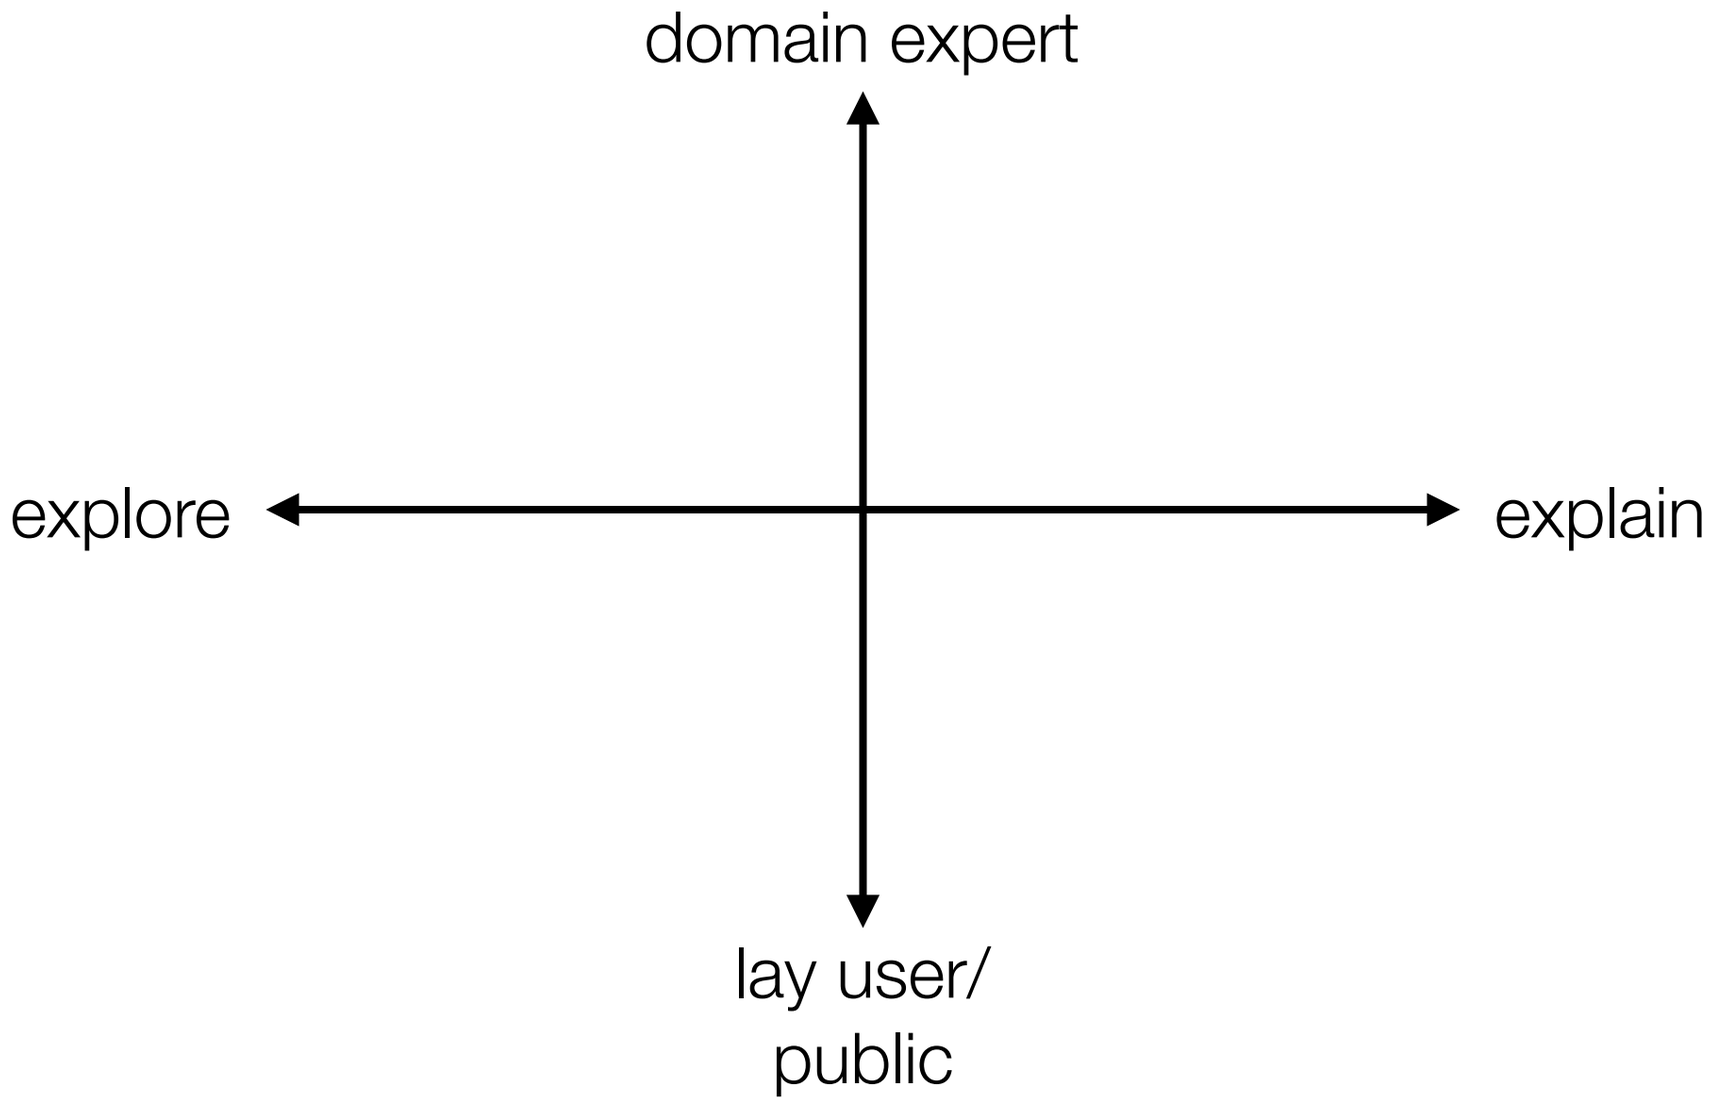

Bret Victor - Ladder of abstraction

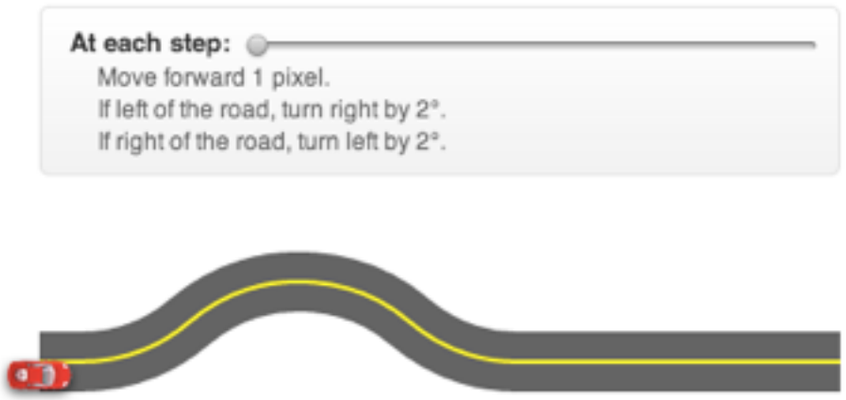

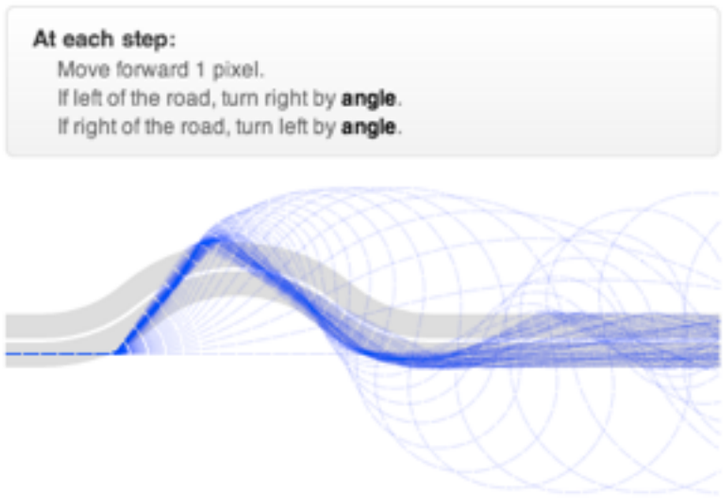

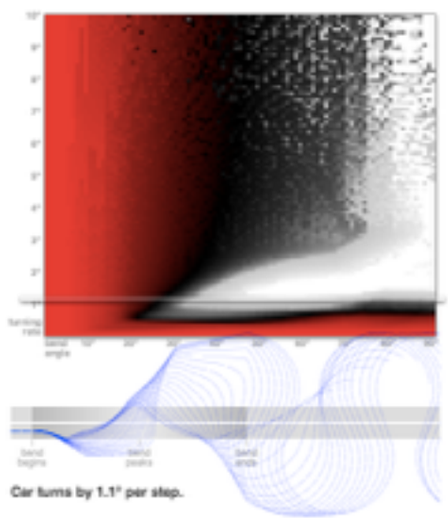

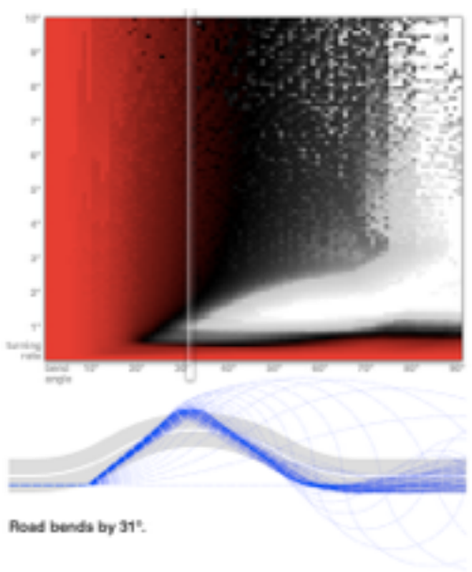

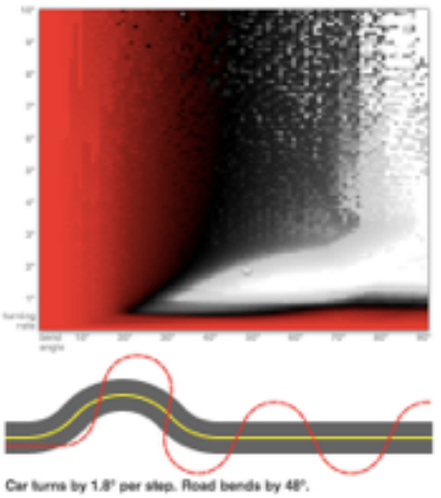

Bret Victor - Ladder of abstraction

literal
lay user/public

abstract
domain expert

Some examples

very (!) limited scattering of visuals, only to indicate the breadth of possibilities

See D3.js examples (https://github.com/mbostock/d3/wiki/Gallery)

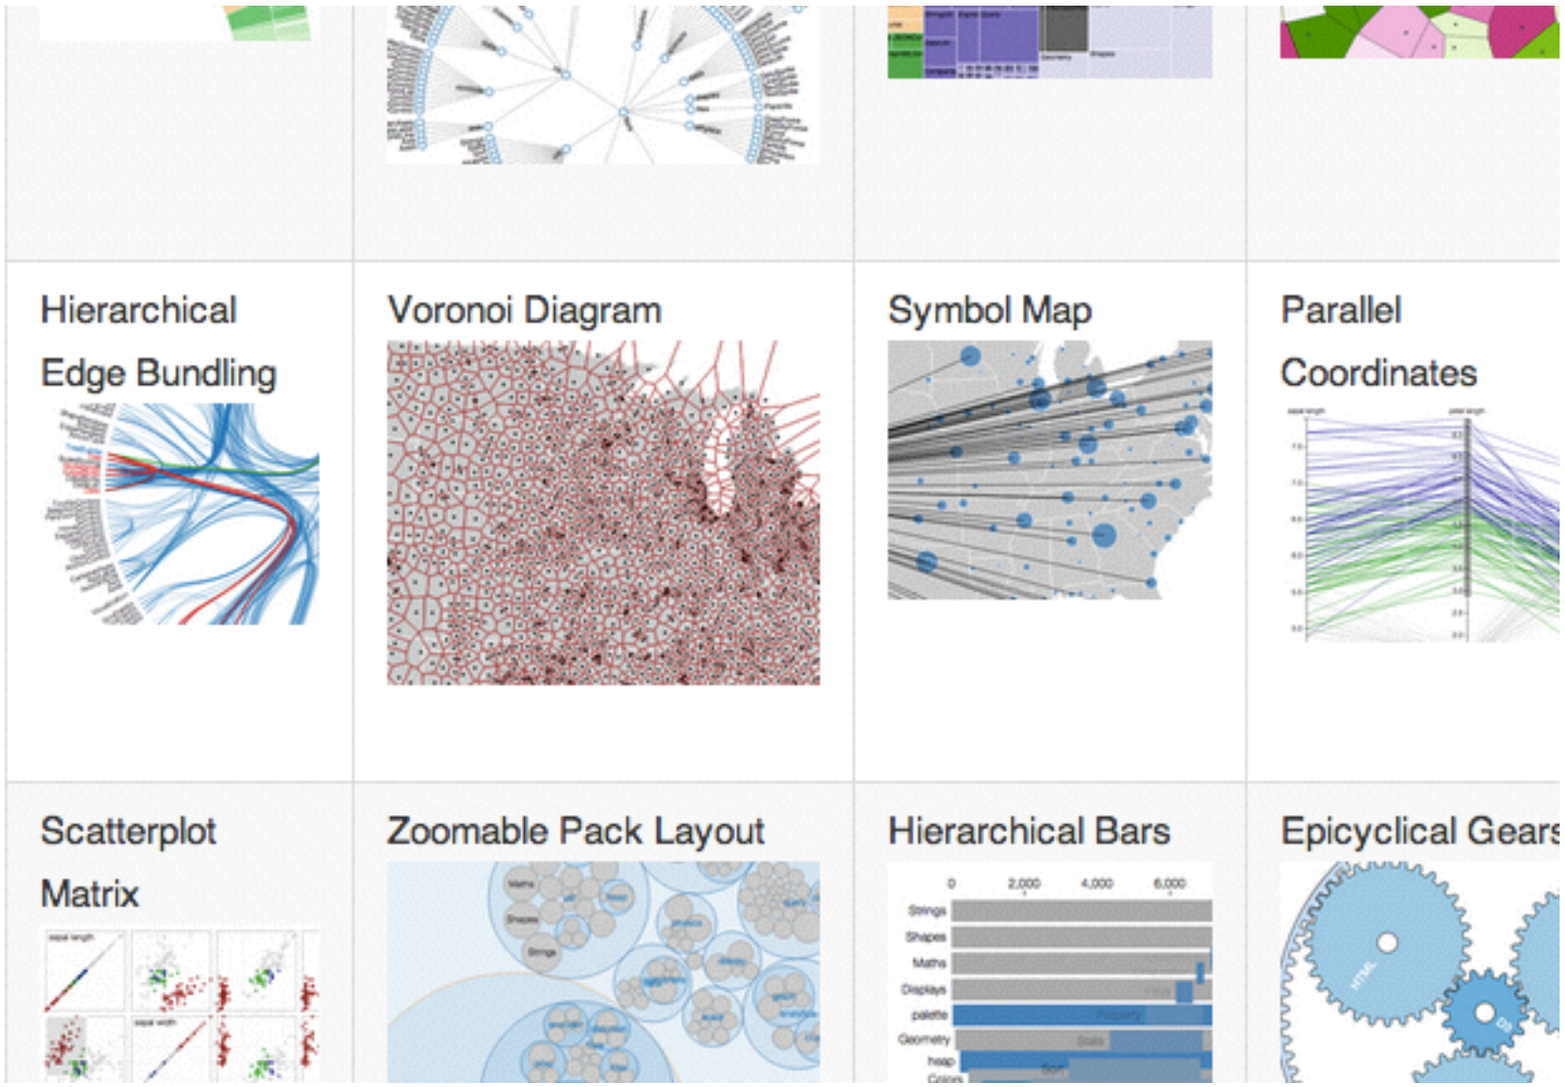

1D, 2D

space-filling curves: 1D -> 2D

e.g. Hilbert curve

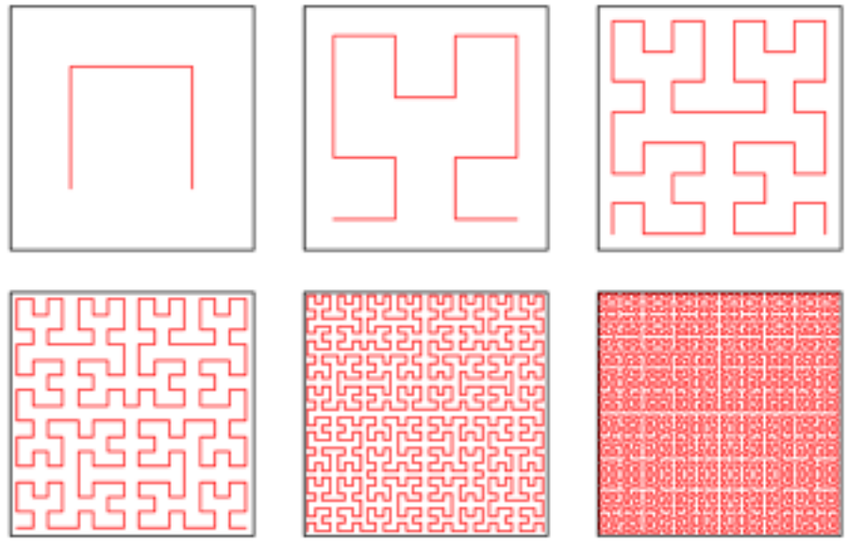

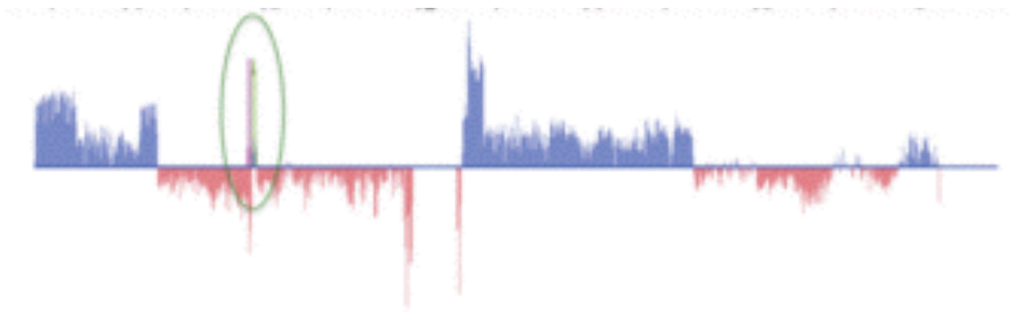

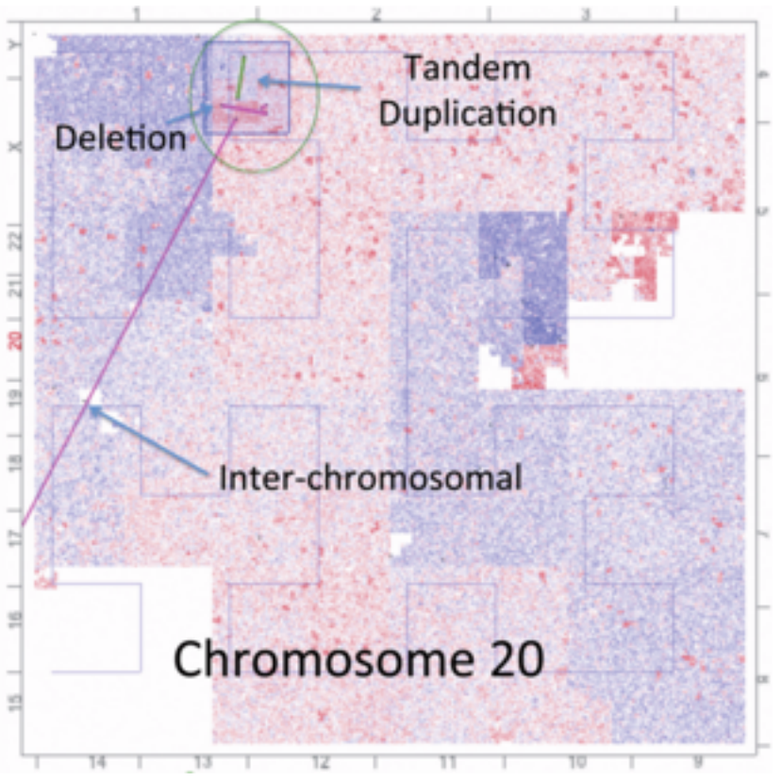

Iris setosa

Iris versicolor

Iris virginica

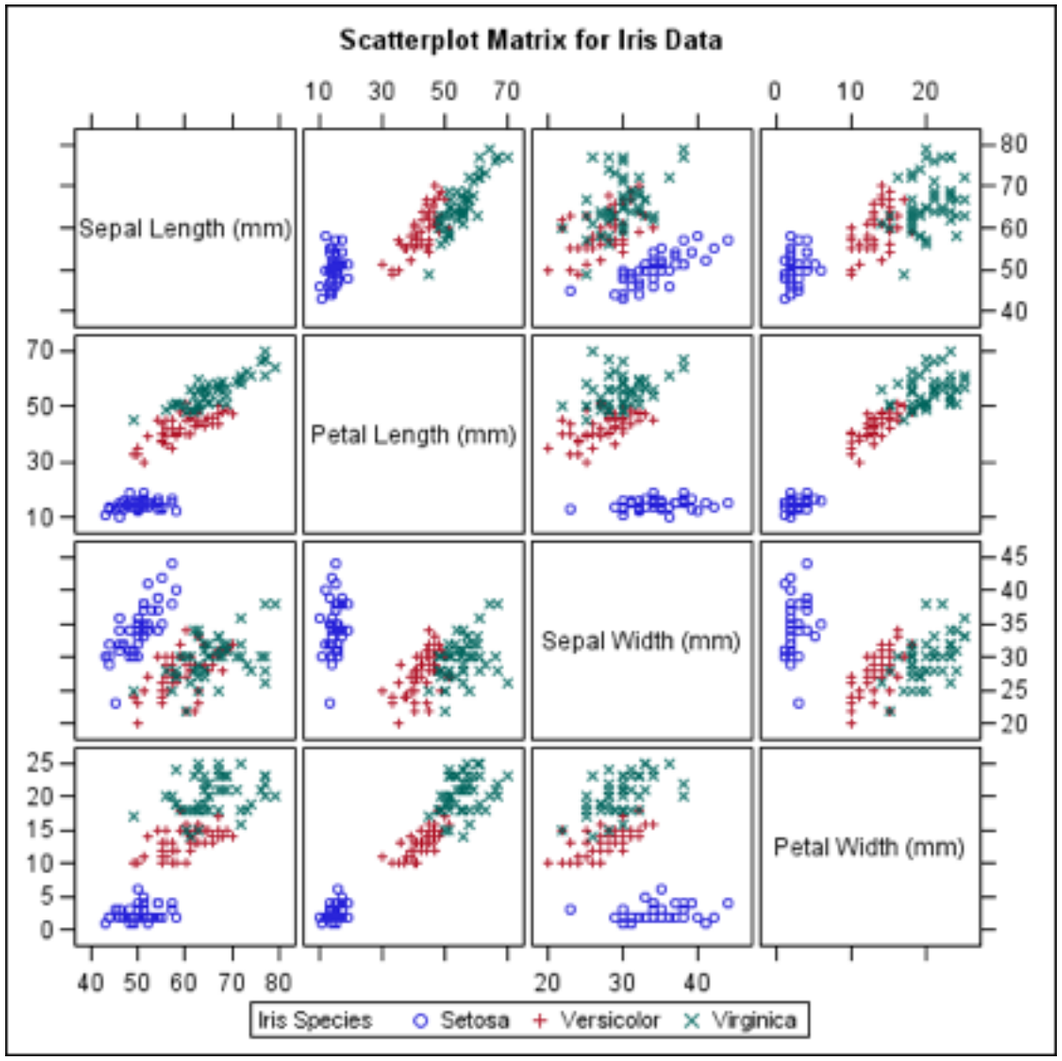

scatterplot matrix

(source: support.sas.com)

Iris setosa

Iris versicolor

Iris virginica

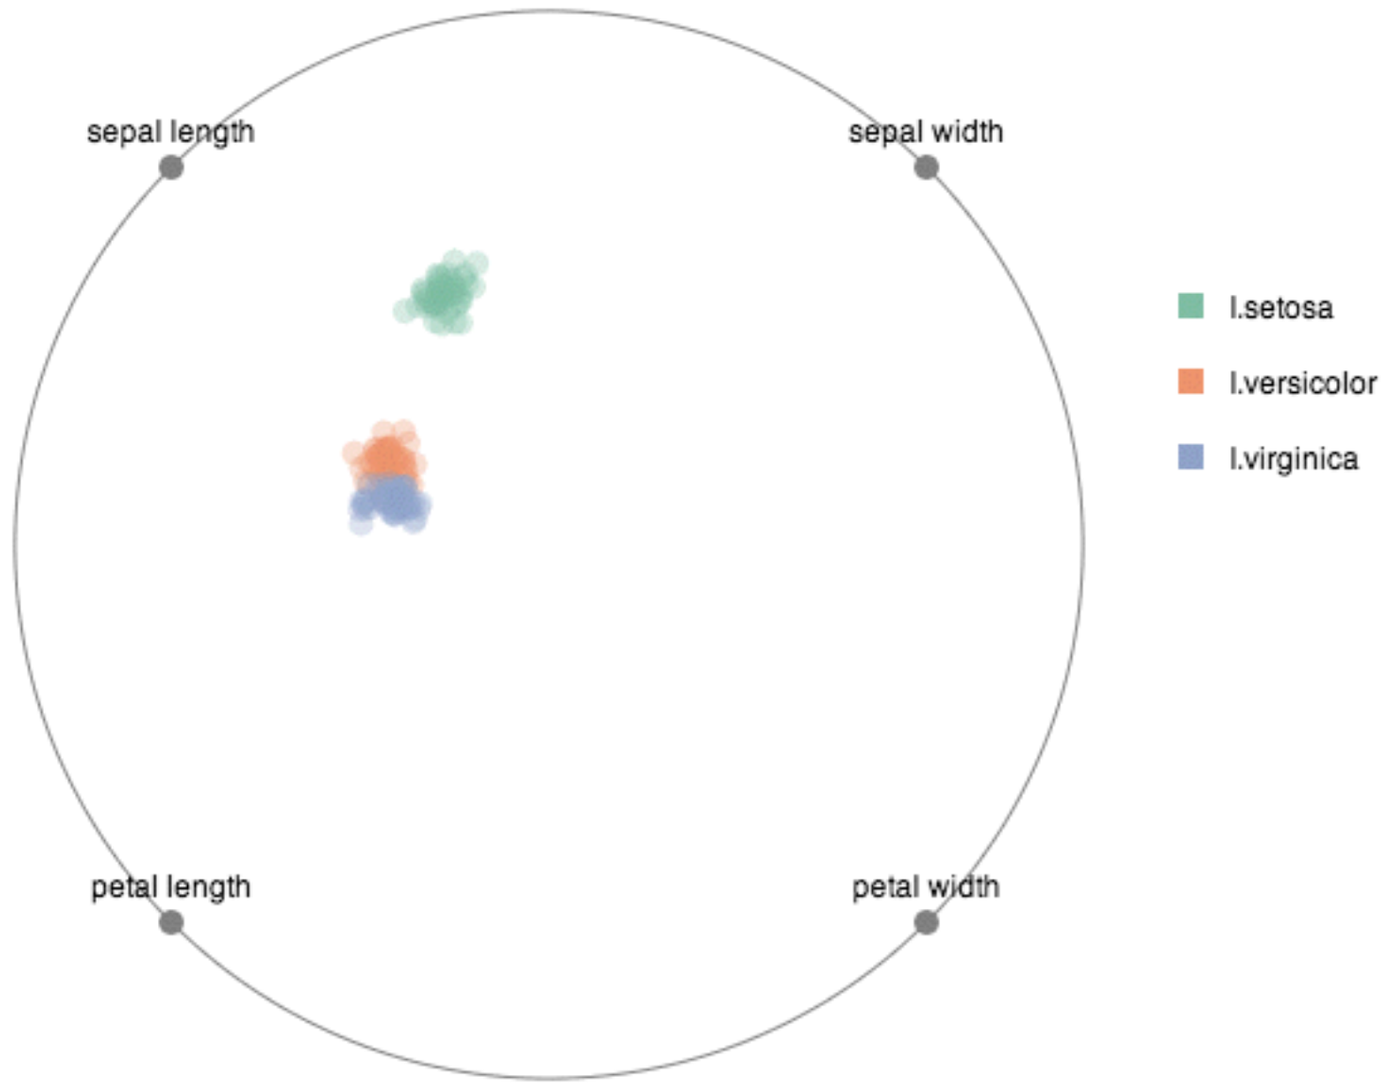

radviz

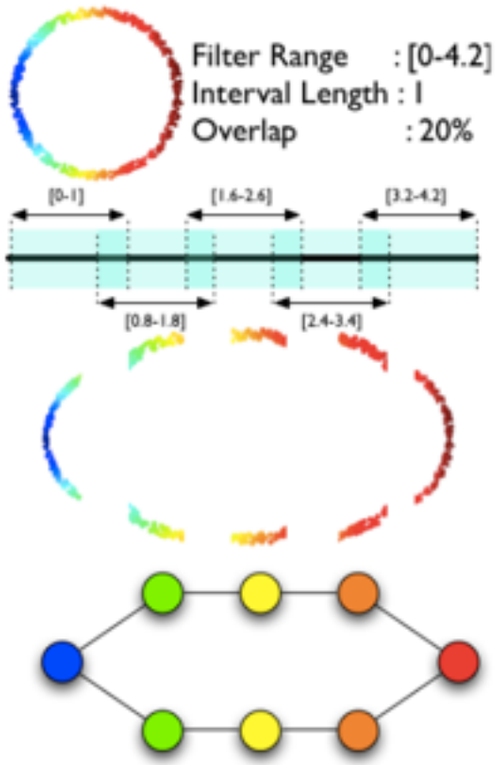

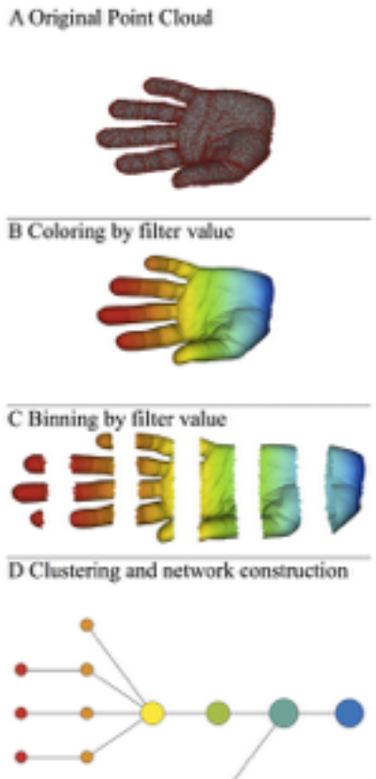

Could they have found this with just clustering or PCA? No.

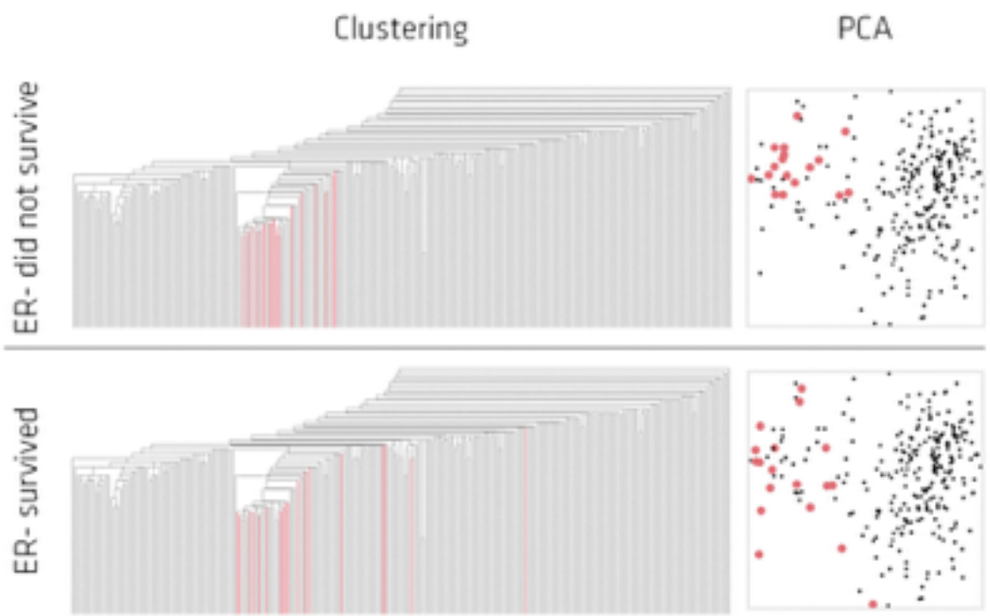

topological data analysis

local clustering in
global context

colour = ESR1 expression
(red = high; blue = low)

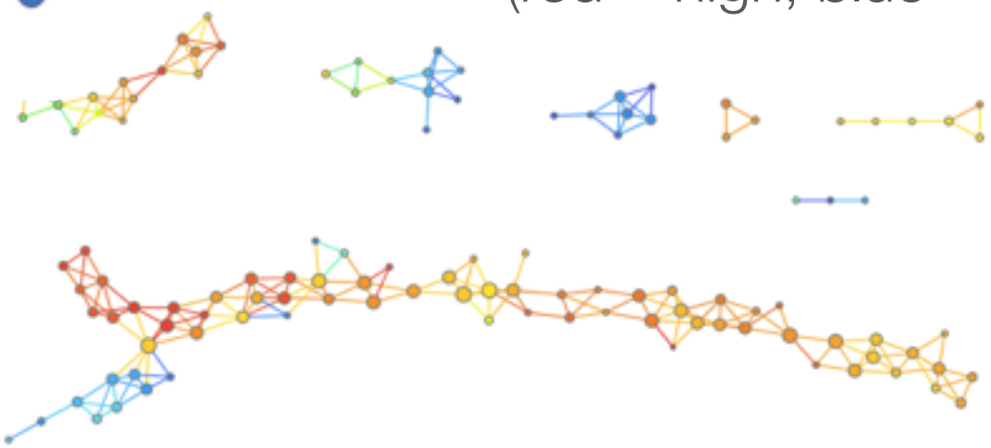

dead

survived

lens = L-infinity centrality
+ event death

Ayasdi, Inc

- parallel coordinates: point in Euclidean space = line in parallel coordinate space and vice versa

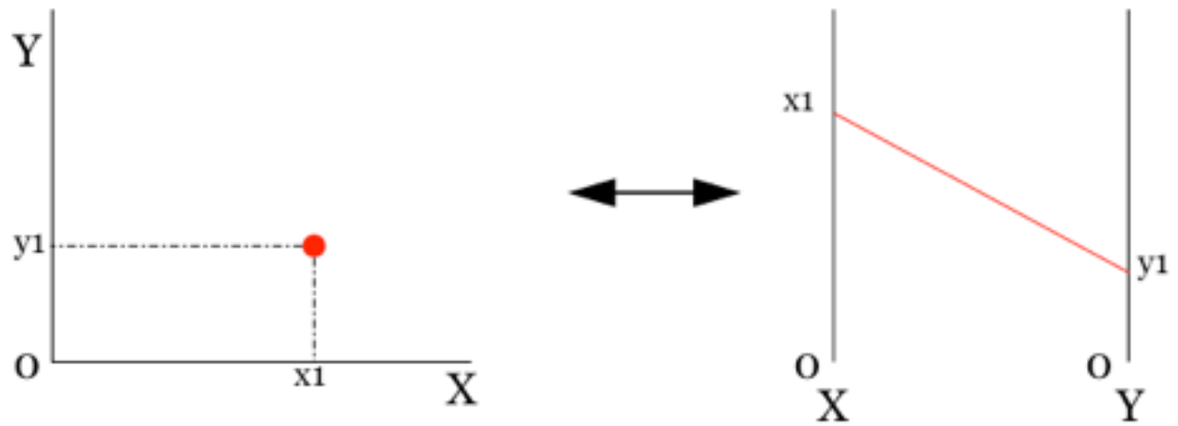

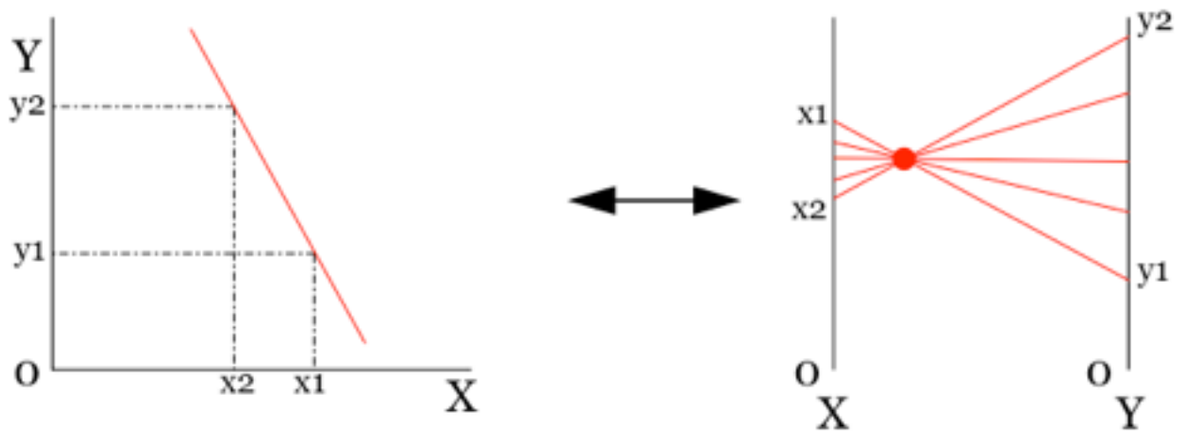

Tour de France, 2011

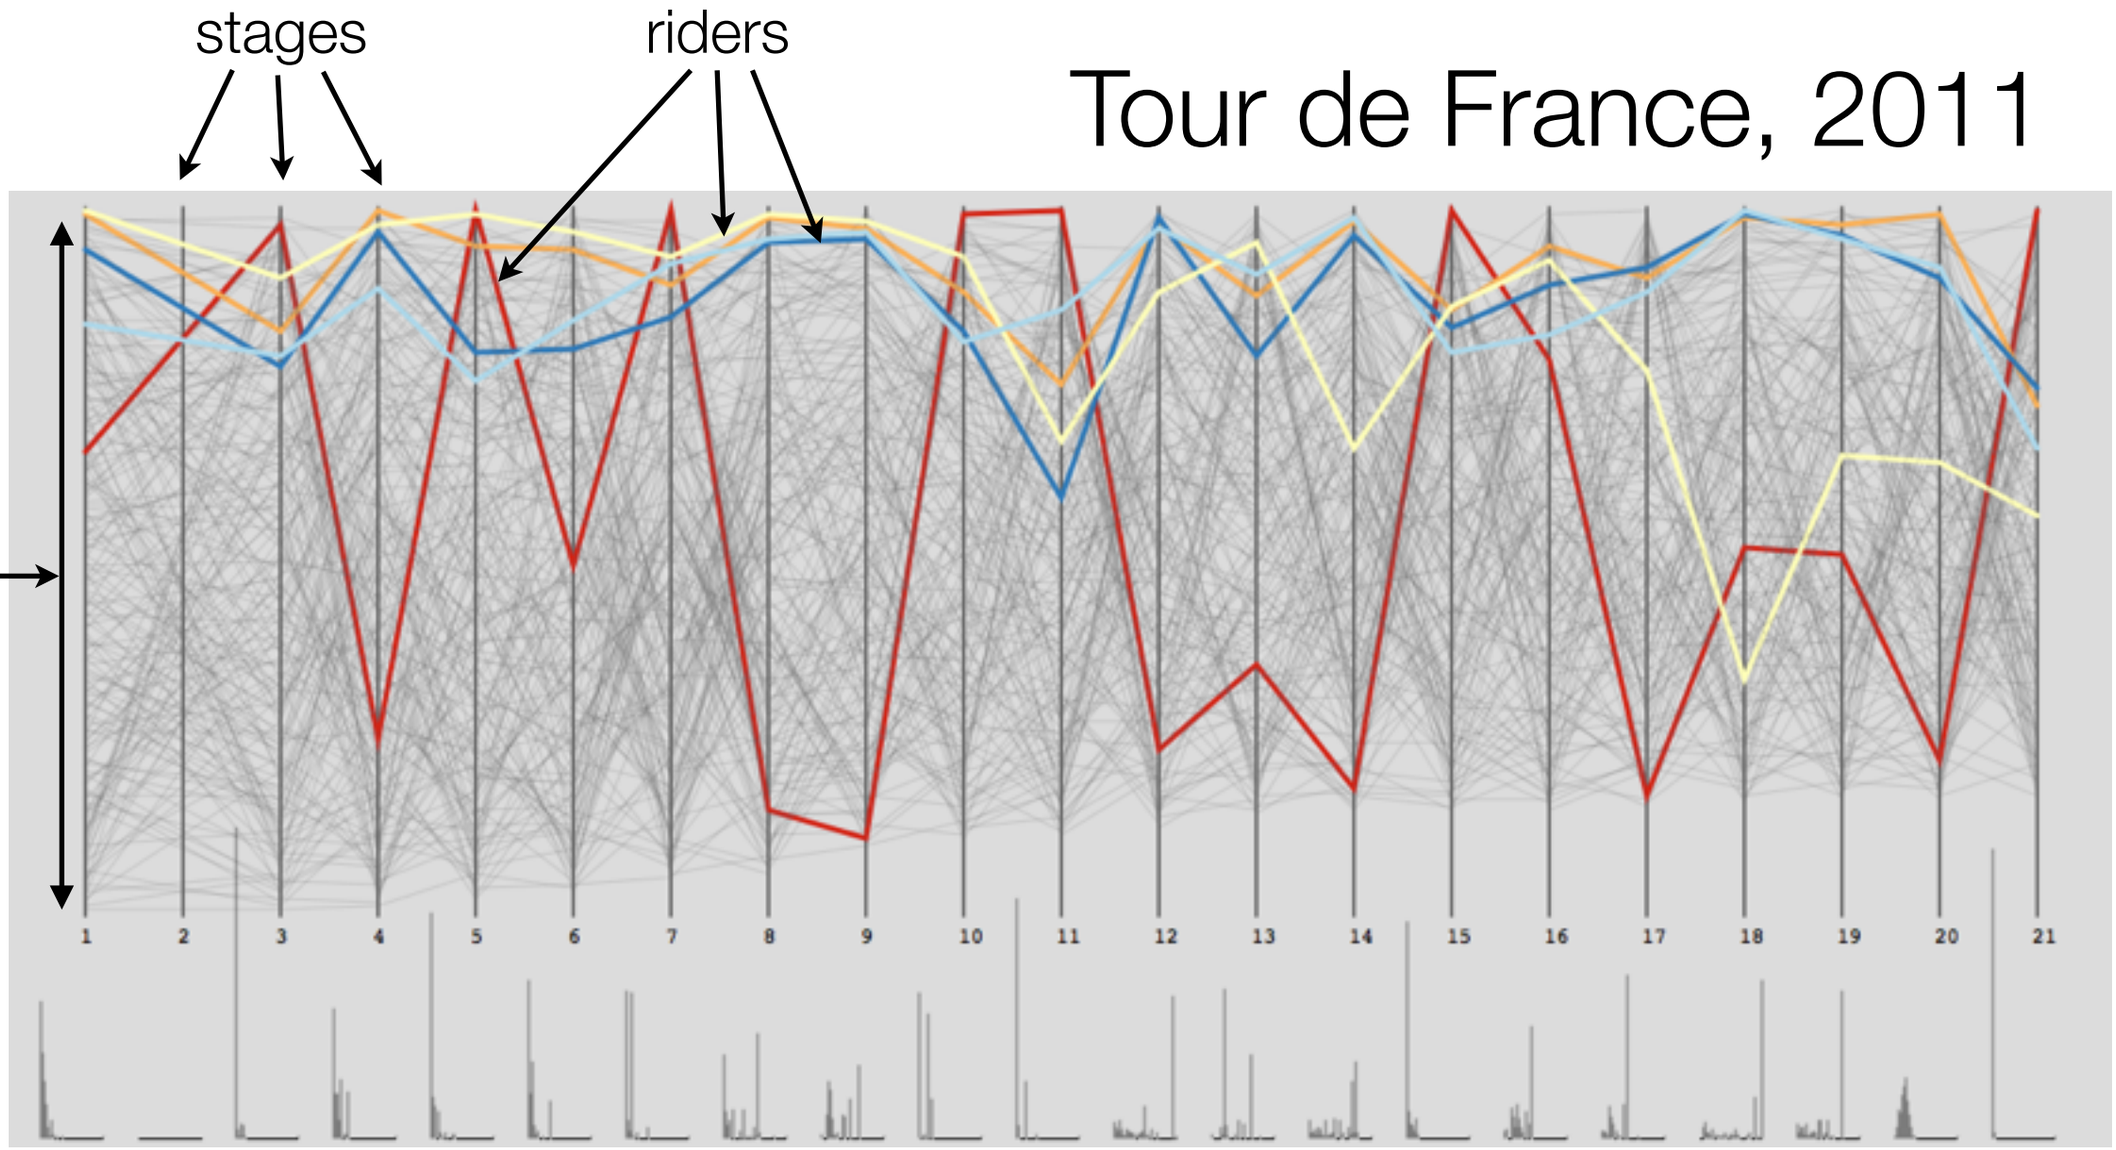

Cavendish: red
Evans: orange
Gilbert: yellow
Andy Schleck: light blue
Frank Schleck: dark blue

time

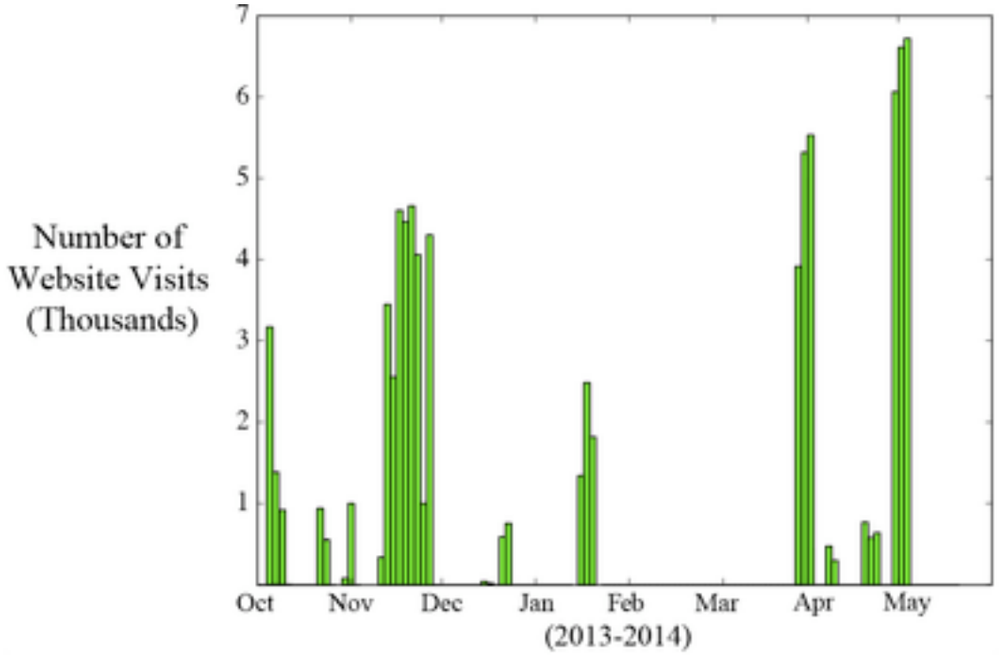

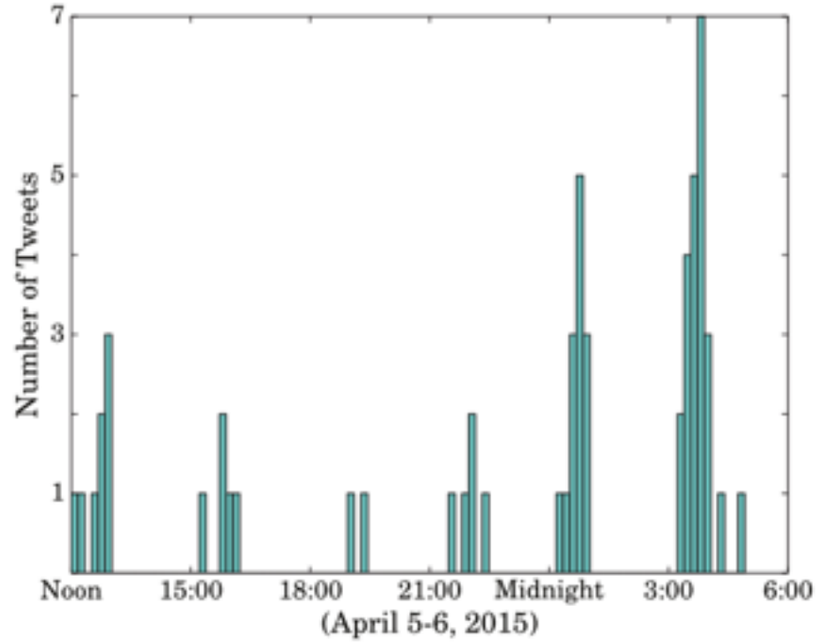

Human

Bot

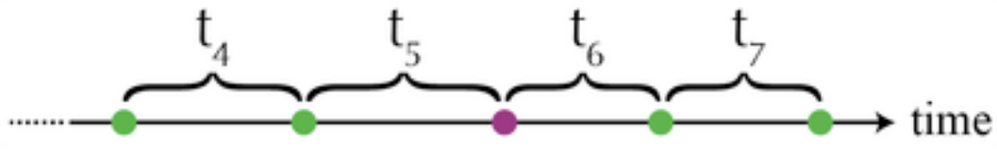

“Time map” (Watson MC, 2015)

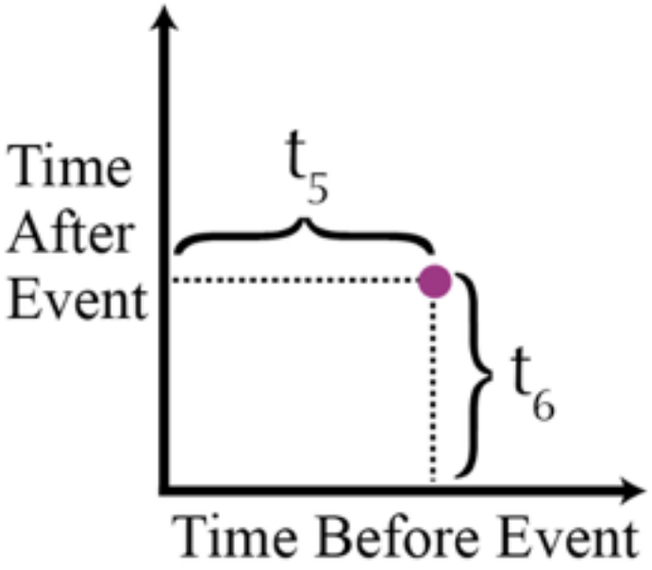

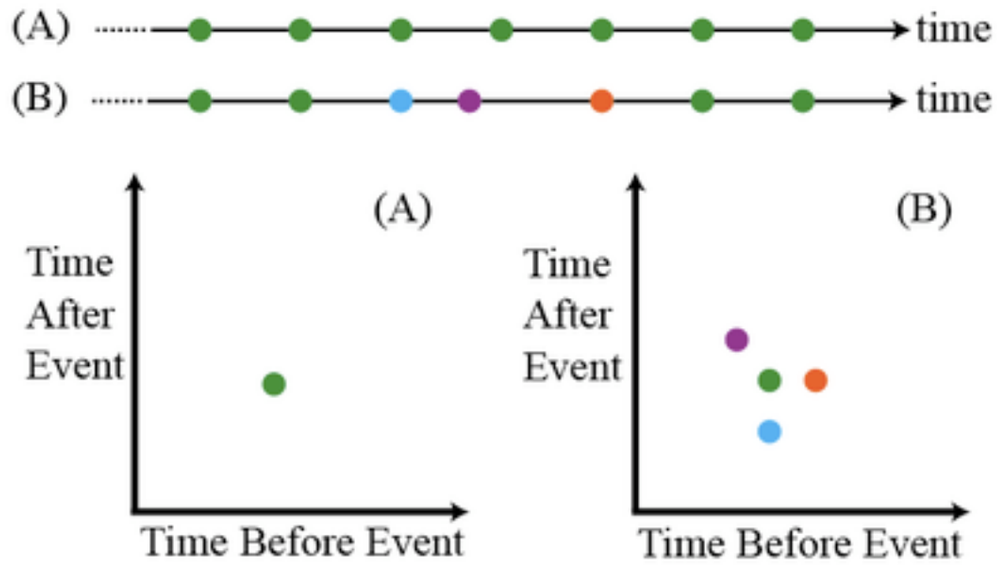

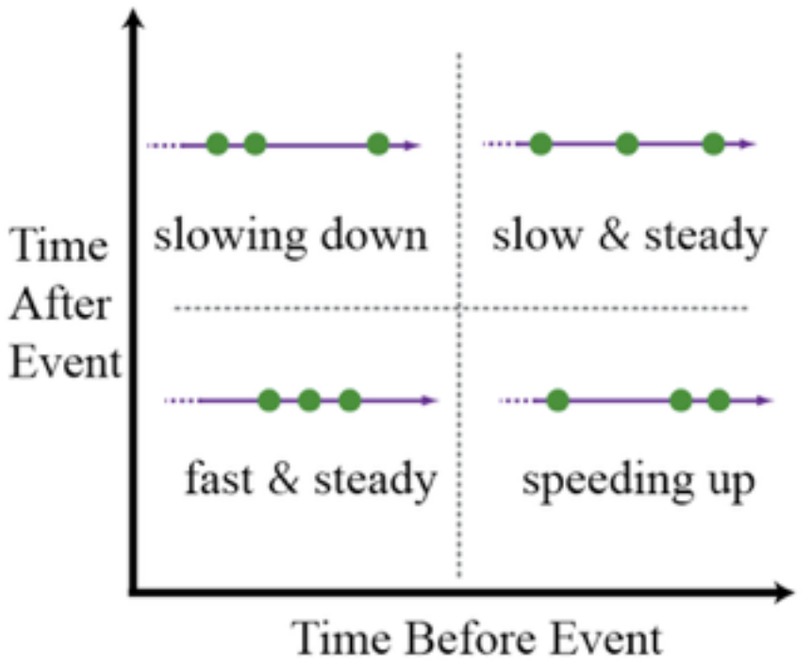

https://districtdatalabs.silvrback.com/time-maps-visualizing-discrete-events-across-many-timescales

time

Tweets

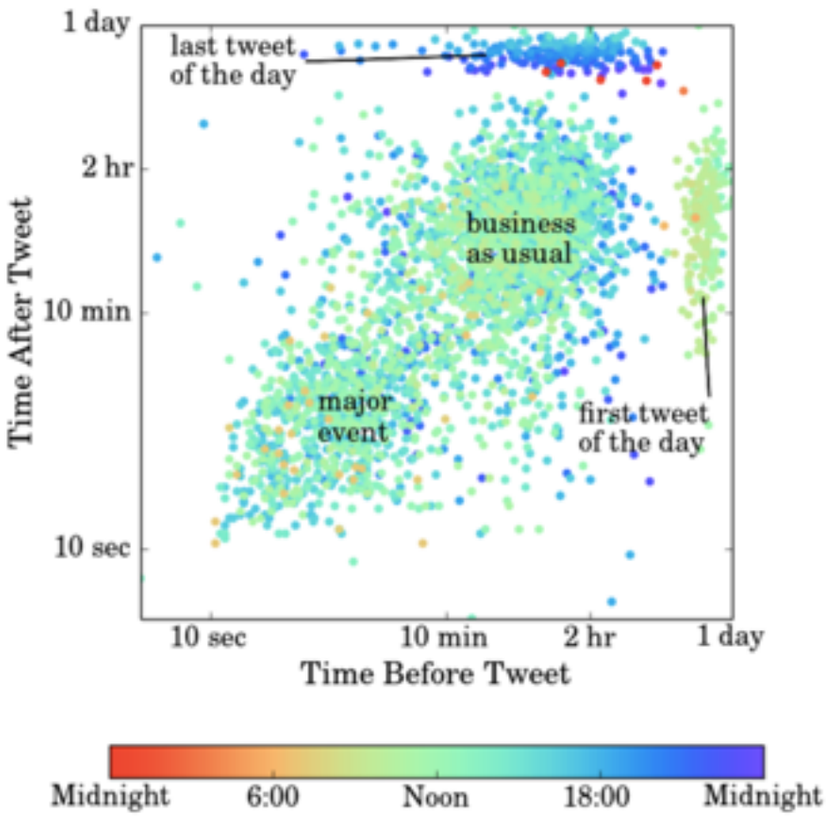

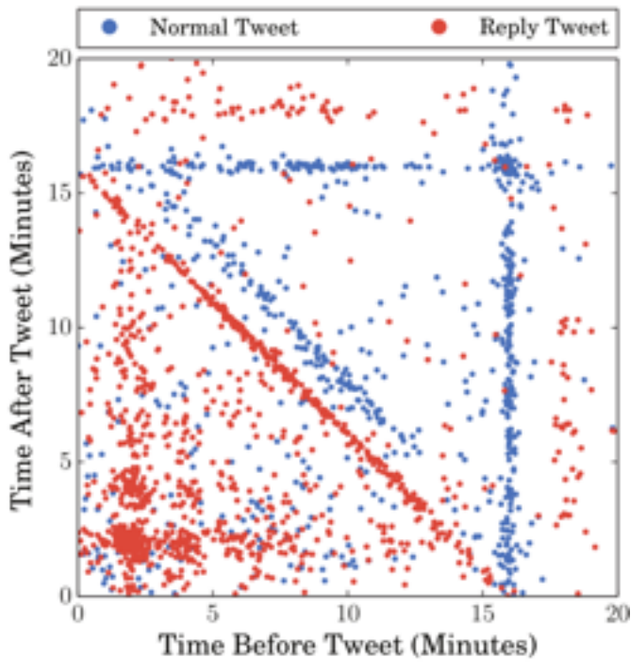

White House

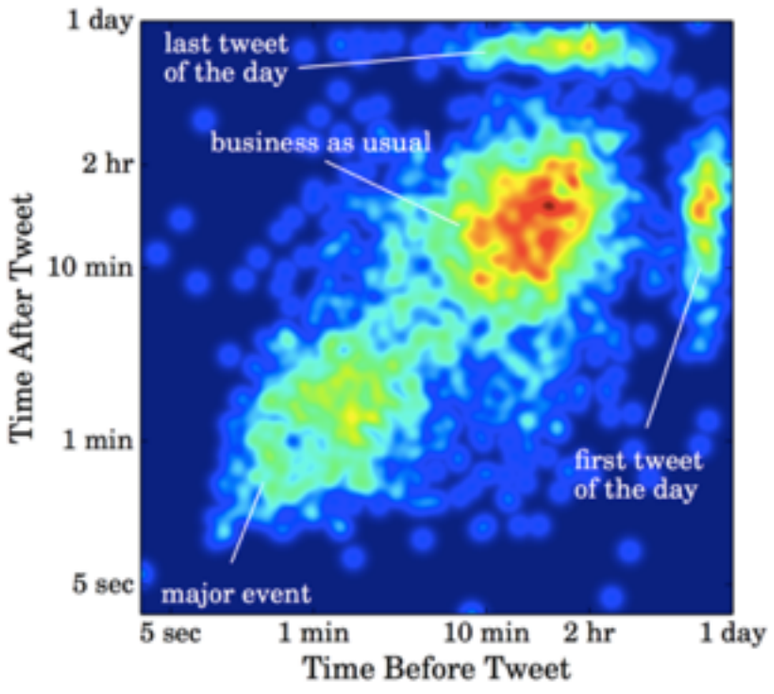

Personal

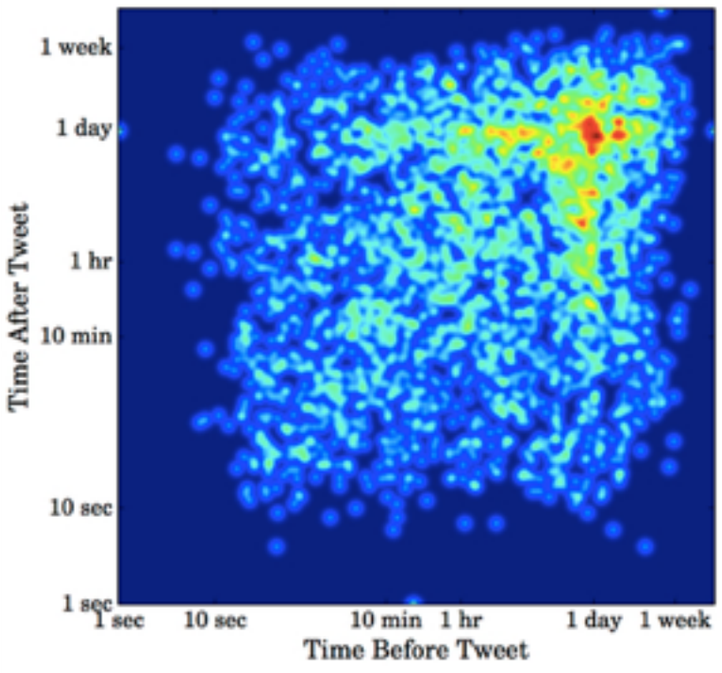

Bot

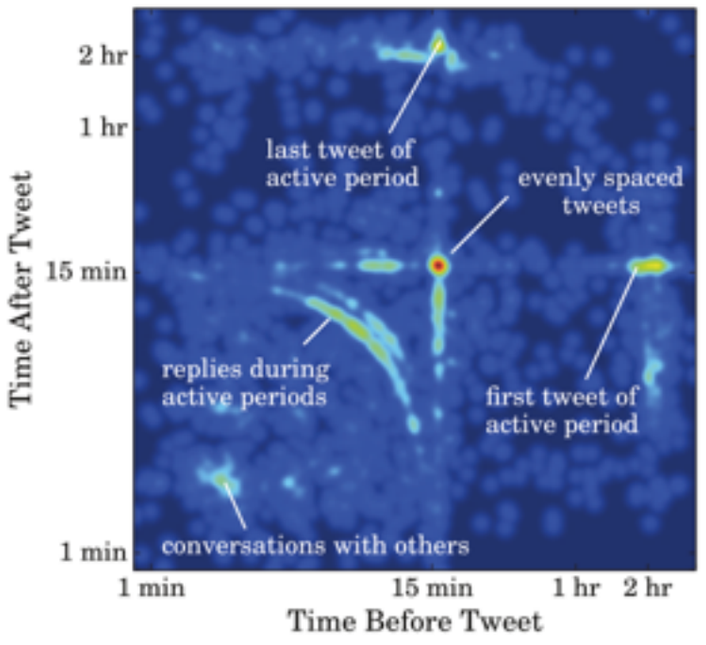

phylogenetic data

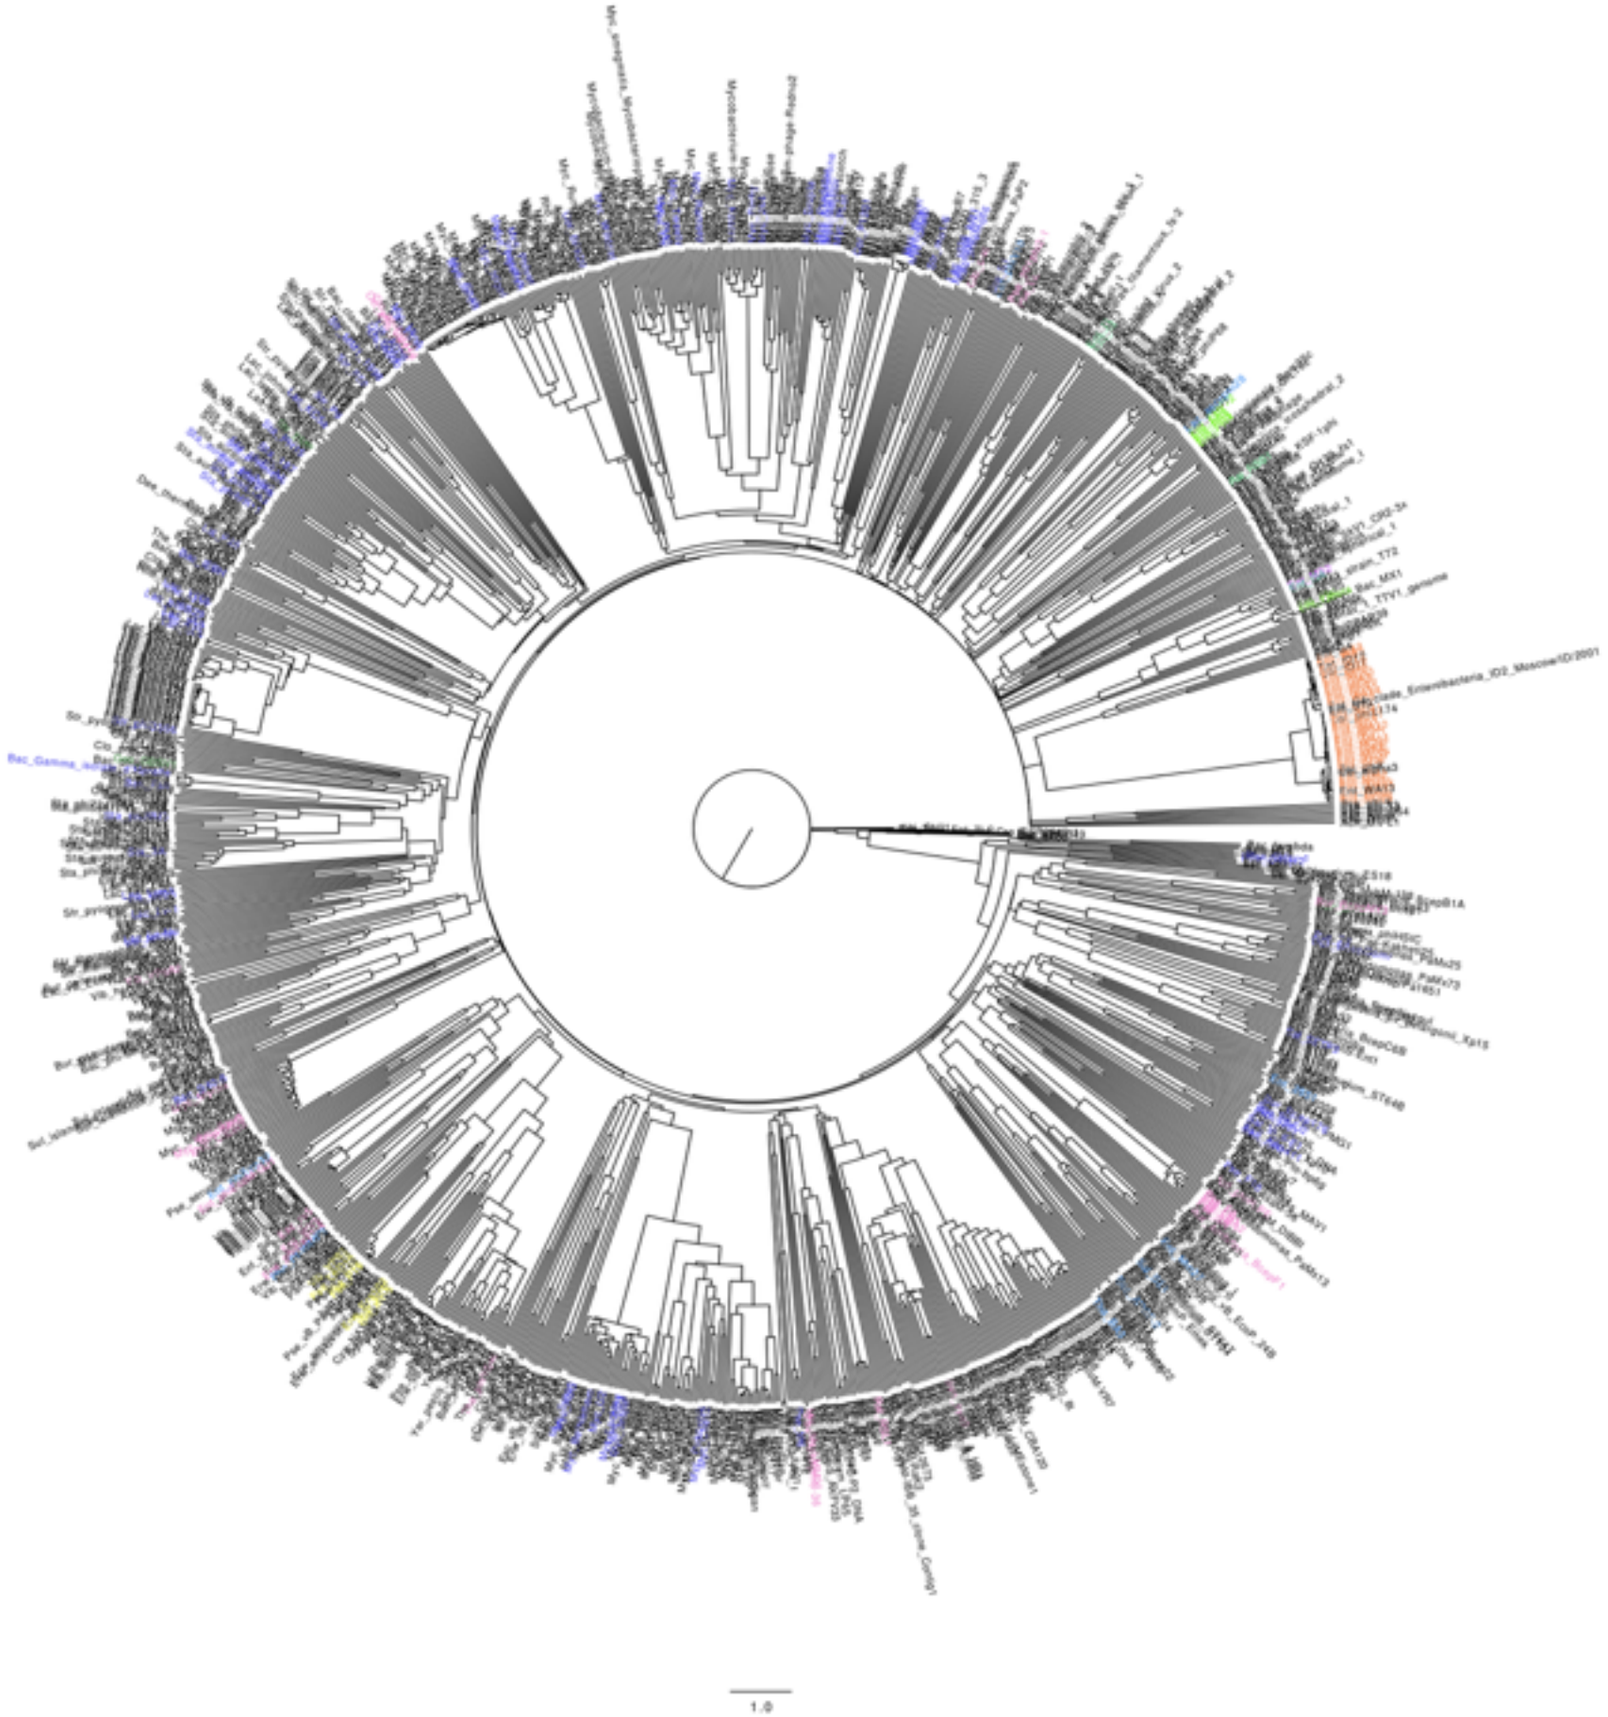

phylogenetic data

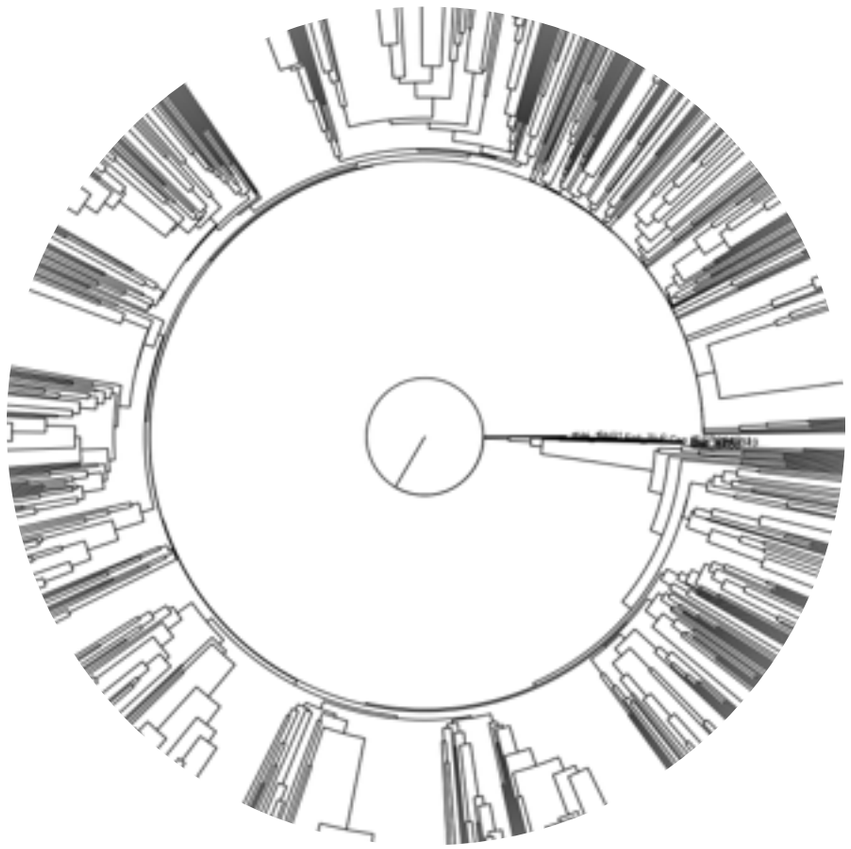

1.0

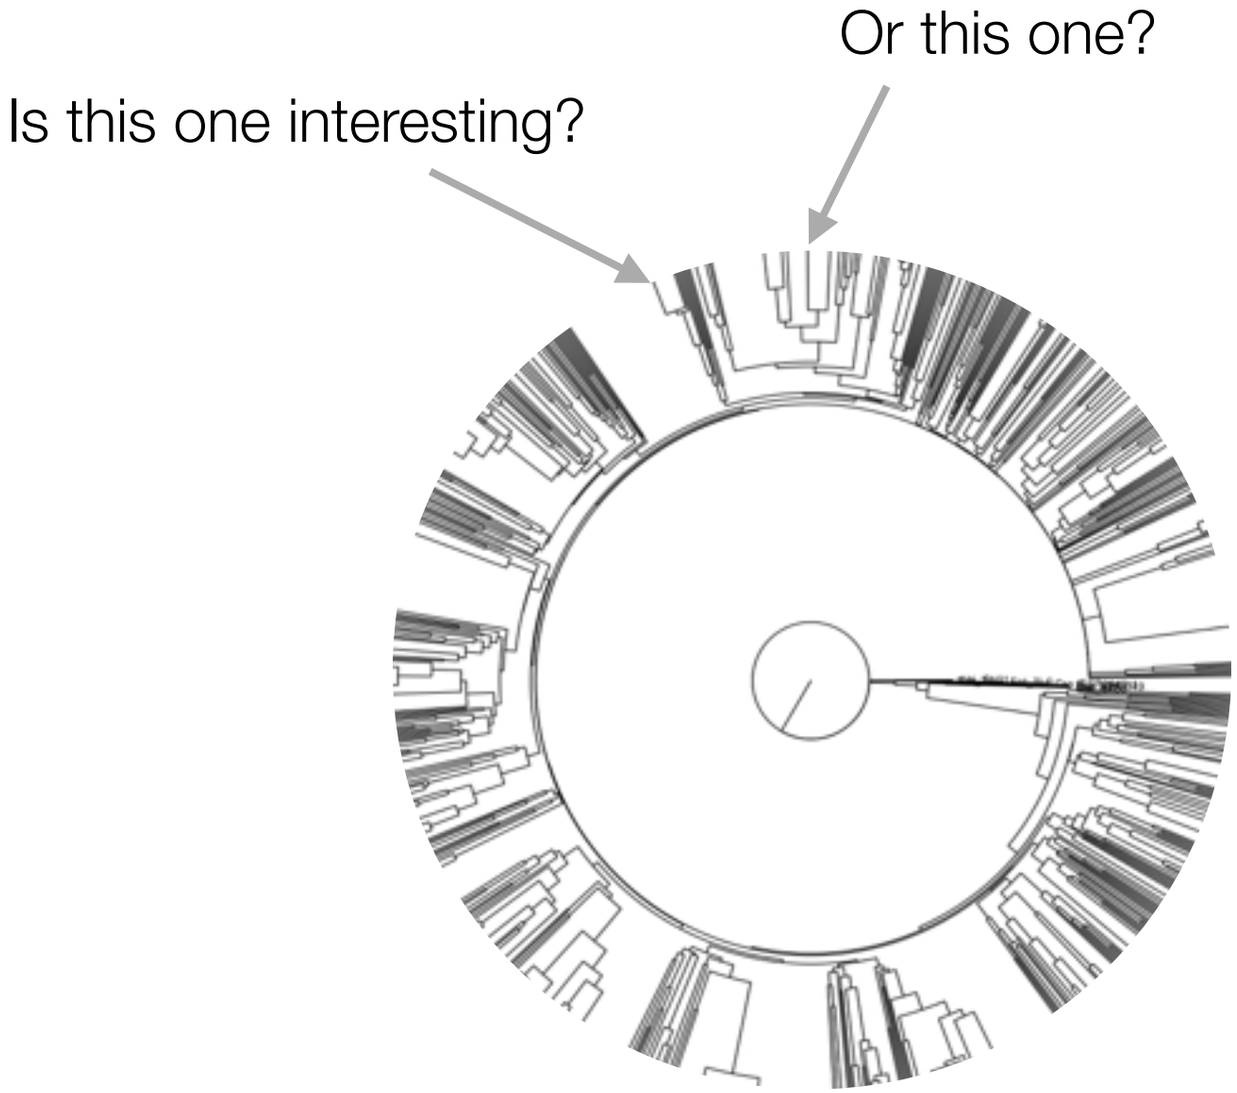
Is this one interesting?
Or this one?
1.0

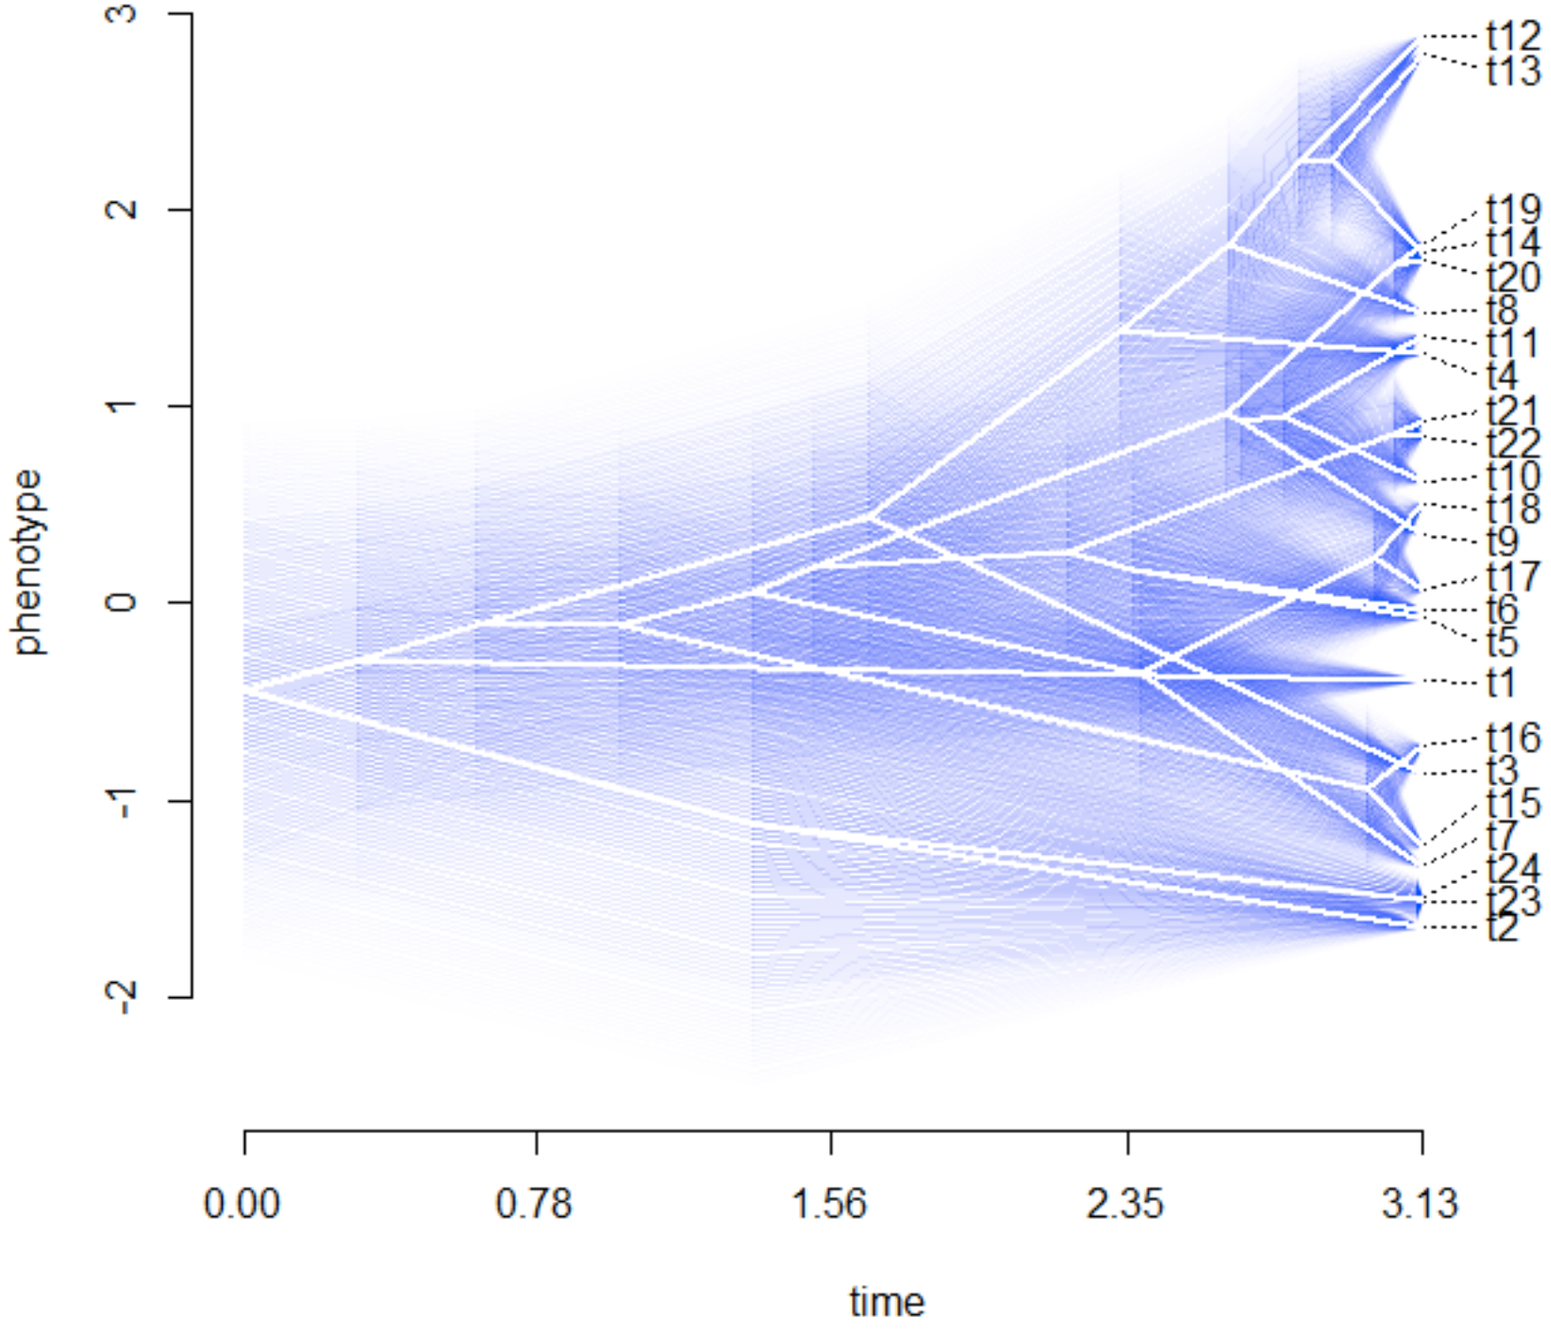
phenotype
3
2
1
0
-1
-2
t12
t13
t19
t14
t20
t8
t11
t4
t21
t22
t10
t18
t9
t17
t6
t5
t1
t16
t3
t15
t7
t24
t23
t2
0.00
0.78
1.56
2.35
3.13
time

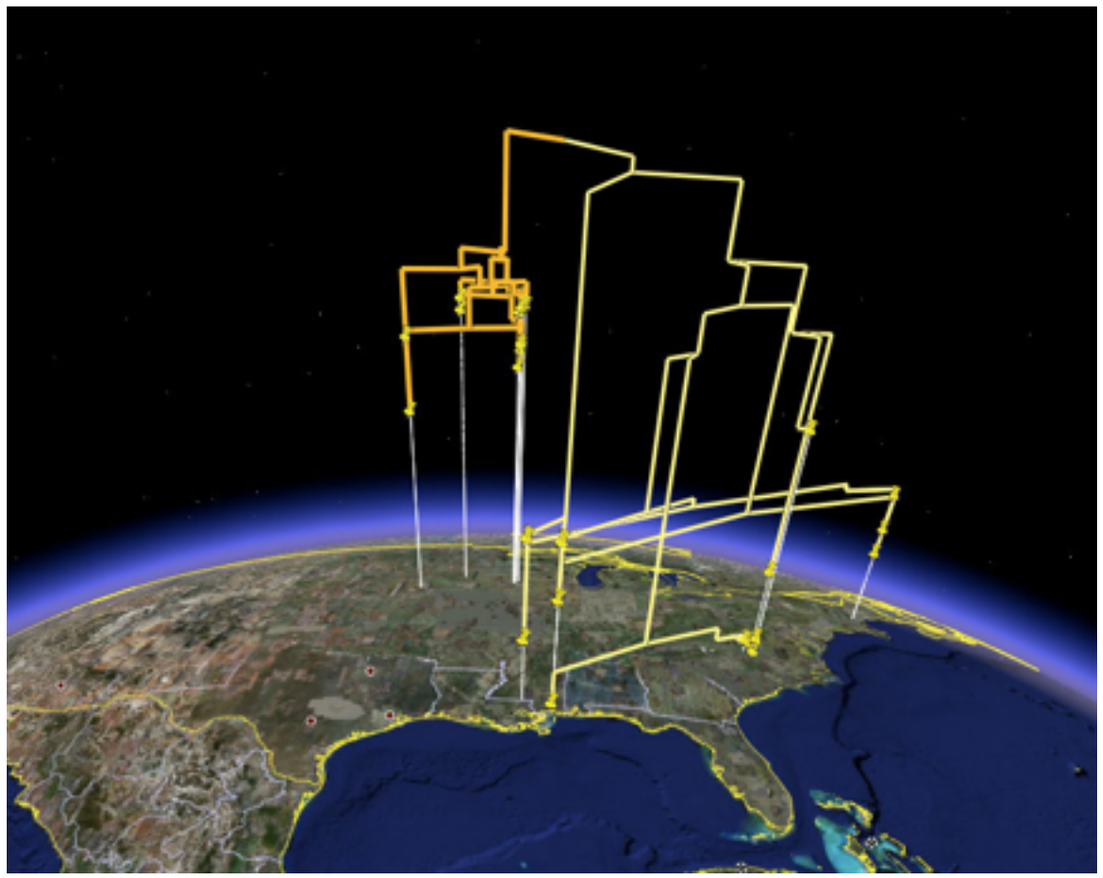

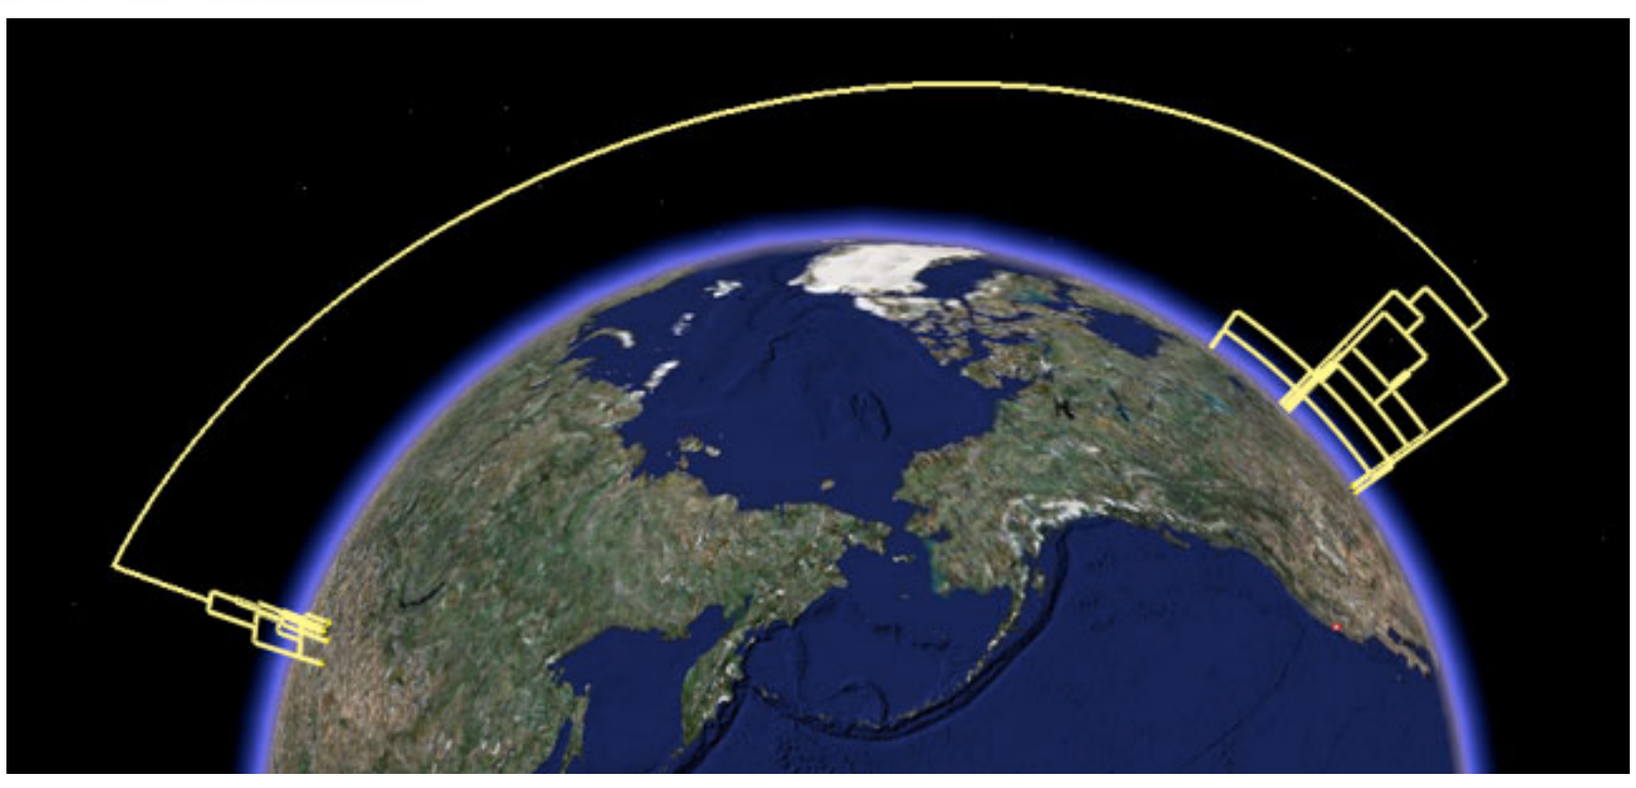

dendrogramix (Blanch et al, 2015)

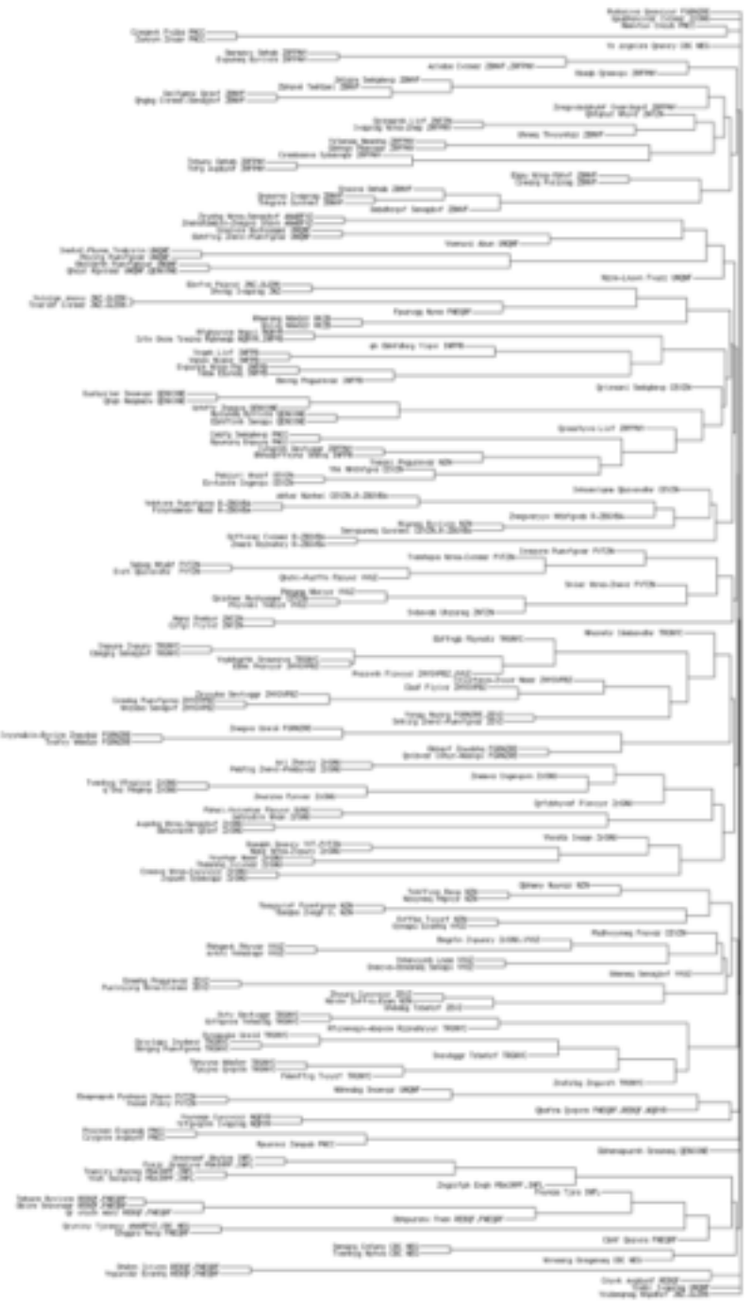

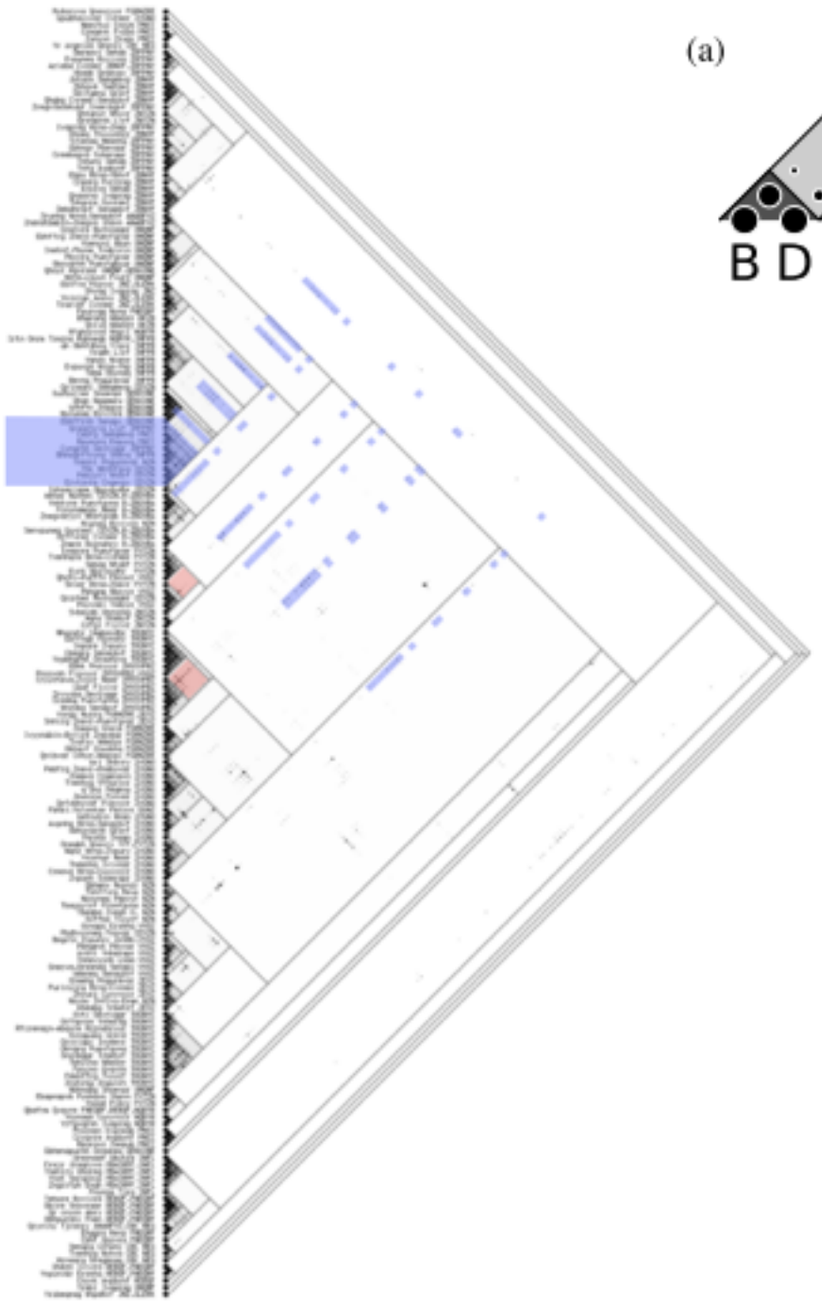

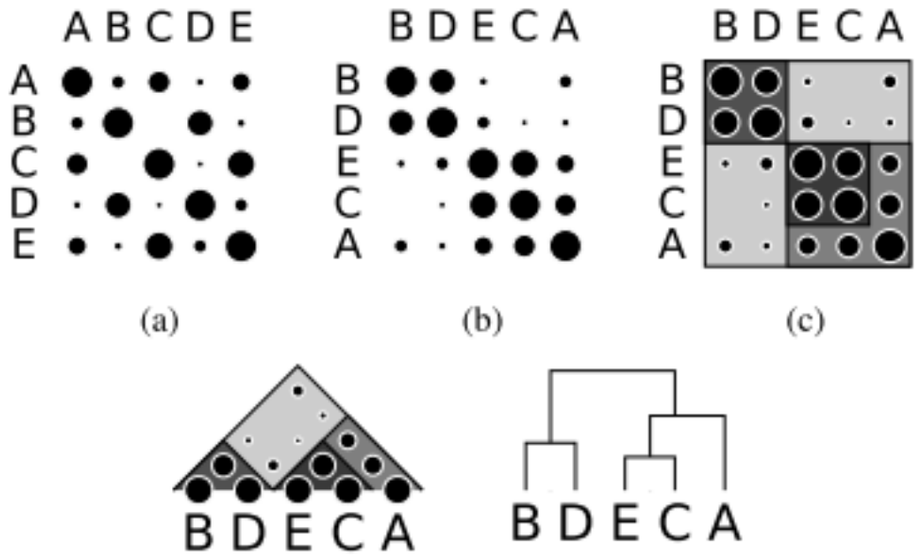

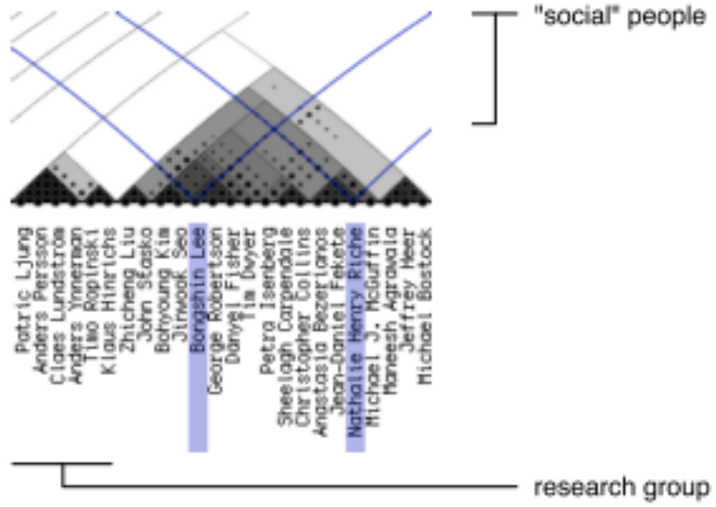

dendsort (R package)

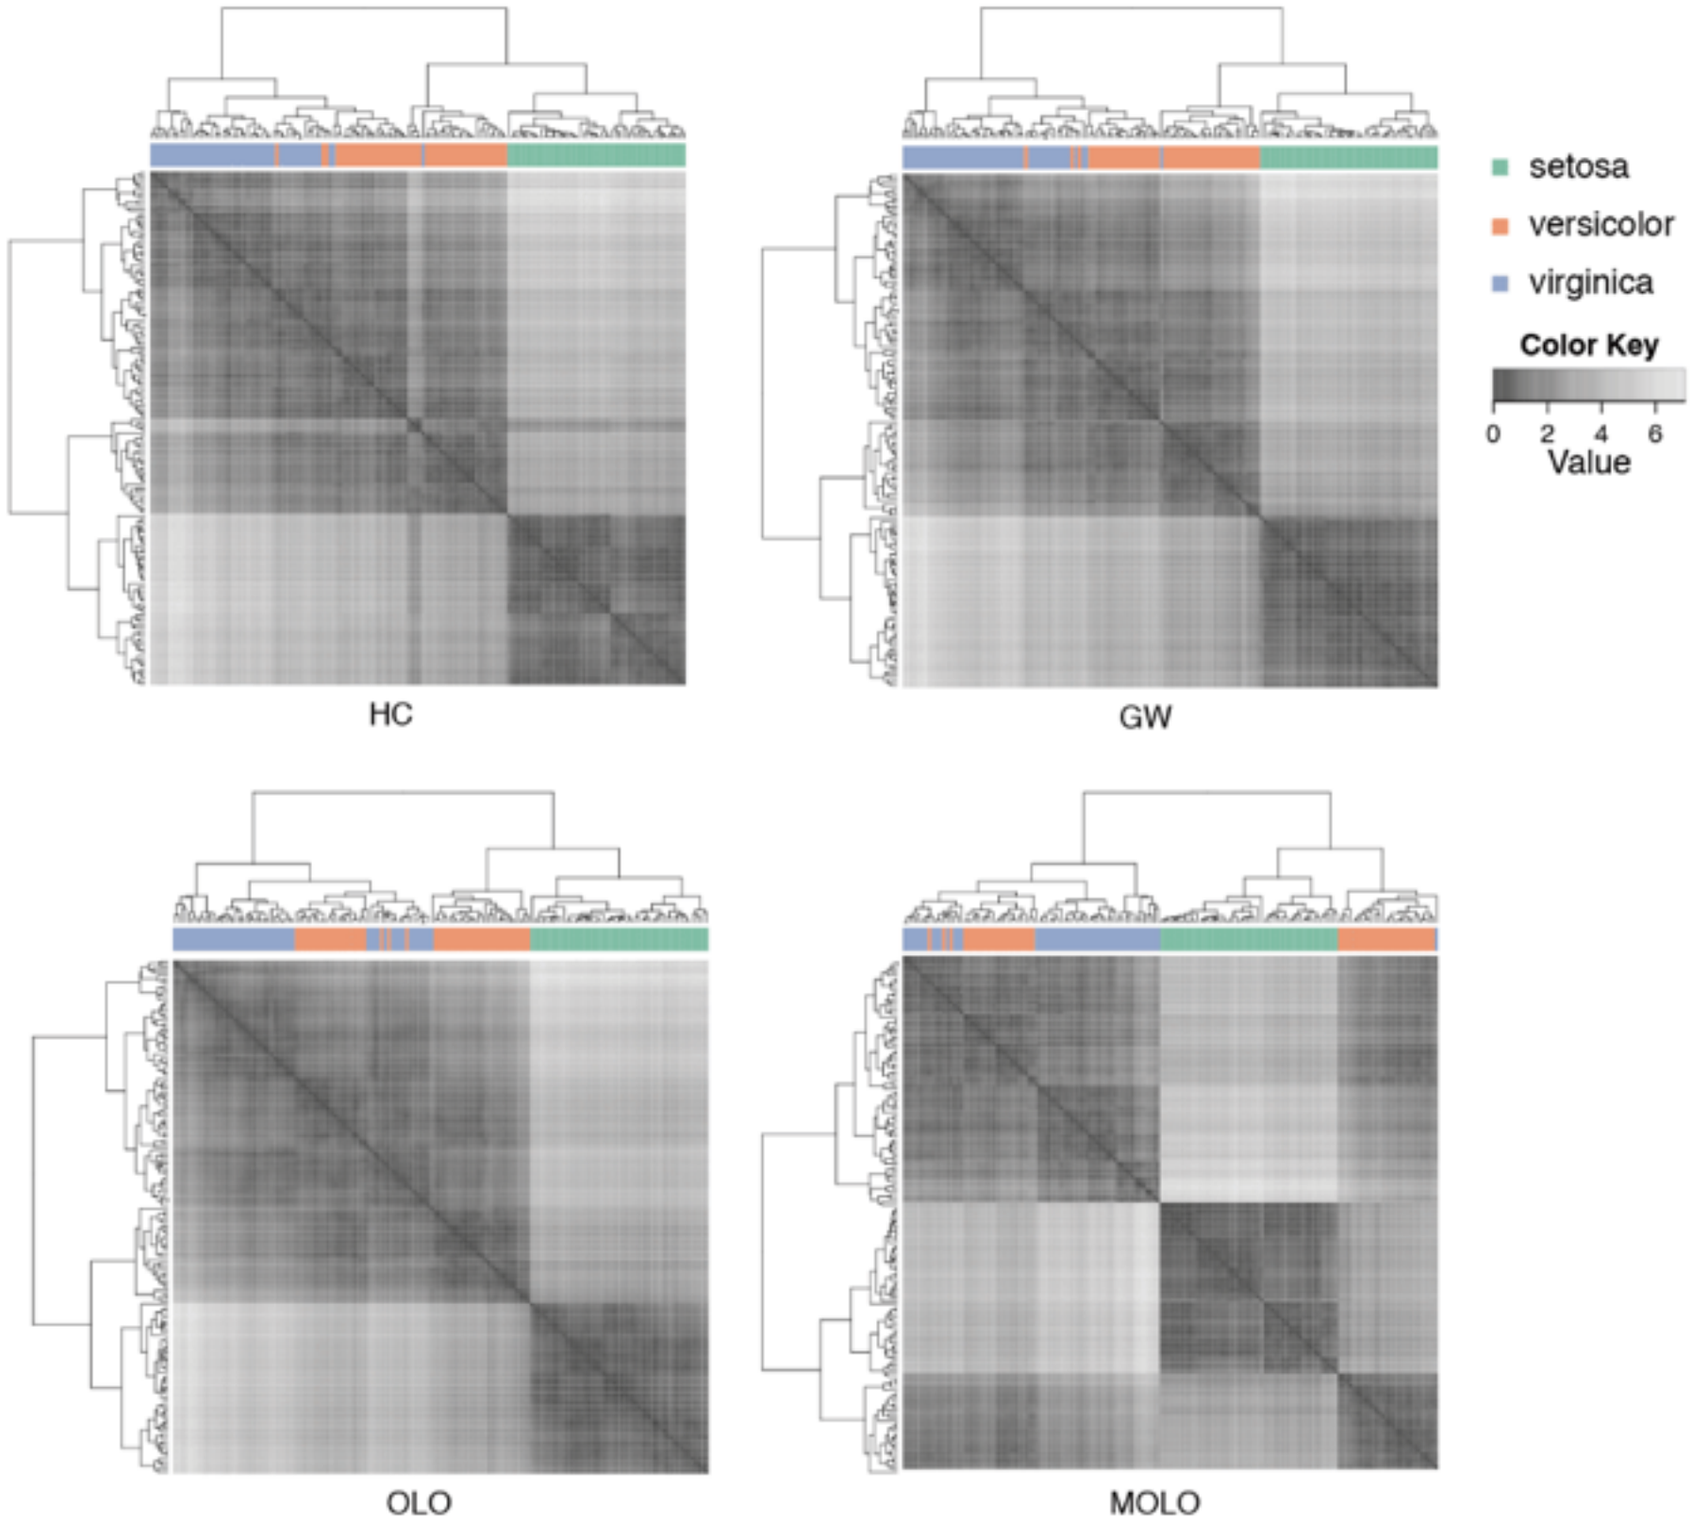

Sakai et al, 2014

networks: node-link diagrams

max 20 nodes…

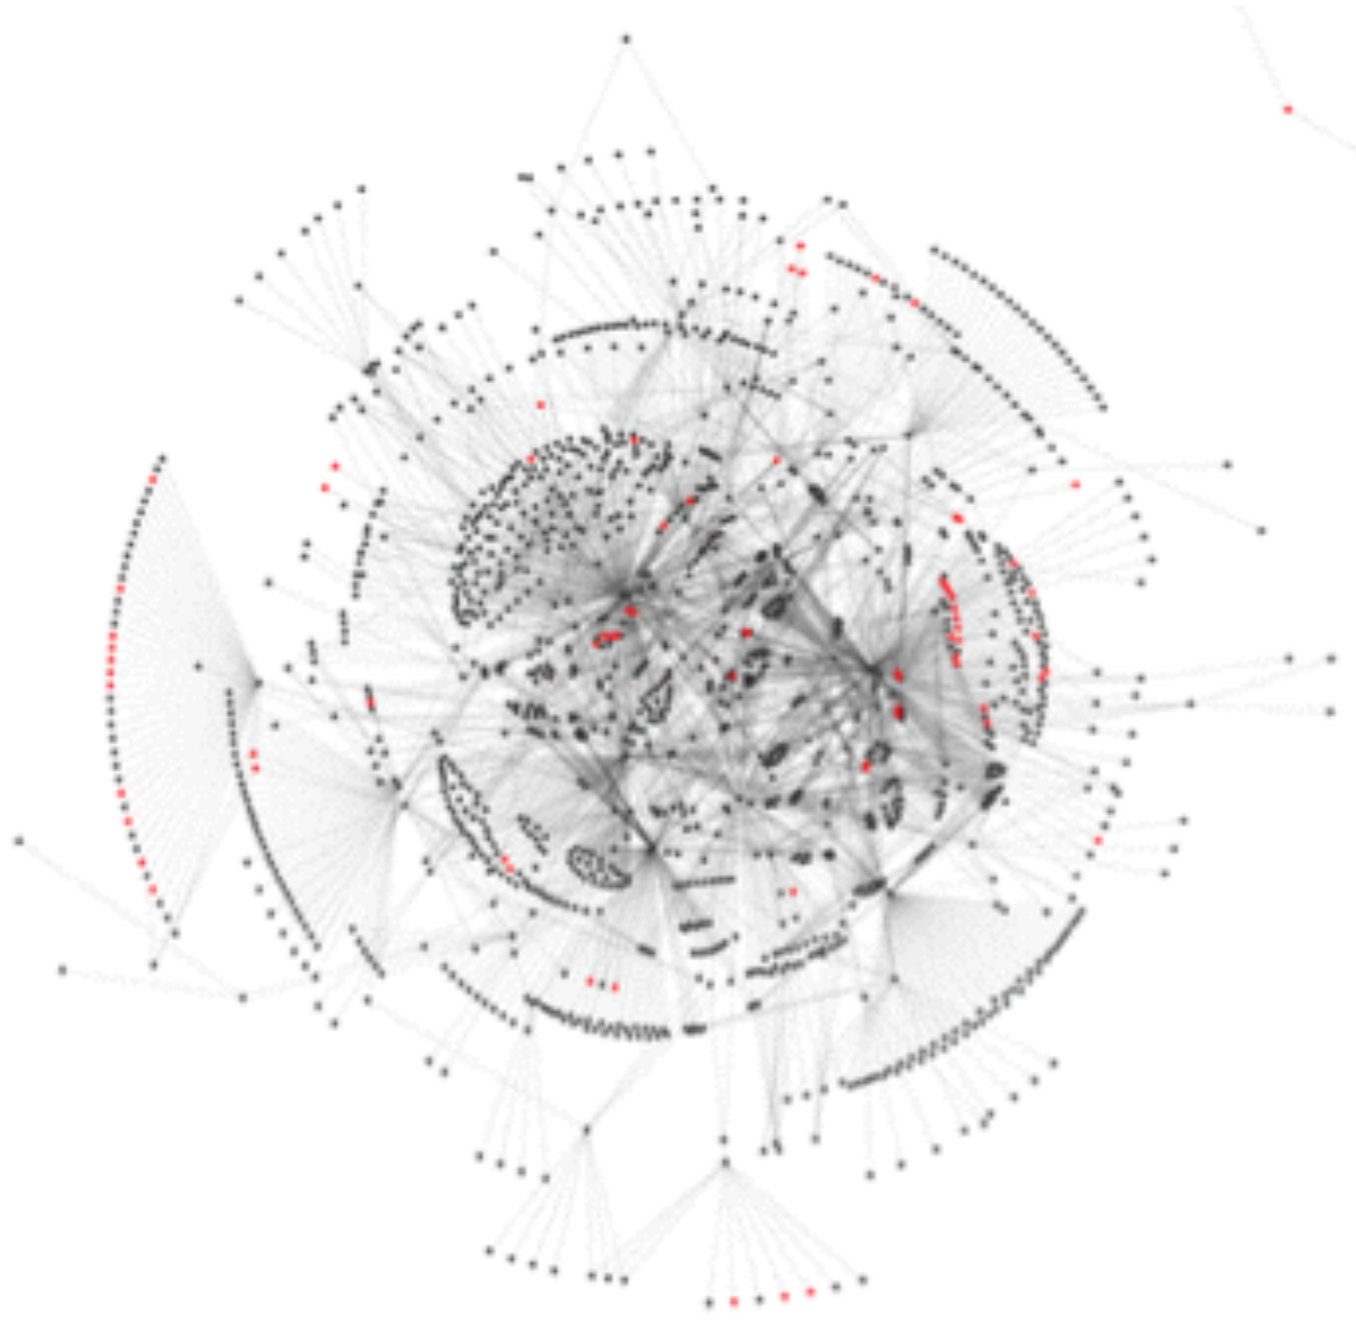

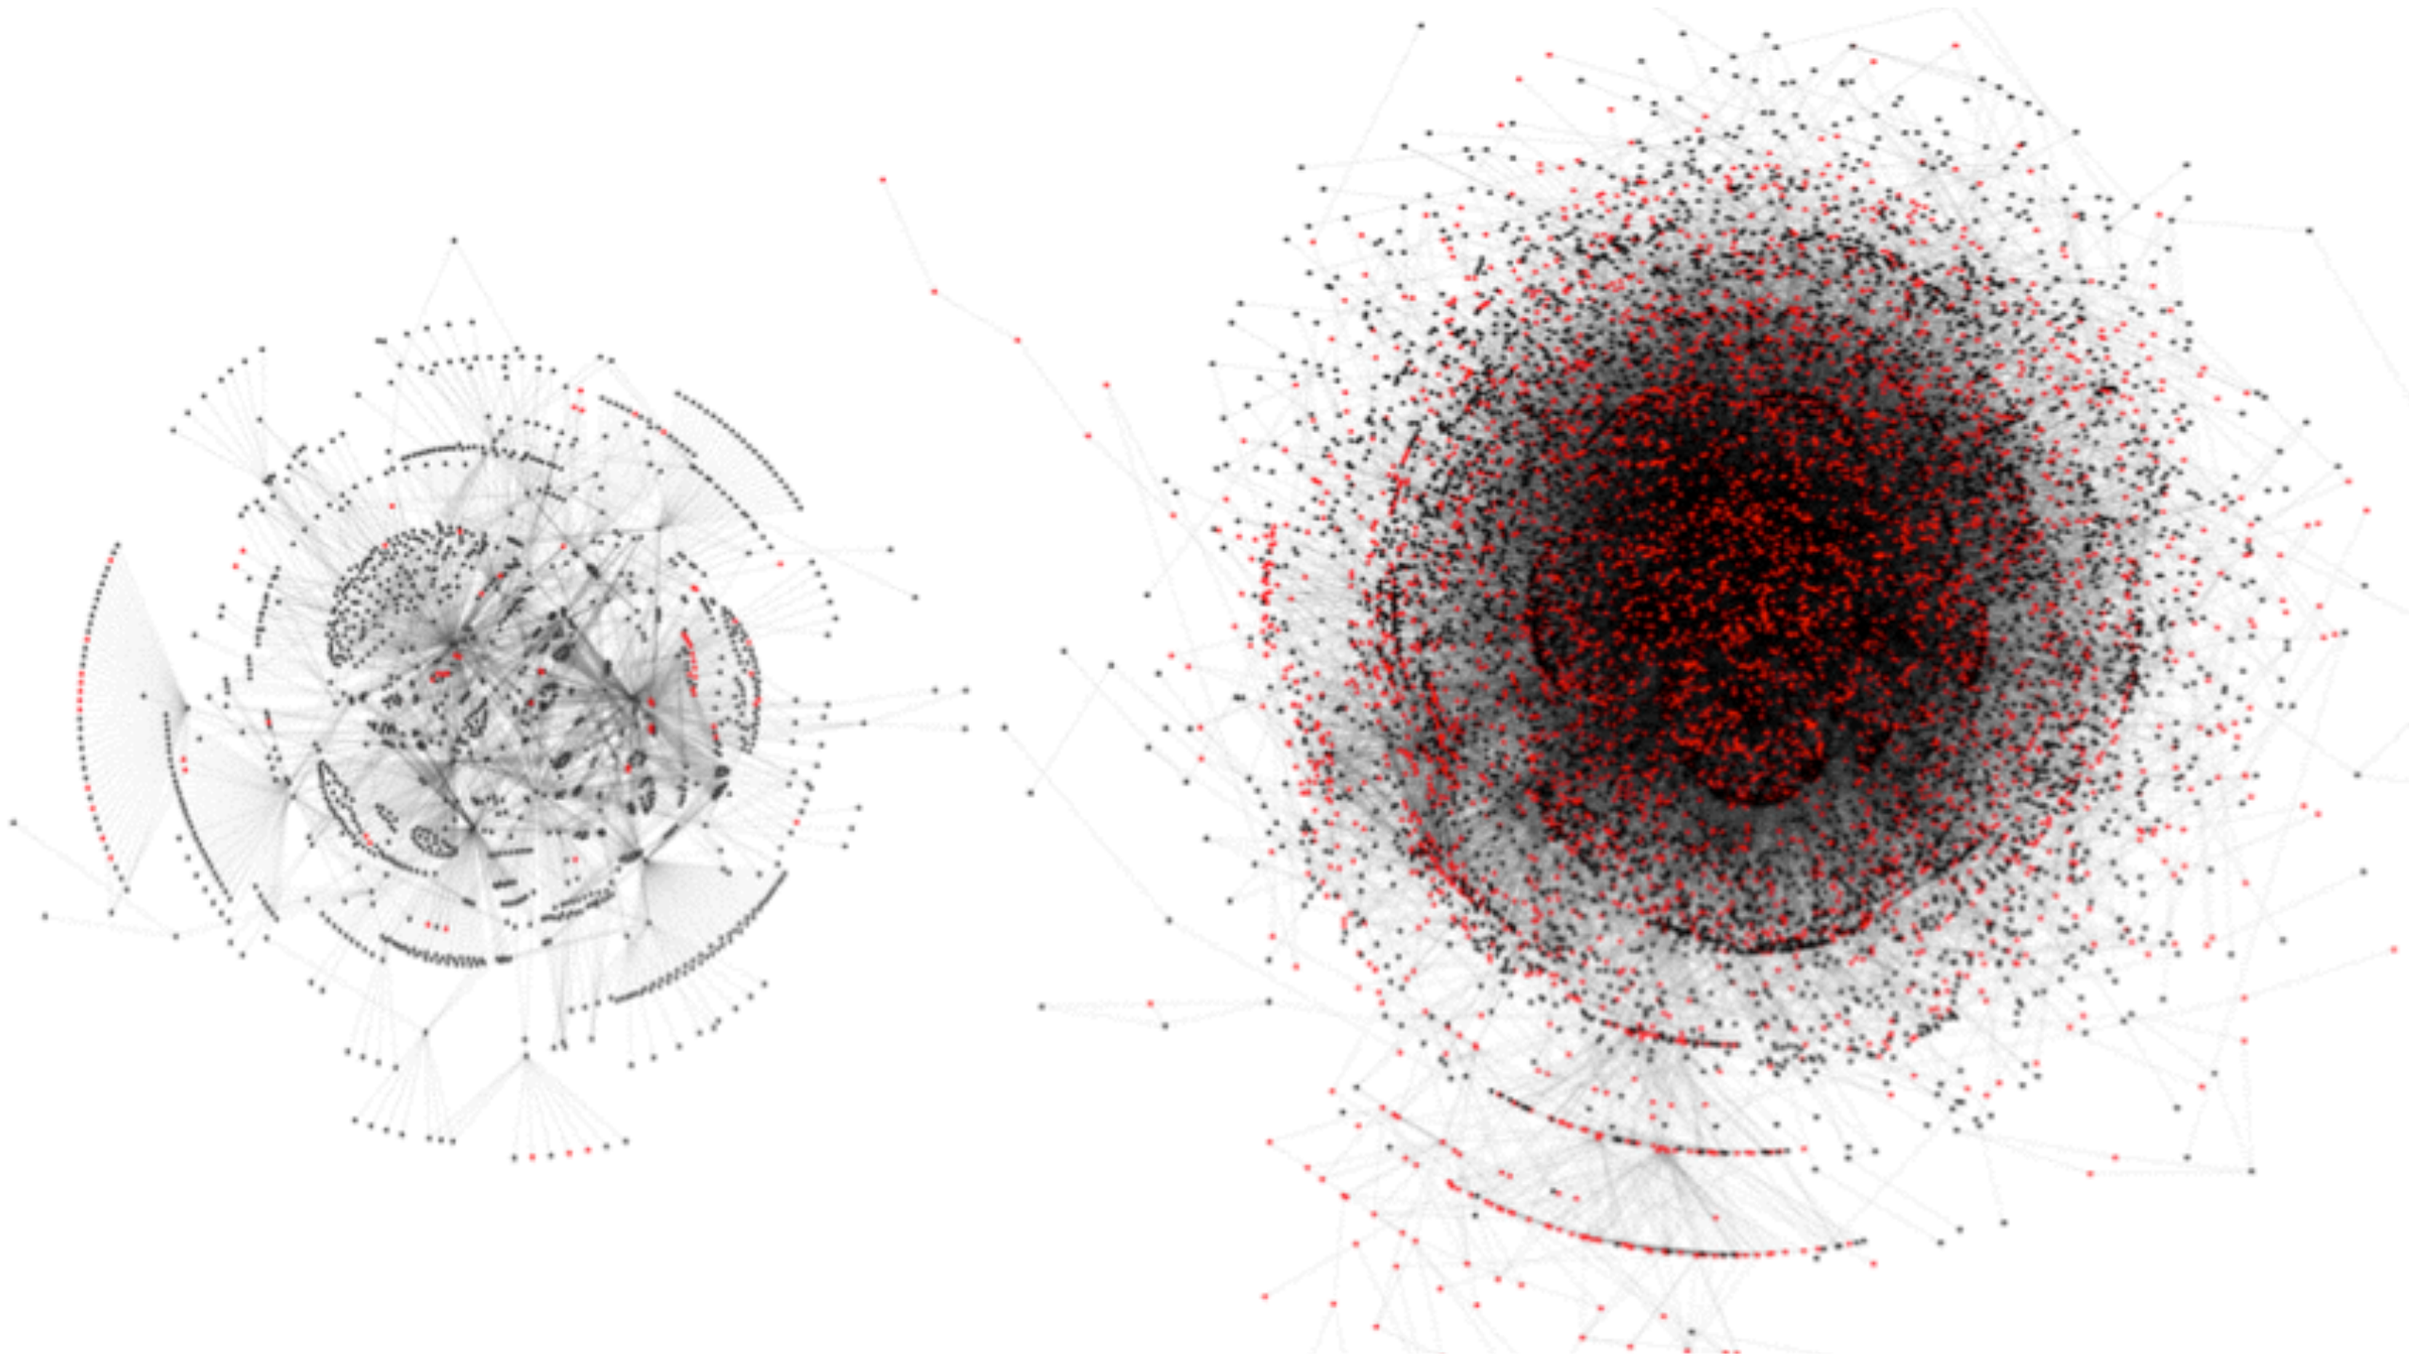

Martin Krzywinski

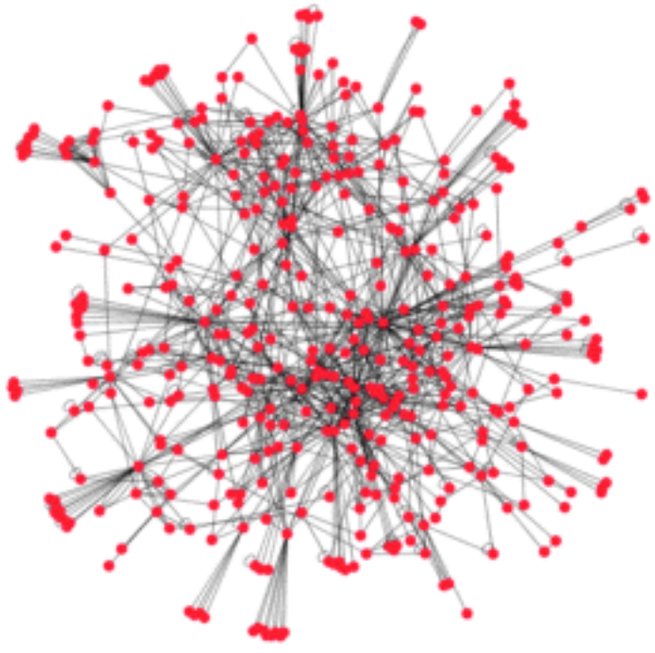

EDGE WEIGHTED
SPRING EMBEDDED

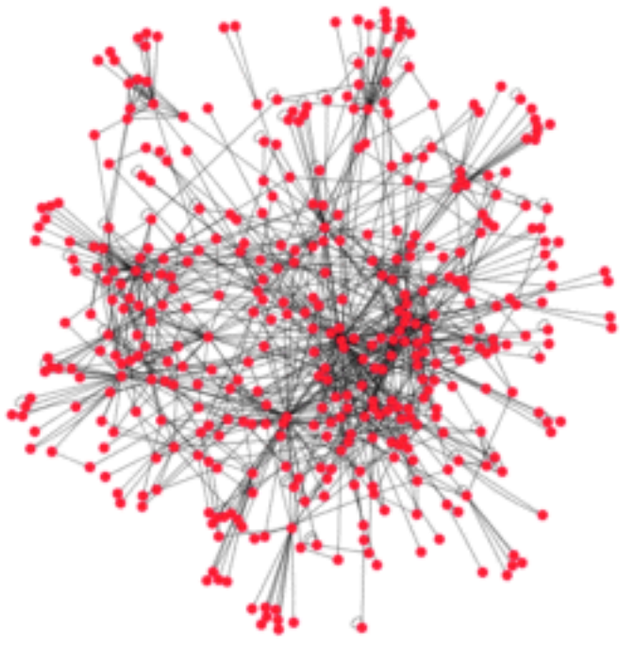

SPRING EMBEDDED

same network

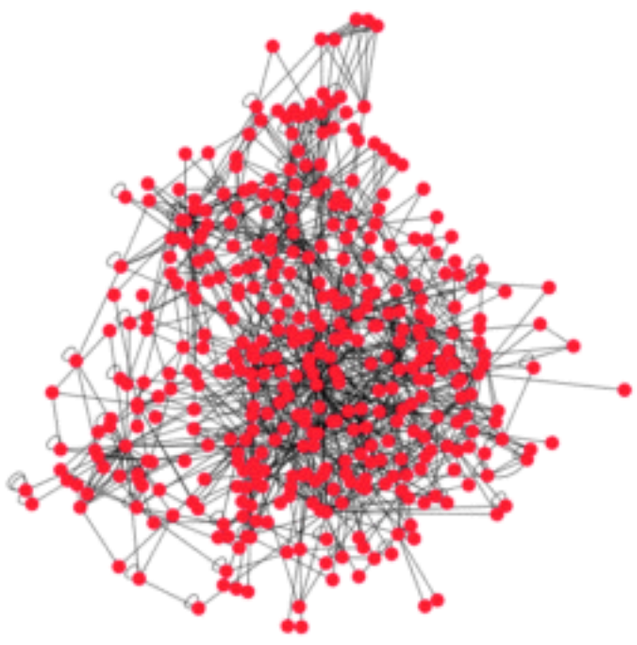

ORGANIC

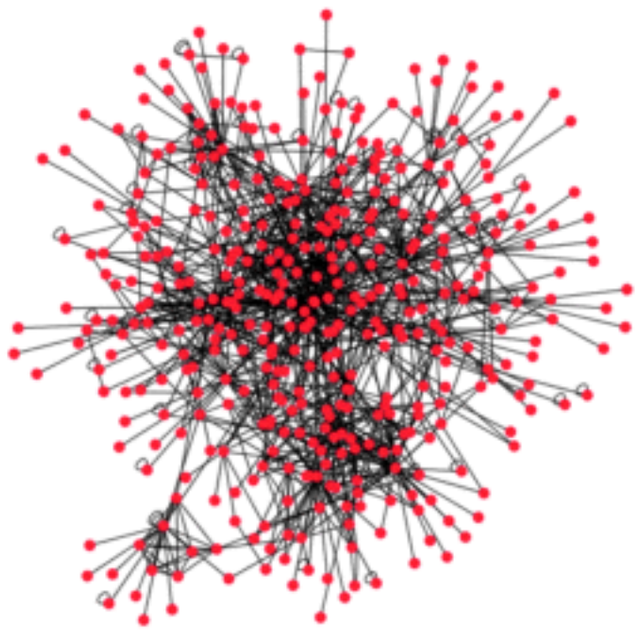

FORCE DIRECTED

- adjacency matrix

- compressed adjacency matrix (Dinkla *et al*, 2012)

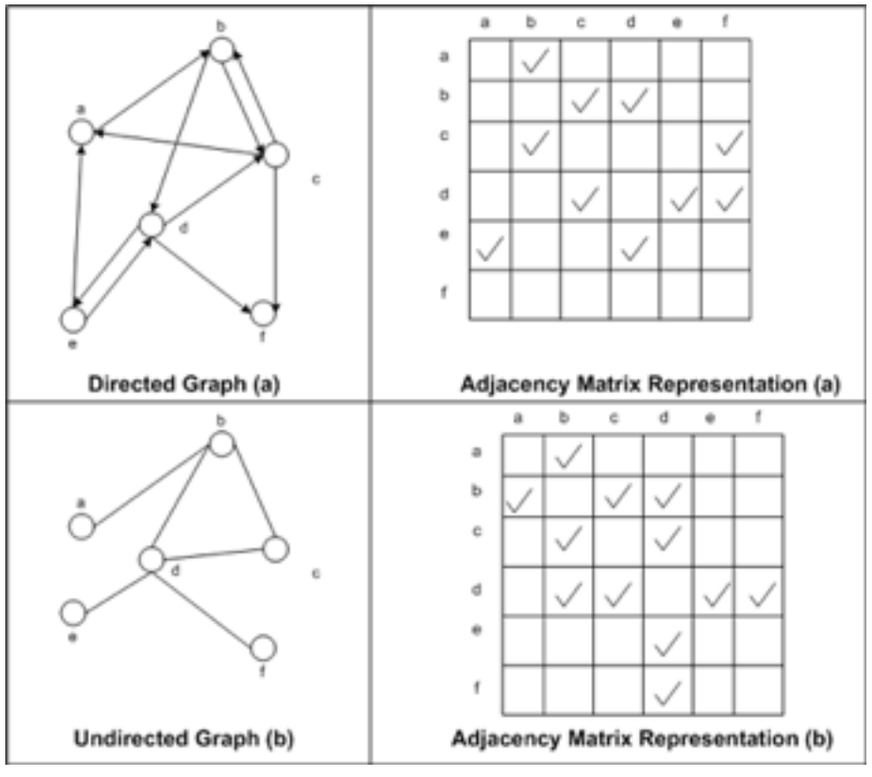

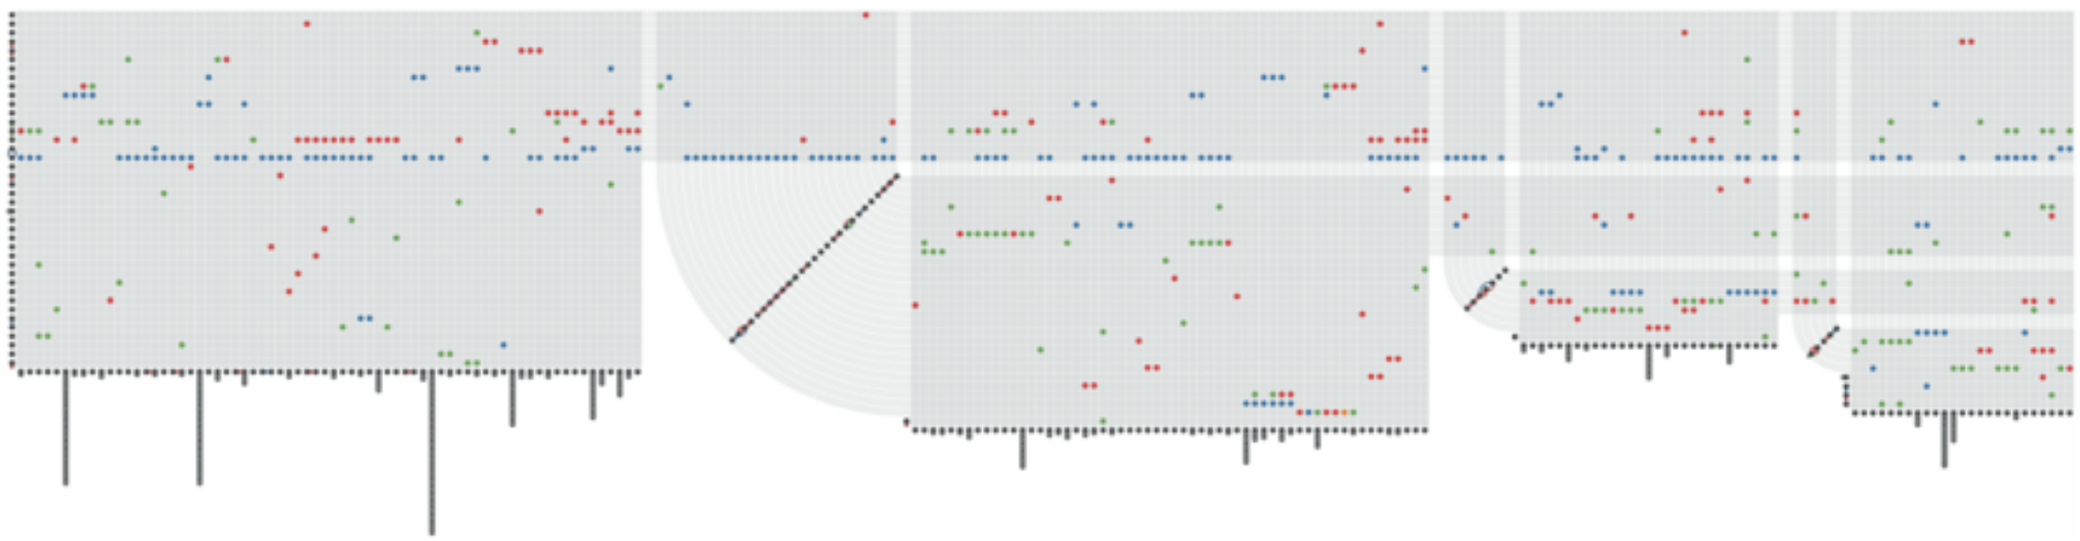

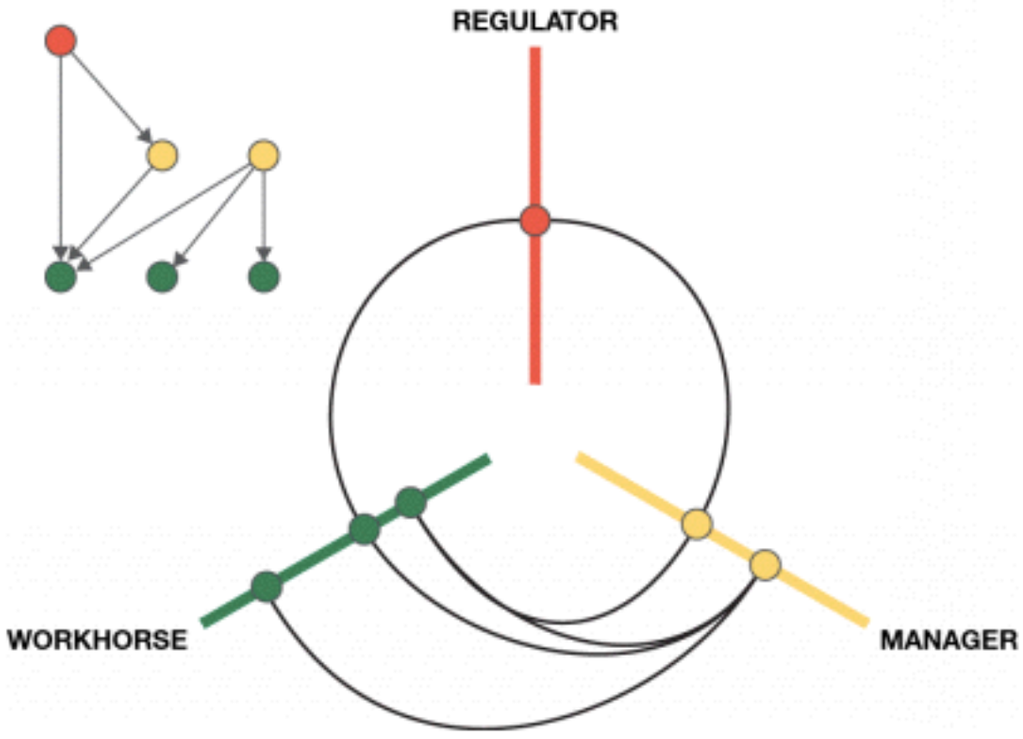
REGULATOR
WORKHORSE
MANAGER

- hive plots

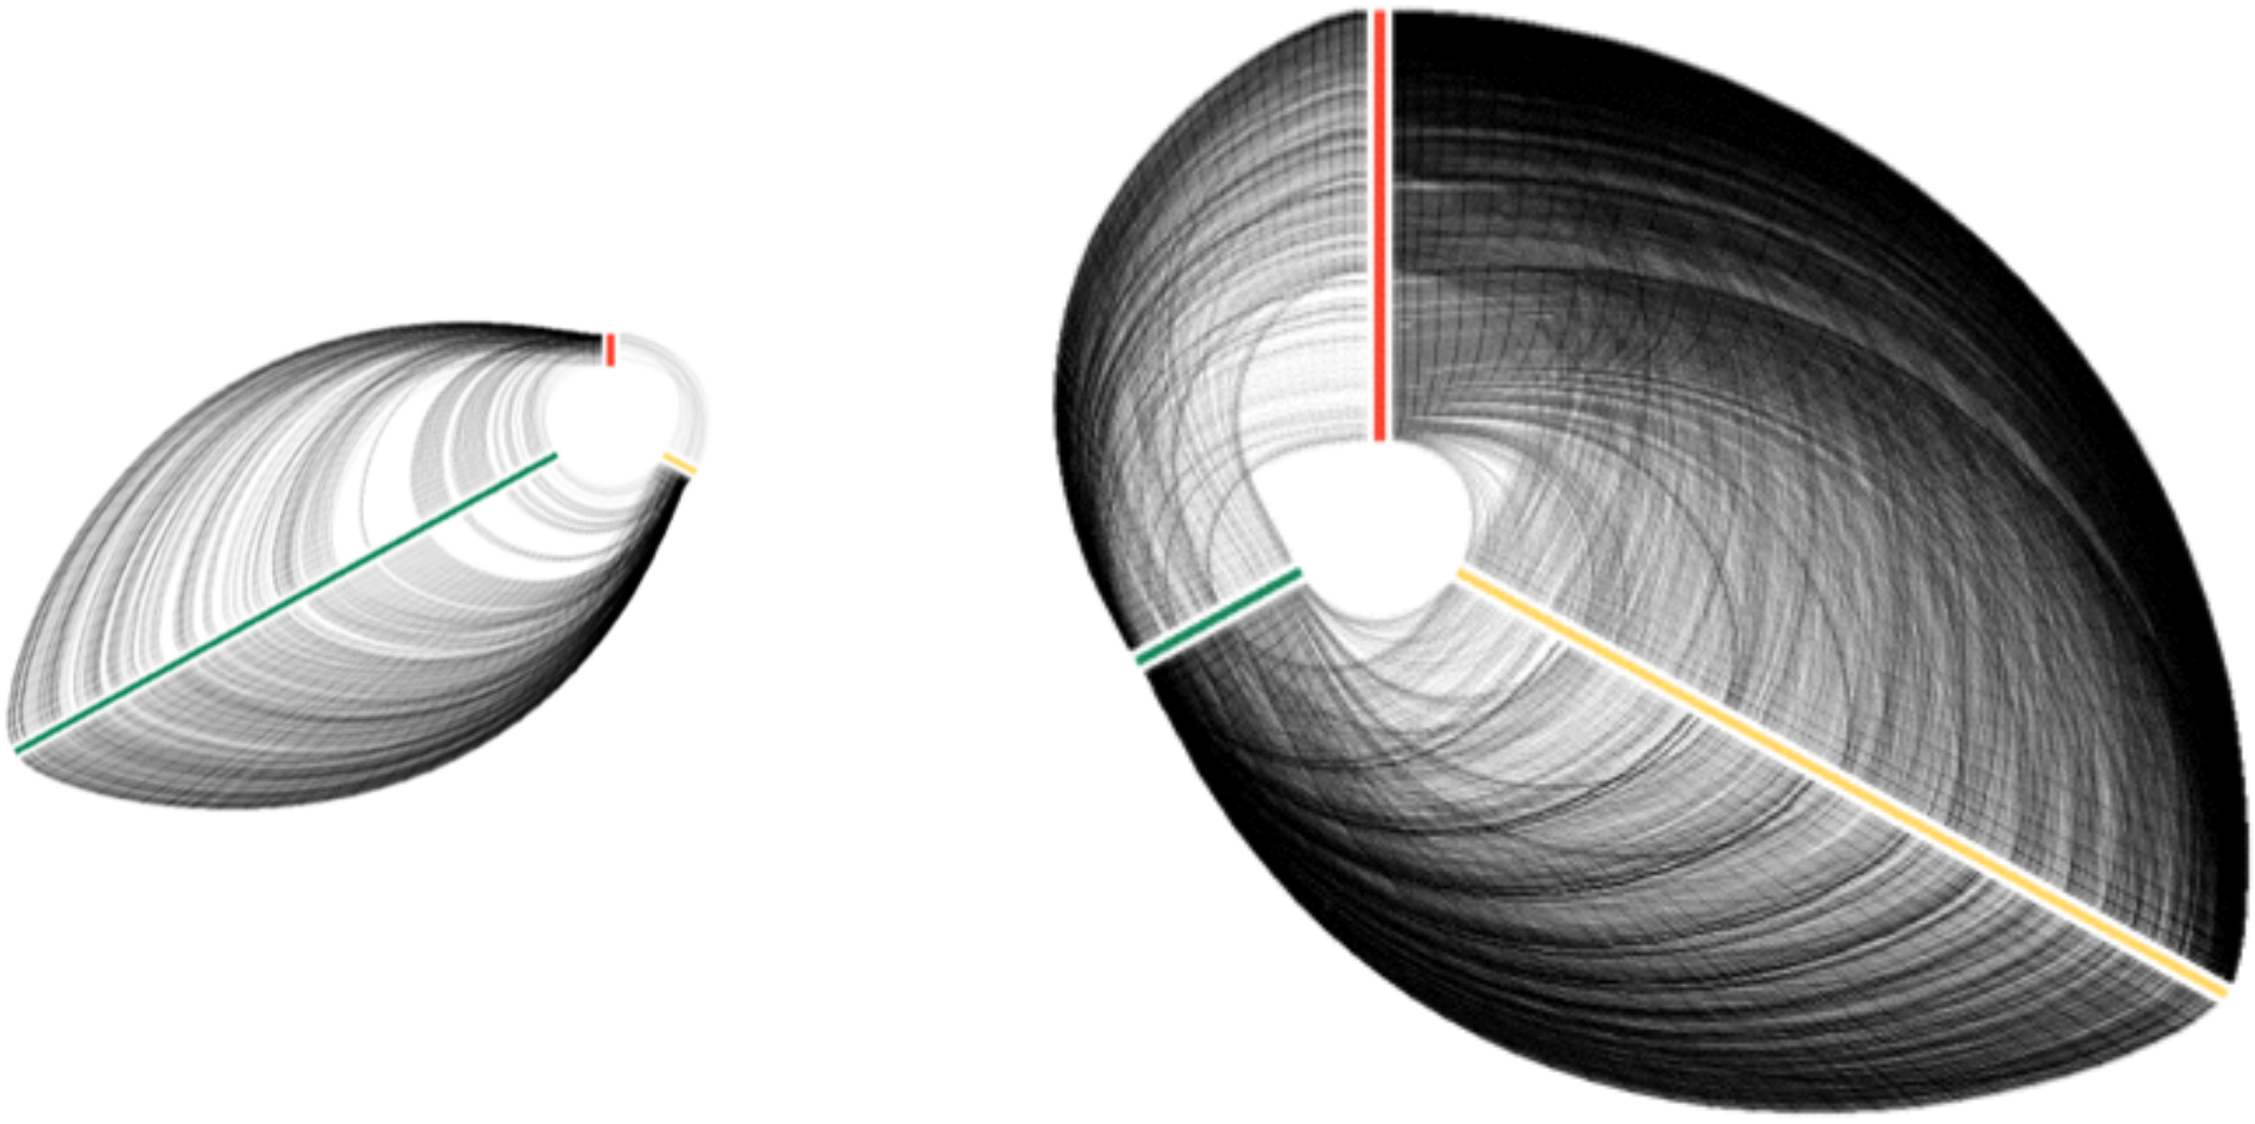

gene interaction networks in different species

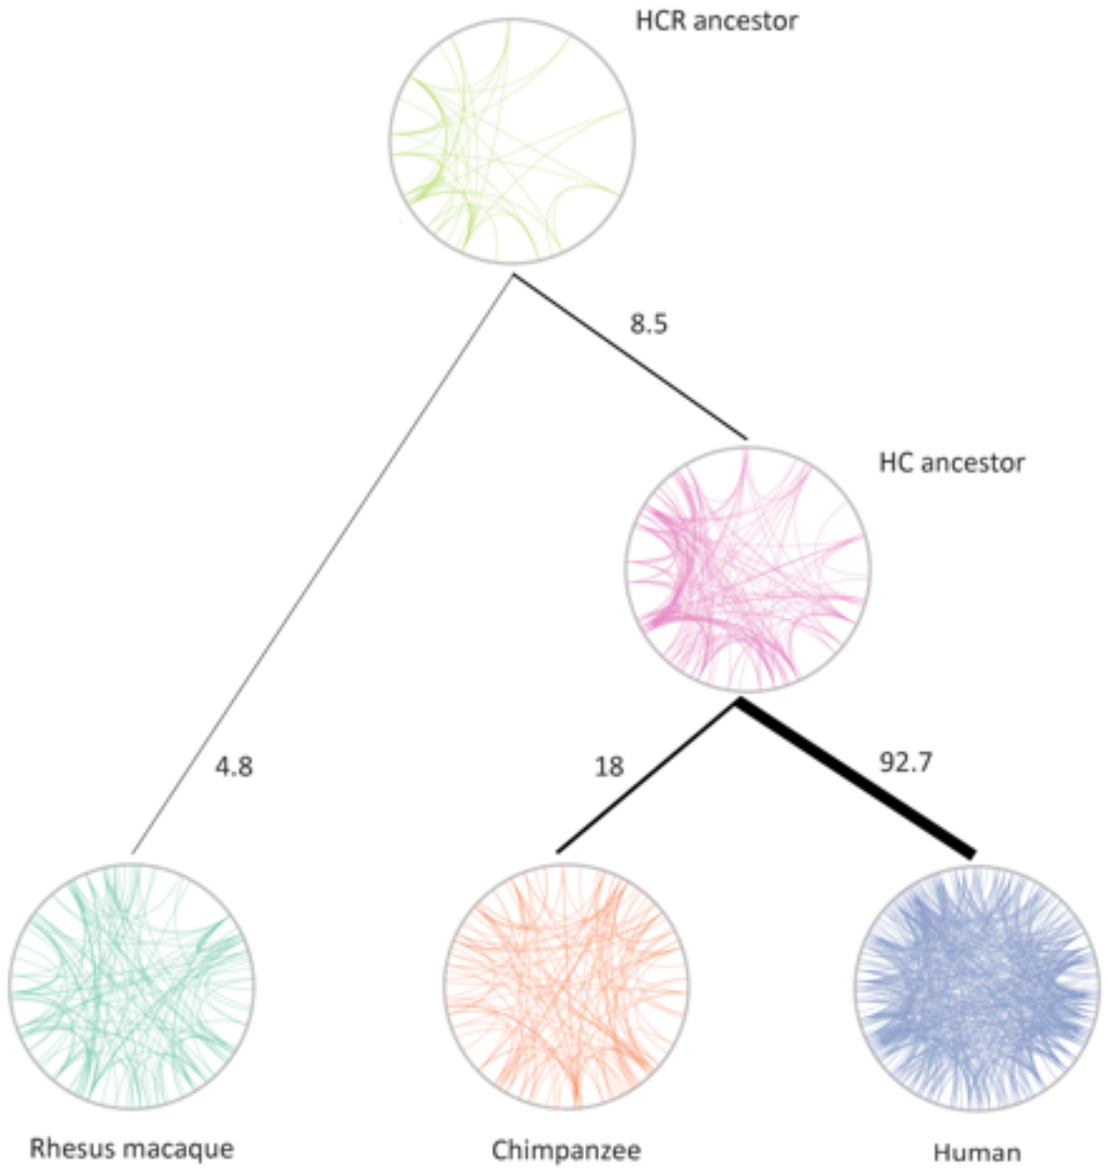

Network comparison gene interaction

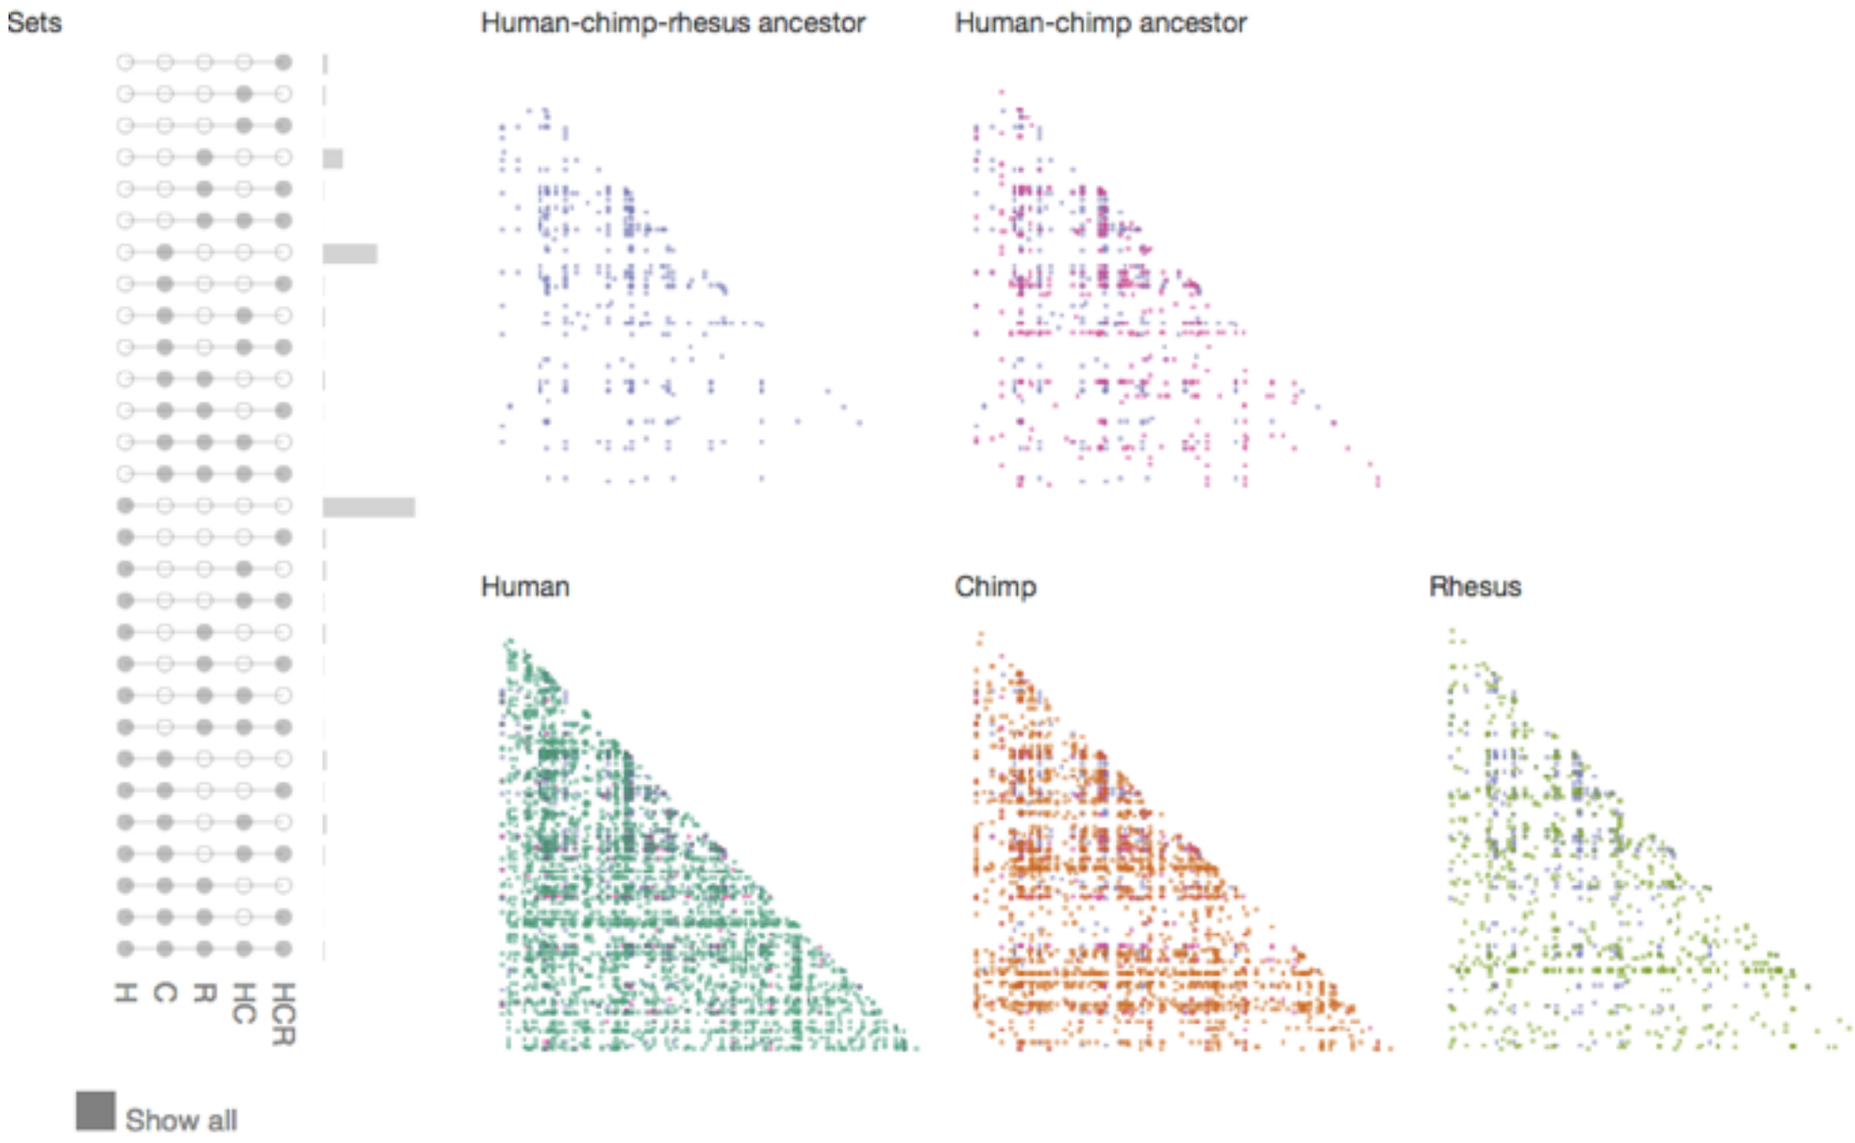

Parallel Edge Splatting (Burch et al, 2011)

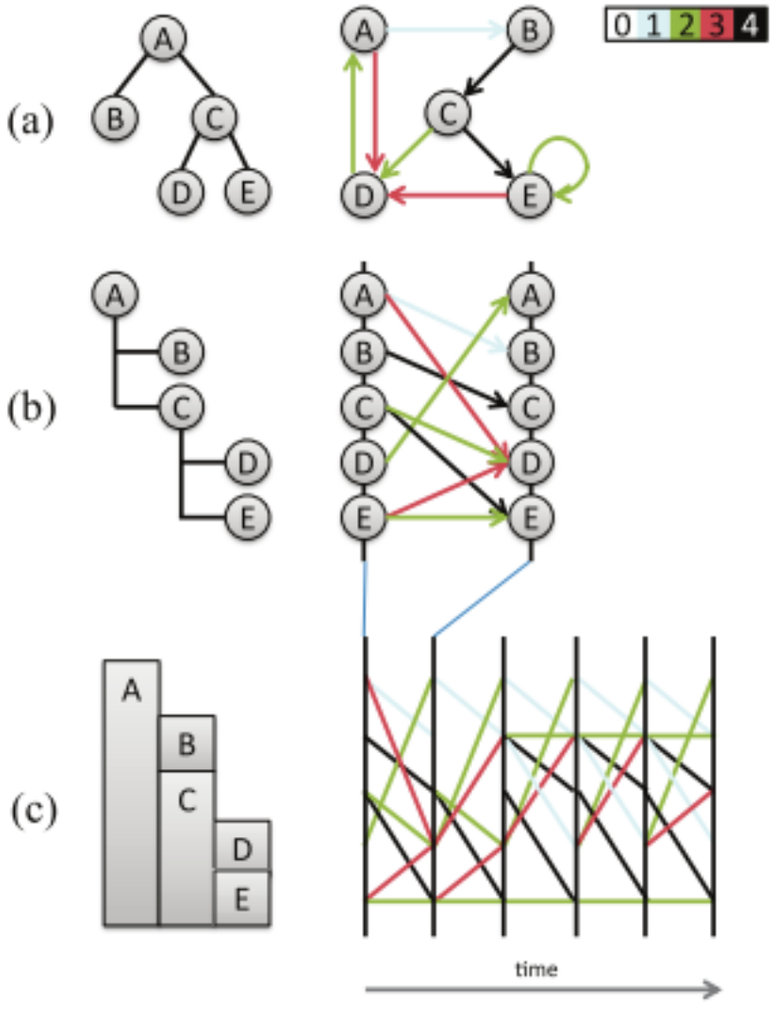

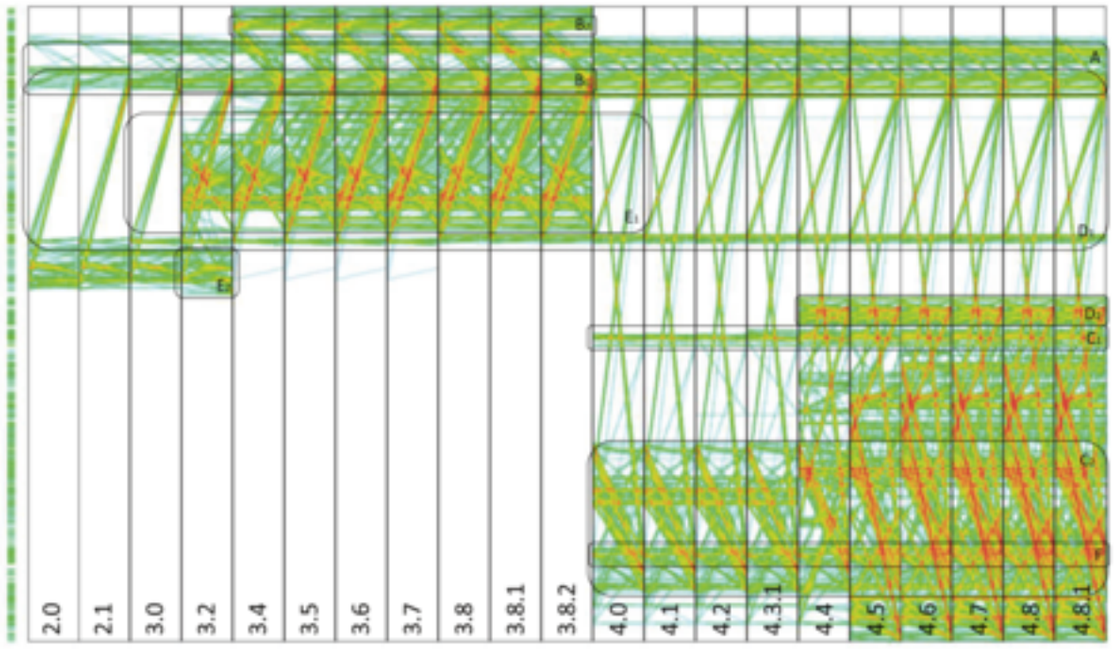

dynamic networks

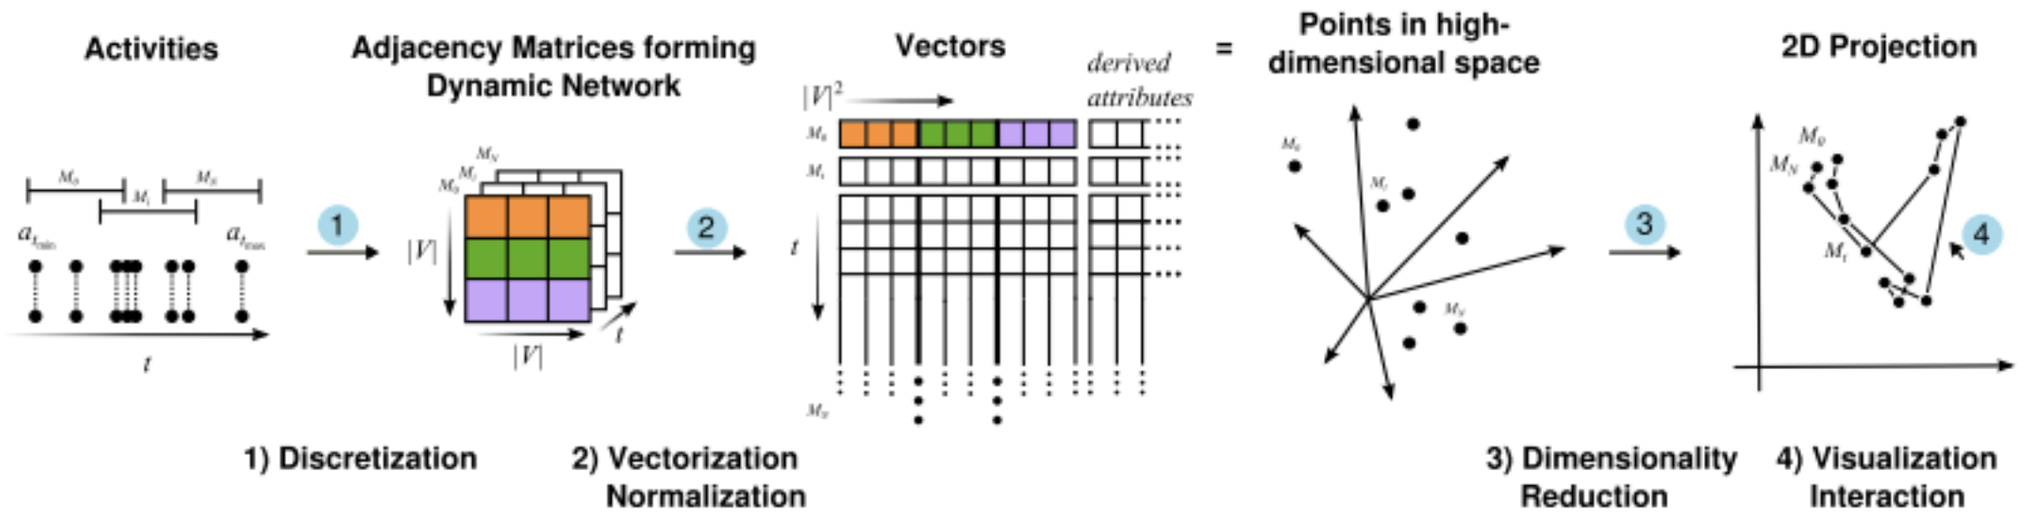

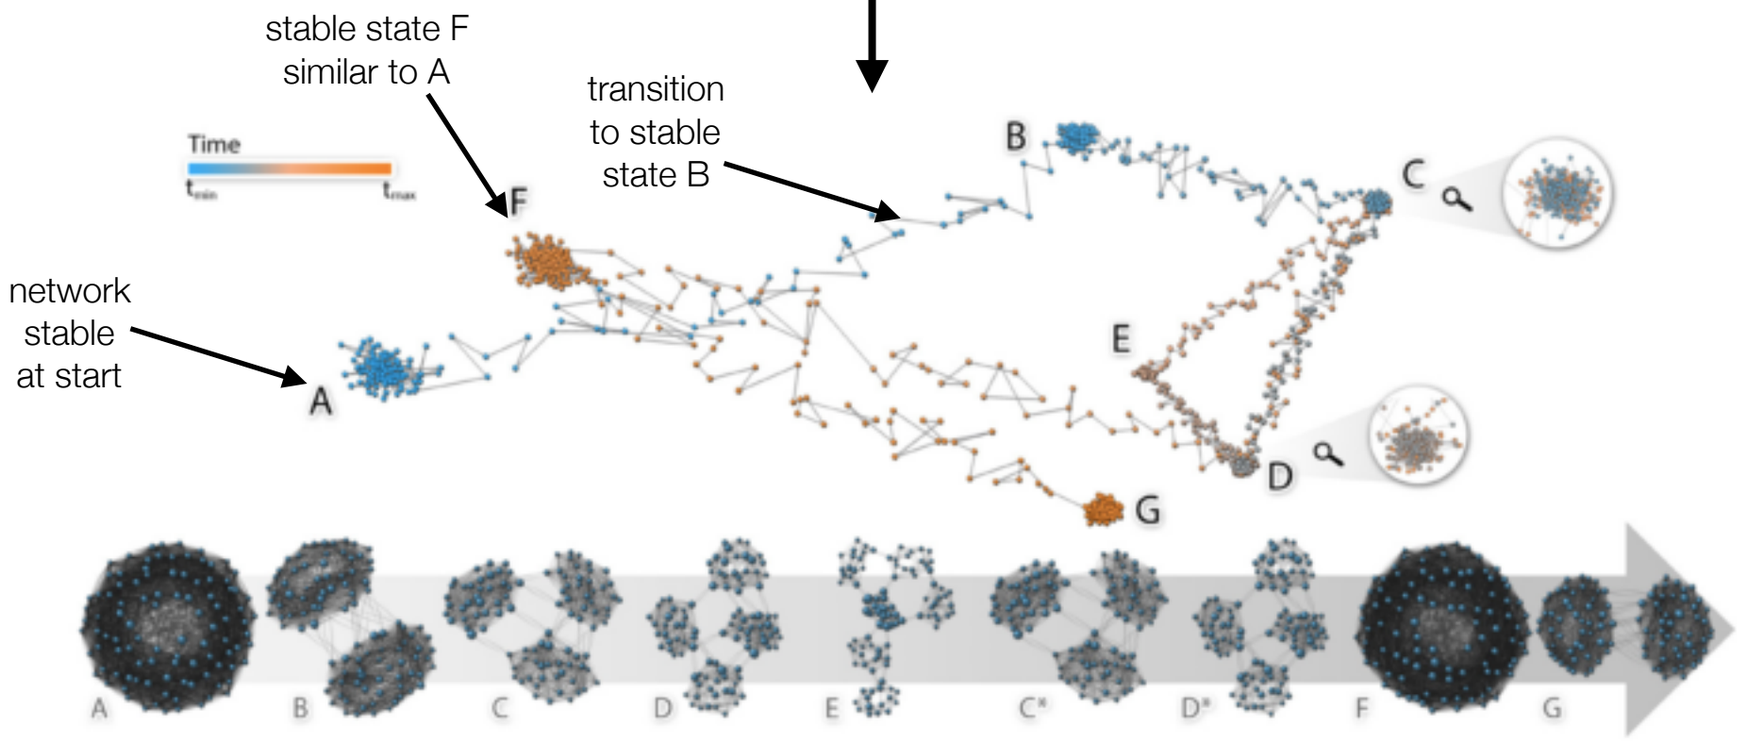

van den Elzen et al, 2015

continuous signal -> what should my histogram bin-sizes be?

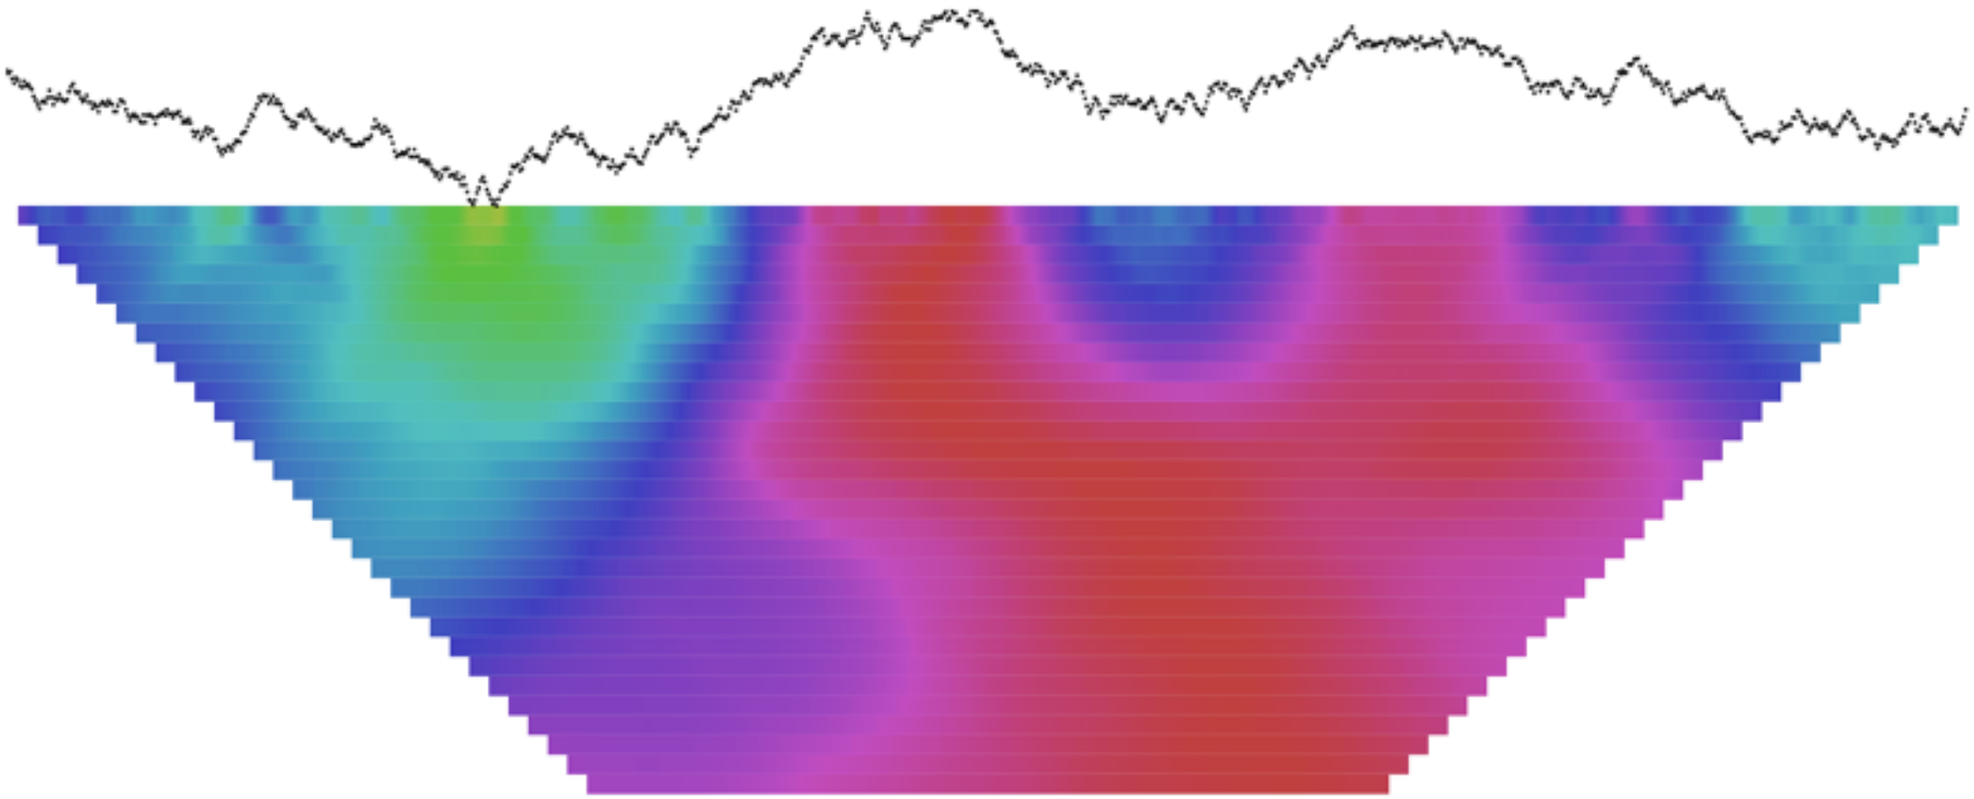

abyssExplorer - making sense of sequence assemblies

Nielsen et al, 2009

Aracari

Bartlett C et al. BMC Bioinformatics (2012)

2-locus eQTL data

Ryo Sakai

Sequence Diversity Diagram

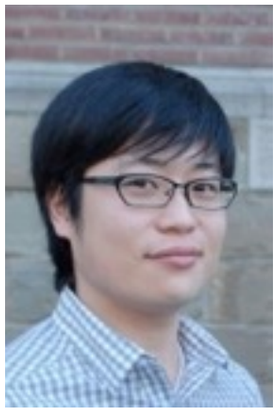

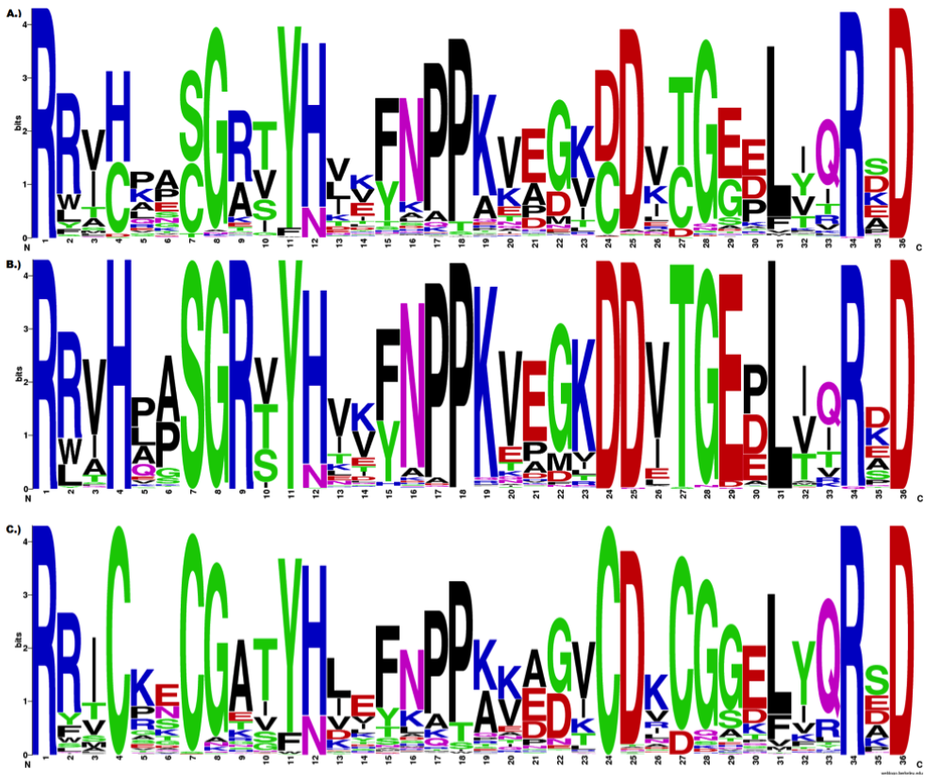

subgroup

similarity

difference

Seagull

sequence diversity diagram

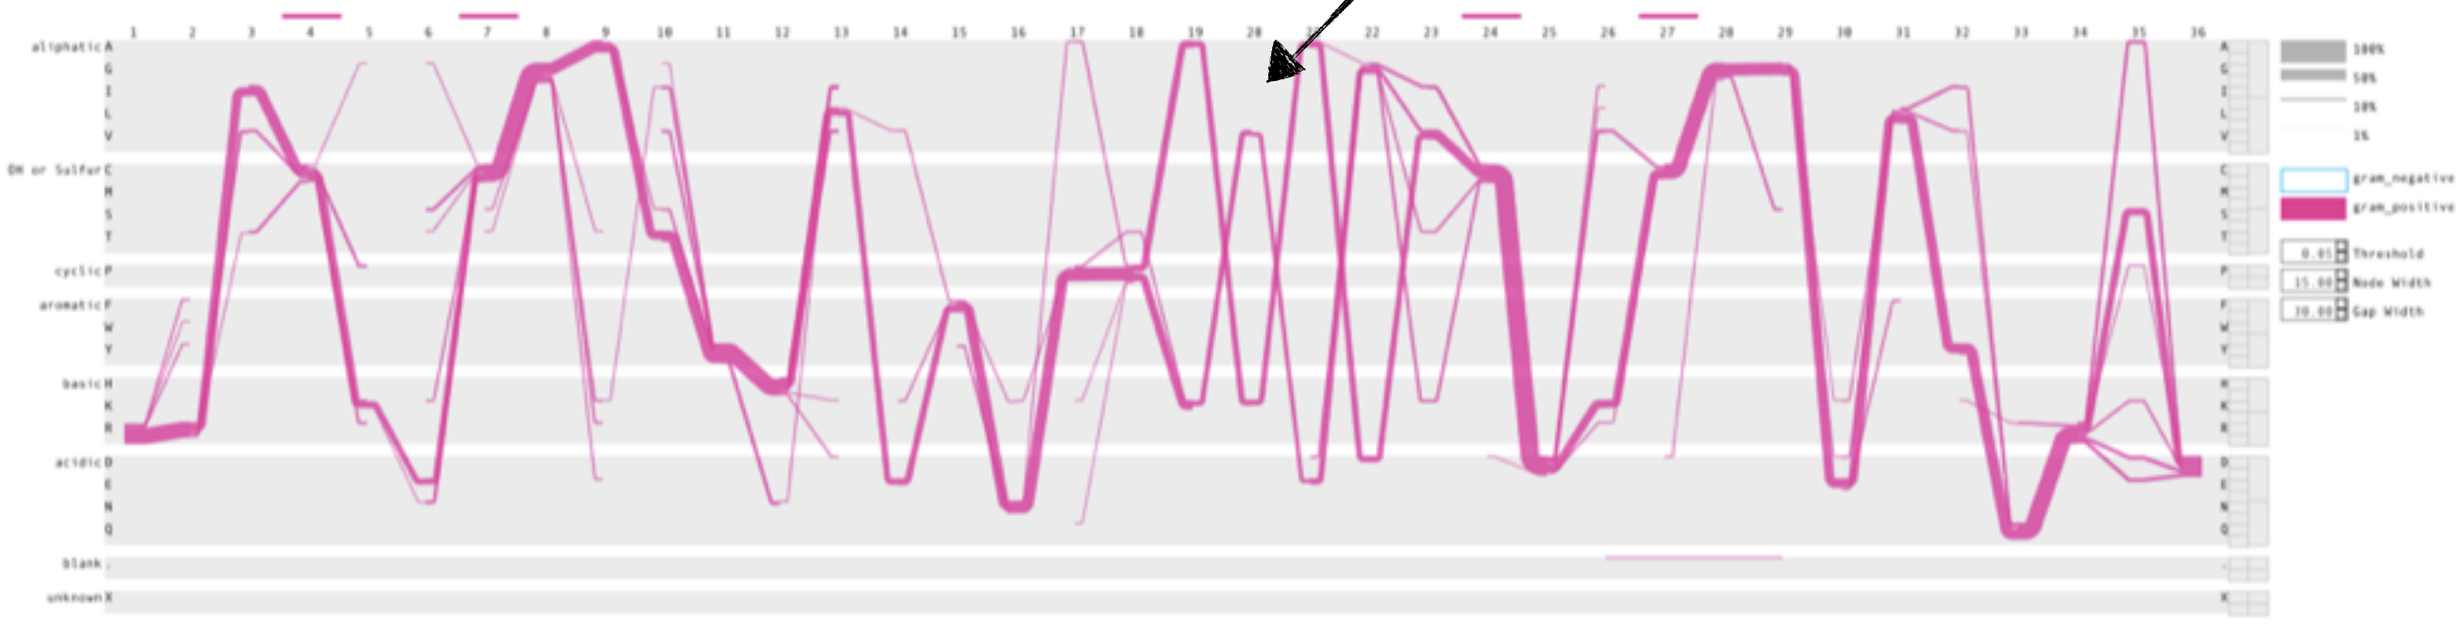

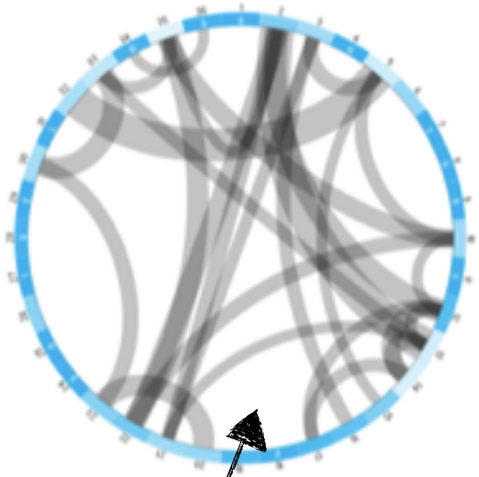

mutual information

3D view

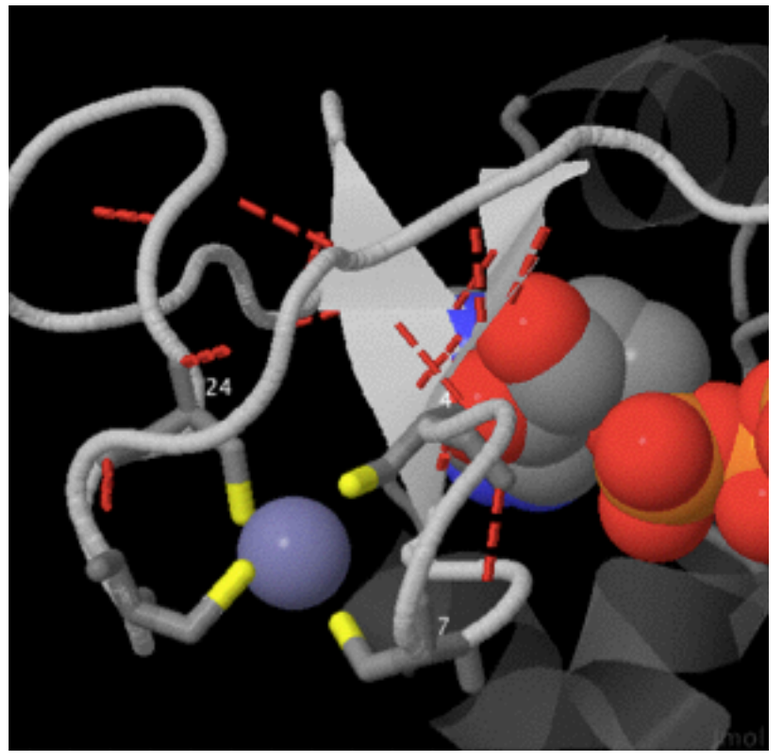

Effect of structural genomic variation

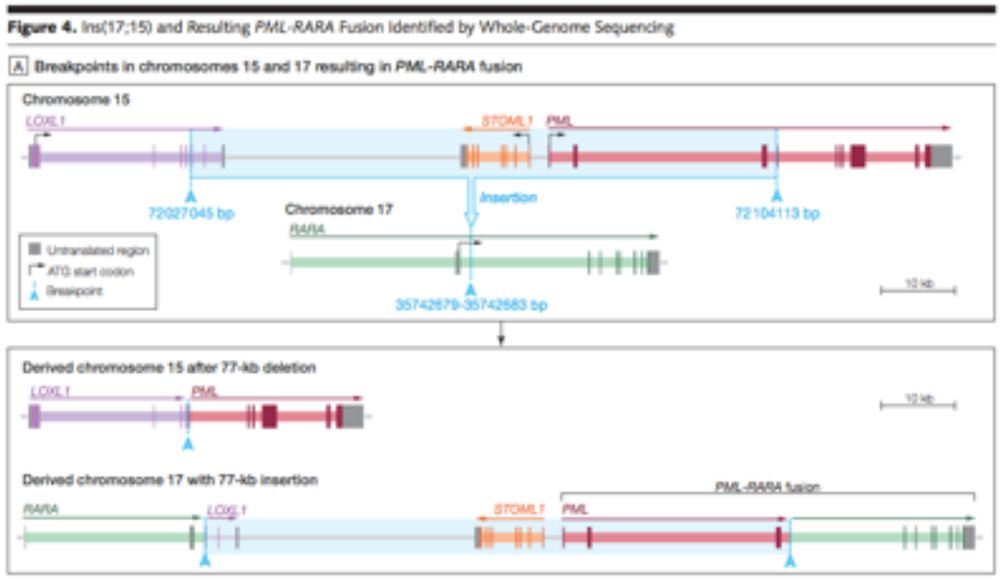

Welch et al, 2011

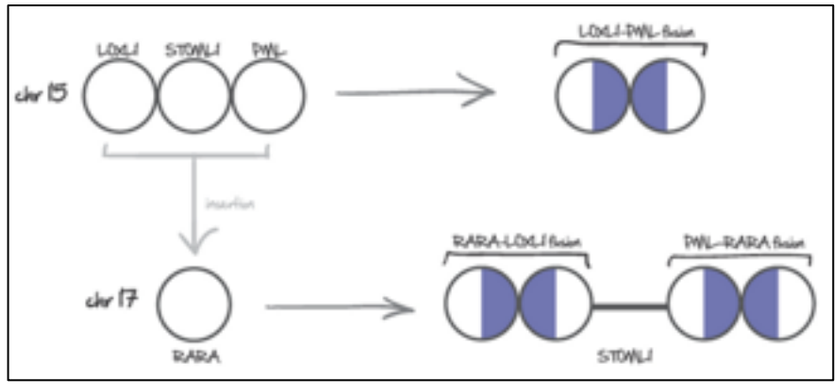

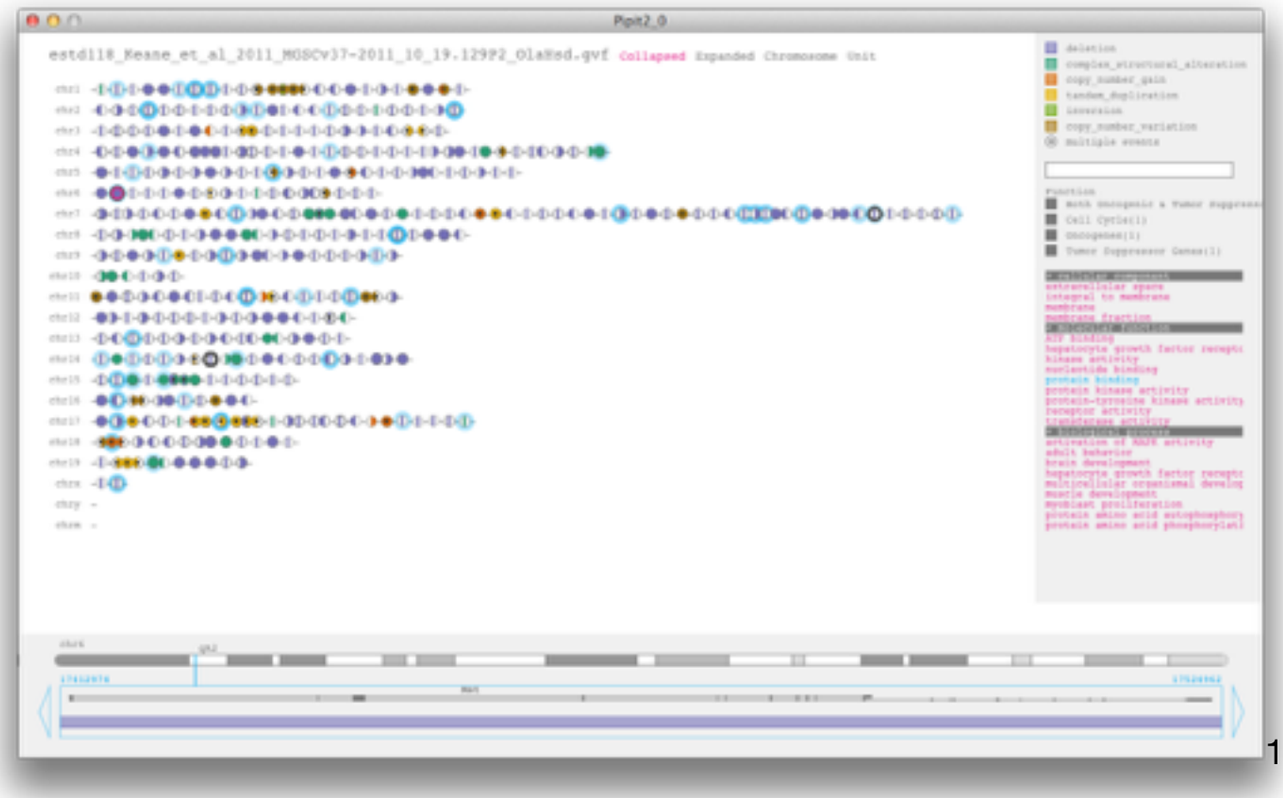

interested in *effect*, not *proof*

collab. Matthieu Moisse & Joke Reumers
University Hospital Leuven
Janssen Pharmaceuticals (J&J)

Overview

A. Why visual analytics?

B. Data visualization

- Data foundations
- Human perception foundations
- Visualization foundations and examples

C. Visualization evaluation

D. Tools of the trade

C. Visualization evaluation

Quantitative evaluation

- spike data with known signal, and record time that it takes for user to find that signal => measure => run statistics “this visualization is better than that one”
- user tasks:
  - identify
  - locate
  - distinguish
  - categorize
  - cluster
  - rank
  - compare
  - associate
  - correlate

Quantitative evaluation

- spike data with known signal, and record time that it takes for user to find that signal => measure => run statistics “this visualization is better than that one”

problem: small sample size

- user tasks:
  - identify
  - locate
  - distinguish
  - categorize
  - cluster
  - rank
  - compare
  - associate
  - correlate

Qualitative evaluation

- very close interaction with domain expert
- let expert use the interactive visualization and try to find out what insights he/she gained from the visualization
- experimenter observation
- think-aloud protocol
- collecting participant opinions

Make sure you measure the right thing

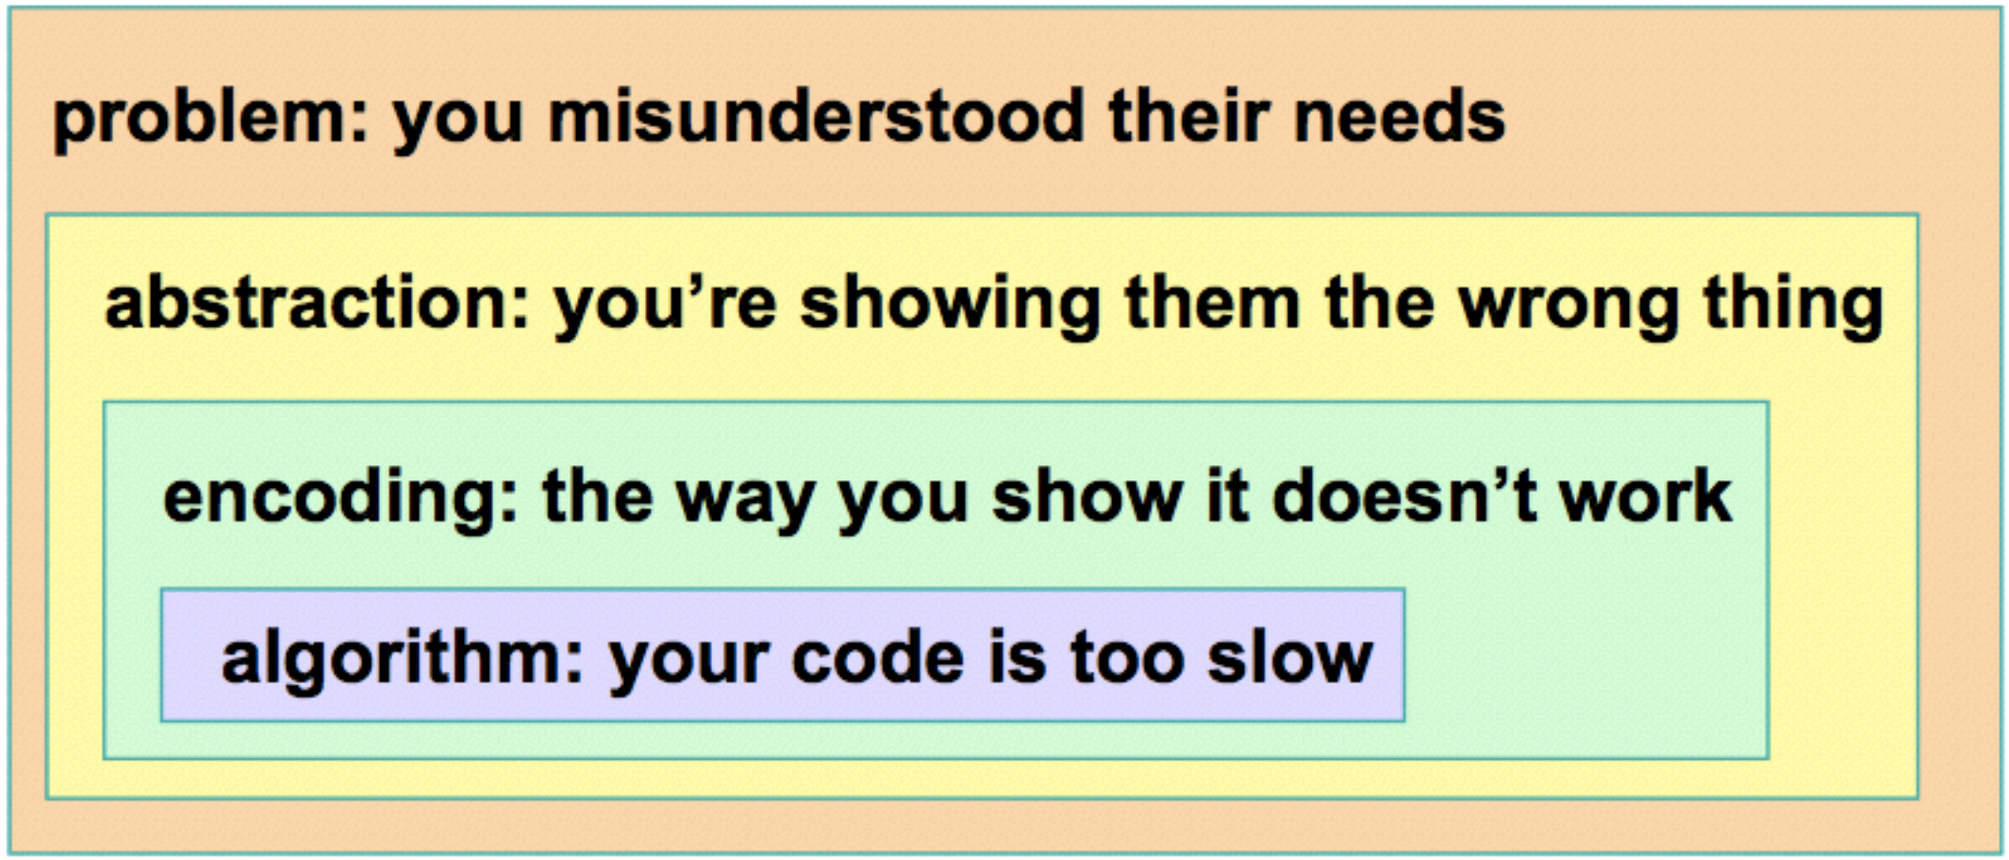

Overview

A. Why visual analytics?

B. Data visualization

- Data foundations
- Human perception foundations
- Visualization foundations and examples

C. Visualization evaluation

D. Tools of the trade

D. Tools of the trade

using
drawing
coding

Using data visualizations

Microsoft Excel

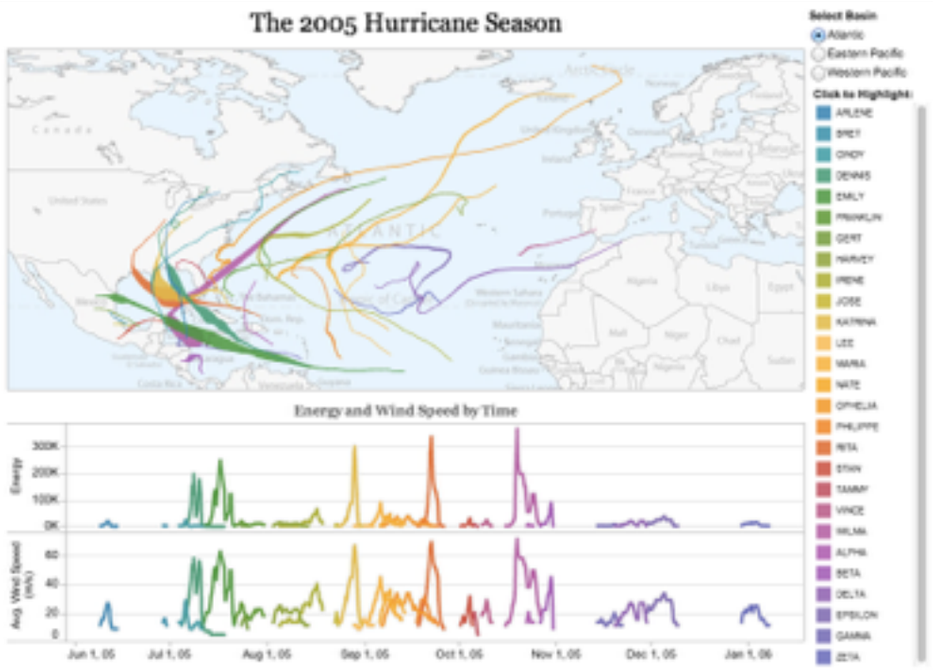

Tableau

Tibco Spotfire

ManyEyes

Drawing data visualizations

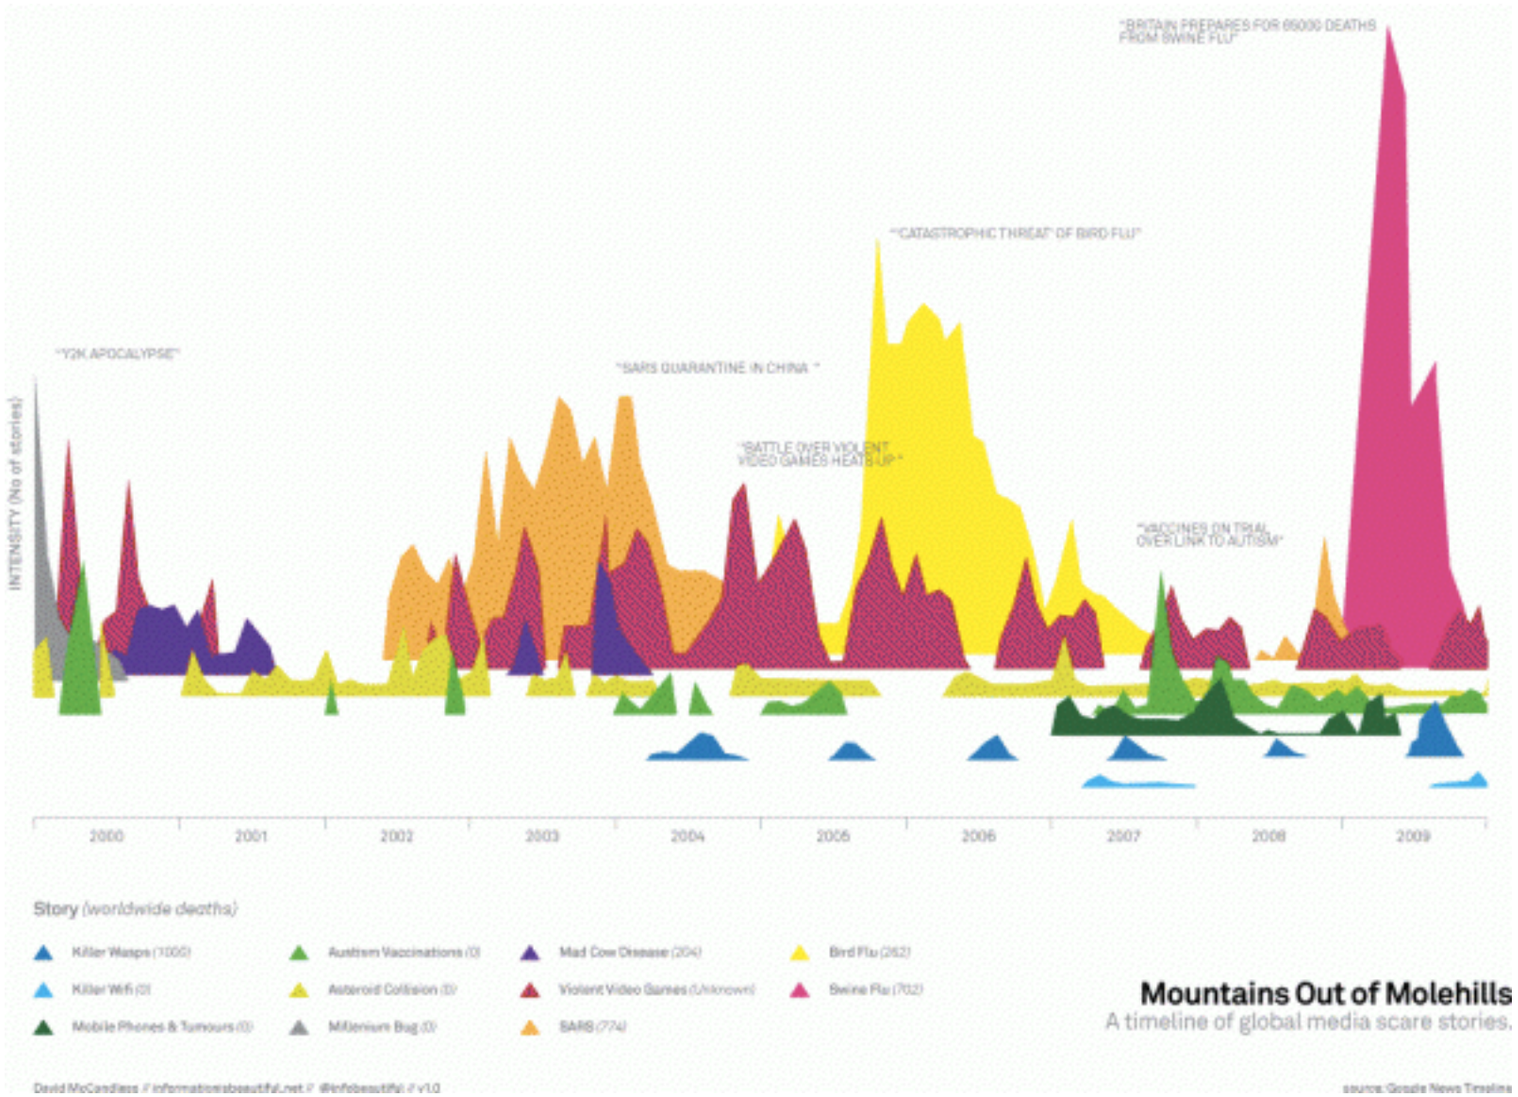

Adobe Illustrator

Coding data visualizations

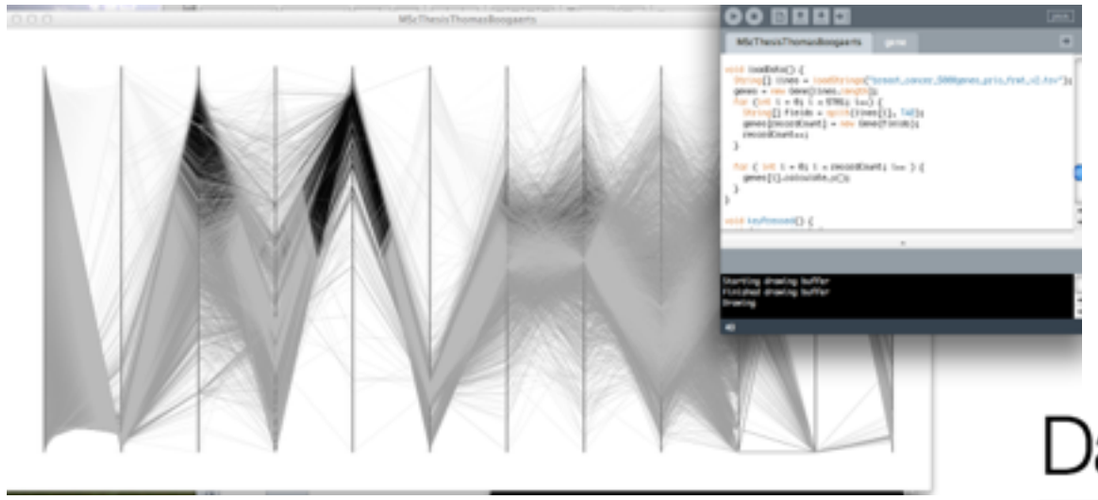

processing.org

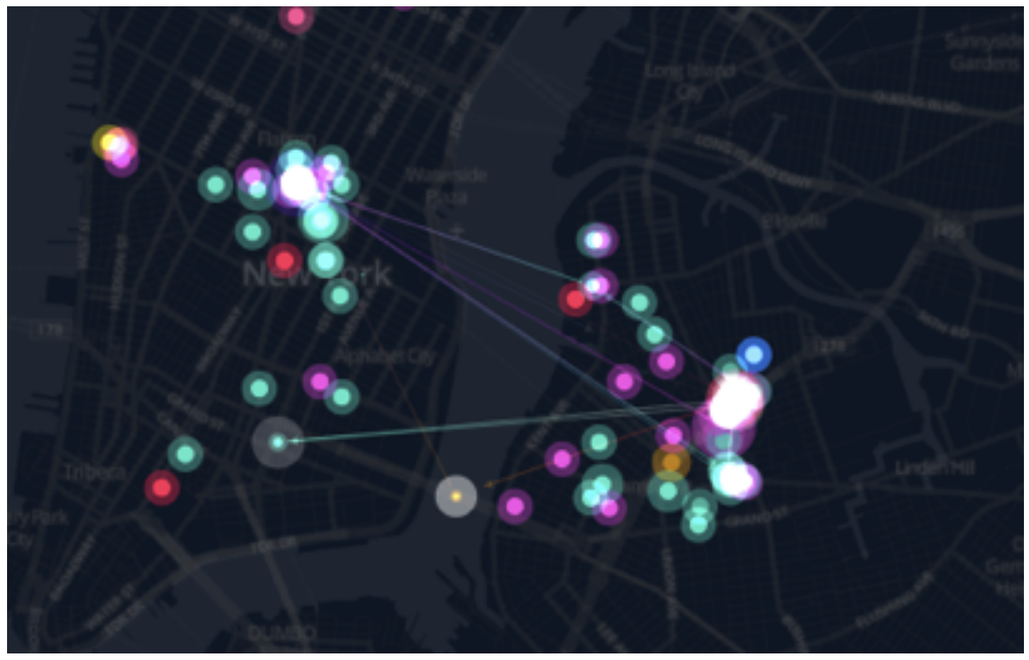

paper.js

Data-Driven Documents

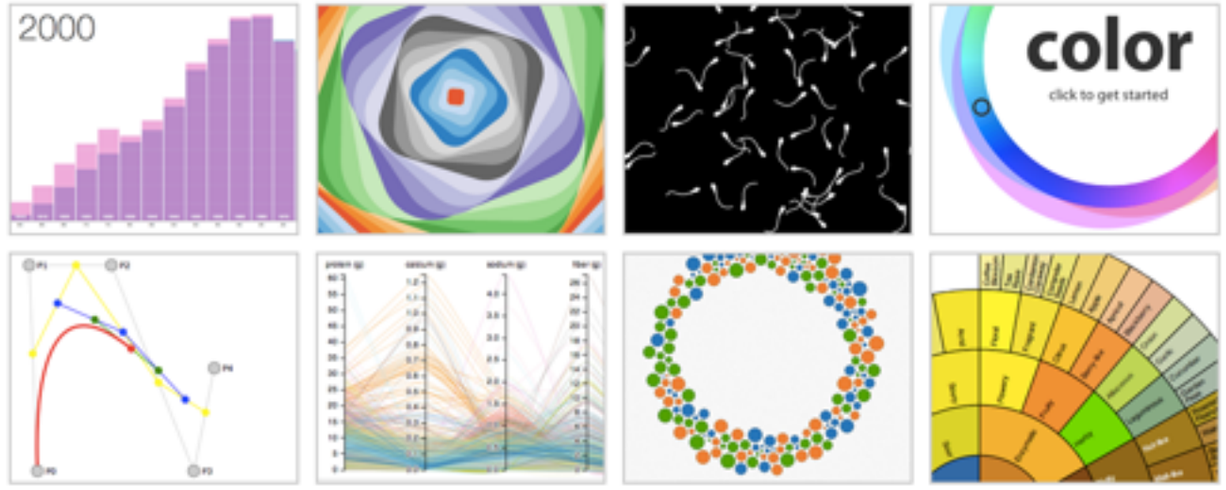

d3.js